
Control Is Motion

Geometry, Feasible Dynamics, and Robust Outcomes

D. Kraft

Chapter 1

Motion, Targets, and Timing

A golf swing is a motion undertaken for an outcome. The club must arrive at contact with a useful position, velocity, orientation, and impact location; the subsequent collision and flight determine what the ball does. The body must produce that delivery through a feasible movement, tolerate relevant variation, and continue into follow-through. Studying only a static pose misses much of this problem. Discarding the outcome would miss its purpose.

“Control is motion” is the organizing perspective of this book. It directs attention to how a trajectory is generated, shaped, and made reliable. It is not a claim that trajectory tracking, optimal control, or orbital stability replaces classical control theory. Those established tools already address moving references, periodic behavior, disturbances, and constrained objectives. The task is to choose and combine them correctly for the question at hand.

The important connections are physical. Earlier applied work changes the state available later. Configuration and velocity change inertial coupling, force transmission, and external loads. Grip and ground constraints restrict compatible motion and affect reactions. Timing changes velocity and effort even along an unchanged configuration path. Feedback and constitutive response alter how disturbances develop. Impact samples the resulting state at an event, which need not occur at a fixed clock time. A rigorous account keeps these connections while distinguishing a model identity from a claim about a golfer.

1.1 Specify What the Motion Should Produce

Let the state be x , the effective input be u , and the declared dynamics be

$$\dot{x} = f(x, u, t), \quad u \in \mathcal{U}(x, t). \quad (1.1)$$

The state must contain enough information for the model to predict its own continuation. For a rigid mechanical model it includes configuration and velocity. A muscle-driven or flexible model may also need activation, tissue, shaft, and other internal states. A list of joint angles alone is not generally a complete dynamic state. Contact modes and their transition rules are part of the model as well.

Several objectives are useful and can coexist. Regulation keeps a quantity near a desired constant value. Trajectory tracking compares the state or output with a time-indexed reference. Path following permits progress along a geometric path with a separately specified timing objective. Terminal or event objectives constrain the state delivered at a particular time or crossing. Periodic tasks may call for orbital stabilization. The names describe different specifications; they are not mutually exclusive schools of control.

For a simplified impact investigation, let the first intended approach crossing be described by

$$h(x(t_e)) = 0, \quad \nabla h(x(t_e))^T f(x(t_e), u(t_e), t_e) \neq 0, \quad (1.2)$$

with an appropriate approach direction and contact mode. The nonzero derivative is a transversality condition: it excludes a grazing crossing where event time can be especially sensitive. Define an output $y_e = H(x(t_e^-))$ from the pre-impact state and a target set $\mathcal{Y}$ of acceptable deliveries. Its components might describe clubhead velocity, face orientation, path direction, and impact location. A ball-outcome model can map those quantities, together with ball and contact properties, into launch and flight predictions.

A terminal requirement $y_e \in \mathcal{Y}$ is compatible with a moving club. It does not demand convergence to an equilibrium with nonzero velocity, nor require that the club accelerate through the collision. The impact model must account for the rapid change in momentum. Follow-through and any subsequent constraints require their own continuation. The event model, output definition, and admissible set determine what constitutes success.

This distinction is practical: two motions can have similar joint paths but different delivery speeds, while different joint paths can deliver similar club states. Neither observation by itself identifies how individual muscles produced the motion.

1.2 A Moving Reference Must Be Feasible

A nominal time-indexed trajectory $x_d(t)$ must have a compatible admissible input $u_d(t)$:

$$\dot{x}_d(t) = f(x_d(t), u_d(t), t), \quad u_d(t) \in \mathcal{U}(x_d(t), t). \quad (1.3)$$

Joint limits, closure conditions, contact forces, and internal-state dynamics must also hold. Drawing a smooth curve in state space does not establish these properties. If a tracking error is expected to remain zero, the closed-loop vector field must reproduce the nominal tangent there.

For a control-affine model $f = f_0 + Gu$, define $r = \dot{x}_d - f_0(x_d, t)$. At a given nominal state, an unconstrained input exists only if $r \in \text{im } G$. When that condition holds,

$$u_d = G^\dagger r + (I - G^\dagger G)z \quad (1.4)$$

expresses the family of solutions in chosen coordinates. The pseudoinverse selects a minimum Euclidean-norm representative; that criterion depends on input scaling and is not a muscle-recruitment law. Input bounds can eliminate all members of the family. A nonlinear input map may require a different solution procedure and still need not have an inverse.

For example, $G = (1, 0)^T$ cannot produce $r = (0, 1)^T$. Its least-squares solution leaves a nonzero residual. Conversely, $G = (1, 1)$ with $r = 1$ admits both $(1, 0)^T$ and $(1/2, 1/2)^T$, among infinitely many choices. Feasibility, uniqueness, and the selection of a preferred input are separate questions.

One useful architecture is $u = u_d + \Delta u$, with corrections designed for the actual tracking-error dynamics. It does not imply a universal feedforward share of measured joint torque. A nominal passive motion can have $u_d = 0$, whereas a strongly forced motion may require substantial nominal input. The split also depends on the chosen reference and policy representation. Establishing a neural feedforward or feedback contribution requires an experimental and physiological model beyond this algebraic split.

1.3 Tracking and Regulation

In a local Euclidean chart, set $e = x - x_d(t)$. Then

$$\dot{e} = f(e + x_d(t), u, t) - \dot{x}_d(t). \quad (1.5)$$

Tracking becomes regulation of the origin in a generally time-varying error system. This change of variables explains why Lyapunov methods and linear-quadratic designs remain useful for moving references. The MIT treatment of finite-horizon LQR and trajectory stabilization provides an established development of these tools. Applying an equilibrium controller without accounting for the changing reference can fail; that is a modeling or design error, not a limitation of all regulation-based analysis.

For a scalar example, take $\dot{x} = x + u$ and $x_d(t) = \sin t$. The nominal input is $u_d = \cos t - \sin t$. Choosing

$$u = u_d - (k + 1)e, \quad k > 0, \quad (1.6)$$

gives $\dot{e} = -ke$ exactly. The state keeps moving while its tracking error converges. If actuator bounds are imposed, this result applies only where the specified input can actually be delivered. Delays, unmodeled dynamics, and measurement error also require their own treatment.

1.4 Configuration Paths, Timing, and State

Write a configuration path as $q_d(s)$, where s is a dimensionless progress coordinate. For a timing law $s(t)$, define $\sigma = \dot{s}$. The chain rule gives

$$\begin{aligned} q(t) &= q_d(s(t)), \\ \dot{q}(t) &= q'_d(s)\sigma, \\ \ddot{q}(t) &= q''_d(s)\sigma^2 + q'_d(s)\dot{\sigma}. \end{aligned} \quad (1.7)$$

Thus the mechanical state is $(q_d(s), q'_d(s)\sigma)$, not a configuration-only curve. Changing σ changes that state and usually the required loads. A prescribed full-state curve $(q_d(s), v_d(s))$ has an additional compatibility condition $q'_d(s)\sigma = v_d(s)$. It cannot generally be traversed at any chosen speed while remaining the same physical state trajectory.

For a mechanical model $M(q)\ddot{q} + h(q, \dot{q}) = B(q)u$, substitution yields

$$B(q_d)u = M(q_d)(q''_d\sigma^2 + q'_d\dot{\sigma}) + h(q_d, q'_d\sigma). \quad (1.8)$$

The right side must lie in the admissible input image. Closure reactions, external forces, and constitutive loads must be included consistently in a more complete model. Geometric path design and timing design are useful subproblems, but the dynamics couples them. Underactuation makes that compatibility especially important; it does not make arbitrary retiming free.

As a reproducible example, take $m\ddot{q} = u$ with mass $m = 1$ kg and applied force u . Let $q_d(s) = \ell s^2$ with $\ell = 1$ m. Use $s = t/T$ on $0 \leq t \leq T$. The same configuration path from 0 to 1 metre gives

$$q = \ell(t/T)^2, \quad \dot{q} = 2\ell t/T^2, \quad u = 2m\ell/T^2. \quad (1.9)$$

At the endpoint, $T = 1$ s gives speed 2 m/s, force 2 N, and kinetic energy 2 J. With $T = 2$ s, the values are 1 m/s, 0.5 N, and 0.5 J. Work over the one-metre displacement equals the kinetic-energy increase in each case. The phase-plane curves are $\dot{q} = 2\sqrt{\ell q}/T$, so they are different state-space curves. This is an illustrative timing calculation, not a measured golf swing.




Figure 1.1: One Configuration Path, Two Timing Laws, and Different State Trajectories. Curves End at the Same Position With Different Velocities.

The example uses a prescribed clock-based phase. Estimating phase from the configuration itself needs care near the stationary start, where $q'_d(0) = 0$. In a swing, a coordinate that reverses at the top of the backswing cannot serve as a globally increasing phase through that reversal without a different construction. A useful phase must be defined on the intended interval, be observable from the chosen state or measurements, and have the required regularity and progress properties.

1.5 Choose a Meaningful Metric and Phase

A raw norm of position and velocity coordinates mixes different units. Even among positions, changing metres to centimetres changes an unweighted norm. One local choice is

$$\|e\|_W^2 = e^T W e, \quad W = W^T \succ 0, \quad (1.10)$$

with weights and scales declared according to the desired interpretation. Under an invertible linear coordinate change $\tilde{e} = S e$, the same quantity uses $\tilde{W} = S^{-T} W S^{-1}$. Keeping W unchanged while changing units generally changes the specification. A physically motivated metric or scaled output error can be useful, but neither is automatically a clinical risk measure or an empirical measure of human effort.

On a configuration manifold, subtraction is a local coordinate operation; rotations, for example, need compatible charts or a chosen local logarithmic error. The mass matrix defines a kinetic-energy metric on configuration velocities, not by itself a unique distance on the entire augmented state including angles, velocities, activation, and other quantities.

For an embedded regular reference curve $\gamma(s)$ in a scaled state chart, one possible local phase is

$$s_*(x) \in \arg \min_s \frac{1}{2}(x - \gamma(s))^T W(x - \gamma(s)). \quad (1.11)$$

At an interior minimizer with constant W ,

$$\gamma'(s_*)^T W(x - \gamma(s_*)) = 0. \quad (1.12)$$

A sufficiently small tubular neighborhood of a suitably regular embedded curve can make the projection unique and smooth. That conclusion is local. At the center of a circle every point on the circle is equally close; self-intersections, nearby branches, zero tangents, or a finite arc's endpoints can invalidate the assumed phase chart. A constrained endpoint minimizer need not satisfy the interior orthogonality equation.

Transverse coordinates need not use nearest-point orthogonality, but must form valid local coordinates together with phase. If the state dimension is d , there are locally $d - 1$ transverse coordinates around a regular one-dimensional curve. This coordinate count is different from the number of actuators and from the dimension of a task-equivalent set.

1.6 Stability, Attraction, and Finite Motion

Lyapunov stability means sufficiently small initial errors remain small over the stated time domain. It does not by itself mean convergence. Asymptotic stability adds attraction as time tends to infinity. Exponential stability adds a quantitative decay bound with specified constants. These distinctions apply to equilibria and have corresponding trajectory and orbital formulations; see the MIT Lyapunov analysis notes.

For a periodic orbit of an autonomous closed-loop system, orbital stability allows phase displacement while controlling distance to the orbit. A nearby solution can converge to the orbit without converging to the same clock phase as a nominal solution. This is useful for periodic motions when phase locking is not the objective. It is not a general license to change the physical speed along a prescribed position-velocity curve.

A golf swing studied only up to an impact event is a finite-duration motion. Its appropriate claims may concern bounded deviation, successful event arrival, and delivered output tolerances. An infinite-time attraction statement needs a defined continuation, such as a periodic task or a subsequent stabilization phase. A finite reference arc is not an invariant orbit once the motion leaves its endpoint. Bounds on a finite interval should not be renamed asymptotic stability without that missing structure.

For example, $\dot{e} = e$ gives $e(t) = e(0)e^t$. On a fixed interval $[0, T]$, choosing $|e(0)| < \varepsilon e^{-T}$ keeps $|e(t)| < \varepsilon$. This finite-time bound holds even though errors amplify and the origin is unstable on the infinite time axis. A useful swing robustness result should therefore report the permitted initial/measurement uncertainty, the amplification or decay, the time or event interval, and the terminal output bound.

Orbital analysis still uses Lyapunov functions. For an explicit example, consider $\dot{x}_1 = x_2$, $\dot{x}_2 = -x_1 + u$. The unit circle is a nominal orbit with $u_d = 0$. Put $r^2 = x_1^2 + x_2^2$ and choose

$$u = kx_2(1 - r^2), \quad V = \frac{1}{4}(r^2 - 1)^2, \quad \dot{V} = -kx_2^2(r^2 - 1)^2 \leq 0, \quad k > 0. \quad (1.13)$$

On a sufficiently small compact annulus around the unit circle, the largest invariant subset of $\dot{V} = 0$ is the circle itself. Points with $x_2 = 0$ away from the circle do not remain there, since $\dot{x}_2 = -x_1 \neq 0$; the origin is excluded by the annulus. Lyapunov and invariance arguments then establish local orbital attraction and stability. They do not establish phase locking or global attraction from the origin. The MIT limit-cycle treatment and Manchester's transverse-dynamics analysis provide the broader framework.

1.7 Verifying a Motion Tube

A time-varying tube can be defined by

$$\mathcal{T}(t) = \{x : V(x, t) \leq \rho(t)\}. \quad (1.14)$$

For example, V may measure weighted tracking error. Selecting $\rho(t)$ to be small near impact expresses a desired tolerance. It does not show that the closed-loop motion can stay inside the tube. Even $\dot{e} = 0$ leaves a shrinking tube when initialized on its boundary: constant error cannot fit inside an arbitrarily decreasing allowed radius.

Suppose the closed-loop dynamics is well posed with admissible inputs, V and ρ are continuously differentiable, and the tube has a regular boundary on the domain of interest. For disturbances d in a declared set $\mathcal{D}$, a robust inward-boundary condition is

$$\partial_t V + \nabla_x V^T f(x, k(x, t), d, t) \leq \dot{\rho}(t) \quad \text{when } V(x, t) = \rho(t), \quad \forall d \in \mathcal{D}. \quad (1.15)$$

Together with the appropriate regularity and initial membership, this supports forward invariance over the specified interval. The derivative must include reference motion and changing metric weights. If V and ρ are indexed by phase, their derivatives involve the actual $\dot{s}$; it cannot silently be replaced by one. Nonsmooth profile corners or hybrid resets need their own one-sided or transition conditions.

A scalar example makes the effort and disturbance limit explicit. Let $\dot{e} = -ke + d$, $|d| \leq \bar{d}$, and require $|e| \leq r(t)$. On either boundary, a sufficient condition is

$$\dot{r} \geq -kr + \bar{d}. \quad (1.16)$$

For $k = 2$, $\bar{d} = 0.1$, and $r(0) = 1$ in declared scaled units, the tight comparison radius is $r(t) = 0.05 + 0.95e^{-2t}$. Constant worst-direction disturbances attain its two boundaries. The nonzero limiting radius follows from the disturbance bound and available correction rate. A drawing that shrinks to zero would not certify this disturbed system.

For an impact specification, also require that every relevant pre-impact boundary state allowed by the verified tube maps into $\mathcal{Y}$, and that the intended event is reached. A phase estimate bounded below by $\dot{s} \geq \alpha_{\min} > 0$ can supply a progress bound on an appropriate interval; distance to a path alone cannot. Actuator saturation, sensor delay, contact feasibility, and uncertainty in the model must be included in the claimed guarantee. A verified model tube is still a statement about that model and its declared uncertainty set, not automatic validation on golfers.

1.8 Input Rank and Reachability

For an n -coordinate mechanical model with $m < n$ effective inputs, rank limits instantaneous input acceleration directions. It does not imply that all feasible states lie on a surface of dimension $2m$. In a state model, drift and input directions evolve as the system moves; their effect over time can provide additional reachable directions. The underactuation chapter develops the associated feasibility and reachability questions.

A concrete counterexample uses two unit masses, one grounded spring, and one coupling spring, all with unit coefficients. In scaled coordinates,

$$\ddot{q} + Kq = bu, \quad K = \begin{pmatrix} 2 & -1 \\ -1 & 1 \end{pmatrix}, \quad b = \begin{pmatrix} 1 \\ 0 \end{pmatrix}. \quad (1.17)$$

There is one applied force. With $x = (q_1, q_2, \dot{q}_1, \dot{q}_2)^T$,

$$\dot{x} = Ax + gu, \quad A = \begin{pmatrix} 0 & I \\ -K & 0 \end{pmatrix}, \quad g = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}. \quad (1.18)$$

The controllability matrix $[g, Ag, A^2g, A^3g]$ has determinant -1 and rank four. With unconstrained signed inputs, this linear system is controllable in its full four-dimensional state space. Input limits restrict the size of reachable sets and time needed; they do not restore the asserted universal $2m$ dimension bound.

A single trajectory is itself a one-dimensional curve, but the union of reachable trajectories and the span of a curve's successive derivatives are different objects. For the same system with $u = 0$ and $x(0) = (1, 0, 0, 0)^T$, the matrix $[Ax, A^2x, A^3x, A^4x]$ at the initial state also has determinant -1 . Thus even an unforced example has four independent time derivatives. At that regular point, smooth arc-length reparameterization preserves their rank. Actuator count therefore does not cap the number of independent Frenet-frame directions at $2m$ either.

The useful golf interpretation is more specific. Coupling can transmit and redistribute the effects of earlier work and present loads. Whether an effective wrist coordinate is unactuated, weakly actuated, or actively constrained is a modeling and measurement question. A rapid release does not prove zero wrist moment, and passive planar dynamics does not guarantee three-dimensional face squaring. Force, joint moment, mechanical power, and muscle activation must be kept distinct when testing such hypotheses.

1.9 Reliability and Performance

A path error, a delivery error, and a ball-dispersion measure are different costs. Choose the one that represents the study, then state how the others enter the model. Both regulation and trajectory problems can use cost functionals over an entire history; neither is restricted to a single instantaneous error value.

For a specified stochastic model and admissible policy π , one possible criterion is

$$\begin{aligned} J_\pi &= \phi(y_e) + \int_0^{t_e} \ell(x(t), u(t)) dt, \\ \min_{\pi} \quad &\mathbb{E}[J_\pi] + \lambda \text{Var}(J_\pi), \quad \lambda \geq 0. \end{aligned} \tag{1.19}$$

The cost and its units, disturbance distribution, event definition, and constraint handling must be declared. The variance weight has corresponding units; it expresses a design preference rather than a universal biological constant. A low variance can accompany a consistently poor mean outcome, so repeatability alone is not sufficient. Expected cost, tail risk, target-set probability, and worst-case guarantees answer different questions. If the intended event is not reached by an allowed horizon, define a failure cost or constraint; silently discarding those trials changes the question.

A candidate policy may use feedback, a learned reference, or both. Comparing policies requires recomputing the motion, loads, event time, and output in each branch. Reusing a measured state or reaction history after the intervention changes the trajectory can invalidate the comparison. Experimental perturbations and held-out predictions can test a model; fitting one observed motion does not establish that the nervous system optimized the proposed criterion.

A practical golf investigation can proceed as follows:

1. Define the delivered output, its tolerance, and the ball/contact model used to connect that output to performance.
2. Declare the mechanical boundary, contact modes, state variables, input meaning, parameter uncertainty, and available measurements.
3. Construct compatible nominal states and inputs. Check closure, timing, actuator limits, work, and contact feasibility before optimizing a path.
4. Model disturbances and measurement/actuation delays. Compare finite-time or event-based output sensitivity and verify any proposed tube condition.
5. Test predictions on independent swings or perturbations, including uncertainty in event timing and the inferred loads.
6. Separate demonstrated mechanical identities, validated predictions, physiological interpretations, and proposed coaching or design hypotheses.

This is the role of the subsequent chapters: curve geometry describes motion; configuration manifolds describe compatible coordinates; transverse analysis and underactuation constrain candidate designs; optimization and robustness analysis evaluate them; learning and experiments test how well they transfer. The later golf case study is a modeling proposal unless supported by a complete specified model, reproducible results, and validation data. The framework's value comes from making the links testable.

1.10 Chapter Exercises

1. Specify an impact-event output set that distinguishes clubhead speed, face orientation, and impact location. Explain why a terminal target does not require stopping there.
2. Derive the error dynamics for the scalar tracking example, then impose an input bound and identify where the proposed law becomes infeasible.
3. Reproduce the one-mass timing example. Compare endpoint velocity, input, work, and the phase-plane curves for $T = 1$ and $T = 2$ seconds.
4. For a prescribed pair $(q_d(s), v_d(s))$, derive the compatibility condition on $\dot{s}$ and give an example where arbitrary retiming fails.
5. Rescale a position coordinate from metres to centimetres. Transport the error metric and verify that the weighted distance remains unchanged.
6. Show why nearest-point phase is not unique at a circle's center and why a finite arc's endpoint needs a separate projection condition.
7. Derive the scalar robust-radius condition and integrate its two worst-direction disturbance cases. Explain the nonzero limiting radius.
8. Verify both rank-four matrices in the spring example. Distinguish actuator rank, reachable-set dimension, and a trajectory's derivative span.
9. Verify the orbital Lyapunov derivative and the invariant-set argument. Explain why this does not prove phase locking or global convergence.
10. Formulate two different reliability objectives for a golf study. State what data would distinguish their predictions and what remains unidentifiable.