

Chapter 4

Orbital Stability and Transverse Linearization

4.1 Choose the Error That Matches the Task

A swing can follow a similar spatial path at a different speed, arrive at impact at a different time, or deliver a different face orientation at the same time. These are different perturbations. Their importance depends on the task and on which variables the path contains. A configuration path omits velocity; a mechanical state-space path includes it. Sliding along the latter is therefore not permission to choose an arbitrary speed along the former. The compatibility conditions from Chapter 1 still apply.

Timed tracking compares $x(t)$ with a feasible reference $x_d(t)$ at the same clock time. Orbital analysis compares a state with an invariant trajectory set. Event-based analysis compares outcomes when a specified event occurs. All three can be useful. A stationary ball may make some absolute start-time shifts irrelevant, while clubhead speed, face orientation, contact order, and timing relative to an external disturbance remain consequential. The choice should follow the outcome, not a general preference for or against time-indexed error.

Lyapunov stability means that sufficiently small initial perturbations remain small. Asymptotic stability additionally requires convergence; exponential stability specifies a rate bound. Lyapunov functions can analyze equilibria, time-varying tracking errors, invariant sets, and periodic orbits. Orbital stability is not an alternative that invalidates Lyapunov analysis. It is a choice of the set whose stability we want to study.

For the unit-speed circle $x_d(t) = (\cos t, \sin t)$, a phase-shifted solution $x_d(t - \varepsilon)$ has same-time distance

$$\|x_d(t - \varepsilon) - x_d(t)\| = 2|\sin(\varepsilon/2)|. \quad (4.1)$$

This remains small with small ε and does not decay. It disproves asymptotic convergence to the scheduled solution, not stability. For a system driven by an explicit clock-dependent input, the shifted curve is not generally a solution under the original input. Either shift the input too, use a phase-dependent autonomous feedback law, or include the forcing phase in the state before drawing an orbital conclusion.

4.2 Invariant Orbits and Finite Segments

Consider a smooth autonomous closed-loop system $\dot{x} = F(x)$ with a regular nonstationary periodic solution $\gamma(t + T) = \gamma(t)$ of minimal period T . Its image $\mathcal{O}$ is an invariant closed

curve. In a neighborhood with a specified distance, orbital stability means

$$\forall \epsilon > 0 \exists \delta > 0 : \quad d(x_0, \mathcal{O}) < \delta \implies d(x(t), \mathcal{O}) < \epsilon \quad \text{for all } t \geq 0. \quad (4.2)$$

Asymptotic orbital stability adds $d(x(t), \mathcal{O}) \rightarrow 0$ locally. Asymptotic phase additionally means convergence to a particular time-shifted nominal solution, $\|x(t) - \gamma(t + \theta_*)\| \rightarrow 0$. These statements require the relevant existence, invariance, and neighborhood conditions; they are not inferred from a visually closed plot of a few coordinates.

A finite downswing segment is generally not invariant: its nominal solution leaves through the endpoint. It cannot satisfy an all-future-time set-stability definition merely because nearby motions remain close until impact. Useful finite-horizon questions instead concern error amplification, constraint satisfaction, disturbance sensitivity, and event outcomes before termination. They need their own bounds. Calling a finite segment an orbit for descriptive purposes should not transfer infinite-time theorems to it.

A nonstationary smooth periodic orbit is regular because its autonomous flow cannot pass through an equilibrium and then leave under uniqueness. Arbitrary reference curves can have stops, self-intersections, or endpoints. A projection into selected measurements can intersect itself even if the full state orbit does not. The state, metric, regularity, and selected phase neighborhood must therefore be specified before constructing transverse coordinates.

The Van der Pol oscillator has an attracting limit cycle for positive damping parameter in its standard form. It is a useful mathematical example. Walking is often modeled with hybrid periodic gaits, and a golf swing is usually modeled as a transient. Neither an analogy with an oscillator nor an attractive orbit in a reduced model establishes a measured human coordination strategy.

4.3 Derive the Local Phase and Transverse Coordinates

Work first in a fixed Euclidean state chart with declared scaling. Let $\gamma(\tau)$ be a regular reference, parameterized by its nominal physical time τ , and let $u_*(\tau)$ satisfy $\gamma'(\tau) = f(\gamma(\tau), u_*(\tau))$. Define $v(\tau) = \|\gamma'(\tau)\| > 0$, unit tangent $T = \gamma'/v$, and a smooth orthonormal normal basis $E(\tau) \in \mathbb{R}^{n \times (n-1)}$, with $E^T E = I$ and $E^T T = 0$. A local coordinate representation is

$$x = \gamma(\tau) + E(\tau)\xi, \quad \xi \in \mathbb{R}^{n-1}. \quad (4.3)$$

The ambient displacement $E\xi$ has n components; the reduced error ξ has $n - 1$. The orthogonal projector $P = EE^T = I - TT^T$ is an $n \times n$ matrix. Projecting with P does not itself produce reduced coordinates.

For nearest-point phase, $T(\tau)^T[x - \gamma(\tau)] = 0$. This defines a smooth local phase only where its derivative with respect to τ is nonzero and the chosen closest point is unique. A compact embedded curve has an appropriate tubular neighborhood, but a curvature radius alone does not establish its global width: distant pieces of the curve can approach one another. Endpoints and branch changes need explicit handling. Frenet frames can fail at zero curvature; a smooth normal frame can often be constructed without that failure.

Differentiate the coordinate representation and project along T and E . For $u = u_*(\tau) + \delta u$, the exact local equations are

$$\dot{\tau} = \frac{T^T f(x, u)}{v - T^T E \xi}, \quad \dot{\xi} = E^T f(x, u) - E^T E' \xi \dot{\tau}. \quad (4.4)$$

Primes here mean derivatives with respect to nominal time τ , and dots mean actual physical time. We used $T^T E' = -T'^T E$. The denominator must remain nonzero; forward progression additionally requires $\dot{\tau} > 0$. At the reference $\dot{\tau} = 1$ and $\dot{\xi} = 0$. This is the nominal-time version of Chapter 2's arc-length equations, with its parameter factors made explicit.

Linearization at fixed phase gives

$$\dot{\xi} = A_{\perp}(\tau)\xi + B_{\perp}(\tau)\delta u + O(\|\xi\|^2 + \|\delta u\|^2), \quad (4.5)$$

$$A_{\perp} = E^T f_x E - E^T E', \quad B_{\perp} = E^T f_u. \quad (4.6)$$

Mixed second-order terms are absorbed into the displayed bound locally. The dimensions are $(n-1) \times (n-1)$ and $(n-1) \times m$, respectively. The frame derivative matters; an arbitrary sequence of independently computed normal bases can introduce sign jumps and fictitious dynamics. In a rotating normal basis this term transforms with the basis, preserving the physical prediction when the full transformation is retained.

The derivation assumes an autonomous plant and feedforward evaluated at the estimated phase. Explicit time dependence, an independent clock-driven input, or a dynamic phase estimator adds variables or coupling terms. For a manifold state use compatible local charts, connections, or group coordinates as in Chapter 3. The fixed-chart formula should not be applied unchanged to an arbitrary position-dependent metric. Under a different path parameter, both the speed and cost integration measure must be transformed consistently.

4.3.1 An Oscillator With a Missing Input Factor

For $\dot{x}_1 = x_2$, $\dot{x}_2 = -x_1 + u$, take $\gamma(\tau) = (\cos \tau, -\sin \tau)$ and $u_* = 0$. Then $T = (-\sin \tau, -\cos \tau)^T$, $E = (-\cos \tau, \sin \tau)^T$, and $x = (1 - \xi)\gamma(\tau)$. The equations become

$$\dot{\xi} = u \sin \tau, \quad \dot{\tau} = 1 - \frac{u \cos \tau}{1 - \xi}, \quad A_{\perp} = 0, \quad B_{\perp} = \sin \tau. \quad (4.7)$$

Choose $u = -k \sin \tau \xi$ with $k > 0$. The exact radial equation is $\dot{\xi} = -k \sin^2 \tau \xi$; along the nominal orbit $\tau = t$ its scalar linearization has one-period multiplier

$$\mu = \exp\left(-k \int_0^{2\pi} \sin^2 t \, dt\right) = e^{-k\pi} < 1. \quad (4.8)$$

This is a local exponential-stability certificate with the usual smoothness and regular-phase assumptions. No averaging is required for the multiplier. The incorrect law $u = -k\xi$ instead gives exponent $-k \int_0^{2\pi} \sin t \, dt = 0$ and multiplier one. The missing sine factor changes the conclusion. For sufficiently small ξ , forward phase remains well-defined; this local argument does not establish global behavior at the origin, where radial phase is undefined, or under torque saturation.

The related exercise $\dot{x}_1 = x_2 + u$, $\dot{x}_2 = -x_1^3$ cannot track the same timed circle by any choice of u : the second equation would require $-\cos t = -\cos^3 t$ for all t . Checking the uncontrolled equation should precede linearization, optimization, or feedback design.

4.4 What Linearization Can Certify

For a smooth invariant orbit, uniform exponential stability of the transverse linearization implies local exponential orbital stability under standard regularity conditions. An unstable transverse multiplier gives local instability. A marginal linearization generally leaves nonlinear behavior undecided. Nonlinear asymptotic stability is not equivalent to asymptotic stability of the linearized system.

For example, in a local annulus use angular phase $\dot{\theta} = 1$ and radial error dynamics either $\dot{\xi} = -\xi^3$ or $\dot{\xi} = +\xi^3$. Both have the same circle and the same transverse linearization $\dot{\xi} = 0$. Yet

$$\xi_-(t) = \frac{\xi_0}{\sqrt{1 + 2\xi_0^2 t}}, \quad \xi_+(t) = \frac{\xi_0}{\sqrt{1 - 2\xi_0^2 t}}. \quad (4.9)$$

The first converges algebraically; the second grows until it leaves the local region. Both linearizations have multiplier one. Higher-order terms decide the distinction. This counterexample also explains why a linear Lyapunov function is a candidate for nonlinear verification, not an automatic proof over a large tube.

4.5 Periodic Orbits, Return Maps, and Hybrid Events

For a periodic transverse linear system $\dot{\xi} = A_\perp(t)\xi$, the state transition matrix satisfies $\partial_t \Phi(t, t_0) = A_\perp(t)\Phi(t, t_0)$, $\Phi(t_0, t_0) = I$. Its one-period map $M_\perp = \Phi(T, 0)$ has Floquet multipliers as eigenvalues. All magnitudes strictly below one give exponential stability of this periodic linear system and local exponential stability of the smooth nonlinear orbit. Any magnitude above one implies instability. For the linear system, semisimple unit-modulus multipliers can permit bounded solutions; for the nonlinear orbit they are inconclusive without more analysis.

A nonstationary autonomous periodic orbit has a unit multiplier in the full variational system because its tangent is a periodic solution of that system. There may be additional unit multipliers. Transverse reduction removes the phase direction, not every possible neutral mode. Periodically forced systems do not automatically have the same neutral phase direction when the forcing clock is held fixed. An augmented phase formulation must state what is varied.

A local Poincare section crosses the flow transversely near a point on a periodic orbit. The local return map has a fixed point there. Its derivative represents the transverse monodromy map up to a consistent change of section coordinates; numerical matrices need not be literally equal under different bases. Their nontrivial multipliers agree. A phase basis used around a cycle must also be identified consistently after one period; a frame mismatch can alter an incorrectly assembled return matrix.

A hypothetical return map with diagonal multipliers 0.7, 0.4, -0.3 contracts those eigenmodes, with the last alternating sign. Along the first eigenmode, ten returns give factor $0.7^{10} \approx 0.028$. This is an illustrative map, not a computed walking or golf result. Nonnormal eigenvectors can produce large norm growth before asymptotic decay. A gain that moves a multiplier cannot be computed from the multiplier alone: the input map is required. For $z_{j+1} = 1.1z_j + bu_j$, $u_j = -kz_j$, the target 0.6 requires $k = 0.5/b$ when $b \neq 0$.

Hybrid motion requires the event and reset derivative as well as continuous variational integration. For a smooth guard $h(x, t) = 0$, reset $x^+ = \Delta(x^-, t)$, and transverse crossing $h_t + h_x f^- \neq 0$, the saltation matrix mapping perturbations at a common nominal time is

$$S = \Delta_x + \frac{(f^+ - \Delta_x f^- - \Delta_t)h_x}{h_t + h_x f^-}. \quad (4.10)$$

This includes the effect of changed event time. A hybrid cycle's linear map composes continuous transitions and these event maps in chronological order, with consistent coordinates. Grazing, simultaneous contacts, or changes in event order can invalidate this smooth local formula.

Even an identity reset does not imply $S = I$. In one dimension, let speed change from one to two when x crosses zero. Then $S = 2$: a small position lead triggers the faster mode earlier and doubles the subsequent position lead. Differentiating only the identity reset would miss this effect. Club-ball contact, ground contact, and changing grip constraints likewise need a stated event model; a smooth downswing Jacobian alone does not describe sensitivity across an impact.

4.6 Finite-Time Amplification and Impact Sensitivity

For a transient, $\Phi(t_2, t_1)$ gives first-order sensitivity under its stated closed-loop dynamics. In a fixed Euclidean norm its largest singular value is a worst-case amplification factor. A value above one means *some* initial direction is amplified, not all directions. A terminal value below one does not exclude earlier growth. Higher-order errors and ongoing disturbances also remain outside that homogeneous linear map.

The stable constant matrix

$$A = \begin{pmatrix} -1 & 10 \\ 0 & -1 \end{pmatrix}, \quad \Phi(t, 0) = e^{-t} \begin{pmatrix} 1 & 10t \\ 0 & 1 \end{pmatrix} \quad (4.11)$$

has both eigenvalues at -1 . Nevertheless $\|\Phi(1, 0)\|_2 > 3.7$, while $\|\Phi(10, 0)\|_2 < 0.005$. A small final error does not prove an intermediate joint or contact constraint was respected. Bounds should cover the whole relevant interval and the actual output, not just a terminal state norm.

With positive-definite endpoint metrics W_1, W_2 , the consistent gain is

$$\rho_W = \|W_2^{1/2} \Phi(t_2, t_1) W_1^{-1/2}\|_2. \quad (4.12)$$

Changing coordinate units without transforming the metrics changes the numerical gain and can change whether it exceeds one. Output sensitivity may matter more than the norm of the full transverse state: some directions barely affect face orientation, while others have large task consequences. For phase-to-phase integration, convert $d\xi/dt$ to $d\xi/ds$ with the phase rate; do not substitute a path coordinate for time in an unchanged ODE.

Impact is usually an event rather than a fixed clock time. Let the nominal event occur at t_e , with $a = h_t + h_x f \neq 0$ and a pre-event fixed-time perturbation $\delta x = \Phi(t_e, t_0) \delta x_0$. To first order,

$$\delta t_e = -\frac{h_x \delta x}{a}, \quad \delta x_e = \delta x + f \delta t_e, \quad \delta y = H_x \delta x + (H_x f + H_t) \delta t_e. \quad (4.13)$$

Here $y = H(x_e, t_e)$ is the event output, and all derivatives are evaluated on the nominal pre-event state. A reset requires its own derivative afterward. The denominator reveals why near-grazing events can have large timing sensitivity even when fixed-time state perturbations look modest.

For constant-speed travel $\dot{q} = v$, $\dot{v} = 0$ toward $q = L$, the fixed-time transition is $\Phi = \begin{pmatrix} 1 & t_e \\ 0 & 1 \end{pmatrix}$. At the event, position perturbation is zero by definition, but velocity perturbation remains and $\delta t_e = -(\delta q_0 + t_e \delta v_0)/v$. A metric that discards event time is appropriate only if the task also discards it. A metric that discards velocity would miss a central contributor to impact performance.

4.7 LQR With the Correct Variables and Cost

For the reduced linear model, define the finite-horizon cost

$$J = \xi(T)^T S_T \xi(T) + \int_0^T (\xi^T Q(t) \xi + \delta u^T R(t) \delta u) dt, \quad Q, S_T \succeq 0, \quad R \succ 0. \quad (4.14)$$

With appropriate continuous bounded coefficients and positive-definite R , the finite-horizon Riccati solution gives

$$-\dot{S} = A_{\perp}^T S + S A_{\perp} - S B_{\perp} R^{-1} B_{\perp}^T S + Q, \quad S(T) = S_T, \quad K = R^{-1} B_{\perp}^T S. \quad (4.15)$$

The input law is

$$u = u_*(\tau) - K(\tau) \xi, \quad \xi = E(\tau)^T [x - \gamma(\tau)]. \quad (4.16)$$

Its correction has m components. Multiplying $K\xi$ by the state basis E is generally dimensionally invalid: that basis maps reduced state error to ambient state displacement, not to actuator input. The negative sign follows the stated gain convention. Finite-horizon optimality for this linear cost is not an unlimited-time stability theorem or a nonlinear optimum.

Evaluating a nominal-time gain schedule at measured phase creates a phase-dependent controller. The exact dynamics include $\dot{\tau}$, so an off-reference nonlinear certificate must differentiate $S(\tau)$ with that actual rate. If the optimization itself uses another phase coordinate s , then $dt = ds/\dot{s}$ changes the dynamics and running cost. A frozen nominal conversion can be a local design approximation, but should not be described as an exact nonlinear optimal-control solution.

The value $V = \xi^T S \xi$ satisfies

$$\dot{V} = -\xi^T (Q + K^T R K) \xi \quad (4.17)$$

along the ideal linear closed loop. If $S > 0$, a constant level set gives an invariant time-varying ellipsoid for that linear system. Its axis radii are $\sqrt{c/\lambda_i(S)}$, but they need not decrease toward the terminal time. For a scalar integrator with $Q = R = 1$, $T = 1$, and $S_T = 1/4$,

$$S(t) = \tanh(1 - t + \operatorname{arctanh}(1/4)). \quad (4.18)$$

This positive solution decreases toward $1/4$, so a fixed-level ellipsoid expands. A larger terminal penalty or state weight does not by itself prove monotone tightening, feasible torque,




Figure 4.1: Transient Growth and Expanding Invariant Sets

or improved impact variance. Those claims require the resulting closed-loop and task calculations.

The left panel evaluates the triangular transition matrix above; the right evaluates the scalar Riccati radius with level one. Both are calculated counterexamples, not measured golf profiles. Together they separate asymptotic stability, intermediate amplification, and the geometric width of a verified set.

Standard full-state time-varying LQR can use nonuniform weights and explicitly represent task-relevant directions. It does not inherently penalize all errors equally. Transverse design is useful when phase freedom matches the task and the phase estimate is reliable. Timed tracking is useful when schedule matters. The comparison should use equal constraints, uncertainty assumptions, and outcome criteria, rather than assume one architecture always expends less effort.

4.8 Verify Nonlinear Tubes and Disturbance Effects

For a candidate tube $V(\tau, \xi) \leq c(\tau)$, verify on its boundary

$$V_\tau \dot{\tau} + V_\xi \dot{\xi} \leq c'(\tau) \dot{\tau} \quad (4.19)$$

under the actual nonlinear controller, admissible disturbances, and input bounds, with the initial condition inside. Also verify phase-coordinate validity, event ordering, and any reset inclusion into the next tube. A positive Riccati matrix is a useful starting candidate, not a replacement for these checks. For $\dot{\xi} = -\xi + \xi^3$, the quadratic $V = \xi^2$ has $\dot{V} = -2\xi^2 + 2\xi^4$: the linearized dissipation holds locally but fails outside $|\xi| = 1$. Available actuator authority alone does not repair an unverified nonlinear inequality.

For additive linear disturbances, the propagation is

$$\xi(t) = \Phi(t, 0)\xi_0 + \int_0^t \Phi(t, s)D(s)d(s)ds. \quad (4.20)$$

Even strong attenuation of initial error does not eliminate disturbances arriving near impact. For white-noise diffusion with covariance intensity Q_d , a linear covariance model satisfies $\dot{\Sigma} = A_{cl}\Sigma + \Sigma A_{cl}^T + DQ_dD^T$. This is a stochastic-model assumption, not a universal description of motor noise. Event-output covariance additionally uses the event sensitivity map, and uncertainty in estimated phase or model parameters must be included.

Classical gain and phase margins concern specified frequency-domain loop models under particular assumptions. An arbitrary time-varying transverse system has no ordinary stationary transfer function to which those margins can be assigned unchanged. Setting $R = rI$ in finite-horizon LQR does not establish universal per-channel half-gain or sixty-degree margins for a time-varying nonlinear swing controller. Appropriate alternatives include verified Lyapunov inequalities over an uncertainty set, induced-gain bounds, delay-aware models, and explicit perturbation studies. Simulation evidence should be distinguished from a worst-case certificate.

4.9 A Reproducible Golf Stability Investigation

Define the measured state and the impact event first. State whether the comparison is at common time, common phase, or each swing's own impact, and which output coordinates determine success. Include clubhead velocity and face orientation with their measurement uncertainties rather than treating proximity in an arbitrary state plot as sufficient evidence of accuracy.

Next construct a feasible nominal motion and identify what is held fixed in the perturbation experiment: prescribed torques, phase-dependent feedback, an estimated biological controller, or a passive model. A nominal trajectory does not identify its surrounding closed-loop vector field. Different feedback laws can generate the same nominal swing and different stability. Inverse dynamics supplies model-dependent net force requirements; it does not alone reveal how those forces change after a perturbation.

Compute the variational dynamics, phase transformation, and event/reset maps for that model. Test numerical derivatives against finite perturbations and check convergence with integration tolerance and perturbation size. Examine weighted singular values throughout the motion, not just eigenvalues or the final gain. Propagate realistic ongoing disturbances and inspect input, contact, and joint constraints. Estimate nonlinear validity by direct tests or a certificate with a stated region.

Only then compare a proposed passive-assistance mechanism with a changed model or controller. Recompute the trajectory and impact event under the intervention. Report whether a difference comes from initial-error attenuation, disturbance rejection, timing, output sensitivity, or altered nominal speed. These aspects can interact: higher speed can improve a performance objective while reducing correction time, and a stable state direction may have little effect on the ball. There is no universal numerical amplification profile through backswing, transition, release, and impact established by geometry alone.

Validate the predicted distinctions on held-out data and report alternative models that fit equally well. A reduction in a model's transverse norm is evidence about that model and metric. Calling it improved repeatability requires the event-output distribution; attributing it to human feedback requires additional identification or intervention evidence. Feedforward

and feedback contributions should be measured under a stated decomposition, not assigned a fixed ninety-to-ten percentage.

4.10 Further Reading

Tedrake’s limit-cycle notes and LQR treatment explain trajectory and orbital analysis within Lyapunov and optimal-control methods. Manchester’s transverse-dynamics work develops local coordinates and nonlinear region verification. Kong and colleagues’ saltation tutorial explains event-time corrections in hybrid linearization. Their assumptions must be matched to the swing model before transferring a stability conclusion.

4.11 Exercises

1. Compute the same-time distance between two phase-shifted unit circles. Distinguish bounded error, convergence to the scheduled solution, and orbital convergence. Explain what changes with a fixed periodic forcing clock.
2. Derive the nominal-time phase and reduced equations from $x = \gamma(\tau) + E(\tau)\xi$. State matrix dimensions and the denominator condition. Convert to an arc-length parameter and check the speed factors.
3. Verify the oscillator’s two feedback laws and their linear multipliers. Simulate small radial perturbations and compare with $e^{-k\pi}$ after one cycle. Check phase positivity and input saturation separately.
4. Compare $\dot{\xi} = -\xi^3$ and $\dot{\xi} = +\xi^3$ around a periodic orbit. Explain why their identical unit multiplier does not settle nonlinear stability.
5. Test feasibility of the proposed circle in $\dot{x}_1 = x_2 + u$, $\dot{x}_2 = -x_1^3$. Identify the violated equation before attempting a feedback or optimization calculation.
6. Compute singular values of the triangular transition matrix throughout $0 \leq t \leq 10$. Compare transient gain with its eigenvalues, then repeat after coordinate scaling with and without transforming the metric.
7. Derive the event-time and output sensitivity for constant-speed travel. Check it by finite differences, then add an identity-reset speed switch and verify the saltation factor two.
8. Verify the scalar Riccati solution with terminal weight 1/4 and plot the fixed-level radius. Explain how invariance and expanding geometric width can coexist. Test the nonlinear cubic remainder on a candidate tube.
9. For $z_{j+1} = 1.1z_j + bu_j$, compute gains producing multiplier 0.6 for $b = 1$ and $b = 2$. Explain why multiplier data alone cannot determine a gain.
10. Design a model-based swing experiment separating initial-error reduction, ongoing noise, phase uncertainty, and impact-output sensitivity. Specify an intervention and held-out evidence that would distinguish competing explanations.