

The Geometry of Motion

Volume 0: The Mathematical Primer

A Foundational Guide to Kinematics, Configuration, and State
Space

Preface

Before we can begin to analyze the geometry of moving physical bodies or command them to navigate the world via control theory, we must establish a common language. The physics of movement happens in the real world, but our understanding of it happens entirely within mathematical spaces.

This primer, **Volume 0** of *The Geometry of Motion*, is designed to bridge the gap between basic calculus/algebra and the rigorous differential geometry required for modern nonlinear control theory. By intuitively introducing State Space, Configuration Manifolds, $SO(3)$ Rotations, and Screw Theory, this text ensures that the "lay user" possesses the foundational armor needed to attack the subsequent volumes on Tangent Spaces (Volume I) and Trajectory Control (Volume II).

The philosophy of this primer maintains our overarching theme: mathematics should be driven by the physical reality of motion, not abstract algebra for its own sake.

Contents

Preface	i
Nomenclature	xii
1 A Primer on Linear Algebra	1
1.1 Historical Context	1
1.2 Vectors, Matrices, and Vector Spaces	2
1.2.1 Vectors	2
1.2.2 Matrices	2
1.2.3 The Axioms of a Vector Space	3
1.2.4 Linear Independence and Span	4
1.2.5 Basis and Dimension	5
1.2.6 Linear Transformations	6
1.2.7 Rank, Null Space, and the Rank-Nullity Theorem	6
1.3 The Dot Product, Cross Product, and Inner Product Spaces	8
1.3.1 The Dot Product (Inner Product)	8
1.3.2 The Cross Product and the Skew-Symmetric Matrix	10
1.4 Eigenvalues, Eigenvectors, and Quadratic Forms	11
1.4.1 The Eigenvalue Equation	11
1.4.2 Physical Interpretation of Eigenvalues	12
1.4.3 The Spectral Theorem for Symmetric Matrices	14
1.4.4 Quadratic Forms and Positive Definiteness	15
1.5 The Singular Value Decomposition (SVD)	16
1.5.1 Derivation of the SVD From the Eigendecomposition	17
1.5.2 Geometric Interpretation of the SVD	17
1.5.3 The Condition Number	18
1.5.4 Truncated SVD and Low-Rank Approximation	18
1.5.5 Instantaneous Rank and Finite-Time Control	19
1.6 Dyads, Dyadics, and Tensor Products	20
1.6.1 What Is a Dyad?	20
1.6.2 From Dyads to Dyadics	21
1.6.3 Physical Meaning in Mechanics	22
1.6.4 The Coordinate-Free Perspective	23
1.6.5 Basis, Metric and Mechanical Duality	23
1.6.6 Connection to the Rest of This Book	23
1.7 Matrix Norms and Conditioning	24
1.7.1 What the Pseudoinverse Does Not Guarantee	25
1.8 Matrix Decompositions for Computation	26
1.8.1 LU Decomposition	26
1.8.2 QR Decomposition	26
1.8.3 Cholesky Decomposition	26
1.9 Worked Example: Eigenvalue and SVD Analysis in Python	27
1.10 Chapter Summary	28

1.11	Further Reading	28
1.12	Exercises	29
2	The Transition to State Space	31
2.1	Introduction: The Physics Abstraction	31
2.2	Historical Context: The Kalman Revolution	32
2.3	Constructing a State Vector	32
2.3.1	Worked Example 1: Simple Pendulum	34
2.3.2	Worked Example 2: DC Motor	34
2.3.3	Worked Example 3: Two-Tank Fluid System	35
2.4	Converting n -th Order ODEs to State Form	35
2.5	State Space Representations	36
2.5.1	Standard Form Differential Equations	36
2.5.2	Linear State Space: The A and B Matrices	37
2.6	Output Equations and Observability	38
2.7	Equilibrium Points and Linearization	38
2.7.1	Equilibrium Point Definition	38
2.7.2	Linearization via Taylor Expansion	38
2.8	Phase Portraits and Stability Classification	39
2.9	Solution of Linear Systems: The Matrix Exponential	40
2.10	Controllability	41
2.11	Observability	42
2.12	The Mass-Spring-Damper System: Complete Analysis	43
2.12.1	Underdamped ($0 < \zeta < 1$)	43
2.12.2	Critically Damped ($\zeta = 1$)	43
2.12.3	Overdamped ($\zeta > 1$)	44
2.12.4	Controllability and Observability	44
2.13	Transfer Function to State Space and Back	44
2.13.1	Discrete-Time State Space	45
2.14	Nonlinear Systems and Phase Portraits	45
2.15	Python Worked Example: State Space Simulation and Phase Portrait	46
2.16	Equilibrium Analysis Summary	47
2.17	Summary	47
2.18	Lyapunov Stability Theory	48
2.19	Further Reading	50
2.20	Exercises	51
3	Configuration Space and Degrees of Freedom	53
3.1	Historical Context: From Lagrange to Riemann	53
3.2	Degrees of Freedom (DOF) and Grübler's Formula	53
3.2.1	Local Freedom and Independent Constraints	53
3.2.2	Grübler's Formula	54
3.2.3	Worked Example 1: Planar Four-Bar Linkage	54
3.2.4	Worked Example 2: Spatial Stewart Platform	55
3.2.5	Worked Example 3: Human Arm Kinematics	56
3.3	The Configuration Space and Topological Manifolds	56
3.4	Charts, Atlases, and Coordinate Patches	56
3.4.1	Worked Example: Two Valid Charts on the Circle	57

3.5	Transition Maps and Smooth Compatibility	57
3.5.1	Worked Example: Circle Transition and Velocity	57
3.6	Tangent Vectors, Curves, and Derivations	58
3.7	Generalized Coordinates and Forward Kinematics	58
3.8	The Jacobian Matrix and the Kinematic Derivative	59
3.8.1	Worked Example: Two-Link Position Jacobian	59
3.8.2	Worked Example: Three-Link Planar Pose	59
3.9	Inverse Kinematics	60
3.9.1	Geometric Two-Link Inverse Kinematics	60
3.9.2	Numerical Inverse Kinematics and Convergence	60
3.10	Reachable and Dexterous Workspace	61
3.11	The Frobenius Theorem and Integrability	61
3.11.1	A Specified Unicycle and Its Lie Bracket	62
3.12	Holonomic and Non-Holonomic Constraints	63
3.12.1	Upright Rolling Wheel	63
3.12.2	Unicycle Versus Bicycle Steering	63
3.12.3	Ball and Moving-Surface Contact	63
3.12.4	Pfaffian Constraints	64
3.13	Python Example: Kinematics and Workspace	64
3.14	From Configuration Geometry to a Golf-Swing Argument	65
3.15	Chapter Summary and Key Principles	66
3.16	Further Reading	66
3.17	Exercises	66
4	Rotations, Quaternions, and SO(3)	68
4.1	Historical Context: From Euler to Hamilton	68
4.2	The Coordinate Frame	69
4.3	Rotation Matrices: From First Principles	69
4.3.1	What Is a Rotation Matrix?	69
4.3.2	Orthonormality Constraints	70
4.3.3	The Three Elementary Rotation Matrices	70
4.4	The Special Orthogonal Group: SO(3)	71
4.4.1	SO(3) as a Group	71
4.4.2	SO(3) as a Smooth Manifold	72
4.5	Euler Angles: Gimbal Lock and Beyond	72
4.5.1	Definition of Euler Angles	72
4.5.2	The ZYX Convention	72
4.5.3	The ZXZ Convention	73
4.5.4	Gimbal Lock: Singularity and Proof	73
4.6	Differential Geometry Primer: Lie Groups and Lie Algebras	74
4.6.1	The Tangent Space and Lie Algebra	74
4.7	Rodrigues' Rotation Formula: Full Derivation	75
4.7.1	The Key Identity: $[\hat{\omega}]^3 = -[\hat{\omega}]$	75
4.7.2	Taylor Expansion and Regrouping	76
4.8	Quaternion Algebra	76
4.8.1	Definition and Representation	76
4.8.2	Quaternion Operations	77
4.8.3	Unit Quaternions and Rotation	77

4.8.4	Double Cover: q and $-q$ Represent the Same Rotation	78
4.9	Conversion Between Quaternions and Rotation Matrices	79
4.9.1	Quaternion to Rotation Matrix	79
4.9.2	Rotation Matrix to Quaternion	79
4.9.3	Rotation Representation Comparison	79
4.10	Spherical Linear Interpolation (SLERP)	80
4.11	Angular Velocity and the Matrix Exponential	81
4.11.1	Angular Velocity in the Spatial Frame	81
4.11.2	Angular Velocity in the Body Frame	81
4.11.3	Proof that Angular Velocity Lives in $\mathfrak{so}(3)$	82
4.12	Axis-Angle Representation and Euler's Rotation Theorem	82
4.13	Worked Example: Quaternions and SLERP in Python	82
4.14	Chapter Summary	86
4.15	Further Reading	87
4.15.1	Foundation Tier	87
4.15.2	Intermediate Tier	88
4.15.3	Advanced Tier	88
5	Twists, Wrenches, and the Screw Axis	89
5.1	Historical Context: From Chasles to Modern Screw Theory	89
5.2	The Unified Four-by-Four Transform $SE(3)$	91
5.2.1	Group Structure of $SE(3)$	91
5.3	Chasles' Theorem and the Screw Axis	91
5.3.1	Geometric Description and Exceptional Cases	92
5.4	The Spatial Twist	92
5.4.1	Screw Representation of a Twist	93
5.4.2	Twists for Standard Joint Types	93
5.4.3	Matrix Representation and Finite Integration	93
5.4.4	The Lie Bracket of Twists	93
5.5	The Adjoint Representation	94
5.5.1	Changing Reference Frames	94
5.6	The Spatial Wrench	94
5.6.1	Wrenches From Forces and Moments	95
5.6.2	Wrench Transformation Under Changes of Frame	95
5.6.3	Physical Interpretation of Wrenches: Forces and Torques Unified	95
5.7	Proof of Chasles' Theorem	95
5.8	Mechanical Work and the Virtual Power Principle	96
5.8.1	Invariance of Power and Its Limits	96
5.9	Reciprocal Screws and Constraint Analysis	96
5.9.1	Two Geometric Screws and the Reciprocal Product	97
5.10	Physical Interpretation of Wrenches	97
5.10.1	Three Fundamental Wrench Types	97
5.10.2	The Wrench–Twist Duality at a Golf Grip	98
5.10.3	Transforming Wrenches Between Frames	98
5.11	Wrench Analysis and Force Polytopes	98
5.12	Worked Example: Planar Four-Bar Statics	99
5.12.1	Setup	99
5.12.2	Twist Space of Each Joint	99

5.12.3	Constraint Wrenches and Separate Free Bodies	99
5.12.4	Equilibrium and the Role of Actuation	100
5.12.5	Geometric Interpretation and Virtual Power	100
5.13	Python Implementation	100
5.14	Summary	101
5.15	Further Reading	101
5.16	Exercises	102
6	Exponential Coordinates and Matrix Logarithms	104
6.1	Historical Context: Sophus Lie and the Theory of Continuous Groups	104
6.2	Lie Algebras and the Tangent Space	105
6.2.1	The Lie Bracket: Non-Commutativity of Motions	106
6.3	Exponential Coordinates and Axis-Angle Representation	107
6.3.1	The Exponential Coordinate Vector	107
6.4	Rodrigues' Formula and the Matrix Exponential for $SO(3)$	107
6.4.1	Verification That the Exponential Map Produces Valid Rotations	109
6.5	The Matrix Logarithm on $SO(3)$	109
6.5.1	Computing the Logarithm: All Cases	110
6.6	The Exponential Map on $SE(3)$	111
6.6.1	Derivation of the $SE(3)$ Exponential Formula	112
6.7	The Matrix Logarithm on $SE(3)$	113
6.7.1	Numerical Stability of the Matrix Logarithm	113
6.8	The Baker-Campbell-Hausdorff Formula	114
6.8.1	The Baker-Campbell-Hausdorff Formula: Higher-Order Terms	115
6.9	Numerical Computation of Matrix Exponentials	116
6.9.1	Padé Approximation	116
6.9.2	Scaling and Squaring	116
6.10	Geodesics on $SO(3)$: Shortest Rotation Paths	116
6.10.1	Geodesics on $SO(3)(3)$: Proof via Riemannian Geometry	117
6.11	SLERP: Spherical Linear Interpolation as Geodesic Interpolation	118
6.12	Python Worked Example: Computing Exponentials and Logarithms	119
6.13	Summary	123
6.14	Further Reading	123
6.15	Exercises	124
7	Recursive Algorithms	126
7.1	Historical Context: The Computational Revolution	126
7.2	The Curse of Dimensionality in Dynamics	127
7.3	Kinematic Tree Representation	128
7.3.1	Parent Arrays and Numbering Conventions	128
7.4	The Adjoint Map ($[Ad_T]$)	128
7.4.1	Derivation of the Adjoint Matrix	129
7.4.2	The Adjoint for Forces (Wrenches)	130
7.5	Recursive Forward Kinematics	130
7.5.1	Recursive Velocity Propagation	130
7.5.2	Algorithm: Forward Kinematics in $O(N)$ Time	131
7.6	Spatial Acceleration and Coriolis Terms	131
7.6.1	The Lie Bracket and Centrifugal Forces	131

7.7	Spatial Inertia	132
7.8	The Recursive Newton-Euler Algorithm (RNEA)	133
7.8.1	The Outward Pass: Velocity and Acceleration Propagation	133
7.8.2	The Inward Pass: Force and Torque Propagation	133
7.8.3	Force Recursion and Extraction of Joint Torques	134
7.9	Computational Complexity Analysis	135
7.9.1	Why $O(N)$ and not $O(N^2)$?	135
7.9.2	Proof of Linearity	136
7.10	Branching Kinematic Trees	136
7.10.1	Handling Multiple Children	136
7.11	Contact Forces and External Disturbances	137
7.11.1	Initialization at Leaf Nodes	137
7.12	RNEA Equivalence to the Lagrangian Equations	137
7.13	Contact Forces and Constraint Handling	138
7.13.1	Bilateral Constraints	138
7.13.2	Unilateral Constraints and Friction	138
7.14	Worked Example: 3-Link Planar Arm	139
7.14.1	Setup	139
7.14.2	Outward Pass: Velocities and Accelerations	139
7.14.3	Inward Pass: Forces and Torques	140
7.15	Validation: RNEA vs. Lagrangian	141
7.15.1	2-Link Planar Arm	141
7.16	Adjoint Matrix Derivation From Geometry	142
7.16.1	Infinitesimal Rotations in SE(3)	142
7.16.2	Adjoint as Conjugation	142
7.17	Python Implementation of RNEA	142
7.18	Summary	144
7.19	Further Reading	145
7.20	Exercises	146
7.20.1	Tier 1 (Foundation)	146
7.20.2	Tier 2 (Intermediate)	146
7.20.3	Tier 3 (Advanced)	146
8	Spatial Vector Algebra and Inertia	148
8.1	Historical Context: From Plücker to Featherstone	148
8.2	Plücker Coordinates and Lines in 3D	149
8.2.1	Connection to Spatial Vectors	150
8.3	Spatial Vectors: Twists and Wrenches	150
8.4	The Spatial Cross Product	151
8.4.1	Motion Cross Product	151
8.4.2	Force Cross Product	151
8.4.3	Why Two Cross Products?	152
8.5	The Spatial Inertia Matrix	152
8.5.1	Definition and Properties	153
8.5.2	Symmetric and Positive Definite	153
8.6	The Parallel Axis Theorem in 6D	154
8.6.1	Derivation From First Principles	154
8.7	Composite Rigid Body Algorithm	154

8.7.1	Concept: Bottom-Up Inertia Assembly	154
8.7.2	Use in Computing the Mass Matrix	155
8.8	Spatial Algebra for Planar Systems	155
8.8.1	3D Vectors for 2D Motion	155
8.8.2	Computational Advantage	156
8.9	Newton-Euler Equations in Spatial Form	156
8.9.1	Spatial Momentum and the Equation of Motion	156
8.10	Connection to Dual Quaternions	157
8.10.1	Brief Overview	157
8.11	Worked Example: Spatial Inertia of a Multi-Link System	157
8.11.1	Building Link Spatial Inertias	158
8.11.2	Composite Inertia	158
8.12	Physical Interpretation of Spatial Inertia Elements	159
8.12.1	Kinetic Energy and the Origin of Spatial Inertia	159
8.13	Python Implementation	160
8.14	Summary	162
8.15	Further Reading	163
8.16	Exercises	164
8.16.1	Tier 1 (Foundation)	164
8.16.2	Tier 2 (Intermediate)	164
8.16.3	Tier 3 (Advanced)	164
9	The Product of Exponentials Formula	166
9.1	Historical Context and the Evolution of Kinematic Formalisms	166
9.2	Abandoning Denavit-Hartenberg	167
9.3	Screws and Twists: A Geometric Foundation	168
9.4	The Forward Kinematics Equation: Space Frame PoE	168
9.4.1	Derivation From First Principles	169
9.5	Body Frame PoE: Motion Relative to the Tool	169
9.5.1	Conversion Between Space and Body Frame	169
9.6	Spatial Jacobian From the Product of Exponentials	170
9.6.1	Direct Numerical Computation	171
9.7	Body Jacobian and the Adjoint Relationship	171
9.8	Singularity Analysis Using the PoE Jacobian	171
9.8.1	Types of Singularities	172
9.8.2	Detection and Handling	172
9.9	Comparison With D-H: A Worked Example	172
9.9.1	D-H Parametrization	172
9.9.2	PoE Parametrization	173
9.9.3	Computational Advantages of PoE	173
9.10	PoE for Prismatic Joints	174
9.10.1	Example: 3D Cartesian Robot	174
9.11	Velocity Kinematics From PoE Differentiation	174
9.12	Worked Example: Complete PoE for a Spatial 3R Robot	175
9.12.1	Home Configuration	175
9.12.2	Screw Axes	175
9.12.3	Numerical Evaluation at a Test Configuration	176
9.13	Inverse Kinematics Using the Jacobian: Newton-Raphson	176

9.13.1	Pose Error Representation	176
9.13.2	Newton-Raphson Update Rule	177
9.13.3	Algorithm and Convergence	177
9.14	Workspace Analysis and Manipulability	177
9.14.1	Redundancy and Null-Space Motion	179
9.15	Python Worked Example	180
9.16	Chapter Summary	181
9.17	Further Reading and References	182
9.18	Exercises	182
10	The Articulated Body Algorithm	184
10.1	Historical Context: Recursive Forward Dynamics	184
10.2	The Forward Dynamics Problem	184
10.2.1	Problem Formulation	184
10.2.2	What the Complexity Comparison Means	185
10.3	Articulated Inertia: The Key Insight	185
10.3.1	A Free Hinge Changes the Apparent Load	185
10.3.2	Mathematical Definition and Frames	186
10.4	Derivation From First Principles	186
10.4.1	Eliminating One Joint Acceleration	186
10.5	The Articulated Body Algorithm: Three Passes	187
10.5.1	Pass 1: Outward Velocity and Bias Calculation	187
10.5.2	Pass 2: Inward Subtree Elimination	188
10.5.3	Pass 3: Outward Acceleration Recovery	188
10.6	Why the Tree Calculation Is Linear-Time	189
10.7	Floating Base: A Six-Degree-of-Freedom Root	189
10.7.1	The Root Articulated Solve	189
10.7.2	Contact Does Not Automatically Fix the Base	189
10.8	Floating-Base Humanoids and Momentum	190
10.8.1	Configuration Coordinates and Generalized Velocity	190
10.8.2	Root Equation With a Known Applied Wrench	190
10.8.3	Centroidal Momentum and Whole-Body Dynamics	190
10.9	Joint Friction and Motor Inertia	190
10.9.1	Viscous Joint Friction	190
10.9.2	An Ideal Reflected Armature Model	191
10.10	ABA and Composite Rigid-Body Methods	191
10.11	Choosing Forward or Inverse Dynamics	191
10.11.1	When the Acceleration Is Prescribed	191
10.11.2	When the Applied Loads Are Specified	191
10.12	Constraint Handling: Contact and Joint Limits	192
10.12.1	Unilateral Contact	192
10.12.2	Joint Limits	192
10.13	Connection to Gaussian Elimination	192
10.14	Worked Example: A Fully Specified Three-Link Arm	192
10.14.1	Configuration and Loads	193
10.14.2	Pass 1: Velocities and Bias	193
10.14.3	Pass 2: Articulated Inertias and Residual Efforts	193
10.14.4	Pass 3 and an Independent Check	194

10.15	Reproducible Python Example	194
10.15.1	Connection to Simulation Frameworks	195
10.16	Contact Forces and Impact Dynamics	195
10.17	Contact Dynamics and Complementarity	196
10.17.1	Non-Penetration and Sustained Contact	196
10.17.2	Friction and Contact Wrenches	196
10.17.3	The Frictionless Acceleration LCP	196
10.17.4	Time-Stepping and Numerical Order	197
10.18	From Coupled Dynamics to a Golf-Swing Argument	197
10.19	Differentiating the Dynamics	198
10.20	Chapter Summary	198
10.21	Further Reading and References	198
10.22	Exercises	199
11	Lagrangian Mechanics	201
11.1	Historical Context: From Newton to Lagrange to Hamilton	201
11.2	The Limitations of Newton	202
11.3	Hamilton's Principle and the Calculus of Variations	203
11.4	Lagrangian and Energy	204
11.5	Derivation of the Standard Form: $\mathbf{M}\ddot{\mathbf{q}} + \mathbf{C}\dot{\mathbf{q}} + \mathbf{G} = \boldsymbol{\tau}$	205
11.6	Properties of the System Matrices	206
11.6.1	The Mass Matrix $\mathbf{M}(\mathbf{q})$	206
11.6.2	The Coriolis/Centripetal Matrix $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$	206
11.6.3	The Gravity Vector $\mathbf{G}(\mathbf{q})$	207
11.7	Energy Conservation	207
11.8	Hamiltonian Mechanics	208
11.9	Noether's Theorem: Symmetry and Conservation Laws	209
11.9.1	Rigorous Proof of Noether's Theorem	210
11.10	Lagrangian for Constrained Systems	211
11.11	D'Alembert's Principle and Virtual Work	211
11.12	Christoffel Symbols and the Geometry of $\mathbf{M}$	212
11.13	Worked Example 1: Double Pendulum	212
11.13.1	System Setup	212
11.13.2	Kinetic Energy	213
11.13.3	Potential Energy	213
11.13.4	Gravity Vector	214
11.13.5	Coriolis/Centripetal Matrix	214
11.14	Worked Example 2: 2-Link Robot Arm With Gravity	214
11.14.1	System Configuration	214
11.14.2	Kinetic Energy	215
11.14.3	Potential Energy	215
11.15	Comparison: Lagrangian vs. Newton-Euler	215
11.16	Symplectic Integrators	216
11.17	Python Worked Example: Double Pendulum Dynamics	217
11.18	Chapter Summary	219
11.19	Further Reading and References	220
11.20	Exercises	220

12 Machine Learning and Neural Networks	223
12.1 Scope of This Chapter	223
12.2 The Limits of Analytical Control	223
12.2.1 Prediction, Inversion and Policy Are Different Maps	224
12.2.2 Why State Choice Matters for Golf	224
12.3 Neural Networks as Function Approximators	224
12.3.1 Universal Approximation and Its Limits	225
12.4 Activation Functions	226
12.5 Loss Functions and What They Estimate	226
12.6 Backpropagation: Chain Rule for Neural Networks	227
12.7 Gradient Descent and Optimization	227
12.8 Markov Decision Processes	228
12.8.1 Return and Value	228
12.8.2 Policy Evaluation Versus Optimality	228
12.9 Policy Gradient Methods	229
12.9.1 Baselines and REINFORCE	229
12.10 Actor-Critic Architecture	230
12.10.1 Multi-Step Returns and Bias-Variance Tradeoffs	230
12.11 Proximal Policy Optimization	230
12.11.1 Stochastic Bounded Actions	231
12.12 Value-Based Methods and Continuous Actions	231
12.13 Exploration Strategies	232
12.14 Worked Example: Pendulum Swing-Up Objective	233
12.15 Physics-Informed Neural Networks	233
12.16 Neural ODEs and the Adjoint Method	234
12.17 Structure-Preserving Neural Networks	235
12.17.1 Hamiltonian Models	235
12.17.2 General Lagrangian Models	235
12.17.3 Natural Mechanical Models and DeLaN	235
12.18 Sim-to-Real Transfer and System Identification	236
12.18.1 What Can the Data Identify?	236
12.19 Python Worked Example: Backpropagation From Scratch	237
12.20 Deep Reinforcement Learning for Robot and Golf Control	238
12.20.1 State Representation and Action Parameterization	238
12.20.2 Reward Shaping and Evidence	238
12.21 Chapter Summary and Further Reading	239
12.22 Exercises	240
Glossary of Important Concepts	242
Further Reading and References	246
Index	248

Nomenclature

General Conventions

- Scalars are lowercase or Greek letters (e.g., m, t, θ).
- Vectors are bold lowercase (e.g., $\mathbf{x}, \mathbf{u}, \mathbf{q}$).
- Matrices are bold uppercase (e.g., $\mathbf{A}, \mathbf{B}, \mathbf{M}, \mathbf{J}$).
- Expected values and sets are denoted by blackboard bold or script (e.g., $\mathbb{E}, \mathcal{S}$).

State and Kinematics

$\mathbf{q} \in \mathbb{R}^n$ Joint configuration vector (generalized coordinates).

$\dot{\mathbf{q}} \in \mathbb{R}^n$ Joint velocity vector (generalized velocities).

$\ddot{\mathbf{q}} \in \mathbb{R}^n$ Joint acceleration vector.

$\mathbf{x} \in \mathbb{R}^{n_x}$ System state vector; commonly $\mathbf{x} = [\mathbf{q}^T, \dot{\mathbf{q}}^T]^T$ for mechanical systems.

$\mathbf{u} \in \mathbb{R}^m$ Control input vector (e.g., joint torques, muscle activations).

$\boldsymbol{\tau} \in \mathbb{R}^n$ Generalized forces/torques acting on the joints.

$t \in \mathbb{R}$ Time.

Inertia and Dynamics

$\mathbf{M}(\mathbf{q})$ Mass (inertia) matrix, symmetric positive definite.

$\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$ Coriolis and centrifugal force matrix.

$\mathbf{g}(\mathbf{q})$ Gravity vector.

$\mathbf{J}(\mathbf{q})$ Jacobian matrix relating joint velocities to task-space velocities ($\dot{\mathbf{p}} = \mathbf{J}(\mathbf{q})\dot{\mathbf{q}}$).

Control and Optimization

$\mathbf{A}_k = \frac{\partial f}{\partial \mathbf{x}}$ Local Jacobian matrix of the state dynamics.

$\mathbf{B}_k = \frac{\partial f}{\partial \mathbf{u}}$ Local Jacobian matrix of the control input.

$V(\mathbf{x})$ or $\mathcal{V}$ Value function or Lyapunov function.

$\mathbf{Q}$ State penalty matrix in LQR/DDP.

R Control penalty matrix in LQR/DDP.

K Feedback gain matrix ($\delta \mathbf{u} = -\mathbf{K}\delta \mathbf{x}$).

γ A trajectory or flow, mapping time and initial conditions to states.

Affine Control and Counterfactuals

$f(\mathbf{x})$ Drift vector field: complete autonomous evolution of the declared effective plant when $\mathbf{u} = \mathbf{0}$.

$G(\mathbf{x})$ Input (control) matrix field mapping control inputs to state derivatives.

$\dot{\mathbf{x}} = f(\mathbf{x}) + G(\mathbf{x})\mathbf{u}$ Control-affine system equations of motion.

$\mathbf{x}^{\text{ZTCF}}(t)$ Forward Zero-Torque Counterfactual trajectory under $\mathbf{u}(t) \equiv \mathbf{0}$.

$\mathbf{a}_{\text{ZVCF}}(\mathbf{q}, \mathbf{z})$ Instantaneous Zero-Velocity Counterfactual acceleration with velocity and declared control set to zero.

$\text{DCR}_{W, \mathcal{U}}(\mathbf{x})$ Drift-Control Ratio in a declared acceleration or task-projected space: $\frac{\|W \mathbf{a}_{\text{drift}}\|_2}{\sup_{\mathbf{u} \in \mathcal{U}} \|W B_a \mathbf{u}\|_2 + \varepsilon}$.

Biomechanics and Motor Control

$a_i \in [0, 1]$ Muscle activation level.

ℓ_i^M Muscle fiber length.

v_i^M Muscle fiber contraction velocity.

ℓ_i^{MT} Musculotendon unit length.

$\mathbf{F}_{\text{muscle}}$ Force generated by muscles.

$\mathbf{W}$ Muscle synergy matrix (weights).

$\mathbf{c}(t)$ Muscle synergy activation vector over time.

Geometry and Manifolds

$SE(3)$ Special Euclidean group of rigid body motions.

$SO(3)$ Special Orthogonal group of 3D rotations.

$\mathfrak{se}(3)$ Lie algebra of $SE(3)$, space of spatial twists (velocities).

$\mathcal{M}$ A smooth manifold.

$T_{\mathbf{x}}\mathcal{M}$ Tangent space at point $\mathbf{x}$.

Chapter 1

A Primer on Linear Algebra

A useful model connects a mathematical description to measurements, forces and motion.

In Plain Language 1.1. Context and Use Case: Linear Algebra for Motion

The Math You Will See: We will cover Vectors ($\mathbf{v}$), Matrices ($\mathbf{M}$), Vector Spaces and their axioms, Linear Independence and Basis, Eigenvalues/Eigenvectors ($\lambda, \mathbf{v}$), the Singular Value Decomposition (SVD), Quadratic Forms ($\mathbf{x}^T \mathbf{M} \mathbf{x}$), Positive Definiteness, and Dyads/Dyadics (the tensor products that underpin inertia, stress, and stiffness).

Why We Need It: Joint velocities, task velocities, forces and measured signals have multiple interacting components. Linear maps describe how those components relate at a specified state; matrix decompositions reveal directions, sensitivity and numerical structure. Predicting a collision or a golf shot additionally requires the dynamics, initial condition, inputs, constraints and contact event. Eigenvalues of one matrix do not answer all of those questions.

All Variables Defined: Check the **Nomenclature** at the start of the book for standard vector and matrix conventions.

1.1 Historical Context

Linear algebra connects the problem of solving simultaneous equations with the geometric study of linear transformations. Its importance in mechanics comes from reusing the same structure for different tasks: resolving velocities, assembling inertias, fitting measurements and propagating local perturbations.

Kalman's 1960 paper *A New Approach to Linear Filtering and Prediction Problems* used state-transition methods to derive a recursive linear estimation scheme and its error-covariance equations [18]. It was a major contribution to estimation, not the first invention of state representation or a claim that every physical system is linear. Controllability and observability make the role of matrix rank precise for declared dynamical models.

Modern simulation and optimization combine kinematic maps, linear solves and nonlinear updates. Solving a linear system within a nonlinear simulation does not make the entire simulated motion linear. The algebra keeps those uses distinct.

1.2 Vectors, Matrices, and Vector Spaces

Before mapping physics, we must understand the fundamental data structures of control theory: vectors and matrices. More importantly, we must understand the abstract *space* in which these objects live, because the space and its metric determine which operations and comparisons are meaningful. Readers seeking a comprehensive undergraduate treatment of these ideas will find Strang's textbook [41] an excellent companion to this chapter.




Figure 1.1: Vector in $\mathbb{R}^3$ and its transformation under a linear map $T(\mathbf{v}) = \mathbf{M}\mathbf{v}$.

1.2.1 Vectors

An element of a vector space is a **vector**; after choosing a finite basis, its coordinates form an ordered array of numbers. It can represent a position in space, a velocity profile, a list of sensor readings, or the activation levels of every muscle in a golfer's forearm. In physics, vectors implicitly assume a coordinate frame (e.g., X, Y, Z). A generic column vector in n -dimensional space is denoted:

$$\mathbf{v} = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} \in \mathbb{R}^n \quad (1.1)$$

It is crucial to distinguish two perspectives on what a vector “is.” The *algebraic* perspective treats a vector as a column of numbers. The *geometric* perspective treats a vector as a directed arrow in space, possessing both magnitude and direction. Both perspectives are useful when the basis and metric are stated. A position is a point in an affine space; displacement and velocity are vectors. Muscle activations can be stored in an array, but their physically admissible bounded, nonnegative set is not itself a real vector space. When we write $\mathbf{v} \in \mathbb{R}^3$, we mean both the triple of numbers (v_1, v_2, v_3) and the arrow pointing from the origin to the point (v_1, v_2, v_3) in three-dimensional Euclidean space.

1.2.2 Matrices

A **matrix** $\mathbf{M}$ is a rectangular grid of numbers arranged in rows and columns. If a vector represents a point or a direction, a matrix acts as a set of instructions on how to *transform*

that point—stretching it, rotating it, projecting it, or mapping it into an entirely different dimensional space. An $n \times m$ matrix maps vectors from $\mathbb{R}^m$ to $\mathbb{R}^n$:

$$\mathbf{M} = \begin{bmatrix} m_{11} & m_{12} & \cdots & m_{1m} \\ m_{21} & m_{22} & \cdots & m_{2m} \\ \vdots & \vdots & \ddots & \vdots \\ m_{n1} & m_{n2} & \cdots & m_{nm} \end{bmatrix} \in \mathbb{R}^{n \times m} \quad (1.2)$$

When we multiply a matrix by a vector ($\mathbf{M}\mathbf{v}$), the result is a *linear combination of the columns of $\mathbf{M}$* , using the elements of $\mathbf{v}$ as weights. Specifically, if $\mathbf{m}_1, \mathbf{m}_2, \dots, \mathbf{m}_m$ denote the columns of $\mathbf{M}$, then:

$$\mathbf{M}\mathbf{v} = v_1\mathbf{m}_1 + v_2\mathbf{m}_2 + \cdots + v_m\mathbf{m}_m \quad (1.3)$$

This “column view” of matrix-vector multiplication is one of the most important conceptual tools in linear algebra. It means that the *range* (or image) of a matrix—the set of all possible outputs $\mathbf{M}\mathbf{v}$ as $\mathbf{v}$ varies over all of $\mathbb{R}^m$ —is precisely the *span* of the columns of $\mathbf{M}$. Throughout this book, when we rotate a robot arm, we are physically applying a matrix to the vectors extending along its links.

1.2.3 The Axioms of a Vector Space

We have described vectors informally as “arrays of numbers.” But the true power of linear algebra lies in the abstract concept of a **vector space**, which captures the essential algebraic properties that make vectors useful—regardless of whether the “vectors” are columns of numbers, functions, or even infinite-dimensional signals.

Definition 1.1. Vector Space

A *vector space* V over the real numbers $\mathbb{R}$ is a set of objects (called vectors) together with two operations—**addition** ($\mathbf{u} + \mathbf{v}$) and **scalar multiplication** ($\alpha\mathbf{v}$, where $\alpha \in \mathbb{R}$)—satisfying the following eight axioms for all $\mathbf{u}, \mathbf{v}, \mathbf{w} \in V$ and all $\alpha, \beta \in \mathbb{R}$:

- (i) **Closure under addition:** $\mathbf{u} + \mathbf{v} \in V$.
- (ii) **Commutativity:** $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$.
- (iii) **Associativity of addition:** $(\mathbf{u} + \mathbf{v}) + \mathbf{w} = \mathbf{u} + (\mathbf{v} + \mathbf{w})$.
- (iv) **Existence of a zero vector:** There exists $\mathbf{0} \in V$ such that $\mathbf{v} + \mathbf{0} = \mathbf{v}$.
- (v) **Existence of additive inverses:** For each $\mathbf{v} \in V$, there exists $-\mathbf{v} \in V$ such that $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$.
- (vi) **Closure under scalar multiplication:** $\alpha\mathbf{v} \in V$.
- (vii) **Distributivity:** $\alpha(\mathbf{u} + \mathbf{v}) = \alpha\mathbf{u} + \alpha\mathbf{v}$ and $(\alpha + \beta)\mathbf{v} = \alpha\mathbf{v} + \beta\mathbf{v}$.
- (viii) **Compatibility:** $\alpha(\beta\mathbf{v}) = (\alpha\beta)\mathbf{v}$, and $1 \cdot \mathbf{v} = \mathbf{v}$.

In Plain Language 1.2. Why Do We Care About Abstract Axioms?

You might wonder: why bother with these formal rules? Because the same mathematical structures appear over and over in radically different contexts. The set $\mathbb{R}^n$ of column vectors satisfies these axioms—but so does the set of all continuous functions $f : [0, T] \rightarrow \mathbb{R}$ (where “addition” means adding functions pointwise). In optimal control, the “vectors” we optimize over are often entire *trajectories*, which are functions of time. By proving a theorem about abstract vector spaces, we prove it simultaneously for column vectors, for functions, and for trajectories. The axioms are the price of generality. The ambient function space may be a vector space, but trajectories satisfying nonlinear dynamics, unilateral contact or bounded actuation generally form a constrained subset that is not closed under addition and scaling.

Example 1.1. Verifying $\mathbb{R}^n$ is a Vector Space

Let us verify that $\mathbb{R}^n$ with standard component-wise addition and scalar multiplication satisfies the axioms.

Given $\mathbf{u} = (u_1, \dots, u_n)$ and $\mathbf{v} = (v_1, \dots, v_n)$ in $\mathbb{R}^n$, and $\alpha \in \mathbb{R}$:

- *Closure under addition:* $\mathbf{u} + \mathbf{v} = (u_1 + v_1, \dots, u_n + v_n)$. Since each $u_i + v_i \in \mathbb{R}$, the sum is in $\mathbb{R}^n$. ✓
- *Zero vector:* $\mathbf{0} = (0, 0, \dots, 0)$. Clearly $\mathbf{v} + \mathbf{0} = \mathbf{v}$. ✓
- *Additive inverse:* $-\mathbf{v} = (-v_1, \dots, -v_n)$. Then $\mathbf{v} + (-\mathbf{v}) = \mathbf{0}$. ✓
- *Closure under scalar multiplication:* $\alpha \mathbf{v} = (\alpha v_1, \dots, \alpha v_n) \in \mathbb{R}^n$. ✓

The remaining axioms (commutativity, associativity, distributivity, compatibility) follow immediately from the corresponding properties of real number arithmetic. Therefore $\mathbb{R}^n$ is a vector space. $\square$

1.2.4 Linear Independence and Span

Definition 1.2. Linear Independence

A set of vectors $\{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k\}$ in a vector space V is called *linearly independent* if the only solution to the equation

$$\alpha_1 \mathbf{v}_1 + \alpha_2 \mathbf{v}_2 + \dots + \alpha_k \mathbf{v}_k = \mathbf{0} \quad (1.4)$$

is the trivial solution $\alpha_1 = \alpha_2 = \dots = \alpha_k = 0$. If a nontrivial solution exists (i.e., at least one $\alpha_i \neq 0$), the vectors are *linearly dependent*.

Intuitively, a set of vectors is linearly independent if none of them can be “built” from the others. In $\mathbb{R}^2$, two nonzero vectors are independent exactly when neither is a scalar multiple of the other. In $\mathbb{R}^3$, three vectors are independent exactly when they do not lie in one plane through the origin.

Example 1.2. Linear Independence in $\mathbb{R}^3$

Consider the vectors:

$$\mathbf{v}_1 = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \quad \mathbf{v}_2 = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \quad \mathbf{v}_3 = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} \quad (1.5)$$

Are these linearly independent? We solve $\alpha_1 \mathbf{v}_1 + \alpha_2 \mathbf{v}_2 + \alpha_3 \mathbf{v}_3 = \mathbf{0}$:

$$\alpha_1 + \alpha_3 = 0 \quad (1.6)$$

$$\alpha_2 + \alpha_3 = 0 \quad (1.7)$$

$$0 = 0 \quad (1.8)$$

The third equation provides no constraint. Setting $\alpha_3 = 1$ yields $\alpha_1 = -1$, $\alpha_2 = -1$. A nontrivial solution exists: $-\mathbf{v}_1 - \mathbf{v}_2 + \mathbf{v}_3 = \mathbf{0}$, confirming that $\mathbf{v}_3 = \mathbf{v}_1 + \mathbf{v}_2$. The vectors are linearly *dependent*. Geometrically, all three vectors lie in the xy -plane, and the third is the diagonal of the rectangle formed by the first two.

1.2.5 Basis and Dimension**Definition 1.3. Basis**

A finite *basis* for a finite-dimensional vector space V is a set of vectors $\mathcal{B} = \{\mathbf{b}_1, \dots, \mathbf{b}_n\}$ that is both linearly independent and spans V (i.e., every vector in V can be written as a linear combination of the basis vectors). The number n of vectors in any basis is called the *dimension* of V , written $\dim(V) = n$.

A fundamental theorem of linear algebra states that *every basis for a given vector space has exactly the same number of elements*. This is why “dimension” is well-defined—it is an intrinsic property of the space, not of any particular choice of basis.

The standard basis for $\mathbb{R}^n$ consists of the unit coordinate vectors $\mathbf{e}_1 = (1, 0, \dots, 0)^T$, $\mathbf{e}_2 = (0, 1, 0, \dots, 0)^T$, through $\mathbf{e}_n = (0, \dots, 0, 1)^T$. Any vector $\mathbf{v} = (v_1, \dots, v_n)^T$ can be written uniquely as $\mathbf{v} = v_1 \mathbf{e}_1 + \dots + v_n \mathbf{e}_n$. The numbers $v_1, \dots, v_n$ are the *coordinates* of $\mathbf{v}$ with respect to this basis.

In Plain Language 1.3. Why Basis Matters for Robotics

Joint angles are coordinates on a configuration space, not basis vectors. A task map $\mathbf{y} = h(\mathbf{q})$ is generally nonlinear and can change dimension. Its Jacobian $J(\mathbf{q}) = Dh(\mathbf{q})$ maps an instantaneous joint velocity to a task velocity, $\dot{\mathbf{y}} = J\dot{\mathbf{q}}$. This is a linear map between tangent spaces at a specified configuration, and may have a nontrivial kernel. A change of basis within one finite-dimensional vector space is instead square and invertible. A seven-joint-to-three-position Jacobian cannot be a change of basis.

1.2.6 Linear Transformations

Definition 1.4. Linear Transformation

A function $T : V \rightarrow W$ between two vector spaces is called a *linear transformation* (or *linear map*) if it preserves vector addition and scalar multiplication:

- (i) $T(\mathbf{u} + \mathbf{v}) = T(\mathbf{u}) + T(\mathbf{v})$ for all $\mathbf{u}, \mathbf{v} \in V$.
- (ii) $T(\alpha \mathbf{v}) = \alpha T(\mathbf{v})$ for all $\alpha \in \mathbb{R}$ and $\mathbf{v} \in V$.

The fundamental connection between linear transformations and matrices is this: *every linear transformation between finite-dimensional vector spaces can be represented by a matrix, and every matrix represents a linear transformation.* Specifically, if $T : \mathbb{R}^m \rightarrow \mathbb{R}^n$ is linear, then there exists a unique $n \times m$ matrix $\mathbf{M}$ such that $T(\mathbf{v}) = \mathbf{M}\mathbf{v}$ for all $\mathbf{v} \in \mathbb{R}^m$. The columns of $\mathbf{M}$ are simply $T(\mathbf{e}_1), T(\mathbf{e}_2), \dots, T(\mathbf{e}_m)$ —the images of the standard basis vectors.

Theorem 1.1. Matrix Multiplication is Linear

Let $\mathbf{M} \in \mathbb{R}^{n \times m}$. Then the function $T(\mathbf{v}) = \mathbf{M}\mathbf{v}$ is a linear transformation from $\mathbb{R}^m$ to $\mathbb{R}^n$.

Proof. Let $\mathbf{u}, \mathbf{v} \in \mathbb{R}^m$ and $\alpha \in \mathbb{R}$. By the distributive property of matrix multiplication:

$$T(\mathbf{u} + \mathbf{v}) = \mathbf{M}(\mathbf{u} + \mathbf{v}) = \mathbf{M}\mathbf{u} + \mathbf{M}\mathbf{v} = T(\mathbf{u}) + T(\mathbf{v}) \quad (1.9)$$

and by the compatibility of matrix multiplication with scalar multiplication:

$$T(\alpha \mathbf{v}) = \mathbf{M}(\alpha \mathbf{v}) = \alpha(\mathbf{M}\mathbf{v}) = \alpha T(\mathbf{v}) \quad (1.10)$$

Both conditions of linearity are satisfied. $\square$

1.2.7 Rank, Null Space, and the Rank-Nullity Theorem

Not all linear transformations preserve full information about the input. Some transformations “collapse” certain directions to zero.

Definition 1.5. Rank

The *rank* of a matrix $\mathbf{M}$, denoted $\text{rank}(\mathbf{M})$, is the number of linearly independent columns (equivalently, the number of linearly independent rows) of $\mathbf{M}$. It equals the dimension of the column space (the range of the corresponding linear transformation).

If a 3×3 matrix has rank 2, it collapses full 3D space into a 2D plane. In robotics, the rank of the Jacobian matrix $\mathbf{J}(\mathbf{q})$ tells you how many independent directions the end-effector can instantaneously move in. If the rank drops below its maximum value, the robot is at a *singularity*—some previously available instantaneous task velocities are no longer in its image. This statement concerns the chosen task and admissible joint velocities; it does not by itself prove finite-time unreachability or loss of dynamical controllability.

Definition 1.6. Null Space (Kernel)

The *null space* (or *kernel*) of a matrix $\mathbf{M}$, denoted $\mathcal{N}(\mathbf{M})$, is the set of all vectors that the matrix maps to zero:

$$\mathcal{N}(\mathbf{M}) = \{\mathbf{x} \in \mathbb{R}^m \mid \mathbf{M}\mathbf{x} = \mathbf{0}\} \quad (1.11)$$

For a redundant arm, $\ker J(\mathbf{q})$ contains instantaneous joint velocities with zero velocity in the chosen hand task. A finite self-motion must keep satisfying the evolving constraint $J(\mathbf{q}(t))\dot{\mathbf{q}}(t) = \mathbf{0}$; a straight step along the initial kernel need not do so. At regular configurations the constant-rank theorem describes the local task-level set. At a singular configuration, even an instantaneous kernel direction need not integrate into a finite constant-task motion.

Definition 1.7. The Determinant (Volumetric Interpretation)

For a square matrix $\mathbf{M} \in \mathbb{R}^{n \times n}$, the *determinant* $\det(\mathbf{M})$ is a scalar whose absolute value measures the factor by which $\mathbf{M}$ scales n -dimensional volume. Imagine a unit hypercube $[0, 1]^n$. Applying $\mathbf{M}$ to every point in that cube warps it into a parallelepiped whose signed volume equals $\det(\mathbf{M})$.

- If $\det(\mathbf{M}) = 1$, volume is perfectly preserved (like a pure rotation in space).
- If $\det(\mathbf{M}) = 0$, volume is crushed completely flat—the matrix is *singular* and loses rank. The transformation is not invertible.
- If $\det(\mathbf{M}) < 0$, orientation is reversed and volume is multiplied by $|\det(\mathbf{M})|$. For example, $\text{diag}(-2, 3)$ reverses orientation and multiplies area by six; only absolute determinant one preserves volume.

Now we arrive at one of the most important structural results in linear algebra:

Theorem 1.2. Rank-Nullity Theorem

For any $n \times m$ matrix $\mathbf{M}$:

$$\underbrace{\text{rank}(\mathbf{M})}_{\text{dimension of range}} + \underbrace{\dim(\mathcal{N}(\mathbf{M}))}_{\text{dimension of null space}} = m \quad (\text{number of columns}) \quad (1.12)$$

Proof. Let $r = \text{rank}(\mathbf{M})$ and let $\{\mathbf{n}_1, \dots, \mathbf{n}_k\}$ be a basis for $\mathcal{N}(\mathbf{M})$, where $k = \dim(\mathcal{N}(\mathbf{M}))$. Extend this to a basis $\{\mathbf{n}_1, \dots, \mathbf{n}_k, \mathbf{w}_1, \dots, \mathbf{w}_{m-k}\}$ for all of $\mathbb{R}^m$ (this extension is always possible by a standard result in linear algebra).

We claim that $\{\mathbf{M}\mathbf{w}_1, \dots, \mathbf{M}\mathbf{w}_{m-k}\}$ is a basis for the column space of $\mathbf{M}$, which would establish $r = m - k$, i.e., $r + k = m$.

Spanning: Any $\mathbf{y}$ in the range of $\mathbf{M}$ can be written as $\mathbf{y} = \mathbf{M}\mathbf{v}$ for some $\mathbf{v} \in \mathbb{R}^m$. Expanding $\mathbf{v}$ in our basis: $\mathbf{v} = \sum_{i=1}^k \alpha_i \mathbf{n}_i + \sum_{j=1}^{m-k} \beta_j \mathbf{w}_j$. Since $\mathbf{M}\mathbf{n}_i = \mathbf{0}$, we get $\mathbf{y} = \sum_{j=1}^{m-k} \beta_j \mathbf{M}\mathbf{w}_j$. So the $\mathbf{M}\mathbf{w}_j$ span the range.

Independence: Suppose

$$\sum_{j=1}^{m-k} \beta_j \mathbf{M} \mathbf{w}_j = 0.$$

Then $M(\sum_j \beta_j \mathbf{w}_j) = 0$, so the vector $\sum_j \beta_j \mathbf{w}_j$ belongs to the kernel. It therefore equals $\sum_{i=1}^k \gamma_i \mathbf{n}_i$ for some scalars γ_i . The combined list of kernel vectors and extension vectors is a basis for $\mathbb{R}^m$, so it is independent. All β_j and γ_i must be zero.

Thus $r = m - k$, establishing $\text{rank}(\mathbf{M}) + \dim(\mathcal{N}(\mathbf{M})) = m$. $\square$

In Plain Language 1.4. The Physical Meaning of Rank-Nullity

For a robot with m independent joint-velocity coordinates and an n -dimensional task:

- The **rank** of the Jacobian tells you the dimension of the instantaneous task-velocity image when all declared joint velocities are admissible.
- The **nullity** (dimension of the null space) tells you how many “extra” internal motions remain after the task is specified.

If a 7-DOF arm ($m = 7$) controls a hand in 3D position space ($n = 3$), and the Jacobian has rank 3, then the nullity is $7 - 3 = 4$. The robot has 4 degrees of internal freedom to reposition its elbow, shoulder, and wrist with zero instantaneous hand-position velocity. This does not fix hand orientation or guarantee a finite self-motion without updating the Jacobian.

1.3 The Dot Product, Cross Product, and Inner Product Spaces

Two fundamental operations exist for combining vectors. Understanding them geometrically is essential for everything from computing kinetic energy to defining stability metrics.

1.3.1 The Dot Product (Inner Product)

The dot product (or inner product) takes two vectors and returns a scalar. It measures how much two vectors “align” with each other.

Definition 1.8. Dot Product

In an orthonormal Euclidean basis, the *dot product* of vectors $\mathbf{a}, \mathbf{b} \in \mathbb{R}^n$ is:

$$\mathbf{a} \cdot \mathbf{b} = \mathbf{a}^T \mathbf{b} = \sum_{i=1}^n a_i b_i \quad (1.13)$$

In $\mathbb{R}^3$, this expands to $a_x b_x + a_y b_y + a_z b_z$.

The geometric interpretation connects the algebraic formula to angles:

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos(\theta) \quad (1.14)$$

where θ is the angle between $\mathbf{a}$ and $\mathbf{b}$, and $\|\mathbf{a}\| = \sqrt{\mathbf{a} \cdot \mathbf{a}}$ is the Euclidean norm (length) of $\mathbf{a}$.

Derivation of Equation (1.14). Consider two vectors $\mathbf{a}$ and $\mathbf{b}$ with an angle θ between them. The vector $\mathbf{a} - \mathbf{b}$ forms the third side of a triangle. By the law of cosines:

$$\|\mathbf{a} - \mathbf{b}\|^2 = \|\mathbf{a}\|^2 + \|\mathbf{b}\|^2 - 2 \|\mathbf{a}\| \|\mathbf{b}\| \cos \theta \quad (1.15)$$

Expanding the left side using the definition of the norm:

$$\|\mathbf{a} - \mathbf{b}\|^2 = (\mathbf{a} - \mathbf{b})^T (\mathbf{a} - \mathbf{b}) \quad (1.16)$$

$$= \mathbf{a}^T \mathbf{a} - 2 \mathbf{a}^T \mathbf{b} + \mathbf{b}^T \mathbf{b} \quad (1.17)$$

$$= \|\mathbf{a}\|^2 - 2 \mathbf{a}^T \mathbf{b} + \|\mathbf{b}\|^2 \quad (1.18)$$

Equating the two expressions:

$$\|\mathbf{a}\|^2 - 2 \mathbf{a}^T \mathbf{b} + \|\mathbf{b}\|^2 = \|\mathbf{a}\|^2 + \|\mathbf{b}\|^2 - 2 \|\mathbf{a}\| \|\mathbf{b}\| \cos \theta \quad (1.19)$$

Canceling $\|\mathbf{a}\|^2 + \|\mathbf{b}\|^2$ from both sides and dividing by -2 :

$$\mathbf{a}^T \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos \theta \quad (1.20)$$

□

Key consequences of the dot product formula:

- If $\theta = 0$ (parallel vectors): $\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\|$ (maximum positive value).
- If $\theta = 90^\circ$ (perpendicular vectors): $\mathbf{a} \cdot \mathbf{b} = 0$. This is the fundamental test for **orthogonality**.
- If $\theta = 180^\circ$ (antiparallel vectors): $\mathbf{a} \cdot \mathbf{b} = -\|\mathbf{a}\| \|\mathbf{b}\|$ (maximum negative value).

In mechanics, $T = \frac{1}{2} \dot{\mathbf{q}}^T M(\mathbf{q}) \dot{\mathbf{q}}$ is a kinetic-energy quadratic form. An optimal-control running cost may be $\ell = \mathbf{x}^T Q \mathbf{x} + \mathbf{u}^T R \mathbf{u}$; a trajectory cost integrates it and may add a terminal cost. Generalized-force projection $\tau_i = \mathcal{F}^T \mathcal{S}_i$ is a dual pairing of a wrench and a joint-motion direction, with consistent frame, reference point and component order. These expressions have different units and meanings; none is automatically a Euclidean angle between like physical quantities.

Orthogonal Projection

A common operation is projecting one vector onto another. For $\mathbf{a} \neq 0$, the **scalar projection** is $\mathbf{a}^T \mathbf{b} / \|\mathbf{a}\|$. The **vector projection** is:

$$\text{proj}_{\mathbf{a}} \mathbf{b} = \frac{\mathbf{a} \cdot \mathbf{b}}{\mathbf{a} \cdot \mathbf{a}} \mathbf{a} = \frac{\mathbf{a}^T \mathbf{b}}{\|\mathbf{a}\|^2} \mathbf{a} \quad (1.21)$$

The residual $\mathbf{b} - \text{proj}_{\mathbf{a}} \mathbf{b}$ is orthogonal to $\mathbf{a}$. This decomposition—splitting a vector into a component along a direction and a component perpendicular to it—is the geometric core of the Gram-Schmidt process and underlies the QR decomposition used in numerical linear algebra.

1.3.2 The Cross Product and the Skew-Symmetric Matrix

The standard cross product used here is defined in oriented three-dimensional Euclidean space. Its result is orthogonal to both inputs. Parallel inputs, or a zero input, give the zero vector, which has no direction.

Definition 1.9. Cross Product

For vectors $\mathbf{a}, \mathbf{b} \in \mathbb{R}^3$, the *cross product* is:

$$\mathbf{a} \times \mathbf{b} = \begin{bmatrix} a_y b_z - a_z b_y \\ a_z b_x - a_x b_z \\ a_x b_y - a_y b_x \end{bmatrix} \quad (1.22)$$

The direction of $\mathbf{a} \times \mathbf{b}$ is determined by the *right-hand rule*: curl the fingers of your right hand from $\mathbf{a}$ toward $\mathbf{b}$; your thumb points in the direction of the cross product. The magnitude satisfies:

$$\|\mathbf{a} \times \mathbf{b}\| = \|\mathbf{a}\| \|\mathbf{b}\| \sin(\theta) \quad (1.23)$$

which equals the area of the parallelogram spanned by $\mathbf{a}$ and $\mathbf{b}$.

The cross product is *anti-commutative*: $\mathbf{a} \times \mathbf{b} = -(\mathbf{b} \times \mathbf{a})$. This has direct physical consequences. If you apply a force $\mathbf{F}$ at a displacement $\mathbf{r}$ from a pivot, the resulting torque is $\boldsymbol{\tau} = \mathbf{r} \times \mathbf{F}$. Reversing the order reverses the torque direction.

Crucially, the cross product can be rewritten as a matrix multiplication using the *skew-symmetric matrix*:

$$[\mathbf{a}_\times] = \begin{bmatrix} 0 & -a_z & a_y \\ a_z & 0 & -a_x \\ -a_y & a_x & 0 \end{bmatrix} \implies \mathbf{a} \times \mathbf{b} = [\mathbf{a}_\times] \mathbf{b} \quad (1.24)$$

Proposition 1. *The skew-symmetric matrix $[\mathbf{a}_\times]$ satisfies:*

- (i) $[\mathbf{a}_\times]^T = -[\mathbf{a}_\times]$ (*skew-symmetry; equivalently, all diagonal entries are zero and the off-diagonal entries are negated across the diagonal*).
- (ii) $\mathbf{a}^T [\mathbf{a}_\times] \mathbf{b} = 0$ for all $\mathbf{b}$ (*the cross product is always perpendicular to both input vectors*).
- (iii) *The eigenvalues of $[\mathbf{a}_\times]$ are $\{0, +i\|\mathbf{a}\|, -i\|\mathbf{a}\|\}$ (purely imaginary, reflecting the rotational nature of the cross product).*

Proof. (i) is immediate from inspection of Equation (1.24): transposing negates every off-diagonal element while the diagonal remains zero.

(ii) We compute $\mathbf{a}^T ([\mathbf{a}_\times] \mathbf{b}) = \mathbf{a}^T (\mathbf{a} \times \mathbf{b})$. Since $\mathbf{a} \times \mathbf{b}$ is perpendicular to $\mathbf{a}$ (by the geometric definition of the cross product), the dot product is zero.

(iii) The characteristic polynomial $\det([\mathbf{a}_\times] - \lambda \mathbf{I}) = -\lambda^3 - \|\mathbf{a}\|^2 \lambda = -\lambda(\lambda^2 + \|\mathbf{a}\|^2)$. The roots are $\lambda = 0$ and $\lambda = \pm i\|\mathbf{a}\|$. $\square$

This skew-symmetric matrix representation is absolutely pivotal for Chapters 4 and 6, where we show that the Lie algebra $\mathfrak{so}(3)$ —the tangent space to the rotation group at the identity—consists entirely of 3×3 skew-symmetric matrices. An angular-velocity vector corresponds to such a matrix after choosing a body or spatial trivialization; a tangent matrix at a general rotation is not itself necessarily skew-symmetric.

The Scalar Triple Product


The scalar triple product $\mathbf{a} \cdot (\mathbf{b} \times \mathbf{c})$ computes the signed volume of the parallelepiped formed by three vectors $\mathbf{a}, \mathbf{b}, \mathbf{c} \in \mathbb{R}^3$. It can be computed as a determinant:

$$\mathbf{a} \cdot (\mathbf{b} \times \mathbf{c}) = \det \begin{bmatrix} a_x & b_x & c_x \\ a_y & b_y & c_y \\ a_z & b_z & c_z \end{bmatrix} \quad (1.25)$$

If this determinant is zero, the three vectors are coplanar and linearly dependent—a direct connection between the cross product and the rank of the matrix formed by stacking the vectors as columns.

1.4 Eigenvalues, Eigenvectors, and Quadratic Forms

When you multiply a general vector by a matrix, the resulting vector typically changes both its magnitude and its direction. Over the complex numbers, every square matrix of positive size has an eigenvector. Real invariant lines need not exist: a planar quarter-turn has eigenvalues $\pm i$ and no real eigenvector. A real eigenvector spans a line mapped into itself; a positive eigenvalue preserves its ray, a negative one reverses it, and zero annihilates it. These invariant directions reveal the deepest structural properties of the linear transformation.



Example: positive eigenvalues preserve eigenvector rays

Figure 1.2: Eigenvectors and a Generic Vector Under $\mathbf{A} = \text{diag}(2, 1/2)$

1.4.1 The Eigenvalue Equation

Definition 1.10. Eigenvalues and Eigenvectors

Given a square matrix $\mathbf{M} \in \mathbb{R}^{n \times n}$, a nonzero vector $\mathbf{v} \in \mathbb{C}^n$ is called an *eigenvector* of $\mathbf{M}$ if there exists a scalar $\lambda \in \mathbb{C}$ such that:

$$\mathbf{M}\mathbf{v} = \lambda\mathbf{v} \quad (1.26)$$

The scalar λ is the corresponding *eigenvalue*. The requirement that $\mathbf{v} \neq \mathbf{0}$ is essential; otherwise the equation would be trivially satisfied for any λ .

Deriving the Characteristic Polynomial

How do we actually find the eigenvalues? We rearrange Equation (1.26):

$$\mathbf{M}\mathbf{v} - \lambda\mathbf{v} = \mathbf{0} \implies (\mathbf{M} - \lambda\mathbf{I})\mathbf{v} = \mathbf{0} \quad (1.27)$$

where $\mathbf{I}$ is the $n \times n$ identity matrix. For a nonzero solution $\mathbf{v}$ to exist, the matrix $(\mathbf{M} - \lambda\mathbf{I})$ must be singular. By our earlier discussion, this means its determinant must vanish:

$$\det(\mathbf{M} - \lambda\mathbf{I}) = 0 \quad (1.28)$$

This is the **characteristic equation**. Expanding the determinant produces a polynomial of degree n in λ , called the **characteristic polynomial**. By the Fundamental Theorem of Algebra, this polynomial has exactly n roots (counted with multiplicity, possibly complex), giving us n eigenvalues.

Example 1.3. Computing Eigenvalues of a 2×2 Matrix

Consider the dynamics matrix of a simple spring system:

$$\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -4 & -5 \end{bmatrix} \quad (1.29)$$

Step 1: Form $\mathbf{A} - \lambda\mathbf{I}$:

$$\mathbf{A} - \lambda\mathbf{I} = \begin{bmatrix} -\lambda & 1 \\ -4 & -5 - \lambda \end{bmatrix} \quad (1.30)$$

Step 2: Compute the determinant and set it to zero:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = (-\lambda)(-5 - \lambda) - (1)(-4) \quad (1.31)$$

$$= \lambda^2 + 5\lambda + 4 \quad (1.32)$$

$$= (\lambda + 1)(\lambda + 4) = 0 \quad (1.33)$$

Step 3: The eigenvalues are $\lambda_1 = -1$ and $\lambda_2 = -4$. Both are negative, meaning this system is **asymptotically stable**—any perturbation decays exponentially back to equilibrium. The mode associated with $\lambda_1 = -1$ decays slowly (time constant $\tau_1 = 1$ second), while the mode associated with $\lambda_2 = -4$ decays four times faster ($\tau_2 = 0.25$ seconds).

Step 4: Find the eigenvectors. For $\lambda_1 = -1$, solve $(\mathbf{A} + \mathbf{I})\mathbf{v}_1 = \mathbf{0}$:

$$\begin{bmatrix} 1 & 1 \\ -4 & -4 \end{bmatrix} \mathbf{v}_1 = \mathbf{0} \implies \mathbf{v}_1 = \begin{bmatrix} 1 \\ -1 \end{bmatrix} \quad (1.34)$$

For $\lambda_2 = -4$, solve $(\mathbf{A} + 4\mathbf{I})\mathbf{v}_2 = \mathbf{0}$:

$$\begin{bmatrix} 4 & 1 \\ -4 & -1 \end{bmatrix} \mathbf{v}_2 = \mathbf{0} \implies \mathbf{v}_2 = \begin{bmatrix} 1 \\ -4 \end{bmatrix} \quad (1.35)$$

1.4.2 Physical Interpretation of Eigenvalues

1. **Inertia tensors and principal axes.** The symmetric rotational inertia tensor about the centre of mass has orthogonal principal axes and corresponding principal moments.

For an isolated rigid body with three distinct principal moments, steady torque-free rotation about the intermediate axis is unstable, whereas the minimum and maximum axes are stable to small rotational perturbations. An eigenvector of inertia alone does not establish stability under applied torques, damping or feedback.

2. **Constant linear dynamics.** For $\dot{\mathbf{x}} = A\mathbf{x}$, all eigenvalues strictly in the open left half-plane imply exponential asymptotic stability. Any eigenvalue with positive real part implies instability. With all real parts nonpositive, stability additionally requires each eigenvalue on the imaginary axis to be semisimple: its Jordan blocks must have size one. Such modes need not decay.
3. **Nonlinear and time-varying systems.** At an equilibrium of a continuously differentiable autonomous system, a Hurwitz linearization establishes local exponential stability, and an eigenvalue with positive real part establishes instability. A zero real part can be inconclusive: $\dot{x} = -x^3$ and $\dot{x} = x^3$ have the same zero linearization and opposite stability. Instantaneous eigenvalues of a time-varying matrix are not a substitute for its state-transition operator.

For example,

$$A_0 = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, \quad e^{tA_0} \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} t \\ 1 \end{pmatrix}$$

is unbounded despite two zero eigenvalues. By contrast, $A = 0$ has bounded constant solutions.

Even negative eigenvalues do not guarantee monotone decrease in a chosen norm. For

$$A = \begin{pmatrix} -1 & 10 \\ 0 & -1 \end{pmatrix}, \quad e^{tA} = e^{-t} \begin{pmatrix} 1 & 10t \\ 0 & 1 \end{pmatrix},$$

the initial vector $(0, 1)^T$ has norm $\sqrt{101}/e > 3.6$ at $t = 1$ before eventually decaying. The transient comes from the nonorthogonal modal structure. Mixed physical units require a declared metric before interpreting its magnitude. For a fixed positive-definite P , the quadratic rate is

$$\frac{d}{dt}(\mathbf{x}^T P \mathbf{x}) = \mathbf{x}^T (A^T P + P A) \mathbf{x}.$$

Asymptotic stability and short-horizon amplification answer different questions in a finite swing.

Example 1.4. Visualizing Eigenvectors: The Stretching Rubber Sheet

A general invertible planar map sends a unit circle to an ellipse. Its output axes are *left singular vectors*, while the corresponding input directions are right singular vectors. They need not be the same directions.

The shear

$$S = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$$

has eigenvalues 1, 1 but singular values $(1 + \sqrt{5})/2$ and $(\sqrt{5} - 1)/2$. It stretches and contracts different directions even though both eigenvalues are one.

For a symmetric matrix one can choose an orthonormal eigenbasis. The image of the unit ball has semiaxis lengths $|\lambda_i|$; a negative eigenvalue also reverses its eigenvector,

and zero gives a collapsed direction. For a positive-definite symmetric matrix the semiaxis lengths equal the positive eigenvalues. This is the special case in which the rubber-sheet picture directly matches the eigenvalue magnitudes.

1.4.3 The Spectral Theorem for Symmetric Matrices

Among all matrices, *symmetric* matrices ($\mathbf{M} = \mathbf{M}^T$) occupy a privileged position. They arise naturally throughout mechanics (mass matrices, stiffness matrices, inertia tensors) and control theory (cost matrices, Lyapunov functions).

Theorem 1.3. The Spectral Theorem

If $\mathbf{M} \in \mathbb{R}^{n \times n}$ is symmetric ($\mathbf{M} = \mathbf{M}^T$), then:

- (i) All eigenvalues of $\mathbf{M}$ are real.
- (ii) Eigenvectors corresponding to distinct eigenvalues are orthogonal.
- (iii) $\mathbf{M}$ can be factored as $\mathbf{M} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$, where $\mathbf{Q}$ is an orthogonal matrix ($\mathbf{Q}^T\mathbf{Q} = \mathbf{I}$) whose columns are the eigenvectors, and $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues.

Proof that eigenvalues of a real symmetric matrix are real. Let $\mathbf{M}$ be real symmetric. Suppose $\mathbf{M}\mathbf{v} = \lambda\mathbf{v}$ with $\mathbf{v} \neq 0$, allowing complex entries. We will show that λ must be real.

Take the complex conjugate transpose (denoted by $*$) of both sides of $\mathbf{M}\mathbf{v} = \lambda\mathbf{v}$:

$$\mathbf{v}^*\mathbf{M}^T = \bar{\lambda}\mathbf{v}^* \quad (1.36)$$

Since $\mathbf{M} = \mathbf{M}^T$ (and $\mathbf{M}$ is real, so $\mathbf{M}^* = \mathbf{M}^T = \mathbf{M}$):

$$\mathbf{v}^*\mathbf{M} = \bar{\lambda}\mathbf{v}^* \quad (1.37)$$

Now multiply the original equation on the left by $\mathbf{v}^*$:

$$\mathbf{v}^*\mathbf{M}\mathbf{v} = \lambda\mathbf{v}^*\mathbf{v} \quad (1.38)$$

And multiply the conjugated equation on the right by $\mathbf{v}$:

$$\mathbf{v}^*\mathbf{M}\mathbf{v} = \bar{\lambda}\mathbf{v}^*\mathbf{v} \quad (1.39)$$

Equating the right-hand sides:

$$\lambda\mathbf{v}^*\mathbf{v} = \bar{\lambda}\mathbf{v}^*\mathbf{v} \quad (1.40)$$

Since $\mathbf{v} \neq \mathbf{0}$, we have $\mathbf{v}^*\mathbf{v} = \sum_i |v_i|^2 > 0$. Dividing both sides by $\mathbf{v}^*\mathbf{v}$ yields $\lambda = \bar{\lambda}$, which is true if and only if λ is real. $\square$

Proof of Orthogonality for Distinct Eigenvalues. For real symmetric $\mathbf{M}$, choose real eigenvectors with distinct eigenvalues:

$$\mathbf{M}\mathbf{v}_1 = \lambda_1\mathbf{v}_1, \quad \mathbf{M}\mathbf{v}_2 = \lambda_2\mathbf{v}_2, \quad \lambda_1 \neq \lambda_2.$$

Compute:

$$\lambda_1(\mathbf{v}_1^T \mathbf{v}_2) = (\mathbf{M}\mathbf{v}_1)^T \mathbf{v}_2 = \mathbf{v}_1^T \mathbf{M}^T \mathbf{v}_2 = \mathbf{v}_1^T \mathbf{M} \mathbf{v}_2 = \mathbf{v}_1^T (\lambda_2 \mathbf{v}_2) = \lambda_2(\mathbf{v}_1^T \mathbf{v}_2) \quad (1.41)$$

Therefore $(\lambda_1 - \lambda_2)(\mathbf{v}_1^T \mathbf{v}_2) = 0$. Since $\lambda_1 \neq \lambda_2$, we must have $\mathbf{v}_1^T \mathbf{v}_2 = 0$. The eigenvectors are orthogonal. $\square$

The spectral decomposition $\mathbf{M} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$ applies an orthogonal coordinate change to eigenvector axes, scales each coordinate by its eigenvalue, and applies the inverse orthogonal change. Negative eigenvalues reverse signs. Unlike the SVD's nonnegative singular scaling, this diagonal scaling retains the eigenvalue signs.

1.4.4 Quadratic Forms and Positive Definiteness

In optimization and energy mechanics, we frequently encounter *quadratic forms*—scalar-valued functions that are quadratic in a vector argument:

Definition 1.11. Quadratic Form

Given a symmetric matrix $\mathbf{M} \in \mathbb{R}^{n \times n}$, the associated *quadratic form* is the scalar function $q : \mathbb{R}^n \rightarrow \mathbb{R}$ defined by:

$$q(\mathbf{x}) = \mathbf{x}^T \mathbf{M} \mathbf{x} = \sum_{i=1}^n \sum_{j=1}^n M_{ij} x_i x_j \quad (1.42)$$

Quadratic forms appear throughout this text series:

- **Kinetic energy:** $T = \frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M}(\mathbf{q}) \dot{\mathbf{q}}$, where $\mathbf{M}(\mathbf{q})$ is the mass matrix (Chapter 11).
- **Optimal control running cost:** $\ell = \mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}$, where $\mathbf{Q}$ penalizes state deviations and $\mathbf{R}$ penalizes control effort.
- **Lyapunov functions:** $V(\mathbf{x}) = \mathbf{x}^T \mathbf{P} \mathbf{x}$, used to prove stability (Volume I).
- **Contraction metrics:** $\delta \mathbf{x}^T \mathbf{M}(x) \delta \mathbf{x}$, measuring squared infinitesimal length on manifolds (Volume I).

Definition 1.12. Positive Definiteness

A symmetric matrix $\mathbf{M}$ is called:

- **Positive definite** ($\mathbf{M} \succ 0$) if $\mathbf{x}^T \mathbf{M} \mathbf{x} > 0$ for all $\mathbf{x} \neq \mathbf{0}$.
- **Positive semi-definite** ($\mathbf{M} \succeq 0$) if $\mathbf{x}^T \mathbf{M} \mathbf{x} \geq 0$ for all $\mathbf{x}$.
- **Negative definite** ($\mathbf{M} \prec 0$) if $\mathbf{x}^T \mathbf{M} \mathbf{x} < 0$ for all $\mathbf{x} \neq \mathbf{0}$.
- **Indefinite** if $\mathbf{x}^T \mathbf{M} \mathbf{x}$ takes both positive and negative values.

Proposition 2. A symmetric matrix $\mathbf{M}$ is positive definite if and only if all of its eigenvalues are strictly positive.

Proof. By the Spectral Theorem (Theorem 1.3), $\mathbf{M} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T$ where $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_n)$ and $\mathbf{Q}$ is orthogonal. For any nonzero $\mathbf{x}$, define $\mathbf{y} = \mathbf{Q}^T \mathbf{x}$. Since $\mathbf{Q}$ is invertible and $\mathbf{x} \neq \mathbf{0}$, we have $\mathbf{y} \neq \mathbf{0}$. Then:

$$\mathbf{x}^T \mathbf{M} \mathbf{x} = \mathbf{x}^T \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^T \mathbf{x} = \mathbf{y}^T \mathbf{\Lambda} \mathbf{y} = \sum_{i=1}^n \lambda_i y_i^2 \quad (1.43)$$

($\Rightarrow$) If $\mathbf{M} \succ 0$, then for each eigenvector $\mathbf{v}_i$ (which is nonzero): $\mathbf{v}_i^T \mathbf{M} \mathbf{v}_i = \lambda_i \|\mathbf{v}_i\|^2 > 0$, so $\lambda_i > 0$.

($\Leftarrow$) If all $\lambda_i > 0$ and $\mathbf{y} \neq \mathbf{0}$, then at least one $y_i \neq 0$, so $\sum_i \lambda_i y_i^2 > 0$. $\square$

In Plain Language 1.5. Why Positive Definiteness Matters for Control

Let the origin be an equilibrium and $V(\mathbf{x}) = \mathbf{x}^T P \mathbf{x}$ with $P \succ 0$. Nonincrease of V along solutions gives a stability bound on a suitable invariant neighborhood; it does not imply convergence. An undamped oscillator has constant positive energy and never comes to rest. A negative-definite derivative, with the usual autonomous-system regularity and domain hypotheses, gives asymptotic convergence; a merely semidefinite derivative needs an additional invariant-set argument. V is a mathematical measure and need not be physical energy. Time-varying or global conclusions require their corresponding uniformity, domain and growth assumptions.

The Cholesky Decomposition

A practical consequence of positive definiteness: any positive definite matrix $\mathbf{M}$ can be uniquely factored as:

$$\mathbf{M} = \mathbf{L}\mathbf{L}^T \quad (1.44)$$

where $\mathbf{L}$ is a lower triangular matrix with strictly positive diagonal entries. This is the **Cholesky decomposition**. It is numerically efficient (roughly half the cost of a general LU factorization) and is the standard method for solving linear systems involving positive definite matrices in physics engines. One usually solves the system rather than explicitly forming its inverse. MuJoCo's documented implementation uses a sparse $\mathbf{L}^T \mathbf{D} \mathbf{L}$ factorization of joint-space inertia and back-substitution; this is related to, but not literally the displayed $\mathbf{L}\mathbf{L}^T$ algorithm.¹

For a real symmetric input in exact arithmetic, Cholesky succeeds with strictly positive pivots precisely when the input is positive definite. In floating-point arithmetic, failure can also reflect rounding near the definiteness boundary, incorrect symmetry/storage assumptions or numerical scaling; inspect the data and residual before interpreting it physically. A routine that reads only one triangle does not by itself verify symmetry of the supplied array.

1.5 The Singular Value Decomposition (SVD)

What if a matrix is not square? What if it does not have a full set of eigenvectors? The **Singular Value Decomposition** (SVD) is the ultimate, universal factorization of *any* matrix $\mathbf{M} \in \mathbb{R}^{n \times m}$, regardless of its shape or rank.

¹<https://mujoco.readthedocs.io/en/stable/computation/index.html>

Theorem 1.4. The Singular Value Decomposition

Every matrix $\mathbf{M} \in \mathbb{R}^{n \times m}$ can be factored as:

$$\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T \quad (1.45)$$

where:

- $\mathbf{U} \in \mathbb{R}^{n \times n}$ is an orthogonal matrix ($\mathbf{U}^T \mathbf{U} = \mathbf{I}$). Its columns are called the *left singular vectors*.
- $\mathbf{\Sigma} \in \mathbb{R}^{n \times m}$ is a diagonal matrix with non-negative entries $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_{\min(n,m)} \geq 0$ on the diagonal. These are the *singular values*.
- $\mathbf{V} \in \mathbb{R}^{m \times m}$ is an orthogonal matrix. Its columns are the *right singular vectors*.

1.5.1 Derivation of the SVD From the Eigendecomposition

The SVD can be derived from the spectral theorem applied to related symmetric matrices. Here is the construction:

Step 1. Form the symmetric matrix $\mathbf{M}^T \mathbf{M} \in \mathbb{R}^{m \times m}$. This matrix is always positive semi-definite because for any $\mathbf{x}$:

$$\mathbf{x}^T (\mathbf{M}^T \mathbf{M}) \mathbf{x} = (\mathbf{M} \mathbf{x})^T (\mathbf{M} \mathbf{x}) = \|\mathbf{M} \mathbf{x}\|^2 \geq 0 \quad (1.46)$$

Step 2. By the Spectral Theorem, $\mathbf{M}^T \mathbf{M}$ has an eigendecomposition:

$$\mathbf{M}^T \mathbf{M} = \mathbf{V} \mathbf{\Lambda} \mathbf{V}^T \quad (1.47)$$

where $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_m)$ with $\lambda_i \geq 0$ (since $\mathbf{M}^T \mathbf{M}$ is positive semi-definite). Order these eigenvalues decreasingly and take their square roots. Only $\min(n, m)$ singular values are listed in the rectangular diagonal $\mathbf{\Sigma}$; when $m > n$, the remaining eigenvalues of $\mathbf{M}^T \mathbf{M}$ are necessarily zero.

Step 3. For each nonzero singular value σ_i , define the corresponding left singular vector:

$$\mathbf{u}_i = \frac{1}{\sigma_i} \mathbf{M} \mathbf{v}_i \quad (1.48)$$

These vectors are orthonormal (a straightforward verification using $\mathbf{M}^T \mathbf{M} \mathbf{v}_i = \sigma_i^2 \mathbf{v}_i$).

Step 4. Extend $\{\mathbf{u}_1, \dots, \mathbf{u}_r\}$ to a complete orthonormal basis for $\mathbb{R}^n$ (where $r = \text{rank}(\mathbf{M})$). Assemble these into the columns of $\mathbf{U}$. Then $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ follows by construction. This is an existence proof, not a recommendation to form $\mathbf{M}^T \mathbf{M}$ numerically: doing so squares the condition number in the full-column-rank case and can obscure small singular values.

1.5.2 Geometric Interpretation of the SVD

The SVD reveals that *any* linear transformation, no matter how complex, can be decomposed into exactly three simple geometric operations:

1. **Reorient the input space by a rotation or reflection** (apply $\mathbf{V}^T$): align the input with the principal directions.

2. **Scale each axis independently** (apply Σ): stretch or shrink along each principal direction.
3. **Reorient the output space by a rotation or reflection** (apply U): reorient the result in the output space.

The image of the unit ball is a solid, possibly degenerate ellipsoid in the output space. Rank loss collapses some dimensions; for a rectangular map, the image of a sphere need not be only an ellipsoid's boundary. The singular values are the semi-axis lengths of this ellipsoid, and the left singular vectors point along the ellipsoid's principal axes.

1.5.3 The Condition Number

For a nonzero, full-rank rectangular matrix, its 2-norm condition number is

$$\kappa_2(M) = \frac{\sigma_{\max}}{\sigma_{\min}},$$

where the denominator is the smallest of the $\min(n, m)$ listed singular values. Rank deficiency is conventionally assigned infinite condition number here. A wide, full-row-rank matrix may have finite condition number and a nontrivial input kernel simultaneously.

For a nonsingular square system $M\mathbf{x} = \mathbf{b}$ with fixed M and nonzero $\mathbf{b}$, the relative perturbation bound is

$$\frac{\|\delta\mathbf{x}\|_2}{\|\mathbf{x}\|_2} \leq \kappa_2(M) \frac{\|\delta\mathbf{b}\|_2}{\|\mathbf{b}\|_2}.$$

Absolute inverse gain is instead $\|M^{-1}\|_2 = 1/\sigma_{\min}$. For $M = 10^{-3}I$, the condition number is one but the absolute inverse gain is 1000. Forward gain is $\sigma_{\max}$, not the condition number. Matrix perturbations and general least-squares problems need additional bounds.

Singular values depend on coordinate units and chosen norms. A task Jacobian mixing translation and rotation, or joint coordinates measured in degrees rather than radians, needs explicit scaling or a metric before its condition number is a physical comparison.

1.5.4 Truncated SVD and Low-Rank Approximation

One of the most powerful applications of the SVD is data compression. Given a matrix $\mathbf{M}$ with singular values $\sigma_1 \geq \sigma_2 \geq \dots$, the best rank- k approximation (in the Frobenius norm) is:

$$\mathbf{M}_k = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^T \quad (1.49)$$

This result (the Eckart-Young-Mirsky theorem) underpins dimensionality reduction techniques like Principal Component Analysis (PCA), which identifies the dominant modes of variation in high-dimensional data such as motion capture recordings of golf swings.

Example 1.5. The SVD of the Jacobian (Manipulability Ellipsoid)

For seven independent joint velocities and a three-dimensional hand-position task, $J \in \mathbb{R}^{3 \times 7}$ maps $\dot{\mathbf{q}}$ to $\mathbf{v}_H$. With the declared budget $\|\dot{\mathbf{q}}\|_2 \leq 1$:

- The columns of U are task-space singular directions. A direction with zero

singular value is absent from the instantaneous velocity image.

- The three singular values give the semiaxes of the image of that joint-speed ball. They are not actual maximum speeds without this budget, its units and actuator feasibility.
- The columns of V , not of V^T , give the corresponding joint-velocity directions. Its remaining columns span the input kernel.

For a positive-definite joint-speed metric H and task metric W , normalize the map as

$$\tilde{J} = W^{1/2} J H^{-1/2}, \quad \dot{\mathbf{q}}^T H \dot{\mathbf{q}} \leq 1.$$

Singular values of this normalized map compare declared physical budgets; unscaled Euclidean values need not survive a change of units.

Virtual work gives $\boldsymbol{\tau} = J^T \mathbf{F}$ and power equality $\boldsymbol{\tau}^T \dot{\mathbf{q}} = \mathbf{F}^T \mathbf{v}_H$. Under the dual torque budget $\boldsymbol{\tau}^T H^{-1} \boldsymbol{\tau} \leq 1$, feasible quasistatic task forces satisfy

$$\mathbf{F}^T J H^{-1} J^T \mathbf{F} \leq 1.$$

At full row rank this gives reciprocal force/velocity semiaxes in consistent normalized coordinates. At rank loss it may be unbounded in a task-force direction supported by ideal constraints. It does not predict unlimited real force: structural strength, joint/contact limits, compliance, actuation, balance and dynamics still matter.

1.5.5 Instantaneous Rank and Finite-Time Control

A full-row-rank 3×7 task Jacobian has a four-dimensional kernel without being singular. More generally, a kinematic singularity is a drop below the maximum rank attainable by the specified mechanism and task; an inherently lower-dimensional image is not itself evidence of a singular configuration. This kinematic redundancy differs from underactuation, which concerns the rank and admissible set of the actuator map in the dynamics.

Initial task-null velocity also need not preserve finite position. For two unit planar links,

$$\mathbf{p}(\mathbf{q}) = \begin{pmatrix} \cos q_1 + \cos(q_1 + q_2) \\ \sin q_1 + \sin(q_1 + q_2) \end{pmatrix},$$

the straight configuration has $J(0) = \begin{pmatrix} 0 & 0 \\ 2 & 1 \end{pmatrix}$. Its kernel contains $(1, -2)^T$, but the straight coordinate step $\mathbf{q}(t) = t(1, -2)^T$ gives $\mathbf{p}(t) = (2 \cos t, 0)^T$, not a stationary hand.

Conversely, a one-dimensional instantaneous input can control more than one state over time. For the double integrator

$$A = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, \quad B = \begin{pmatrix} 0 \\ 1 \end{pmatrix},$$

B has rank one while $[B, AB]$ has rank two. Its controllability Gramian is

$$W_c(T) = \begin{pmatrix} T^3/3 & T^2/2 \\ T^2/2 & T \end{pmatrix}, \quad \det W_c(T) = T^4/12 > 0 \quad (T > 0).$$

This unconstrained linear result does not remove finite torque bounds or contact requirements. It shows why instantaneous rank, finite-time reachability, observation and a successful golf delivery must be evaluated using their respective maps and assumptions.

Observation has the same temporal distinction. If position alone is measured, $C = [1, 0]$ has an instantaneous kernel containing velocity. Yet the observability matrix $\mathcal{O} = [C^T, (CA)^T]^T = I_2$ has full rank: position measurements over time reveal velocity in this ideal model. A kernel of C is not necessarily an unobservable dynamical subspace.

1.6 Dyads, Dyadics, and Tensor Products

The matrix decompositions we have studied so far—eigendecomposition, SVD—break a matrix into pieces that reveal its geometric action. In this section we meet a complementary idea: building matrices up from the simplest possible pieces, called *dyads*. This viewpoint is not merely algebraic; it is the language that mechanics has used for over a century to describe inertia, stress, and stiffness without ever choosing a coordinate frame [31].

1.6.1 What Is a Dyad?

Definition 1.13. Dyad (Outer Product)

Given two vectors $\mathbf{a} \in \mathbb{R}^m$ and $\mathbf{b} \in \mathbb{R}^n$, the **dyad** (or **outer product**) is the $m \times n$ matrix

$$\mathbf{D} = \mathbf{a} \mathbf{b}^T.$$

This matrix has rank at most one: every column is a scalar multiple of $\mathbf{a}$, and every row is a scalar multiple of $\mathbf{b}^T$.

A dyad encodes *two* directions simultaneously. When it acts on a vector $\mathbf{v}$,

$$\mathbf{D} \mathbf{v} = (\mathbf{a} \mathbf{b}^T) \mathbf{v} = \mathbf{a} \langle \mathbf{b}, \mathbf{v} \rangle,$$

it first *senses* how much $\mathbf{v}$ aligns with $\mathbf{b}$ (via the inner product $\langle \mathbf{b}, \mathbf{v} \rangle$) and then *responds* along the direction $\mathbf{a}$, scaled by that alignment.

In Plain Language 1.6. Dyads: Direction of Action and Direction of Sensitivity

The dyad $\mathbf{a} \mathbf{b}^T$ measures the component of an input along $\mathbf{b}$ and produces a response along $\mathbf{a}$. It has rank one if both vectors are nonzero, and rank zero otherwise. This is an algebraic response channel, not automatically a passive spring.

For an unstressed conservative spring along a unit direction $\mathbf{n}$, stored energy $\frac{1}{2} k (\mathbf{n}^T \mathbf{u})^2$ gives restoring force $-k \mathbf{n} \mathbf{n}^T \mathbf{u}$. Its stiffness is symmetric. A generic cross-direction dyad is nonsymmetric and cannot by itself be the Hessian of a smooth elastic potential. For example, the field $(0, x)$ does one unit of work around the positively oriented unit square; a conservative spring would do zero net closed-path work.

Example 1.6. A Rank-1 Linear Response

Let $\mathbf{a} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$ (vertical) and $\mathbf{b} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$ (horizontal). Then

$$\mathbf{D} = \mathbf{a}\mathbf{b}^T = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}.$$

Applied to a displacement $\mathbf{v} = \begin{bmatrix} 3 \\ 5 \\ 2 \end{bmatrix}$:

$$\mathbf{D}\mathbf{v} = \begin{bmatrix} 0 \\ 3 \\ 0 \end{bmatrix}.$$

The dyad extracts the horizontal component (3) and produces a purely vertical response—exactly the one-channel behaviour described above.

1.6.2 From Dyads to Dyadics

A nonzero dyad has rank one. Physical stiffness, inertia and stress maps may have full rank or may have null directions. We get there by *summing* dyads.

Definition 1.14. Dyadic (Tensor Product Expansion)

A **dyadic** is a finite sum of dyads:

$$\mathbf{T} = \sum_{k=1}^r \mathbf{a}_k \mathbf{b}_k^T.$$

Every $m \times n$ matrix can be written in this form with $r \leq \min(m, n)$ terms.

This is not a new fact—it is the column–row factorisation you already know from the SVD. If $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$, then

$$\mathbf{A} = \sum_{k=1}^r \sigma_k \mathbf{u}_k \mathbf{v}_k^T,$$

which is a dyadic with r terms (the rank of $\mathbf{A}$). The SVD merely chooses the *orthogonal* dyad decomposition; infinitely many other decompositions exist.

Abstractly, a linear map $V \rightarrow W$ belongs to $W \otimes V^*$: its input slot is a covector. In standard Euclidean coordinates the inner product identifies $\mathbb{R}^n$ with its dual, so $\mathbf{a}\mathbf{b}^T$ can be written using two vector coordinate arrays and identified with an elementary tensor in $\mathbb{R}^m \otimes \mathbb{R}^n$. Under general coordinate changes one must retain which slots are vectors and which are covectors.

1.6.3 Physical Meaning in Mechanics

The power of the dyadic viewpoint is that it names the two “slots” of a matrix: input direction and output direction. Three central objects in mechanics are dyadics.

The Inertia Tensor. About the centre of mass, the rotational inertia of a rigid body maps angular velocity to angular momentum:

$$\mathbf{L} = \mathbf{I}\boldsymbol{\omega}. \quad (1.50)$$

Here $\boldsymbol{\omega}$ is the “input” (the axis and rate of spin) and $\mathbf{L}$ is the “output” (the direction and magnitude of angular momentum). For a generic direction in a body with unequal principal moments, $\mathbf{L}$ and $\boldsymbol{\omega}$ need not be parallel; they are parallel along every principal axis even without an isotropic inertia; the inertia tensor is the dyadic that encodes this directional coupling.

The Stress Tensor. The Cauchy stress tensor $\boldsymbol{\sigma}$ maps a surface-normal direction $\hat{\mathbf{n}}$ to the traction (force per unit area) $\mathbf{t}$ acting on that surface:

$$\mathbf{t} = \boldsymbol{\sigma} \hat{\mathbf{n}}.$$

Input: the orientation of a small area element. Output: the force that the material on one side exerts on the other.

The Stiffness Matrix. For a small-displacement, unstressed truss, let $\mathbf{u}$ collect all nodal displacements and $B_e \mathbf{u}$ be member e ’s axial extension. The external balancing force is $\mathbf{f} = K \mathbf{u}$, with

$$K = \sum_e k_e B_e^T B_e, \quad k_e = \frac{E_e A_e}{L_e}.$$

For a member between two three-dimensional nodes, its nonzero incidence blocks are $B_e = [-\mathbf{n}_e^T, \mathbf{n}_e^T]$. Thus one member contributes a rank-one *nodal* matrix, not merely a 3×3 direction matrix. A single anchored-node star is a special case in which the free-node stiffness reduces to $\sum_e k_e \mathbf{n}_e \mathbf{n}_e^T$.

Unrestrained rigid translations give null modes of the assembled matrix. Prestress, geometric stiffness, constraints and nonlinear deformation require additional terms or a tangent linearization. Summing member directions alone does not encode the connectivity of a general truss.

In Plain Language 1.7. Why Engineers Say “Tensor”

When a mechanical engineer refers to the “inertia tensor” or “stress tensor,” the usual rigid-body inertia and classical Cauchy stress have symmetric 3×3 representations. Symmetry of Cauchy stress here assumes the classical angular-momentum balance without independent couple stresses; not every tensor or linear response is symmetric. The word *tensor* signals that the object transforms in a specific, predictable way under changes of coordinate frame. It is not just a bag of nine numbers; it is a physical machine that takes one vector in and produces another vector out, regardless of how you choose your axes.

1.6.4 The Coordinate-Free Perspective

A dyadic is a geometric object that lives in the space of linear maps $\mathbb{R}^n \rightarrow \mathbb{R}^n$. Choosing a basis gives it a matrix representation, but the underlying map does not depend on that choice.

Example 1.7. Inertia Tensor of a Point Mass

A point mass m at displacement $\mathbf{r}$ from a chosen reference origin contributes inertia tensor

$$\mathbf{I} = m(\|\mathbf{r}\|^2 \mathbf{1} - \mathbf{r} \mathbf{r}^T), \quad (1.51)$$

where $\mathbf{1}$ is the 3×3 identity. Notice that no basis vectors appear— $\mathbf{r} \mathbf{r}^T$ is a dyad formed from the position vector with itself, and $\mathbf{1}$ is the identity map. This expression is *coordinate-free*: it is valid in every frame simultaneously.

When we change from a body-fixed frame $\{B\}$ to a world frame $\{W\}$ via a rotation $\mathbf{R} \in \text{SO}(3)$, the matrix representation of the inertia tensor transforms by the **congruence transformation**:

$$\mathbf{I}^W = \mathbf{R} \mathbf{I}^B \mathbf{R}^T. \quad (1.52)$$

The congruence form $\mathbf{R}(\cdot)\mathbf{R}^T$ rotates both the input slot and the output slot of the dyadic. It preserves symmetry, positive-definiteness, and—for rotations—eigenvalues (the principal moments of inertia). This formula changes orientation about the same reference point; changing the origin also requires the parallel-axis contribution.

1.6.5 Basis, Metric and Mechanical Duality

If old velocity coordinates and new coordinates satisfy $\mathbf{v} = S\tilde{\mathbf{v}}$ for an invertible S , preservation of kinetic energy and power requires

$$\begin{aligned} \tilde{M} &= S^T M S, & \tilde{\boldsymbol{\tau}} &= S^T \boldsymbol{\tau}, \\ \tilde{\mathbf{v}}^T \tilde{M} \tilde{\mathbf{v}} &= \mathbf{v}^T M \mathbf{v}, & \tilde{\boldsymbol{\tau}}^T \tilde{\mathbf{v}} &= \boldsymbol{\tau}^T \mathbf{v}. \end{aligned}$$

This congruence law describes the kinetic bilinear form and its vector-to-covector map. A linear endomorphism instead has the similarity law $\tilde{A} = S^{-1}AS$. For an orthogonal S , the two transformation patterns can look alike, which can hide the distinction.

Congruence preserves definiteness but generally changes ordinary matrix eigenvalues. Physical comparisons in nonorthogonal or differently scaled coordinates require the corresponding metric, not raw eigenvalues or a dot product chosen only because the arrays have equal length. This distinction connects velocity transmission through J to force transmission through J^T , and prevents coordinate-dependent numerical norms from being labeled energy or effort without a model.

1.6.6 Connection to the Rest of This Book

The dyadic point of view recurs throughout our study of rigid-body mechanics:

- **Screw axes and wrenches (Chapter 5).** A wrench combines a force and a moment into a single six-dimensional object. The spatial inertia is a 6×6 map from twist to

spatial momentum, not directly to wrench. A wrench is related to momentum rate, including the transport terms required by a moving coordinate frame.

- **Spatial algebra (Chapter 8).** With twist ordered as $(\boldsymbol{\omega}, \mathbf{v}_O)$ and centre-of-mass displacement $\mathbf{c}$ from the reference origin O , the spatial inertia is

$$\mathcal{I}_O = \begin{bmatrix} I_O & m[\mathbf{c}]_{\times} \\ -m[\mathbf{c}]_{\times} & mI_3 \end{bmatrix}.$$

Here I_O is rotational inertia about O , not about the centre of mass unless the origins coincide. The matrix maps twist to angular momentum about O and linear momentum, and is SPD when every nonzero rigid-body twist has positive kinetic energy. In a body-fixed frame attached to the rigid body, $\mathbf{w} = \mathcal{I}_O \dot{\mathbf{v}} + \mathbf{v} \times^* (\mathcal{I}_O \mathbf{v})$, where the force cross-product operator includes the required frame transport. The momentum map itself contains no acceleration.

- **The mass matrix (Chapter 11).** The joint-space mass matrix $\mathbf{M}(\mathbf{q})$ that appears in the manipulator equation $\mathbf{M}\ddot{\mathbf{q}} + \mathbf{c} = \boldsymbol{\tau}$ is a configuration-dependent dyadic on joint-velocity space. Each link contributes a term built from Jacobians and spatial inertias through $M(\mathbf{q}) = \sum_i J_i^T \mathcal{I}_i J_i$, with each twist Jacobian and inertia expressed about matching origins in matching frames. This is the kinetic-energy counterpart of assembling nodal stiffness through a compatibility map.
- **The articulated-body algorithm (Chapter 10).** Featherstone's algorithm propagates *articulated-body inertias* through a kinematic tree. These are 6×6 dyadics that account for the coupling between a link's motion and the reactions of all its descendants.

Understanding that all these objects are dyadics—linear maps built from sums of outer products—provides a unifying thread: the same algebraic structure appears whether we are analysing a truss, spinning a gyroscope, or computing the forward dynamics of a humanoid robot.

1.7 Matrix Norms and Conditioning

Having decomposed matrices via the eigendecomposition and SVD, we now turn to a fundamental practical question: *how large is a matrix, and how sensitive is a linear system to perturbations?* Matrix norms quantify the first question; the *condition number* answers the second.

Definition 1.15. Frobenius Norm $\|\mathbf{A}\|_F$

The **Frobenius norm** of a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ is

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n a_{ij}^2} = \sqrt{\text{tr}(\mathbf{A}^T \mathbf{A})} = \sqrt{\sigma_1^2 + \sigma_2^2 + \cdots + \sigma_{\min(m,n)}^2}. \quad (1.53)$$

The last equality follows directly from the SVD. The Frobenius norm treats the matrix as a long vector and computes its Euclidean length.

Definition 1.16. Spectral Norm $\|\mathbf{A}\|_2$

The **spectral norm** (also called the *operator 2-norm* or *induced 2-norm*) is

$$\|\mathbf{A}\|_2 = \max_{\|\mathbf{x}\|_2=1} \|\mathbf{A}\mathbf{x}\|_2 = \sigma_{\max}(\mathbf{A}). \quad (1.54)$$

This is the maximum factor by which $\mathbf{A}$ can stretch a unit vector—exactly the length of the longest semi-axis of the image ellipsoid from the SVD’s geometric interpretation (Section 1.5).

Definition 1.17. Condition Number $\kappa(\mathbf{A})$

For a matrix $\mathbf{A}$ with $\sigma_{\min} > 0$, the **condition number** with respect to the 2-norm is

$$\kappa(\mathbf{A}) = \frac{\sigma_{\max}}{\sigma_{\min}}. \quad (1.55)$$

When $\sigma_{\min} = 0$ (i.e., $\mathbf{A}$ is singular), we define $\kappa(\mathbf{A}) = \infty$. A well-conditioned matrix has $\kappa \approx 1$; an ill-conditioned matrix has $\kappa \gg 1$.

1.7.1 What the Pseudoinverse Does Not Guarantee

For $J = U\Sigma V^T$, the Moore–Penrose pseudoinverse is $J^+ = V\Sigma^+U^T$, with each nonzero singular value inverted. The command $J^+\mathbf{v}$ is the minimum-Euclidean-norm least-squares solution. It exactly realizes $\mathbf{v}$ only when $\mathbf{v} \in \text{im } J$; its notion of minimum norm depends on the declared units and metric.

The pseudoinverse does not remove near-singular amplification. With $J = \text{diag}(1, 10^{-3})$ and desired velocity $(1, 1)^T$, it requests joint velocity $(1, 1000)^T$. Error in the weak task direction is amplified; error in every direction need not be. A condition-number threshold alone also misses an overall loss of scale such as $J = \epsilon I$, whose condition number stays one as $\epsilon \rightarrow 0$.

A damped least-squares alternative solves

$$\min_{\dot{\mathbf{q}}} \|J\dot{\mathbf{q}} - \mathbf{v}\|_2^2 + \lambda^2 \|\dot{\mathbf{q}}\|_2^2, \quad \lambda > 0,$$

and applies singular-direction gains

$$\frac{\sigma_i}{\sigma_i^2 + \lambda^2} \leq \frac{1}{2\lambda}.$$

This bounds inverse gain by accepting a task residual. It does not enforce joint-speed limits, contact constraints or closed-loop stability by itself. Truncated SVD likewise requires an explicit tolerance for treating small measured singular values as zero; numerical rank is not automatically exact physical rank.

For a golf model, a small singular value can explain why a chosen hand-velocity correction demands a large joint-velocity change at one configuration. Predicting actual motion additionally requires the inertial equations, actuator limits and controller. Neither shaking nor a beneficial release follows from the condition number alone.

1.8 Matrix Decompositions for Computation

The eigendecomposition and SVD reveal geometric structure. This section surveys three additional decompositions whose primary purpose is *computational efficiency* when solving the linear systems that arise throughout dynamics and control.

1.8.1 LU Decomposition

Any invertible matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can (with row permutations) be factored as

$$\mathbf{PA} = \mathbf{LU}, \quad (1.56)$$

where $\mathbf{P}$ is a permutation matrix, $\mathbf{L}$ is unit lower triangular, and $\mathbf{U}$ is upper triangular. Solving $\mathbf{Ax} = \mathbf{b}$ then reduces to two triangular solves (forward and back substitution), each costing $O(n^2)$ after the $O(n^3)$ factorization. When multiple right-hand sides share the same $\mathbf{A}$, LU is the workhorse: factor once, solve many times.

1.8.2 QR Decomposition

For any $\mathbf{A} \in \mathbb{R}^{m \times n}$ with $m \geq n$,

$$\mathbf{A} = \mathbf{QR}, \quad (1.57)$$

where $\mathbf{Q} \in \mathbb{R}^{m \times m}$ is orthogonal and $\mathbf{R} \in \mathbb{R}^{m \times n}$ is upper trapezoidal: an upper-triangular $n \times n$ block above zero rows. QR is the preferred method for solving least-squares problems $\min_{\mathbf{x}} \|\mathbf{Ax} - \mathbf{b}\|_2$ because, unlike the normal equations $\mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$, it avoids squaring the condition number ($\kappa_2(\mathbf{A}^T \mathbf{A}) = \kappa_2(\mathbf{A})^2$ for full column rank). A backward-stable QR solve avoids that additional conditioning penalty. Rank-deficient problems need an appropriate rank-revealing or SVD treatment, particularly when a minimum-norm solution is required.

1.8.3 Cholesky Decomposition

If $\mathbf{M} \in \mathbb{R}^{n \times n}$ is symmetric positive definite (SPD), then there exists a unique lower-triangular matrix $\mathbf{L}$ with positive diagonal entries such that

$$\mathbf{M} = \mathbf{LL}^T. \quad (1.58)$$

Dense Cholesky costs roughly half the leading-order flops of dense LU and is backward stable under the usual floating-point assumptions for a sufficiently well-resolved SPD input. No pivoting is needed in exact arithmetic. Near-singularity, underflow, overflow and loss of definiteness through rounding remain numerical concerns.

Mathematical Principle 1.1. Cholesky and the Mass Matrix

In independent generalized-velocity coordinates, a model with strictly positive kinetic energy for every nonzero admissible velocity has an SPD mass matrix. Redundant embedding coordinates, massless idealizations or singular coordinate charts can violate that premise. With ideal constraints one may instead solve a reduced system or an indefinite saddle-point system; Cholesky does not directly apply to every augmented matrix.

For the unconstrained SPD solve $M\ddot{\mathbf{q}} = \boldsymbol{\tau} - \mathbf{h} - \mathbf{g}$, Cholesky or a structure-exploiting factorization is appropriate. An articulated-body algorithm can solve tree dynamics recursively without explicitly forming the full dense matrix. Choose the algorithm for the model and sparsity rather than treating one factorization as universal [32].

1.9 Worked Example: Eigenvalue and SVD Analysis in Python

Let us bring together the theory developed in this chapter with a concrete computational example. The following Python code constructs a 3×3 symmetric matrix (representing a simplified 3D inertia tensor), computes its eigendecomposition and SVD, and verifies the key properties we have proven.

```

1 import numpy as np
2
3 # ---- Define a symmetric positive definite matrix ----
4 # This could represent a simplified 3D inertia tensor
5 M = np.array([
6     [5.0, 1.0, 0.5],
7     [1.0, 4.0, 0.3],
8     [0.5, 0.3, 3.0]
9 ])
10
11 # Verify symmetry
12 assert np.allclose(M, M.T), "Matrix is not symmetric!"
13
14 # ---- Eigendecomposition ----
15 eigenvalues, eigenvectors = np.linalg.eigh(M) # eigh for symmetric matrices
16 print("Eigenvalues:", np.round(eigenvalues, 4))
17 print("Eigenvectors (columns):\n", np.round(eigenvectors, 4))
18
19 # Positive principal moments plus their triangle inequalities
20 # are required for a physical three-dimensional inertia tensor.
21 assert eigenvalues[-1] <= eigenvalues[0] + eigenvalues[1]
22
23 # Verify: all eigenvalues positive => positive definite
24 assert all(eigenvalues > 0), "Matrix is not positive definite!"
25 print("Matrix is POSITIVE DEFINITE (all eigenvalues > 0)")
26
27 # Verify: eigenvectors are orthogonal
28 QTQ = eigenvectors.T @ eigenvectors
29 print("Q^T Q (should be identity):\n", np.round(QTQ, 10))
30
31 # Verify spectral decomposition: M = Q * Lambda * Q^T
32 M_reconstructed = eigenvectors @ np.diag(eigenvalues) @ eigenvectors.T
33 print("Reconstruction error:", np.linalg.norm(M - M_reconstructed))
34
35 # ---- Singular Value Decomposition ----
36 U, sigma, VT = np.linalg.svd(M)
37 print("\nSingular values:", np.round(sigma, 4))
38 print("Eigenvalues (sorted descending):", np.round(sorted(eigenvalues,
39     reverse=True), 4))
39 print("For symmetric PD matrices, singular values = eigenvalues (sorted)")
40
41 # ---- Condition Number ----

```

```

42 kappa = sigma[0] / sigma[-1]
43 print(f"\nCondition number: {kappa:.4f}")
44 print("(Close to 1 means well-conditioned; >>1 means ill-conditioned)")
45
46 # ---- Cholesky Decomposition ----
47 L = np.linalg.cholesky(M)
48 print("\nCholesky factor L:\n", np.round(L, 4))
49 print("Verify L @ L^T = M, error:", np.linalg.norm(M - L @ L.T))

```

Listing 1.1: Eigendecomposition and SVD of a Symmetric Matrix

1.10 Chapter Summary

Mathematical Principle 1.2. Key Takeaways from Chapter 1

1. A **vector space** is defined by eight axioms guaranteeing closure, commutativity, and compatibility of addition and scalar multiplication. $\mathbb{R}^n$ is the prototypical example.
2. The **rank** of a matrix counts independent columns; the **null space** captures the “lost” dimensions. The **Rank-Nullity Theorem** connects them: rank + nullity = number of columns.
3. The **dot product** measures alignment between vectors and defines orthogonality. The **cross product** in $\mathbb{R}^3$ produces a perpendicular vector and can be written as a *skew-symmetric matrix* multiplication—a representation central to Lie Algebra ($\mathfrak{so}(3)$).
4. **Eigenvalues** characterize asymptotic stability of a constant continuous-time linear system, with Jordan structure needed on the imaginary-axis boundary. Eigenvectors of a symmetric inertia tensor give principal axes; its eigenvalues give principal moments. The **Spectral Theorem** guarantees that symmetric matrices have real eigenvalues and orthogonal eigenvectors.
5. **Positive definiteness** ($\mathbf{M} \succ 0$) guarantees a quadratic form has a strict minimum—the cornerstone of Lyapunov stability proofs and energy-based control design.
6. The **SVD** decomposes a matrix into orthogonal transformations and nonnegative scaling. Applied to a Jacobian with declared input and output metrics, it describes instantaneous velocity transmission. Relative conditioning, absolute inverse gain and finite-time controllability answer different questions.

1.11 Further Reading

The material in this chapter draws on a long tradition of textbooks at the intersection of linear algebra and applied mathematics. For a comprehensive undergraduate treatment that emphasizes geometric intuition alongside computational technique, Strang [41] remains

the standard reference; his “four fundamental subspaces” perspective directly informs the rank-nullity discussion presented here.

In the robotics context, Chapter 2 of Murray, Li, and Sastry [31] develops the specific linear-algebraic tools—rotation matrices, skew-symmetric operators, and the Lie bracket—that recur throughout the rest of this textbook. Readers who wish to pursue the numerical aspects of linear algebra, particularly iterative solvers and the role of condition numbers in optimization, will find Nocedal and Wright [32] an invaluable companion.

The dyadic and tensor-product viewpoint introduced in Section 1.6 connects directly to continuum mechanics, where stress and strain are second-order tensors built from exactly the same outer-product construction. Students interested in these connections should consult any graduate text on continuum mechanics or tensor analysis.

Finally, the eigenvalue structures and matrix exponentials introduced in this chapter are the first hints of *Lie theory*. Hall [17] provides a mathematically rigorous bridge from linear algebra to Lie groups and Lie algebras, foreshadowing the rotation groups ($SO(3)$) and rigid-body transformation groups ($SE(3)$) that appear in Chapters 4 and 5.

1.12 Exercises

Computational Exercises

1. **(Rank and Null Space)** Given the matrix

$$\mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

compute the rank, determinant, and a basis for the null space. Verify the Rank-Nullity Theorem. *Hint: row reduce to echelon form.*

2. **(Eigenvalue Computation)** For the matrix $\mathbf{A} = \begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix}$, compute the eigenvalues and eigenvectors by hand. Verify that $\mathbf{A}\mathbf{v}_i = \lambda_i\mathbf{v}_i$ for each pair.
3. **(SVD by Hand)** Compute the SVD of $\mathbf{M} = \begin{bmatrix} 3 & 0 \\ 0 & -2 \end{bmatrix}$. What are the singular values, and how do they relate to the eigenvalues?
4. **(Cholesky)** Verify that $\mathbf{M} = \begin{bmatrix} 4 & 2 \\ 2 & 3 \end{bmatrix}$ is positive definite by: (a) checking all eigenvalues are positive, (b) computing the Cholesky factorization $\mathbf{LL}^T$.
5. **(Python)** Write a Python script that generates a random 5×5 symmetric matrix $\mathbf{M} = \mathbf{A}^T\mathbf{A} + \epsilon\mathbf{I}$ with $\epsilon > 0$ (guaranteeing positive definiteness in exact arithmetic), computes its eigendecomposition, SVD, and Cholesky factorization, and verifies all three decompositions reproduce the original matrix within numerical tolerance.

Conceptual Exercises

6. **(Vector Spaces)** Is the set of all 2×2 symmetric matrices a vector space under standard matrix addition and scalar multiplication? Prove or disprove by checking the axioms. What is its dimension?

7. **(Physical Interpretation)** A robot Jacobian $\mathbf{J} \in \mathbb{R}^{3 \times 5}$ has singular values $\sigma_1 = 2.0$, $\sigma_2 = 0.8$, $\sigma_3 = 0.01$. Describe the shape of the velocity manipulability ellipsoid. Is the robot near a singularity? What does the condition number tell you about the reliability of inverse kinematics computations?
8. **(Positive Definiteness in Control)** Explain why strictly positive kinetic energy for every nonzero independent generalized velocity implies an SPD mass matrix. How can redundant coordinates, a massless idealization or a singular coordinate chart instead produce a zero eigenvalue? Distinguish this from an ordinary massive rigid body having a physically accessible motion with zero kinetic energy.
9. **(Skew-Symmetric Matrices)** Prove that for any skew-symmetric matrix $\mathbf{S}$ ($\mathbf{S}^T = -\mathbf{S}$), the matrix $e^{\mathbf{S}} = \mathbf{I} + \mathbf{S} + \frac{1}{2!}\mathbf{S}^2 + \cdots$ is an orthogonal matrix (i.e., $(e^{\mathbf{S}})^T e^{\mathbf{S}} = \mathbf{I}$). *Hint: use the fact that $(e^{\mathbf{S}})^T = e^{\mathbf{S}^T} = e^{-\mathbf{S}}$.* This result is the mathematical foundation of Chapter 6.

Proof-Based Exercises

10. **(Rank-Nullity Application)** Prove that if $\mathbf{A} \in \mathbb{R}^{m \times n}$ with $m < n$, then $\mathbf{A}$ must have a nontrivial null space (i.e., $\mathcal{N}(\mathbf{A}) \neq \{\mathbf{0}\}$). Interpret this result for a robot with more independent joint coordinates than task-space dimensions. Explain why this kinematic redundancy does not establish underactuation, and why a full-row-rank task Jacobian can have a nontrivial kernel.
11. **(Trace and Eigenvalues)** Prove that the trace of a matrix (sum of diagonal elements) equals the sum of its eigenvalues: $\text{tr}(\mathbf{M}) = \sum_{i=1}^n \lambda_i$. Count eigenvalues with algebraic multiplicity, over $\mathbb{C}$ if necessary. *Hint: use trace invariance under similarity and a complex Schur triangularization; do not assume that every matrix is diagonalizable.*
12. **(Determinant and Eigenvalues)** Prove that the determinant of a matrix equals the product of its eigenvalues: $\det(\mathbf{M}) = \prod_{i=1}^n \lambda_i$. *Hint: evaluate the characteristic polynomial at $\lambda = 0$.*
13. **(Uniqueness of SVD)** The sorted singular values are unique. Even for a distinct positive singular value, a corresponding left/right pair can have both signs reversed. For a repeated positive singular value, both bases can undergo the same orthogonal change within their paired subspaces. Demonstrate this with I_2 ; also explain why bases for the left and right zero-singular-value subspaces can be chosen independently.
14. **(Schur Complement and Positive Definiteness)** Let $\mathbf{M} = \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{C} \end{bmatrix}$ be a symmetric block matrix with $\mathbf{A} \succ 0$. Prove that $\mathbf{M} \succ 0$ if and only if $\mathbf{C} - \mathbf{B}^T \mathbf{A}^{-1} \mathbf{B} \succ 0$ (the *Schur complement* condition). This result is used extensively in the Articulated Body Algorithm (Chapter 10).

Chapter 2

The Transition to State Space

You cannot control a machine by staring at the machine. You control it by staring at the mathematical space that completely defines it.

2.1 Introduction: The Physics Abstraction

In Plain Language 2.1. Context & Use Case: Why We Banish 3D Physics

The Math You Will See: We will build the **State Vector** ($\mathbf{x}$), a vertical stack of numbers containing positions ($\mathbf{q}$) and velocities ($\dot{\mathbf{q}}$). We will formulate the universal control equation $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u}, t)$, and its linear cousin $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$.

Why We Need It: When a robotic arm swings a bat, the shoulder, elbow, and wrist all move in weird, overlapping arcs in 3D physics space (X, Y, Z). This is impossibly hard to calculate directly. Instead, we map the entire robot into an abstract mathematical realm called "State Space". In State Space, the complex multi-joint swinging robot is represented by a *single, infinitesimally small dot* moving along a single curved line. This abstraction makes solving incredibly complex control problems possible.

All Variables Defined: $\mathbf{x}$ is the state, $\mathbf{u}$ is the control input (e.g., motor torque), and f is the system dynamics mapping. See the **Nomenclature** at the start of the book for a full reference.

When a layman pictures a robotic arm throwing a ball or a drone correcting its pitch in a crosswind, they visualize physical bodies acting in 3-dimensional Euclidean space.

But control theorists and kinematic engineers almost never look at the physical space. They work in a wholly abstract, higher-dimensional mathematical realm known as **State Space**. State space is arguably the most fundamental translation required in control theory. If you do not understand state space, nothing else in the subsequent volumes of *The Geometry of Motion* will make sense.

2.2 Historical Context: The Kalman Revolution

In Plain Language 2.2. Historical Context: From Transfer Functions to State Space

In the 1950s, control engineers exclusively used **transfer functions**, which relate outputs to inputs in the frequency domain via Laplace transforms. Transfer functions are well suited for single-input single-output (SISO) systems and for analyzing stability via Nyquist diagrams and Bode plots. However, they fail catastrophically for multi-input multi-output (MIMO) systems and provide no insight into what is happening *inside* the system.

In 1960, **Rudolf E. Kalman** introduced the state-space representation of linear dynamical systems and the structural concepts of **controllability** and **observability** [18]. The state-space formulation made the analysis of multi-input multi-output systems systematic, opened the door to optimal linear-quadratic control, and led within a few years to the recursive estimation problem now known as the Kalman filter. Within a decade, state space had become the standard working language of quantitative control engineering.

The fundamental advantage of state-space methods over transfer functions is *transparency* [19]. A transfer function $G(s) = \frac{Y(s)}{U(s)}$ tells you how the input drives the output, but it hides the internal structure. A state-space model $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ exposes the entire flow of energy and information through the system. This transparency makes it possible to:

- Design controllers that regulate internal state, not just final outputs
- Detect when a system is uncontrollable or unobservable
- Handle systems with multiple inputs and multiple outputs naturally
- Compute optimal control laws analytically
- Estimate hidden state variables from measurements (Kalman filter)

The rest of this chapter builds the state-space framework from first principles. By the end, the reader should understand why Kalman's formulation became the standard and why state space is the working representation for almost every modern control problem.

2.3 Constructing a State Vector

Let us build a state space from the ground up using the simplest dynamical system possible: a 1-dimensional point mass (like a train on a straight track).

We want to predict what the train will do. To do this, we need to know its position, q . Is q the "state" of the system?

No. If we only know that the train is at $q = 100$ meters, we cannot predict where it will be 1 second from now. It could be moving right, moving left, or stationary.

Therefore, to complete our understanding, we must also know its velocity, $\dot{q}$ (the time-derivative of position). If we know $(q, \dot{q})$, we know its exact location and its exact speed.

State-Space System Architecture




Figure 2.1: Block diagram of a continuous-time linear state-space system. The input $\mathbf{u}$ drives the dynamics, producing state changes. The state $\mathbf{x}$ is obtained by integrating the state derivative. The output $\mathbf{y}$ is a linear combination of state and input.

Do we need to know acceleration ($\ddot{q}$) as part of the state? According to Newton's Second Law: $F = m\ddot{q}$. Acceleration is not an independent property of the train; it is completely driven by the external forces applied at that instant. Therefore, if we know the forces (the Control Inputs, $\mathbf{u}$), we can immediately compute the acceleration. The state itself is fully defined by purely the position and velocity.

Definition 2.1. The State Vector $\mathbf{x}$

The state vector $\mathbf{x}$ is the minimum set of mathematical variables completely describing the condition of a dynamic system at a specific point in time.

For our train, it is a 2-dimensional column vector:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} q \\ \dot{q} \end{bmatrix} = \begin{bmatrix} \text{position} \\ \text{velocity} \end{bmatrix} \quad (2.1)$$

For a general mechanical system with n configuration degrees of freedom $\mathbf{q} \in \mathbb{R}^n$, the state is the pair of configuration and velocity, $\mathbf{x} = (\mathbf{q}, \dot{\mathbf{q}}) \in \mathcal{X} \subset \mathbb{R}^{2n}$. More generally, the state space may be any manifold whose dimension equals the number of first-order ODEs required to predict the future: for a mechanical system that is $2n$; for an electromechanical system it is typically larger (see the DC-motor example below, where the inductor current adds an extra state).

If you take a photograph of a system, you capture its position $\mathbf{q}$. But a state vector $\mathbf{x}$ is the "ultimate photograph"—it captures the position *and* the velocity. With $\mathbf{x}$ and the laws of physics, the entire future of the system is deterministic.

2.3.1 Worked Example 1: Simple Pendulum

Example 2.1. State Vector for a Simple Pendulum

A pendulum of length L and mass m is suspended from a pivot. Let θ denote the angle from vertical, and τ denote an applied torque at the pivot. The equation of motion is:

$$mL^2\ddot{\theta} + mgL\sin(\theta) = \tau \quad (2.2)$$

To predict the pendulum's future, what do we need?

- The angle θ (where is it pointing?)
- The angular velocity $\dot{\theta}$ (is it spinning?)
- The torque τ is the control input (provided at the moment we want to predict)

Thus the state vector is:

$$\mathbf{x} = \begin{bmatrix} \theta \\ \dot{\theta} \end{bmatrix} \quad (2.3)$$

This is a 2-dimensional state space, even though the pendulum physically swings in a 1-dimensional arc in our visual space.

2.3.2 Worked Example 2: DC Motor

Example 2.2. State Vector for a DC Motor

A DC motor is driven by input voltage V_{in} . Let θ be the rotor angle, and let i be the current flowing through the motor windings.

The governing equations are:

$$L \frac{di}{dt} = V_{\text{in}} - Ri - K_e \dot{\theta} \quad (2.4)$$

$$J \frac{d\dot{\theta}}{dt} = K_t i - b\dot{\theta} \quad (2.5)$$

where:

- L is the motor inductance, R is resistance
- K_e is the back-EMF constant, K_t is the torque constant
- J is the rotor inertia, b is viscous friction

The state vector must include all variables whose time-derivatives depend on the inputs. These are:

$$\mathbf{x} = \begin{bmatrix} i \\ \theta \\ \dot{\theta} \end{bmatrix} \quad (2.6)$$

This is a 3-dimensional state space. Notice we include θ (rotor position) even though it does not directly appear in the governing equations—we need it to fully specify the motor's configuration, and its time-derivative $\dot{\theta}$ does appear. As noted in Definition 2.1, the state space dimension (3) exactly matches the number of first-order ODEs needed to predict the system (one electrical + two mechanical), despite there being only 1 mechanical DOF.

2.3.3 Worked Example 3: Two-Tank Fluid System

Example 2.3. State Vector for Two Coupled Tanks

Two tanks, each with cross-sectional area A , are connected by a valve. Let h_1 and h_2 be the fluid levels. Flow into tank 1 is Q_{in} , and flow between tanks is proportional to the height difference:

$$Q_{1 \rightarrow 2} = C \sqrt{|h_1 - h_2|} \quad (2.7)$$

where C is a flow coefficient, and outflow from tank 2 is:

$$Q_{\text{out}} = C_{\text{out}} \sqrt{h_2} \quad (2.8)$$

The height rates of change are:

$$A \frac{dh_1}{dt} = Q_{\text{in}} - Q_{1 \rightarrow 2} \quad (2.9)$$

$$A \frac{dh_2}{dt} = Q_{1 \rightarrow 2} - Q_{\text{out}} \quad (2.10)$$

To predict the system, we need the current heights of both tanks. The state vector is:

$$\mathbf{x} = \begin{bmatrix} h_1 \\ h_2 \end{bmatrix} \quad (2.11)$$

This is a 2-dimensional state space. The control input is the inflow rate Q_{in} .

2.4 Converting n -th Order ODEs to State Form

Theorem 2.1. Conversion of n -th Order ODE to First-Order State Form

Any n -th order differential equation of the form

$$y^{(n)} + a_{n-1}y^{(n-1)} + \cdots + a_1\dot{y} + a_0y = u(t) \quad (2.12)$$

can be converted to a system of n first-order ODEs via state-vector substitution.

Proof. Define the state variables as consecutive derivatives of y :

$$x_1 = y \quad (2.13)$$

$$x_2 = \dot{y} \quad (2.14)$$

$$x_3 = \ddot{y} \quad (2.15)$$

$$\vdots \quad (2.16)$$

$$x_n = y^{(n-1)} \quad (2.17)$$

Then:

$$\dot{x}_1 = x_2 \quad (2.18)$$

$$\dot{x}_2 = x_3 \quad (2.19)$$

$$\vdots \quad (2.20)$$

$$\dot{x}_{n-1} = x_n \quad (2.21)$$

$$\dot{x}_n = u(t) - a_0 x_1 - a_1 x_2 - \cdots - a_{n-1} x_n \quad (2.22)$$

where the final equation is obtained by solving the original ODE for $y^{(n)} = \dot{x}_n$:

$$\dot{x}_n = u(t) - a_0 x_1 - a_1 x_2 - \cdots - a_{n-1} x_n \quad (2.23)$$

In matrix form, this is $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}u$ where $\mathbf{A}$ is the companion matrix:

$$\mathbf{A} = \begin{bmatrix} 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \\ -a_0 & -a_1 & -a_2 & \cdots & -a_{n-1} \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 1 \end{bmatrix} \quad (2.24)$$

This representation is unique up to change of state variables (i.e., state transformation). $\square$

2.5 State Space Representations

Because the state vector $\mathbf{x}$ contains n variables, it naturally lives in an n -dimensional geometry: the State Space, denoted as $\mathcal{X} \subset \mathbb{R}^n$.

An astonishing realization for beginners is that when a dynamic system (like an airplane or a robotic arm) moves through physical space over time, its entire sprawling motion is represented as a single, unified **1-dimensional curve** snaking through n -dimensional state space. This curve is called the system's **Trajectory** $\mathcal{T}$.

2.5.1 Standard Form Differential Equations

The immense power of state space is that it allows us to convert complex, higher-order physical differential equations into simple, first-order vector math.

Consider a simple mass-spring-damper system governed by a second-order differential equation:

$$m\ddot{q} + c\dot{q} + kq = u \quad (2.25)$$

where q is position, m is mass, c is damping, k is stiffness, and u is our control force (e.g., a motor pushing the mass).

By defining our state vector as $\mathbf{x} = [x_1, x_2]^T = [q, \dot{q}]^T$, we can rewrite this physical equation as purely mathematical first-order derivatives. First, isolate the highest derivative ($\ddot{q}$):

$$\ddot{q} = \frac{1}{m}(u - c\dot{q} - kq) \quad (2.26)$$

Now, substitute the state variables ($x_1 = q, x_2 = \dot{q}$):

$$\dot{x}_1 = x_2 \quad (2.27)$$

$$\dot{x}_2 = \frac{1}{m}(u - cx_2 - kx_1) \quad (2.28)$$

This is called the **State Space Representation**. It is universally written in nonlinear control literature as:

Mathematical Principle 2.1. The Fundamental Equation of Control Theory

Every forced control system can be written as a first-order vector differential equation mapping the current state and current control commands to the velocity of the state vector:

$$\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u}, t) \quad (2.29)$$

This vector field f dictates how the system "flows" through the abstract geometry of the state space.

2.5.2 Linear State Space: The A and B Matrices

When the system is perfectly linear (like the mass-spring-damper above), the function $f(\mathbf{x}, \mathbf{u}, t)$ can be split into two matrices: the Dynamics Matrix (**A**) and the Input Matrix (**B**).

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu} \quad (2.30)$$

If we package our mass-spring equations into matrices, we get:

$$\begin{bmatrix} \dot{x}_1 \\ \dot{x}_2 \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{k}{m} & -\frac{c}{m} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{1}{m} \end{bmatrix} u \quad (2.31)$$

- **The A Matrix:** $\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -k/m & -c/m \end{bmatrix}$. This matrix defines the "drift" of the system.

If you turn off all motors ($u = 0$), the **A** matrix dictates exactly how the system will coast to a stop, oscillate, or explode. Its eigenvalues completely define the system's natural stability.

- **The B Matrix:** $\mathbf{B} = \begin{bmatrix} 0 \\ 1/m \end{bmatrix}$. This matrix defines how your control inputs u actually inject energy into the state. Notice that the top value is 0, meaning our motor cannot instantly teleport the position (x_1); it can only directly push the velocity (x_2).

2.6 Output Equations and Observability

Often, we cannot measure the entire state vector. For example, in a DC motor, we can measure the voltage and current flowing through it, but we may not directly measure the angular position θ . The **output equation** describes what we actually measure:

Definition 2.2. Output Equation

For a linear system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$, the output equation is:

$$\mathbf{y} = \mathbf{C}\mathbf{x} + \mathbf{D}\mathbf{u} \quad (2.32)$$

where:

- $\mathbf{y} \in \mathbb{R}^p$ is the measurement vector (what we can observe)
- $\mathbf{C} \in \mathbb{R}^{p \times n}$ is the output matrix
- $\mathbf{D} \in \mathbb{R}^{p \times m}$ is the feedthrough matrix (often zero)

The question "Can we infer the full state $\mathbf{x}$ from the measurements $\mathbf{y}$?" is precisely the question of **observability**, which we formalize later in Section 2.11.

2.7 Equilibrium Points and Linearization

2.7.1 Equilibrium Point Definition

Definition 2.3. Equilibrium Point

An **equilibrium point** (or fixed point) $\mathbf{x}^* \in \mathbb{R}^n$ of the system $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u}^*)$ is a point where the vector field is zero:

$$f(\mathbf{x}^*, \mathbf{u}^*) = \mathbf{0} \quad (2.33)$$

At an equilibrium, the state does not change; the system is in a "steady state."

For linear systems $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$, if $\mathbf{A}$ is invertible and $\mathbf{u}^* = \mathbf{0}$, the unique equilibrium is $\mathbf{x}^* = \mathbf{0}$. For nonlinear systems, there may be multiple equilibria.

2.7.2 Linearization via Taylor Expansion

To understand the local behavior near an equilibrium, we linearize via Taylor expansion. Suppose we have an equilibrium $\mathbf{x}^*$ at control input $\mathbf{u}^*$. Consider a small perturbation $\delta\mathbf{x} = \mathbf{x} - \mathbf{x}^*$ and $\delta\mathbf{u} = \mathbf{u} - \mathbf{u}^*$.

Expand f in a Taylor series around $(\mathbf{x}^*, \mathbf{u}^*)$:

$$f(\mathbf{x}^* + \delta\mathbf{x}, \mathbf{u}^* + \delta\mathbf{u}) = f(\mathbf{x}^*, \mathbf{u}^*) + \mathbf{J}_f \delta\mathbf{x} + \mathbf{J}_u \delta\mathbf{u} + O(\|\delta\mathbf{x}\|^2, \|\delta\mathbf{u}\|^2) \quad (2.34)$$

where $\mathbf{J}_f$ is the **Jacobian matrix** of f with respect to the state:

$$\mathbf{J}_f = \frac{\partial f}{\partial \mathbf{x}} \Big|_{(\mathbf{x}^*, \mathbf{u}^*)} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_n}{\partial x_1} & \frac{\partial f_n}{\partial x_2} & \cdots & \frac{\partial f_n}{\partial x_n} \end{bmatrix}_{(\mathbf{x}^*, \mathbf{u}^*)} \quad (2.35)$$

and similarly:

$$\mathbf{J}_u = \frac{\partial f}{\partial \mathbf{u}} \Big|_{(\mathbf{x}^*, \mathbf{u}^*)} \quad (2.36)$$

Since $f(\mathbf{x}^*, \mathbf{u}^*) = \mathbf{0}$, the linearized dynamics are:

$$\dot{\delta \mathbf{x}} \approx \mathbf{A}_c \delta \mathbf{x} + \mathbf{B}_c \delta \mathbf{u} \quad (2.37)$$

where $\mathbf{A}_c = \mathbf{J}_f$ and $\mathbf{B}_c = \mathbf{J}_u$. This is a linear system describing the small-signal behavior around the equilibrium.

Example 2.4. Pendulum Linearization

Consider the pendulum from Example 2.3.1 with $\tau = 0$ (no control). The equation is:

$$\ddot{\theta} + \frac{g}{L} \sin(\theta) = 0 \quad (2.38)$$

State vector: $\mathbf{x} = [\theta, \dot{\theta}]^T$. The vector field is:

$$f(\mathbf{x}) = \begin{bmatrix} x_2 \\ -\frac{g}{L} \sin(x_1) \end{bmatrix} \quad (2.39)$$

Equilibria: $f(\mathbf{x}^*) = \mathbf{0}$ when $x_2 = 0$ and $\sin(x_1) = 0$, so $\mathbf{x}^* = [k\pi, 0]^T$ for $k \in \mathbb{Z}$.

Bottom equilibrium ($\theta = 0$): The Jacobian is:

$$\mathbf{J}_f = \begin{bmatrix} 0 & 1 \\ -\frac{g}{L} \cos(\theta) & 0 \end{bmatrix} \Big|_{\theta=0} = \begin{bmatrix} 0 & 1 \\ -\frac{g}{L} & 0 \end{bmatrix} \quad (2.40)$$

The linearized system near the bottom is $\ddot{\theta} = -\frac{g}{L}\theta$, a harmonic oscillator.

Top equilibrium ($\theta = \pi$):

$$\mathbf{J}_f \Big|_{\theta=\pi} = \begin{bmatrix} 0 & 1 \\ \frac{g}{L} & 0 \end{bmatrix} \quad (2.41)$$

The linearized system is $\ddot{\theta} = \frac{g}{L}\theta$, which has exponentially growing solutions—an unstable saddle.

2.8 Phase Portraits and Stability Classification

For a 2D system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$, the phase portrait is the collection of all trajectories in the (x_1, x_2) plane. The behavior depends entirely on the eigenvalues of $\mathbf{A}$.

Theorem 2.2. Stability Classification from Eigenvalues

For a 2D linear system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$ with eigenvalues λ_1, λ_2 :

- **Stable Node:** Both eigenvalues are real and negative ($\lambda_1, \lambda_2 < 0$). Trajectories spiral into the origin monotonically.
- **Unstable Node:** Both eigenvalues are real and positive ($\lambda_1, \lambda_2 > 0$). Trajectories spiral away from the origin.
- **Saddle Point:** One eigenvalue is positive, one is negative. Trajectories approach along the stable manifold (eigenvector of $\lambda < 0$) and diverge along the unstable manifold (eigenvector of $\lambda > 0$).
- **Stable Spiral (Focus):** Eigenvalues are complex conjugates with negative real part ($\lambda = \alpha \pm i\beta$, $\alpha < 0$). Trajectories spiral inward.
- **Unstable Spiral:** Eigenvalues are complex conjugates with positive real part. Trajectories spiral outward.
- **Center:** Eigenvalues are purely imaginary ($\lambda = \pm i\omega$). Trajectories form closed orbits (periodic motion).

The origin is **asymptotically stable** if and only if all eigenvalues have negative real part. The origin is **marginally stable** if the largest real part is zero. The origin is **unstable** if any eigenvalue has positive real part.

2.9 Solution of Linear Systems: The Matrix Exponential

For a linear time-invariant (LTI) system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$ with initial condition $\mathbf{x}(0) = \mathbf{x}_0$, the solution is:

Theorem 2.3. Solution via Matrix Exponential

The unique solution is:

$$\mathbf{x}(t) = e^{\mathbf{A}t} \mathbf{x}_0 \quad (2.42)$$

where $e^{\mathbf{A}t}$ is the **matrix exponential**, defined as:

$$e^{\mathbf{A}t} = \sum_{k=0}^{\infty} \frac{(\mathbf{A}t)^k}{k!} = \mathbf{I} + \mathbf{A}t + \frac{(\mathbf{A}t)^2}{2!} + \frac{(\mathbf{A}t)^3}{3!} + \dots \quad (2.43)$$

Proof. Consider the scalar ODE $\dot{x} = ax$ with initial condition $x(0) = x_0$. The solution is $x(t) = e^{at}x_0$.

By analogy, define $e^{\mathbf{A}t}$ as the solution to:

$$\frac{d}{dt}e^{\mathbf{A}t} = \mathbf{A}e^{\mathbf{A}t}, \quad e^{\mathbf{A} \cdot 0} = \mathbf{I} \quad (2.44)$$

The Taylor series ansatz:

$$e^{\mathbf{A}t} = \sum_{k=0}^{\infty} \frac{(\mathbf{A}t)^k}{k!} \quad (2.45)$$

satisfies this differential equation because:

$$\frac{d}{dt} \sum_{k=0}^{\infty} \frac{(\mathbf{A}t)^k}{k!} = \sum_{k=1}^{\infty} \frac{k(\mathbf{A}t)^{k-1} \mathbf{A}}{k!} = \sum_{k=1}^{\infty} \frac{(\mathbf{A}t)^{k-1} \mathbf{A}}{(k-1)!} = \mathbf{A} \sum_{j=0}^{\infty} \frac{(\mathbf{A}t)^j}{j!} = \mathbf{A} e^{\mathbf{A}t} \quad (2.46)$$

with $e^{\mathbf{A} \cdot 0} = \mathbf{I}$. Now verify that $\mathbf{x}(t) = e^{\mathbf{A}t} \mathbf{x}_0$ solves the original system:

$$\dot{\mathbf{x}}(t) = \frac{d}{dt} (e^{\mathbf{A}t} \mathbf{x}_0) = \mathbf{A} e^{\mathbf{A}t} \mathbf{x}_0 = \mathbf{A} \mathbf{x}(t) \quad \checkmark \quad (2.47)$$

□

Practical Computation: For small-dimensional systems, the matrix exponential can be computed via diagonalization. If $\mathbf{A} = \mathbf{P} \mathbf{\Lambda} \mathbf{P}^{-1}$ where $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_n)$, then:

$$e^{\mathbf{A}t} = \mathbf{P} e^{\mathbf{\Lambda}t} \mathbf{P}^{-1} = \mathbf{P} \text{diag}(e^{\lambda_1 t}, \dots, e^{\lambda_n t}) \mathbf{P}^{-1} \quad (2.48)$$

For systems with input, $\dot{\mathbf{x}} = \mathbf{A} \mathbf{x} + \mathbf{B} \mathbf{u}(t)$:

$$\mathbf{x}(t) = e^{\mathbf{A}t} \mathbf{x}_0 + \int_0^t e^{\mathbf{A}(t-\tau)} \mathbf{B} \mathbf{u}(\tau) d\tau \quad (2.49)$$

This is the superposition of free response (first term) and forced response (second term).

2.10 Controllability

The key question in control is: *Can we steer the system from any initial state to any desired final state using our inputs?*

Definition 2.4. Controllability

A linear system $\dot{\mathbf{x}} = \mathbf{A} \mathbf{x} + \mathbf{B} \mathbf{u}$ is **controllable** if for any initial state $\mathbf{x}_0$ and any final state $\mathbf{x}_f$, there exists a finite time $T > 0$ and an input sequence $\mathbf{u}(t)$ for $t \in [0, T]$ that drives the system from $\mathbf{x}_0$ to $\mathbf{x}_f$.

Theorem 2.4. Kalman Controllability Rank Condition

The system $\dot{\mathbf{x}} = \mathbf{A} \mathbf{x} + \mathbf{B} \mathbf{u}$ is controllable if and only if the **controllability matrix**

$$\mathbf{C}_{\text{rank}} = [\mathbf{B} \quad \mathbf{A}\mathbf{B} \quad \mathbf{A}^2\mathbf{B} \quad \dots \quad \mathbf{A}^{n-1}\mathbf{B}] \quad (2.50)$$

has full rank: $\text{rank}(\mathbf{C}_{\text{rank}}) = n$.

Sketch. The controllability matrix $\mathbf{C}_{\text{rank}}$ has dimensions $n \times nm$. Each column $\mathbf{A}^k \mathbf{B}$ represents the "reach" of the input after k applications of the dynamics. The rank condition ensures that the span of these columns covers all n dimensions of state space.

The intuition is from the Cayley-Hamilton theorem: any power $\mathbf{A}^k$ for $k \geq n$ can be expressed as a linear combination of $\mathbf{A}^0, \mathbf{A}^1, \dots, \mathbf{A}^{n-1}$. Thus, checking columns up to $\mathbf{A}^{n-1} \mathbf{B}$ is sufficient.

A rigorous proof uses the Gram matrix approach: controllability is equivalent to the controllability Gram matrix

$$\mathbf{W}_c = \int_0^T e^{\mathbf{A}\tau} \mathbf{B} \mathbf{B}^T e^{\mathbf{A}^T \tau} d\tau \quad (2.51)$$

being positive definite, which is equivalent to $\mathbf{C}_{\text{rank}}$ having full rank. $\square$

Physical Interpretation: If the system is not controllable, it means some mode or degree of freedom cannot be directly influenced by the available inputs. For a DC motor whose position is θ and current is i , if the only measurement is current and there is no direct control over electrical resistance, the position might be uncontrollable.

2.11 Observability

Observability is the dual of controllability: *Can we infer the full state from measurements alone?*

Definition 2.5. Observability

A linear system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$, $\mathbf{y} = \mathbf{C}\mathbf{x} + \mathbf{D}\mathbf{u}$ is **observable** if the initial state $\mathbf{x}(0)$ can be uniquely determined from knowledge of the input $\mathbf{u}(t)$ and the output $\mathbf{y}(t)$ over a finite time interval $[0, T]$.

Theorem 2.5. Kalman Observability Rank Condition

The system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$, $\mathbf{y} = \mathbf{C}\mathbf{x} + \mathbf{D}\mathbf{u}$ is observable if and only if the **observability matrix**

$$\mathbf{O} = \begin{bmatrix} \mathbf{C} \\ \mathbf{CA} \\ \mathbf{CA}^2 \\ \vdots \\ \mathbf{CA}^{n-1} \end{bmatrix} \quad (2.52)$$

has full rank: $\text{rank}(\mathbf{O}) = n$.

The observability matrix has dimensions $np \times n$ where p is the number of measurements. Each row $\mathbf{CA}^k$ represents the k -th time derivative of the measurement, weighted by the dynamics. If these rows span all n dimensions, we can reconstruct the state.

2.12 The Mass-Spring-Damper System: Complete Analysis

Let us apply the concepts from this chapter to the mass-spring-damper system with stiffness k , damping c , and mass m :

$$m\ddot{q} + c\dot{q} + kq = u \quad (2.53)$$

State vector: $\mathbf{x} = [q, \dot{q}]^T$.

Matrices:

$$\mathbf{A} = \begin{bmatrix} 0 & 1 \\ -k/m & -c/m \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} 0 \\ 1/m \end{bmatrix} \quad (2.54)$$

Characteristic polynomial:

$$\det(\mathbf{A} - \lambda\mathbf{I}) = \det \begin{bmatrix} -\lambda & 1 \\ -k/m & -c/m - \lambda \end{bmatrix} = \lambda^2 + \frac{c}{m}\lambda + \frac{k}{m} \quad (2.55)$$

Eigenvalues:

$$\lambda_{1,2} = -\frac{c}{2m} \pm \sqrt{\left(\frac{c}{2m}\right)^2 - \frac{k}{m}} = -\frac{c}{2m} \pm \sqrt{\frac{c^2}{4m^2} - \frac{k}{m}} \quad (2.56)$$

Define the natural frequency $\omega_0 = \sqrt{k/m}$ and damping ratio $\zeta = \frac{c}{2\sqrt{km}}$. Then:

$$\lambda_{1,2} = -\zeta\omega_0 \pm i\omega_0\sqrt{1 - \zeta^2} \quad (2.57)$$

(assuming $\zeta < 1$ for the underdamped case; see below for other cases).

2.12.1 Underdamped ($0 < \zeta < 1$)

The discriminant is negative, so eigenvalues are complex: $\lambda = -\zeta\omega_0 \pm i\omega_d$ where $\omega_d = \omega_0\sqrt{1 - \zeta^2}$ is the damped frequency.

The mass oscillates while decaying. The solution is:

$$q(t) = e^{-\zeta\omega_0 t} \left(q_0 \cos(\omega_d t) + \frac{\dot{q}_0 + \zeta\omega_0 q_0}{\omega_d} \sin(\omega_d t) \right) \quad (2.58)$$

The trajectory in phase space is a **stable spiral**, spiraling into the origin.

2.12.2 Critically Damped ($\zeta = 1$)

The eigenvalues coincide: $\lambda_1 = \lambda_2 = -\omega_0$. The matrix $\mathbf{A}$ is not diagonalizable; it has a Jordan block.

The solution is:

$$q(t) = (q_0 + (\dot{q}_0 + \omega_0 q_0)t)e^{-\omega_0 t} \quad (2.59)$$

This is the fastest non-oscillatory return to equilibrium. The trajectory does *not* spiral; it approaches the origin monotonically.

2.12.3 Overdamped ($\zeta > 1$)

The eigenvalues are real and distinct:

$$\lambda_1 = -\zeta\omega_0 + \omega_0\sqrt{\zeta^2 - 1}, \quad \lambda_2 = -\zeta\omega_0 - \omega_0\sqrt{\zeta^2 - 1} \quad (2.60)$$

Both are negative (since $\zeta > 0$). The solution is:

$$q(t) = C_1 e^{\lambda_1 t} + C_2 e^{\lambda_2 t} \quad (2.61)$$

The slower eigenvalue λ_1 dominates for large times. The trajectory is a **stable node**—no oscillation, just exponential decay. The system returns to rest more slowly than the critically damped case.

2.12.4 Controllability and Observability

The controllability matrix is:

$$\mathbf{C}_{\text{rank}} = [\mathbf{B} \quad \mathbf{AB}] = \begin{bmatrix} 0 & 1/m \\ 1/m & -c/m^2 \end{bmatrix} \quad (2.62)$$

Its determinant is $\det(\mathbf{C}_{\text{rank}}) = 0 - \frac{1}{m^2} = -\frac{1}{m^2} \neq 0$. Thus $\mathbf{C}_{\text{rank}}$ has full rank and the system is **controllable**. We can steer position and velocity to any desired values using the input force u .

If we measure position only, $\mathbf{C} = [1, 0]$, the observability matrix is:

$$\mathbf{O} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad (2.63)$$

which has rank 2. The system is **observable**. We can infer both position and velocity from position measurements alone (by differentiating or using a Kalman filter).

2.13 Transfer Function to State Space and Back

The transfer function $G(s) = \frac{Y(s)}{U(s)}$ relates input and output in the Laplace domain. For a linear system with zero initial conditions, the transfer function is:

$$G(s) = \mathbf{C}(s\mathbf{I} - \mathbf{A})^{-1}\mathbf{B} + \mathbf{D} \quad (2.64)$$

This is derived from the state-space equations by taking Laplace transforms:

$$s\mathbf{X}(s) = \mathbf{A}\mathbf{X}(s) + \mathbf{B}U(s) \quad (2.65)$$

$$\mathbf{Y}(s) = \mathbf{C}\mathbf{X}(s) + \mathbf{D}U(s) \quad (2.66)$$

Solve for $\mathbf{X}(s)$:

$$(s\mathbf{I} - \mathbf{A})\mathbf{X}(s) = \mathbf{B}U(s) \quad \Rightarrow \quad \mathbf{X}(s) = (s\mathbf{I} - \mathbf{A})^{-1}\mathbf{B}U(s) \quad (2.67)$$

Substitute into the output equation:

$$\mathbf{Y}(s) = [\mathbf{C}(s\mathbf{I} - \mathbf{A})^{-1}\mathbf{B} + \mathbf{D}] U(s) = G(s)U(s) \quad (2.68)$$

Conversion from transfer function to state space is not unique—there are many valid state representations. The most common choice is the controllable canonical form (the companion matrix form from Section 2.4).

2.13.1 Discrete-Time State Space

Every modern robot controller is implemented on a digital processor that reads sensors and writes actuator commands at a fixed sample rate $f_s = 1/T_s$. Between samples the controller holds its output constant (zero-order hold). We therefore need a *discrete-time* counterpart of the continuous state-space model.

Definition 2.6. Discrete-Time State-Space Model

A linear discrete-time system is

$$\mathbf{x}[k+1] = \mathbf{A}_d \mathbf{x}[k] + \mathbf{B}_d \mathbf{u}[k], \quad (2.69)$$

$$\mathbf{y}[k] = \mathbf{C} \mathbf{x}[k] + \mathbf{D} \mathbf{u}[k], \quad (2.70)$$

where $k \in \{0, 1, 2, \dots\}$ indexes the sample instants $t_k = k T_s$.

Exact Discretization (Zero-Order Hold). Given the continuous system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ and a sampling period T_s , the exact discrete-time matrices are

$$\mathbf{A}_d = e^{\mathbf{A}T_s}, \quad \mathbf{B}_d = \int_0^{T_s} e^{\mathbf{A}\tau} \mathbf{B} d\tau = \mathbf{A}^{-1}(e^{\mathbf{A}T_s} - \mathbf{I})\mathbf{B}, \quad (2.71)$$

where the closed-form expression for $\mathbf{B}_d$ holds when $\mathbf{A}$ is invertible. The matrix exponential $e^{\mathbf{A}T_s}$ connects directly to Theorem 2.3: the free response over one sample period is simply propagated by $\mathbf{A}_d$.

Theorem 2.6. Discrete-Time Stability Criterion

The equilibrium $\mathbf{x}^* = \mathbf{0}$ of the autonomous discrete-time system $\mathbf{x}[k+1] = \mathbf{A}_d \mathbf{x}[k]$ is asymptotically stable if and only if every eigenvalue of $\mathbf{A}_d$ lies strictly inside the unit circle:

$$|\lambda_i(\mathbf{A}_d)| < 1 \quad \forall i. \quad (2.72)$$

This is the discrete analogue of requiring all eigenvalues of $\mathbf{A}$ to have negative real parts in continuous time (Theorem 2.2).

In Plain Language 2.3. All Robot Controllers Are Discrete

All modern robot controllers run in discrete time at sample rates of 1–10 kHz. The continuous-time model $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ is the physicist's description of reality; the discrete-time model $\mathbf{x}[k+1] = \mathbf{A}_d \mathbf{x}[k] + \mathbf{B}_d \mathbf{u}[k]$ is what actually executes on the real-time processor. A control design that ignores this distinction—for example, by using too low a sample rate so that eigenvalues migrate outside the unit circle—will destabilise an otherwise stable system.

2.14 Nonlinear Systems and Phase Portraits

While we have focused on linear systems, real physical systems are often nonlinear. For a nonlinear system $\dot{\mathbf{x}} = f(\mathbf{x})$, the phase portrait can be qualitatively understood by:

1. Finding equilibrium points by solving $f(\mathbf{x}^*) = \mathbf{0}$
2. Linearizing around each equilibrium: $\delta \dot{\mathbf{x}} \approx \mathbf{J}_f(\mathbf{x}^*)\delta \mathbf{x}$
3. Classifying the equilibrium by its eigenvalues (Theorem 2.8)
4. Plotting the vector field $\dot{\mathbf{x}} = f(\mathbf{x})$ to visualize global behavior

The linearization is valid in a neighborhood of the equilibrium. For large perturbations, we must rely on numerical simulation or other techniques (Lyapunov theory, bifurcation analysis).

2.15 Python Worked Example: State Space Simulation and Phase Portrait

Consider the underdamped mass-spring-damper system with *illustrative numerical values*: $m = 1$ kg, $c = 0.5$ N · s/m, $k = 2$ N/m, with initial conditions $q_0 = 1$ m, $\dot{q}_0 = 0$.

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.integrate import odeint

# Parameters
m, c, k = 1.0, 0.5, 2.0
A = np.array([[0, 1], [-k/m, -c/m]])
B = np.array([[0], [1/m]])

# System dynamics
def dynamics(x, t, u=0):
    return A @ x + B.flatten() * u

# Initial condition
x0 = np.array([1.0, 0.0])
t = np.linspace(0, 10, 1000)

# Simulate
x = odeint(dynamics, x0, t)

# Phase portrait
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

# Time evolution
ax1.plot(t, x[:, 0], label='Position q(t)')
ax1.plot(t, x[:, 1], label='Velocity dq/dt(t)')
ax1.set_xlabel('Time (s)')
ax1.set_ylabel('State')
ax1.legend()
ax1.grid()
```

```

ax1.set_title('Time Evolution')

# Phase portrait
ax2.plot(x[:, 0], x[:, 1], 'b-', linewidth=2)
ax2.plot(x[0, 0], x[0, 1], 'go', markersize=10, label='Initial')
ax2.plot(x[-1, 0], x[-1, 1], 'ro', markersize=10, label='Final')
ax2.set_xlabel('Position q')
ax2.set_ylabel('Velocity dq/dt')
ax2.grid()
ax2.legend()
ax2.set_title('Phase Portrait (Stable Spiral)')

plt.tight_layout()
plt.show()

# Eigenvalues
eigenvalues, eigenvectors = np.linalg.eig(A)
print(f"Eigenvalues: {eigenvalues}")
print(f"Damping ratio zeta = {c / (2*np.sqrt(k*m)):.3f}")
print(f"Natural freq omega_0 = {np.sqrt(k/m):.3f} rad/s")

```

This code simulates the system and plots both the time evolution and the phase portrait. For the underdamped case ($\zeta = 0.177$), you should see a spiral converging to the origin with oscillation.

2.16 Equilibrium Analysis Summary

For any equilibrium point $\mathbf{x}^*$:

- Compute the Jacobian: $\mathbf{A} = \left. \frac{\partial f}{\partial \mathbf{x}} \right|_{\mathbf{x}^*}$
- Find eigenvalues of $\mathbf{A}$
- Classify the equilibrium according to Theorem 2.8
- If all eigenvalues have negative real part: **asymptotically stable**
- If any eigenvalue has positive real part: **unstable**
- If the largest real part is zero (center or boundary case): **marginally stable**

2.17 Summary

Mathematical Principle 2.2. Key Takeaways: State Space Fundamentals

1. **State Vector:** The state $\mathbf{x}$ is the minimum information needed to predict the system's future. For an n -th order ODE, the state is n -dimensional.
2. **Kalman's Revolution (1960):** State-space representations replace transfer

functions, enabling control of MIMO systems, design of optimal controllers, and systematic analysis of internal system structure.

3. Standard Forms: Any system can be written as $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u}, t)$ (nonlinear) or $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ (linear). The $\mathbf{A}$ matrix encodes natural stability; the $\mathbf{B}$ matrix encodes control authority.

4. ODE Reduction: Any n -th order ODE reduces to n first-order state equations via companion form.

5. Equilibrium and Linearization: Equilibria are found by solving $f(\mathbf{x}^*, \mathbf{u}^*) = \mathbf{0}$. Local behavior is determined by the Jacobian $\mathbf{J}_f = \partial f / \partial \mathbf{x}$.

6. Stability from Eigenvalues: Eigenvalues of $\mathbf{A}$ determine phase-portrait structure (node, spiral, saddle, center). Negative real parts mean stability.

7. Matrix Exponential: The solution $\mathbf{x}(t) = e^{\mathbf{A}t}\mathbf{x}_0 + \int_0^t e^{\mathbf{A}(t-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau$ decomposes into free response and forced response.

8. Controllability (Kalman Rank Test): The system is controllable iff $\text{rank}([\mathbf{B}, \mathbf{A}\mathbf{B}, \dots, \mathbf{A}^{n-1}\mathbf{B}]) = n$. If not controllable, some state variables cannot be influenced by inputs.

9. Observability (Dual of Controllability): The system is observable iff $\text{rank}([\mathbf{C}^T, \mathbf{A}^T\mathbf{C}^T, \dots, (\mathbf{A}^T)^{n-1}\mathbf{C}^T]^T) = n$. If not observable, some state variables cannot be inferred from measurements.

10. Duality: A system is controllable iff its dual (with $\mathbf{A}^T, \mathbf{C}^T, \mathbf{B}^T$) is observable. This duality underpins Kalman filter design.

2.18 Lyapunov Stability Theory

The eigenvalue-based stability analysis of Theorem 2.2 is powerful but inherently *local*: it applies only in a neighbourhood of an equilibrium where the linearization is valid. For nonlinear systems, we need a theory that can certify stability over large—possibly global—regions of state space. **Lyapunov's direct method** provides exactly this by replacing spectral analysis with an energy-like argument.

Linearization tells us about local stability. Lyapunov theory works globally.

Definition 2.7. Lyapunov Function $V(\mathbf{x})$

Consider the autonomous system $\dot{\mathbf{x}} = f(\mathbf{x})$ with an equilibrium at $\mathbf{x}^* = \mathbf{0}$. A continuously differentiable function $V : \mathbb{R}^n \rightarrow \mathbb{R}$ is called a **Lyapunov function candidate** if

1. $V(\mathbf{0}) = 0$ (zero at the equilibrium),
2. $V(\mathbf{x}) > 0$ for all $\mathbf{x} \neq \mathbf{0}$ in a neighbourhood of the origin (*positive definite*).

Its time derivative along trajectories of the system is

$$\dot{V}(\mathbf{x}) = \frac{\partial V}{\partial \mathbf{x}} f(\mathbf{x}) = \nabla V(\mathbf{x})^T f(\mathbf{x}). \quad (2.73)$$

Theorem 2.7. Lyapunov's Direct Method (Second Method)

Let $\mathbf{x}^* = \mathbf{0}$ be an equilibrium of $\dot{\mathbf{x}} = f(\mathbf{x})$ and let $V(\mathbf{x})$ be a Lyapunov function candidate in a domain $\mathcal{D}$ containing the origin.

1. If $\dot{V}(\mathbf{x}) \leq 0$ in $\mathcal{D}$, the equilibrium is **stable** (in the sense of Lyapunov).
2. If $\dot{V}(\mathbf{x}) < 0$ for all $\mathbf{x} \neq \mathbf{0}$ in $\mathcal{D}$, the equilibrium is **asymptotically stable**.
3. If the conditions above hold with $\mathcal{D} = \mathbb{R}^n$ and, in addition, $V(\mathbf{x}) \rightarrow \infty$ as $\|\mathbf{x}\| \rightarrow \infty$ (*radially unbounded*), the equilibrium is **globally asymptotically stable**.

The power of the theorem lies in the fact that we never need to solve the differential equation—we only need to find a single scalar function V with the right properties.

Example 2.5. Pendulum with Damping

Consider a simple pendulum of mass m and length ℓ with viscous damping $b > 0$:

$$m\ell^2\ddot{\theta} + b\dot{\theta} + mg\ell \sin \theta = 0. \quad (2.74)$$

The state is $\mathbf{x} = (\theta, \dot{\theta})^T$ and the downward equilibrium is $\mathbf{x}^* = \mathbf{0}$. Choose the total mechanical energy as a Lyapunov candidate:

$$V(\theta, \dot{\theta}) = \underbrace{\frac{1}{2} m\ell^2 \dot{\theta}^2}_{\text{kinetic}} + \underbrace{mg\ell (1 - \cos \theta)}_{\text{potential}}. \quad (2.75)$$

Clearly $V(\mathbf{0}) = 0$ and $V > 0$ for $(\theta, \dot{\theta}) \neq \mathbf{0}$ with $|\theta| < \pi$. Differentiating along trajectories:

$$\dot{V} = m\ell^2 \dot{\theta} \ddot{\theta} + mg\ell \sin \theta \dot{\theta} = \dot{\theta}(m\ell^2 \ddot{\theta} + mg\ell \sin \theta) = \dot{\theta}(-b\dot{\theta}) = -b\dot{\theta}^2 \leq 0. \quad (2.76)$$

Since $\dot{V} \leq 0$, the equilibrium is stable. Because $\dot{V} = 0$ only on the set $\{\dot{\theta} = 0\}$ and no trajectory other than the equilibrium can stay there forever (by LaSalle's invariance principle), the equilibrium is in fact **asymptotically stable**.

Physically, the damper dissipates energy at rate $b\dot{\theta}^2$, and the total energy V is a non-increasing “altitude” that the system rolls down until it reaches the bottom of the energy landscape [19].

In Plain Language 2.4. Lyapunov Functions Are Like Altitude

Lyapunov functions are like altitude—if you can prove you are always going downhill, you must eventually reach the valley. The function $V(\mathbf{x})$ measures how “high” the state is above the equilibrium, and the condition $\dot{V} \leq 0$ guarantees that the system never climbs. Finding a good Lyapunov function is an art: energy is often the first candidate, but there are systematic construction methods for certain system classes [39].

Mathematical Principle 2.3. Lyapunov Theory as Foundation

Lyapunov stability theory is the foundation of virtually all stability proofs in modern control. Every adaptive controller, every robust controller, and every learning-based stability certificate in Chapters 11–12 ultimately relies on constructing a Lyapunov (or Lyapunov-like) function and showing that its derivative is non-positive along closed-loop trajectories.

2.19 Further Reading

The state-space framework presented in this chapter traces its origins to Kalman’s foundational 1960 paper [18], which introduced the concepts of controllability and observability that remain central to modern control engineering. That single paper launched an entire field; readers who wish to appreciate its full scope should read the original.

For nonlinear extensions of the state-space ideas developed here, Slotine and Li [39] provide an accessible treatment oriented toward robotic systems, covering sliding-mode control, adaptive control, and Lyapunov-based design—all formulated in state space. A more mathematically rigorous treatment of nonlinear systems theory, including stability in the sense of Lyapunov, input-output stability, and singular perturbation methods, can be found in Khalil [19].

Spong, Hutchinson, and Vidyasagar [40] develop state-space models specifically for robotic manipulators and mobile robots, making explicit the connection between the equations of motion derived from Lagrangian or Newton–Euler methods and the $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u})$ formalism of this chapter.

Finally, readers should note the deep connection between this chapter and Chapter 11 (Lagrangian Mechanics): the Euler–Lagrange equations provide the equations of motion that, once cast in first-order form, populate the state-space models studied here. Understanding both perspectives—the energy-based derivation of dynamics and the state-space representation of those dynamics—is essential for modern robotics.

2.20 Exercises

Computational Exercises

1. For the DC motor from Example 2.3.2, suppose $L = 0.01$ H, $R = 1$ Ω , $K_e = 0.02$ V·s/rad, $K_t = 0.02$ N·m/A, $J = 0.001$ kg·m², $b = 0.001$ N·m·s/rad. Write the system in the form $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}V_{\text{in}}$. Compute the eigenvalues of $\mathbf{A}$ numerically and classify the open-loop stability.
2. For the two-tank system (Example 2.3.3), assume $A = 0.1$ m², $C = 0.05$, $C_{\text{out}} = 0.03$. Starting from $h_1(0) = 0.5$ m, $h_2(0) = 0.2$ m, with constant inflow $Q_{\text{in}} = 0.01$ m³/s, simulate the system for 100 seconds and plot $h_1(t)$ and $h_2(t)$.
3. Consider the mass-spring-damper with $m = 2$ kg, $c = 4$ N·s/m, $k = 3$ N/m. Compute the damping ratio ζ . Is the system underdamped, critically damped, or overdamped? Plot the phase portrait for initial conditions $(q, \dot{q}) = (1, 0)$ and $(0.5, 1)$.
4. For the pendulum system (Example 2.3.1) with $L = 1$ m, $g = 9.81$ m/s², write the nonlinear state equations. Linearize around both the upright ($\theta = 0$) and inverted ($\theta = \pi$) equilibria. Classify each equilibrium.

Conceptual Exercises

5. Explain in your own words why knowledge of position alone is insufficient to specify the state of a system. Use a simple example (e.g., a mass on a spring) to illustrate why velocity must also be included.
6. Why did Rudolf Kalman's 1960 paper on state-space representations represent a breakthrough compared to classical transfer-function methods? What specific limitations of transfer functions does state space overcome?
7. Consider a 3-dimensional system with eigenvalues $\lambda_1 = -2$, $\lambda_2 = 1 + 2i$, $\lambda_3 = 1 - 2i$. Sketch a qualitative phase portrait. What will happen to a trajectory starting near the origin?
8. A system is controllable but not observable. What does this mean physically? Can you still design a feedback controller, and if so, what information would you need?
9. The Kalman rank condition for controllability checks that $\text{rank}([\mathbf{B}, \mathbf{A}\mathbf{B}, \dots, \mathbf{A}^{n-1}\mathbf{B}]) = n$. Why is it sufficient to check only up to $\mathbf{A}^{n-1}\mathbf{B}$ and not higher powers?

Proof-Based Exercises

10. Prove that if $\mathbf{A}$ is diagonalizable with eigenvalues $\lambda_1, \dots, \lambda_n$, then the matrix exponential is:

$$e^{\mathbf{A}t} = \mathbf{P} \text{diag}(e^{\lambda_1 t}, \dots, e^{\lambda_n t}) \mathbf{P}^{-1} \quad (2.77)$$

where $\mathbf{P}$ is the matrix of eigenvectors.

11. Derive the Jacobian matrix for the pendulum system $\ddot{\theta} + (g/L)\sin(\theta) = \tau$ at the bottom equilibrium ($\theta = 0$). Show that the linearized system is $\ddot{\theta} = -(g/L)\theta$ and classify the equilibrium's stability.

12. Show that for a 2D system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$ with $\mathbf{A} = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$, the eigenvalues are $\lambda = \frac{(a+d) \pm \sqrt{(a-d)^2 + 4bc}}{2}$. Under what conditions are the eigenvalues (i) both real, (ii) complex conjugates?
13. Prove that the controllability matrix $\mathbf{C}_{\text{rank}} = [\mathbf{B}, \mathbf{A}\mathbf{B}, \dots, \mathbf{A}^{n-1}\mathbf{B}]$ has full rank if and only if there is no eigenvector of $\mathbf{A}$ orthogonal to $\mathbf{B}^T$.
14. Using the variation of constants formula, show that the solution to $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}(t)$ with $\mathbf{x}(0) = \mathbf{x}_0$ is:

$$\mathbf{x}(t) = e^{\mathbf{A}t}\mathbf{x}_0 + \int_0^t e^{\mathbf{A}(t-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \quad (2.78)$$

15. (Challenge) Prove that a system is observable (can reconstruct initial state from output) if and only if the observability matrix $\mathbf{O} = [\mathbf{C}^T, \mathbf{A}^T\mathbf{C}^T, \dots, (\mathbf{A}^T)^{n-1}\mathbf{C}^T]^T$ has full rank. Hint: Use the dual system.

Chapter 3

Configuration Space and Degrees of Freedom

A golfer can place the clubhead at the same point using different elbow, shoulder and torso configurations. Those postures need not permit the same next motion, require the same joint torques, or produce the same impact. Configuration space provides the language for distinguishing these questions. It describes the modeled arrangement of the bodies; kinematics maps that arrangement to a task; dynamics determines how forces change the motion.

This chapter develops local coordinates, constraint counts, tangent vectors, forward and inverse kinematics, workspace and non-holonomic motion. The examples are specified mechanisms, not measurements of an anatomical golfer. Keeping that boundary explicit lets us transfer the mathematics without transferring unjustified performance claims.

3.1 Historical Context: From Lagrange to Riemann

Generalized coordinates let mechanics describe a constrained arrangement without carrying every particle coordinate independently. Manifold theory explains why a useful local coordinate description need not extend over the entire configuration space. Modern robotics combines these ideas with maps between joint arrangements and end-effector tasks [31, 27]. A coordinate system is a representation; it does not determine the underlying physical freedom.

Nor does curvature define a manifold: Euclidean space is a manifold, and a torus can carry an intrinsically flat metric despite its familiar curved embedding in three dimensions. Topology, smooth structure and metric answer different questions. A model's mass matrix later supplies one physically meaningful metric; the configuration manifold itself does not select it.

3.2 Degrees of Freedom (DOF) and Grübler's Formula

3.2.1 Local Freedom and Independent Constraints

At a regular configuration, the number of degrees of freedom is the local dimension of the configuration space. A free rigid body has three translational and three rotational freedoms. Six local coordinates suffice near a regular chart point, but there is no globally nonsingular, unique three-angle description of all orientations. A rotation matrix or unit quaternion stores more scalar entries with constraints; that storage does not add physical freedoms.

For an ambient configuration manifold of dimension D , suppose smooth equations $\phi(q) = 0$ have locally constant derivative rank r and define a regular level set. Then

$$\dim \mathcal{Q} = D - r, \quad T_q \mathcal{Q} = \ker D\phi(q). \quad (3.1)$$

This follows from the constant-rank theorem: locally, r independent coordinates are fixed by the equations and $D-r$ remain free. Counting written equations instead of their independent rank can overcount constraints.

At a singular point the linearized kernel can be misleading. The equation $x^2 + y^2 = 0$ has the single real solution $(0, 0)$, while its derivative there is zero and its kernel is all of $\mathbb{R}^2$. A two-dimensional linearized kernel does not make the solution set locally two-dimensional. Closed linkages may similarly have singular branches or instantaneous motions that do not extend to finite motion.

Two welded floating bodies form one rigid assembly with six freedoms. A hinge between two otherwise free spatial bodies leaves six assembly freedoms plus one relative rotation. Fixing one body to ground removes the six assembly freedoms. These counts concern the chosen mechanical model, not the number of motors or available control inputs.

3.2.2 Grübler's Formula

Let n count moving rigid links, excluding a fixed ground link. Let joint i permit d_i relative freedoms within a planar or spatial model. Under independent regular joint constraints, a planar count is

$$\text{DOF} = 3n - \sum_i c_i, \quad c_i = 3 - d_i, \quad (3.2)$$

or, for j joints,

$$\text{DOF} = 3n - 3j + \sum_{i=1}^j d_i. \quad (3.3)$$

Each moving planar body begins with two translations and one rotation. A revolute or prismatic joint removes two independent relative freedoms. Subtracting their independent constraints gives the formula. For spatial bodies, replace each 3 by 6:

$$\text{DOF} = 6n - \sum_{i=1}^j (6 - d_i). \quad (3.4)$$

With $N = n + 1$ total links including ground, both forms are $\lambda(N - 1 - j) + \sum_i d_i$, where $\lambda = 3$ or 6. This is a useful count, not a proof that a particular linkage has independent constraints. Redundant constraints, special geometry and singular configurations require a rank or geometric analysis [27].

3.2.3 Worked Example 1: Planar Four-Bar Linkage

Three moving links and four revolute joints give $3(3) - 4(2) = 1$ regular freedom. Write a closure equation using absolute link angles and a fixed ground displacement de_x :

$$ae(\theta_1) + be(\theta_2) - ce(\theta_3) - de_x = 0, \quad e(\theta) = (\cos \theta, \sin \theta)^T. \quad (3.5)$$

These are two scalar equations in three angles. Their derivative normally has rank two. The input θ_1 determines the other angles locally only when the two-column derivative with

respect to (θ_2, θ_3) is invertible. Its determinant is $-bc\sin(\theta_3 - \theta_2)$, so that choice of input coordinate fails when the coupler and output link are parallel.

Different assembly branches can share an input angle. The branch must be specified and followed continuously; one motor does not make the global closure solution unique. At some singular postures the full closure derivative also loses rank, and finite mobility needs more than a first-order count.




Figure 3.1: Four-Bar Closure on One Regular Assembly Branch. Ground Contains Two Fixed Pivots; the Input Angle Is Only a Local Coordinate Where the Selected Closure Derivative Is Nonsingular.

3.2.4 Worked Example 2: Spatial Stewart Platform

For an ideal six-leg UPS platform, each leg has a universal joint, a prismatic joint and a spherical joint, with $2 + 1 + 3$ freedoms. There are $N = 14$ bodies including ground, $n = 13$ moving bodies and $j = 18$ joints. Thus the regular independent-constraint count is

$$6(14 - 1 - 18) + 6(2 + 1 + 3) = 6. \quad (3.6)$$

This is a mobility count for the mechanism with variable leg lengths. Local control of all platform velocities also requires the actuator-length derivative to have rank six.

If base anchors are a_i , platform-frame anchors are b_i , and the platform pose is (p, R) , the required lengths are directly

$$\ell_i = \|p + Rb_i - a_i\|. \quad (3.7)$$

No coupled root-finding is needed to obtain lengths from a specified pose. The reverse problem, pose from lengths, may have multiple solutions. Joint-angle limits, leg interference and stroke limits must still be checked. A singular length-to-pose relation is not diagnosed by merely recounting bodies.

3.2.5 Worked Example 3: Human Arm Kinematics

A common reduced arm model assigns three shoulder rotations, one elbow flexion angle, one forearm pronation/supination angle and two wrist angles. That gives seven coordinates relative to a fixed shoulder base. Forearm rotation must not be relabeled as a third rotation of a ball-and-socket wrist. Scapular motion, shoulder translation, compliant tissues and hand joints change a more detailed model's configuration.

For a six-dimensional hand-pose task, a seven-coordinate model has a one-dimensional local inverse-kinematics family near a reachable regular point where the task derivative has rank six. This conclusion follows locally from the regular-level-set theorem. It does not say that every target is reachable or that every posture admits the same redundancy. Joint limits, singularities and grasp constraints can change the admissible set.

With both hands holding a rigid club, the arms and torso form a closed kinematic loop. Counting two unconstrained seven-coordinate arms and adding a club would ignore the grip constraints. Compliant grips require a different model. In either case, configuration redundancy is distinct from muscle-force redundancy and from the ability to alter a selected impact outcome.

3.3 The Configuration Space and Topological Manifolds

The configuration space $\mathcal{Q}$ is the set of modeled arrangements satisfying the selected geometric conditions. Its regular unconstrained regions are often smooth manifolds. Joint limits can introduce boundaries and corners; collision exclusions remove configurations; singular loop closures can create sets that are not manifolds. A contact-mode change can alter which equations apply. Therefore “every physical configuration space is a smooth manifold” is too strong.

A topological n -manifold without boundary is a Hausdorff, second-countable space in which every point has a neighborhood homeomorphic to an open subset of $\mathbb{R}^n$. Hausdorff means distinct points have disjoint neighborhoods; second-countable means a countable basis suffices. A manifold with boundary uses a local half-space at its boundary instead.

The circle $S^1 = \{(x, y) : x^2 + y^2 = 1\}$ has dimension one. The sphere S^2 has dimension two, including at its poles; longitude fails at the poles, not the manifold. Two independent unrestricted revolute angles give $S^1 \times S^1$, a torus. With restricted joint ranges, the admissible set can instead be a rectangle or another subset of that torus. Tracking cable wind-up or accumulated revolutions introduces additional model information; an unwrapped angle is not automatically the physical orientation itself.

For a free spatial rigid body, position and orientation give $\mathbb{R}^3 \times SO(3)$ as a manifold. Its group multiplication is the semidirect product $SE(3)$, because rotating one frame also changes how translations compose. The orientation space $SO(3)$ is three-dimensional and compact. Unit quaternions give a two-to-one representation: opposite quaternion points represent the same orientation.

3.4 Charts, Atlases, and Coordinate Patches

A chart is a homeomorphism $\varphi : U \rightarrow \varphi(U) \subset \mathbb{R}^n$, where both U and its image are open in their respective spaces. An atlas is a collection of charts covering the manifold. A smooth atlas has smooth transition maps on every overlap; the associated maximal compatible atlas

defines the smooth structure. Here “smooth” means C^∞ , although other differentiability classes can be specified.

3.4.1 Worked Example: Two Valid Charts on the Circle

No single chart covers all of S^1 . A hypothetical homeomorphism would map the compact circle onto a nonempty compact open subset of $\mathbb{R}$, which does not exist. The half-open interval representation $[0, 2\pi)$ is a bijective labeling of circle points, but its inverse-angle assignment jumps across the cut and its image is not open. Boundedness alone is not the problem: a bounded open interval can be a valid chart image.

Let $U_N = S^1 \setminus \{(0, 1)\}$ and $U_S = S^1 \setminus \{(0, -1)\}$. Stereographic coordinates are

$$u = \varphi_N(x, y) = \frac{x}{1 - y}, \quad v = \varphi_S(x, y) = \frac{x}{1 + y}. \quad (3.8)$$

Both images are all of $\mathbb{R}$. Their explicit inverses are

$$\varphi_N^{-1}(u) = \frac{(2u, u^2 - 1)}{1 + u^2}, \quad \varphi_S^{-1}(v) = \frac{(2v, 1 - v^2)}{1 + v^2}. \quad (3.9)$$

Substitution verifies unit norm and recovers the original coordinate. Each point other than its chart’s omitted pole has one coordinate. The two chart domains cover the circle and overlap away from both poles.

In contrast, projecting to x on a region $y \geq -\epsilon$ is generally not injective: for a small positive $\eta < \epsilon$, both $(\sqrt{1 - \eta^2}, \eta)$ and $(\sqrt{1 - \eta^2}, -\eta)$ belong to the region and share x . A smooth formula is not a valid chart unless it is also a homeomorphism onto an open image.

3.5 Transition Maps and Smooth Compatibility

For overlapping charts, the transition map is $\varphi_j \circ \varphi_i^{-1}$ on $\varphi_i(U_i \cap U_j)$. It changes coordinates of the same point. It need not preserve their numerical values, Euclidean lengths or component velocities.

3.5.1 Worked Example: Circle Transition and Velocity

The circle inverses give, on $u \neq 0$,

$$v = \frac{1}{u}, \quad u = \frac{1}{v}, \quad \dot{v} = -\frac{\dot{u}}{u^2}. \quad (3.10)$$

These transition maps are smooth on their open, disconnected overlap images $\mathbb{R} \setminus \{0\}$. Their singularity at zero lies outside the overlap: zero is the opposite chart’s omitted pole.

For example, $u = 1$ corresponds to $(1, 0)$ and $v = 1$. If $\dot{u} = 2$, then $\dot{v} = -2$. These represent the same physical tangent vector. Differentiating the inverse map also gives the circle’s induced line element

$$ds^2 = \frac{4}{(1 + u^2)^2} du^2. \quad (3.11)$$

The same speed is obtained in the v chart after transforming both the metric and the velocity. Comparing raw coordinate rates without this transformation can create an apparent difference that is purely representational.

3.6 Tangent Vectors, Curves, and Derivations

The tangent space $T_q\mathcal{Q}$ is an n -dimensional vector space attached to a particular configuration. It describes unconstrained instantaneous directions on the smooth configuration manifold. Additional non-holonomic restrictions may allow only a subspace of these directions; an active unilateral boundary may instead admit a one-sided cone.

Two smooth curves through q represent the same tangent vector when their coordinate derivatives agree at zero in one chart. The chain rule for a transition map then makes them agree in every chart. The tangent vector is the equivalence class, while $\dot{q}$ is its coordinate representation.

Alternatively, a tangent vector is an $\mathbb{R}$ -linear derivation v on smooth real functions satisfying

$$v(fg) = v(f)g(q) + f(q)v(g). \quad (3.12)$$

A curve γ supplies $v(f) = \frac{d}{dt}|_0 f(\gamma(t))$. Conversely, in coordinates z centered at q , a smooth function locally has the form

$$f(z) = f(0) + \sum_i z_i h_i(z), \quad h_i(0) = \partial_i f(0). \quad (3.13)$$

The product rule gives $v(f) = \sum_i v(z_i) \partial_i f(0)$. Constants have zero derivative, and a derivation depends only on the function near q ; smooth cutoff functions justify using local coordinate functions. The coordinate curve $z(t) = t(v(z_1), \dots, v(z_n))$ realizes this derivation for sufficiently small t . Thus curve classes and derivations are isomorphic, with n independent components.

The tangent bundle is the disjoint union

$$T\mathcal{Q} = \bigsqcup_{q \in \mathcal{Q}} T_q\mathcal{Q}. \quad (3.14)$$

Its elements are pairs (q, v_q) and it has dimension $2n$. Locally its coordinates transform as $(z, \dot{z}) \mapsto (\tau(z), D\tau(z)\dot{z})$. In general it is not a global product $\mathcal{Q} \times \mathbb{R}^n$.

For an unconstrained second-order mechanical model, configuration and velocity can form the state in $T\mathcal{Q}$. With a rank- k homogeneous allowed-velocity distribution, the constrained mechanical state lies in that distribution and has dimension $n + k$, under the regularity assumptions. Muscle activation, tendon deformation, flexible-shaft modes, controller memory or contact-mode variables add states. Consequently “state always has twice the configuration dimension” is not a general modeling rule.

3.7 Generalized Coordinates and Forward Kinematics

Generalized coordinates parameterize configurations locally; they need not all be motor angles or all be independently actuated. Prismatic displacement, a passive hinge angle and a floating-base coordinate are examples. A redundant constrained representation can also be computationally useful, even though it is not a minimal chart.

For a fixed-length arm model, joint coordinates and a base pose determine Cartesian segment positions. A shoulder angle alone is not enough for a general spatial shoulder: its orientation requires the chosen multi-coordinate representation. The anatomical elbow is approximated as a hinge only within a declared model.

Forward kinematics is a map $f : \mathcal{Q} \rightarrow \mathcal{Y}$ to a selected task space. Position may lie in $\mathbb{R}^3$, planar pose in $SE(2)$, and spatial pose in $SE(3)$. These targets are different maps, with potentially different derivative ranks.

For planar links of positive lengths L_1, L_2 and relative joint angles,

$$\begin{aligned} x &= L_1 \cos q_1 + L_2 \cos(q_1 + q_2), \\ y &= L_1 \sin q_1 + L_2 \sin(q_1 + q_2). \end{aligned} \quad (3.15)$$

The final-link orientation is $q_1 + q_2$ modulo 2π . The position equations by themselves do not include it as a controlled task coordinate.

3.8 The Jacobian Matrix and the Kinematic Derivative

The intrinsic derivative is $Df_q : T_q\mathcal{Q} \rightarrow T_{f(q)}\mathcal{Y}$. In local coordinates its matrix has entries

$$J_{ij}(q) = \frac{\partial f_i}{\partial q_j}, \quad \dot{y} = J(q)\dot{q}. \quad (3.16)$$

In particular, the bottom-left entry is $\partial f_m / \partial q_1$, not $\partial f_m / \partial q_n$. Changing nonsingular coordinate charts changes the matrix but preserves its rank. Singular coordinate representations can introduce an apparent singularity that is not a loss of physical task freedom.

For a spatial pose, Euler-angle derivatives are not generally the angular-velocity vector. A geometric Jacobian maps the declared generalized velocities to a declared twist, with its frame, reference point and angular/linear ordering specified. A six-coordinate pose derivative and a six-component twist Jacobian are related through the relevant chart map, not interchangeable by notation alone.

3.8.1 Worked Example: Two-Link Position Jacobian

Differentiating the two-link equations gives

$$J_p = \begin{bmatrix} -L_1 \sin q_1 - L_2 \sin(q_1 + q_2) & -L_2 \sin(q_1 + q_2) \\ L_1 \cos q_1 + L_2 \cos(q_1 + q_2) & L_2 \cos(q_1 + q_2) \end{bmatrix}, \quad (3.17)$$

with $\det J_p = L_1 L_2 \sin q_2$. For $L_1 = 2, L_2 = 1, q = (0, 0)$,

$$J_p = \begin{bmatrix} 0 & 0 \\ 3 & 1 \end{bmatrix}. \quad (3.18)$$

Only vertical endpoint velocity is available to first order. The arm still has two joint configuration freedoms. Motion in its position-Jacobian null space can alter joint posture, and second-order endpoint changes can appear even when the first-order velocity is zero.

3.8.2 Worked Example: Three-Link Planar Pose

Let $\alpha_i = \sum_{j=1}^i q_j$. A three-link planar pose is

$$(x, y) = \sum_{i=1}^3 L_i (\cos \alpha_i, \sin \alpha_i), \quad \theta = \alpha_3 \pmod{2\pi}. \quad (3.19)$$

On a local angle chart, a compact expression for every Jacobian column is

$$J_{1j} = -\sum_{i=j}^3 L_i \sin \alpha_i, \quad J_{2j} = \sum_{i=j}^3 L_i \cos \alpha_i, \quad J_{3j} = 1. \quad (3.20)$$

Subtracting consecutive columns shows $\det J = L_1 L_2 \sin q_2$. Full rank gives a locally invertible planar pose map. It does not establish global uniqueness or a six-dimensional spatial pose capability.

3.9 Inverse Kinematics

Inverse kinematics asks for configurations satisfying $f(q) = y_d$ within the admissible set. Solutions may be absent, isolated, multiple or continuous. A mathematical solution may violate a joint limit or collision condition. Even an admissible configuration does not supply a dynamically feasible path to it.

3.9.1 Geometric Two-Link Inverse Kinematics

For a target (x_d, y_d) with $r^2 = x_d^2 + y_d^2$, the forward equations imply

$$r^2 = L_1^2 + L_2^2 + 2L_1 L_2 \cos q_2. \quad (3.21)$$

Therefore, defining $c_2 = (r^2 - L_1^2 - L_2^2)/(2L_1 L_2)$, a reachable target requires $|c_2| \leq 1$. The two generic branches are

$$\begin{aligned} s_2 &= \pm \sqrt{1 - c_2^2}, & q_2 &= \text{atan2}(s_2, c_2), \\ q_1 &= \text{atan2}(y_d, x_d) - \text{atan2}(L_2 s_2, L_1 + L_2 c_2). \end{aligned} \quad (3.22)$$

Angles are understood modulo 2π before applying physical joint ranges. The sign of the cosine numerator follows directly from the squared vector sum; using the opposite sign produces the wrong elbow angle.

For $L_1 = L_2 = 1$ and target $(1, 1)$, the branches are $(0, \pi/2)$ and $(\pi/2, -\pi/2)$. Substitution recovers the same position but different final-link orientations. At ordinary inner or outer boundaries, the two branches coalesce modulo angle periodicity. At the equal-link origin, $q_2 = \pi$ and any q_1 works, so the solution is a continuous family and $\text{atan2}(0, 0)$ must not be treated as a unique shoulder solution.

3.9.2 Numerical Inverse Kinematics and Convergence

For local Euclidean task coordinates, set $e(q) = y_d - f(q)$. Linearizing gives $J(q)\delta q \approx e(q)$. A common iteration is

$$q_{k+1} = q_k + \alpha_k J(q_k)^\dagger e(q_k). \quad (3.23)$$

For an invertible square Jacobian, solve the linear system rather than explicitly forming an inverse. Classical Newton iteration has local quadratic convergence near a nonsingular root under suitable smoothness and sufficiently close initialization, with full steps eventually accepted. That statement is not a universal guarantee for singular, rectangular, damped or constrained iterations.

The pseudoinverse minimizes the linearized residual, and among tied solutions chooses the smallest Euclidean coordinate step. Both choices depend on scaling and metrics. A regularized alternative solves

$$\min_{\delta q} \|W^{1/2}(J\delta q - e)\|^2 + \lambda^2 \|R^{1/2}\delta q\|^2, \quad (3.24)$$

with declared positive-definite task and joint weights. It trades residual against step size; it does not make an unreachable target reachable. Use meaningful stopping tolerances, step limits and admissibility checks. On an orientation manifold, use a consistent local error or retraction rather than subtracting arbitrary wrapped angles across a cut.

3.10 Reachable and Dexterous Workspace

Let $\mathcal{Q}_{\text{adm}}$ be the admissible configuration set and p the position map. The reachable position workspace is $p(\mathcal{Q}_{\text{adm}})$. For a required orientation set $\mathcal{R}_{\text{req}}$, define the dexterous position workspace by

$$\begin{aligned} p_d \in \mathcal{W}_{\text{dex}} \iff & \text{for every } R_d \in \mathcal{R}_{\text{req}}, \\ & \text{there exists } q \in \mathcal{Q}_{\text{adm}} \text{ such that} \\ & (p(q), R(q)) = (p_d, R_d). \end{aligned} \quad (3.25)$$

The configuration may differ for each orientation. Existence of one full-rank Jacobian at a position is a local differential property, not this global orientation quantifier. A determinant is also undefined for a rectangular Jacobian.

For two unrestricted positive-length planar links, vector geometry gives

$$|L_1 - L_2| \leq \|p\| \leq L_1 + L_2. \quad (3.26)$$

Every intermediate radius occurs as q_2 varies, and q_1 rotates the resultant through all directions. Unequal lengths give a closed annulus; equal lengths give a closed disk. With $L_1 = 2, L_2 = 1$, its area is $\pi(3^2 - 1^2) = 8\pi$ in squared length units. This analytic argument establishes coverage; a finite sample cloud only illustrates it.

For an equal-link planar 2R arm, the folded origin is reachable with every final-link orientation by varying q_1 while keeping $q_2 = \pi$. Yet its position Jacobian has rank one. At a generic nonzero point, the two inverse branches give only two final-link orientations, even when the position Jacobian is full rank. Thus, for the all-planar-orientations requirement and unrestricted equal links, the dexterous workspace is just the origin. The example directly separates dexterity in this defined sense from local position manipulability.

The magnitude $|\det J_p| = L_1 L_2 |\sin q_2|$ measures local area scaling for the selected coordinate metrics. It is not a geometric distance to singularity or a condition number. Singular values describe directional gains; their ratio describes conditioning after choosing units and metrics. Two matrices with equal determinant can have very different conditioning, such as I_2 and $\text{diag}(100, 0.01)$. A colored determinant plot must say what it measures.

3.11 The Frobenius Theorem and Integrability

A smooth constant-rank distribution Δ assigns a smoothly varying k -dimensional subspace $\Delta(q) \subset T_q \mathcal{Q}$. It is integrable when it is tangent to a foliation by k -dimensional leaves. The

regular Frobenius theorem states that this occurs exactly when the distribution is involutive:

$$X, Y \text{ smooth sections of } \Delta \implies [X, Y] \text{ is a section of } \Delta. \quad (3.27)$$

Use the convention $[X, Y] = DY X - DX Y$, equivalent to the commutator of derivations. The brackets can leave the original distribution when it is nonintegrable. They need not span the entire ambient tangent space.

For example, on $\mathbb{R}^4$ with coordinates (x, y, z, w) , take

$$X_1 = \partial_x, \quad X_2 = \partial_y + x\partial_z, \quad [X_1, X_2] = \partial_z. \quad (3.28)$$

The rank-two distribution is nonintegrable, but every generated motion keeps w fixed. Its Lie closure has rank three, not four. Nonintegrability alone cannot imply full configuration reachability.

3.11.1 A Specified Unicycle and Its Lie Bracket

For an ideal unicycle with independently commanded signed forward speed and heading rate, let $q = (x, y, \theta)$ and

$$X_1 = (\cos \theta, \sin \theta, 0)^T, \quad X_2 = (0, 0, 1)^T. \quad (3.29)$$

The no-lateral-slip restriction is

$$-\sin \theta \dot{x} + \cos \theta \dot{y} = 0, \quad (3.30)$$

so the allowed-velocity distribution has rank two in a three-dimensional configuration space. Its bracket is

$$[X_1, X_2] = (\sin \theta, -\cos \theta, 0)^T. \quad (3.31)$$

These three vectors span the tangent space everywhere. For the smooth driftless system with reversible controls around zero, the bracket-generating condition yields local controllability; appropriate connected, obstacle-free models admit connections by concatenated admissible motions [31, 5]. This conclusion uses more than failure of Frobenius.

The four-step sequence forward, rotate, backward, reverse-rotate, each for duration ϵ with unit magnitudes, produces

$$\Delta q = \epsilon^2 [X_1, X_2] + O(\epsilon^3). \quad (3.32)$$

Every segment obeys the velocity constraint. The net sideways displacement is not instantaneous sideways slipping. Its order also matters: a small maneuver does not provide a first-order sideways velocity channel. Finite time, speed limits, collision clearance and lack of reverse motion can change the practical or local reachability question.

For two points connected by a positive-length rigid rod, $\|r_1 - r_2\|^2 = L^2$ is a regular position constraint in six-dimensional particle space. Its tangent constraint is integrable and the selected length level set has dimension five. The differentiated velocity equation alone preserves whichever separation was selected initially; an initial position constraint chooses the physical leaf.

3.12 Holonomic and Non-Holonomic Constraints

A holonomic equality constraint can be written $\phi(q, t) = 0$. At regular points it defines the corresponding position set or time-dependent family. Differentiation gives $D_q\phi\dot{q} + \partial_t\phi = 0$. Time-dependent constraints are generally affine in velocity. Joint-limit and separation inequalities need their own admissibility and active-set treatment; an inequality is not simply another equality that always subtracts a dimension.

A genuinely non-holonomic velocity restriction cannot locally be replaced by a family of position equalities of the same codimension. It restricts paths in the ambient configuration space. Which positions a controlled system can reach still depends on all its vector fields, their Lie closure, control bounds, drift and obstacles.

3.12.1 Upright Rolling Wheel

For an ideal upright wheel of radius $r > 0$, use configuration (x, y, θ, φ) , with heading θ and signed rolling angle φ . Choose the sign of φ so positive rolling advances along heading. No slip then requires

$$\dot{x} = r\dot{\varphi} \cos \theta, \quad \dot{y} = r\dot{\varphi} \sin \theta. \quad (3.33)$$

These are two independent velocity constraints on a four-dimensional configuration. The factors have units of speed. Replacing the signed rolling relation by an unspecified length ratio times speed gives the wrong units and loses the direction of rotation. An actual coin that can lean needs additional configuration and contact dynamics; this is the upright-wheel idealization.

3.12.2 Unicycle Versus Bicycle Steering

The independently driven unicycle has two controls:

$$\dot{x} = v \cos \theta, \quad \dot{y} = v \sin \theta, \quad \dot{\theta} = \omega. \quad (3.34)$$

A kinematic bicycle with wheelbase $L > 0$ and steering angle δ instead uses $\dot{\theta} = v \tan \delta / L$. Its admissible $(v, \dot{\theta})$ pairs depend on steering bounds, and it cannot turn in place with finite steering angle. If steering rate is the input, δ becomes a state. Neither model is accurately described as a one-input unicycle while simultaneously treating speed and steering as independent inputs.

3.12.3 Ball and Moving-Surface Contact

Let C be the ball center and P its contact point, all expressed in the same inertial frame. No slip equates velocities of the contacting material points:

$$v_C + \omega \times r_{CP} = v_{\text{surface}}(P, t). \quad (3.35)$$

Here r_{CP} is the vector from center to contact, not the contact's absolute position. For a sphere of radius r on a fixed plane with unit normal n toward the center, $r_{CP} = -rn$, giving $v_C = r\omega \times n$. Rotation about n is not excluded by no slip at an ideal point; a no-twist or finite-patch law would add a condition.

Maintained sphere-plane contact leaves two center-position coordinates plus three orientation freedoms, a five-dimensional configuration manifold. Tangential no slip removes

two velocity directions, leaving rank three and an eight-dimensional constrained mechanical state under this idealization. A tilted fixed plane changes gravitational dynamics, not the algebra of the contact-velocity equality. A moving plate requires its actual surface velocity and maintained-contact geometry. Friction and normal-force feasibility must be checked separately.

3.12.4 Pfaffian Constraints

A homogeneous Pfaffian constraint is $A(q)\dot{q} = 0$. Its independent rows are covectors that annihilate allowed velocities. If the rank is r on an n -dimensional manifold, $\ker A$ has dimension $n - r$. Calling velocities “orthogonal to the rows” uses the Euclidean metric of a chosen coordinate representation; the intrinsic statement is covector-vector pairing equal to zero.

For the unicycle, $A = (-\sin\theta, \cos\theta, 0)$ has one independent row and a two-dimensional kernel. A holonomic constraint’s differentiated equation can also be written in Pfaffian form. Thus the appearance of $A\dot{q} = 0$ alone does not prove non-holonomy; integrability is the additional question.

3.13 Python Example: Kinematics and Workspace

The repository module implements the stated planar chain, two-link inverse branches, circle charts and exact constant-input unicycle step. Its tests compare Jacobians with finite differences, close both inverse branches through the forward map, verify chart transitions and recover the finite-motion bracket sign. No numerical sample replaces the geometric proofs.

```

1 import numpy as np
2 from src.tools.configuration_examples import (
3     planar_pose, planar_jacobian, inverse_two_link,
4     circle_point, circle_coordinate, unicycle_step,
5 )
6
7 lengths = np.ones(2)
8 angles = np.array([np.pi / 4, np.pi / 3])
9 pose = planar_pose(lengths, angles)
10 print(pose)
11 print(np.linalg.det(planar_jacobian(lengths, angles)[:2]))
12 for branch in inverse_two_link(lengths, pose[:2]):
13     np.testing.assert_allclose(
14         planar_pose(lengths, branch)[:2], pose[:2]
15     )
16     print(branch)
17 point = circle_point(0.7, 1)
18 print(point, circle_coordinate(point, -1))
19 state = np.zeros(3)
20 for speed, turn in [(1, 0), (0, 1), (-1, 0), (0, -1)]:
21     state = unicycle_step(state, speed, turn, 0.01)
22 print(state / 0.01*2)

```

For the published unit-link example at $q = (\pi/4, \pi/3)$, the endpoint is approximately $(0.448287736, 1.673032607)$, with final-link orientation $7\pi/12$. The position determinant is

$\sqrt{3}/2$. The two-link inverse reproduces the position on both branches; comparing their orientations illustrates why a position task is not a pose task.




Figure 3.2: A Sampled Unit-Link Position Workspace and Its Local Area Scaling. The Exact Reachable Set Is a Disk; Determinant Magnitude Is Not Dexterous Workspace or Distance to Singularity.

The figure samples an unrestricted unit-link arm and colors configurations by $|\det J_p|$. The analytic reachable set is the disk of radius two. Joint-angle samples are not uniformly distributed in workspace, and overlapping points can hide branch information. This is an area-scaling illustration, not a dexterous-workspace plot, a singularity-distance map or evidence of a human swing strategy.

3.14 From Configuration Geometry to a Golf-Swing Argument

A defensible analysis first names the modeled bodies, joint coordinates, admissible contact and grip modes, and the task. Clubhead position, face orientation, path direction, impact speed and contact location are different outputs. Their Jacobians and tolerances differ. A change invisible to one task may matter to another.

At a fixed regular posture, $\delta y = J\delta q$ describes first-order geometric sensitivity. A null-space displacement can leave that task unchanged to first order while changing joint leverage, mass distribution, available muscle force or contact feasibility. It need not preserve the task after a finite displacement; following a level set requires integrating a compatible direction field or resolving the nonlinear constraint.

Kinematic freedom is also not torque authority. Ground reactions and two-hand grip forces couple the body and club. Joint torques act through the mass matrix and the constraint equations; muscle activation and tendon dynamics restrict how those torques can change. A desired endpoint velocity in the Jacobian image may demand unavailable joint

rates or forces. Conversely, a singular instantaneous map need not eliminate all finite-motion options.

Finally, club delivery is evaluated along a trajectory, often at an impact event. The same posture reached with different velocities, activation states or shaft deformation can produce different outcomes. A claim that a geometric feature improves delivery must specify a feasible intervention, propagate the controlled dynamics and evaluate the selected outcome under measurement uncertainty. This connects configuration geometry to dynamics and control while keeping their distinct evidential roles clear.

3.15 Chapter Summary and Key Principles

Configuration dimension is local freedom under a stated model. Charts represent it locally; topology and constraints determine where those charts work. Tangent vectors transform with chart derivatives, and state requires all variables needed to predict motion. Kinematic maps and their ranks depend on the selected task. Inverse solutions and workspace require global admissibility checks beyond local rank.

Frobenius tests integrability of a regular distribution. Reachability additionally requires the controlled vector fields and their admissible inputs. Geometric freedom, instantaneous velocity authority and dynamically feasible golf delivery are related through explicit maps and equations, not interchangeable descriptions of the same quantity.

3.16 Further Reading

Lynch and Park [27] distinguish configuration topology, independent-constraint counts, task space and workspace. Their official task-space supplement [26] explains the orientation requirement in dexterous workspace. The numerical inverse-kinematics supplement [25] develops the pseudoinverse's local linear role; the mobile-robot controllability supplement [24] connects bracket rank to admissible control sets.

Murray, Li and Sastry [31], Siciliano and colleagues [38], and Bullo and Lewis [5] provide broader treatments of robot kinematics and geometric mechanics/control. The explicit chart inverses, triangle derivation and nonintegrability counterexample here can be checked directly. A source citation does not remove the need to state regularity, coordinate conventions or the physical model used in an application.

3.17 Exercises

1. Explain the six local freedoms of a free spatial rigid body. Why do a rotation matrix's nine entries and a unit quaternion's four entries not increase its three rotational freedoms?
2. Count the regular mobility of a planar four-bar with three moving links and four revolute joints. Differentiate its two closure equations. State when the input angle is a valid local coordinate and why assembly branches can share its value.
3. Prove that no single chart covers the circle using compactness and openness of a chart image. Explain why the angle labeling on $[0, 2\pi)$ is not such a chart. Give a bounded open interval that is a valid possible chart image.

4. Show that $T_{(x,y)}S^1 = \{(a,b) : xa + yb = 0\}$ has dimension one. Compare this tangent space with the derivative kernel of $x^2 + y^2 = 0$ at its isolated solution.
5. For two unit links at $q = (0,0)$, compute position and position Jacobian. Identify the available first-order endpoint velocity direction and a nonzero joint velocity giving zero endpoint velocity.
6. A standard planar slider-crank has three moving links, three revolute joints and one prismatic joint. Sketch its connectivity and obtain the regular count. Explain why the mobility count alone does not uniquely identify these joint types.
7. For two unit links at $q = (\pi/2, 0)$, compute the singular position Jacobian and its image. Describe the missing radial velocity direction. Does the mechanism lose a joint configuration coordinate?
8. Derive the inner radius, outer radius and area of the unrestricted planar workspace with $L_1 = 2, L_2 = 1$. State what changes when the lengths are equal or the joint angles are limited.
9. Verify both stereographic inverses, derive $v = 1/u$ on the overlap and transform the velocity. Recover the same circle speed using the induced metric in either chart.
10. Write the ideal unicycle's Pfaffian constraint and count its kernel dimension. Compute $[X_1, X_2]$ using the stated convention. Explain which additional control hypotheses permit a local controllability conclusion, and test the four-step finite maneuver numerically.
11. Prove the reachable-radius bounds for arbitrary positive L_1, L_2 and show every intermediate radius occurs. For equal links, determine the all-planar-orientations dexterous workspace and explain why it differs from the nonsingular position workspace.
12. For a planar arm with lengths $(1, 1, 0.5)$ and target $(x, y, \theta) = (2, 0.5, \pi/4)$, subtract the final-link vector to obtain the two-link wrist target. Solve both branches geometrically. Compare with a safeguarded Newton iteration starting at $(\pi/6, \pi/6, \pi/6)$, reporting residuals and any rejected steps rather than assuming convergence.
13. For a sphere of radius r on a fixed horizontal plane, write its five-dimensional configuration, no-slip velocity equation and rank-three allowed-velocity space. Explain the eight-dimensional constrained mechanical state and how no twist or a moving plate would change the model.
14. Complete the curve/derivation equivalence proof using the local expansion $f(z) = f(0) + \sum_i z_i h_i(z)$. Explain why the equivalence is independent of chart choice and why tangent vectors at different base points are not automatically interchangeable.
15. For specified Stewart platform anchors, derive the six lengths directly from a requested pose. Then differentiate them with respect to translational velocity and angular velocity, identifying each row's moment arm and direction. State the rank test for local platform control and distinguish direct length calculation from solving pose from lengths.

Chapter 4

Rotations, Quaternions, and $SO(3)$

The tragedy of Euclidean space is that it works perfectly until you spin.

In Plain Language 4.1. Context & Use Case: The Problem with Spinning

The Math You Will See: We introduce **Rotation Matrices** ($\mathbf{R}$), the **Special Orthogonal Group** $SO(3)$, and **Quaternions** ($\mathbf{q}$). **Why We Need It:** Moving a robot arm up, down, left, or right is easy; you just add numbers to X, Y, Z . But calculating what happens when a drone performs a backflip while spinning like a top is a mathematical nightmare. The most intuitive way to track rotation (Euler angles like Roll, Pitch, Yaw) literally breaks the laws of math in certain positions, causing flight controllers to crash. To solve this, we must map rotations using completely different objects: 3×3 matrices and 4-dimensional Quaternions. **All Variables Defined:** $\mathbf{R} \in SO(3)$ is a rotation matrix, and $\mathbf{q} \in \mathbb{R}^4$ is a quaternion.

4.1 Historical Context: From Euler to Hamilton

The mathematics of 3D rotation has been refined over centuries by some of history's greatest mathematicians. In **1775**, Leonhard Euler published his groundbreaking *Euler's Rotation Theorem*, establishing that any finite rotation in 3D space can be described as a single rotation about a fixed axis by some angle—a foundational geometric insight that remains central to modern rotation calculus.

However, Euler's representation (axis-angle form) requires solving nonlinear equations and computing trigonometric functions. In **1843**, the Irish mathematician **William Rowan Hamilton** invented **Quaternions** ($\mathcal{H}$), an extension of complex numbers to four dimensions. Quaternions elegantly encode rotations as multiplication operations on 4-tuples, bypassing the singularities that plague other representations. Almost simultaneously, **Olinde Rodrigues** (1795–1853) derived the famous **Rodrigues rotation formula**, expressing how a vector rotates about an axis—a result that connects Euler's axis-angle theorem to matrix algebra.

These three frameworks—rotation matrices, quaternions, and axis-angle representations—are *mathematically equivalent* but offer different computational and conceptual advantages. Modern robotics and computer graphics employ all three, often switching between them fluidly.

4.2 The Coordinate Frame

When we model movement or control robotic systems, we are not looking at the physical object itself; we are tracking a **Coordinate Frame** rigidly attached to the object, often denoted as the Body Frame ($\{b\}$).




Figure 4.1: Spatial and body frames: The spatial frame is fixed in the world; the body frame is rigidly attached to a moving object. The relationship between them is encoded by a rotation matrix $\mathbf{R}$ and a position vector $\mathbf{p}$.

To know where the Body Frame $\{b\}$ is relative to the absolute stationary universe—the Spatial (or World) Frame ($\{s\}$)—we must define two mathematical properties:

1. The **translation** vector $\mathbf{p} \in \mathbb{R}^3$ pointing from the origin of $\{s\}$ to the origin of $\{b\}$.
2. The **rotation** of the axes of $\{b\}$ relative to the axes of $\{s\}$.

Unlike the straightforward (X, Y, Z) translation vector, rotation is mathematically subtle and surprisingly difficult to parameterize.

4.3 Rotation Matrices: From First Principles

4.3.1 What Is a Rotation Matrix?

To represent a rotation, we examine the three orthonormal basis vectors (axes) of the body frame $\{b\}$: the unit vectors $\hat{\mathbf{x}}_b$, $\hat{\mathbf{y}}_b$, and $\hat{\mathbf{z}}_b$, each expressed in the coordinates of the spatial frame $\{s\}$. Stack these three column vectors side-by-side:

$$\mathbf{R} = \begin{bmatrix} | & | & | \\ \hat{\mathbf{x}}_b & \hat{\mathbf{y}}_b & \hat{\mathbf{z}}_b \\ | & | & | \end{bmatrix} \in \mathbb{R}^{3 \times 3} \quad (4.1)$$

This is a **Rotation Matrix**. Each column is a unit vector, and any two columns are orthogonal to each other. This orthonormality is the key property.

4.3.2 Orthonormality Constraints

Theorem 4.1. Orthonormality Conditions

A matrix $\mathbf{R} \in \mathbb{R}^{3 \times 3}$ is a valid rotation matrix if and only if:

1. **Orthogonality:** $\mathbf{R}^T \mathbf{R} = \mathbf{I}$, which means $\mathbf{R}^{-1} = \mathbf{R}^T$.
2. **Determinant Constraint:** $\det(\mathbf{R}) = +1$.

Proof. Let $\mathbf{R} = [\mathbf{c}_1 \mid \mathbf{c}_2 \mid \mathbf{c}_3]$, where $\mathbf{c}_i$ are the columns. The product $\mathbf{R}^T \mathbf{R}$ yields a 3×3 matrix whose (i, j) entry is $\mathbf{c}_i^T \mathbf{c}_j$. For this to equal $\mathbf{I}$:

- Diagonal: $\mathbf{c}_i^T \mathbf{c}_i = 1$ for each i (each column is unit length).
- Off-diagonal: $\mathbf{c}_i^T \mathbf{c}_j = 0$ for $i \neq j$ (columns are mutually orthogonal).

These are exactly 6 independent constraints on the 9 matrix entries, leaving $9 - 6 = 3$ degrees of freedom.

The determinant condition $\det(\mathbf{R}) = 1$ distinguishes *proper rotations* (right-handed orientation preserving) from reflections, which would have $\det(\mathbf{R}) = -1$. $\square$

4.3.3 The Three Elementary Rotation Matrices

Any rotation in 3D can be decomposed into rotations about the principal axes. We derive the three elementary rotation matrices.

Theorem 4.2. Rotation About the x -Axis

A rotation by angle θ about the x -axis, with the rotation viewed as right-hand rule (counterclockwise when looking from positive x toward the origin), is given by:

$$\mathbf{R}_x(\theta) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta \\ 0 & \sin \theta & \cos \theta \end{bmatrix} \quad (4.2)$$

Geometric Interpretation:

- **Column 1:** $[1, 0, 0]^T$ is the x -axis, unchanged by rotation about itself.
- **Column 2:** $[0, \cos \theta, \sin \theta]^T$ is where the original y -axis rotates to.
- **Column 3:** $[0, -\sin \theta, \cos \theta]^T$ is where the original z -axis rotates to.

Derivation. The x -axis unit vector $\hat{\mathbf{e}}_x = [1, 0, 0]^T$ is invariant. The y and z axes lie in the yz -plane, which undergoes a 2D rotation:

$$\begin{bmatrix} y' \\ z' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} y \\ z \end{bmatrix}.$$

Embedding this in 3D yields equation (4.2). $\square$

Theorem 4.3. Rotation About the y -Axis

A rotation by angle θ about the y -axis is:

$$\mathbf{R}_y(\theta) = \begin{bmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{bmatrix} \quad (4.3)$$

Theorem 4.4. Rotation About the z -Axis

A rotation by angle θ about the z -axis is:

$$\mathbf{R}_z(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (4.4)$$

Each elementary matrix satisfies $\mathbf{R}^T \mathbf{R} = \mathbf{I}$ and $\det(\mathbf{R}) = 1$ by construction (they encode orthonormal bases).

4.4 The Special Orthogonal Group: $SO(3)$

4.4.1 $SO(3)$ as a Group

Definition 4.1. The Group $SO(3)(3)$

The **Special Orthogonal Group** $SO(3)(3)$ [31] is the set of all 3×3 real matrices satisfying:

$$SO(3)(3) = \{\mathbf{R} \in \mathbb{R}^{3 \times 3} : \mathbf{R}^T \mathbf{R} = \mathbf{I}, \det(\mathbf{R}) = 1\} \quad (4.5)$$

Theorem 4.5. $SO(3)(3)$ is a Group

The set $SO(3)(3)$ forms a mathematical group under matrix multiplication.

Proof. We verify the four group axioms:

Closure: If $\mathbf{R}_1, \mathbf{R}_2 \in SO(3)(3)$, then:

$$(\mathbf{R}_1 \mathbf{R}_2)^T (\mathbf{R}_1 \mathbf{R}_2) = \mathbf{R}_2^T \mathbf{R}_1^T \mathbf{R}_1 \mathbf{R}_2 = \mathbf{R}_2^T \mathbf{I} \mathbf{R}_2 = \mathbf{I}, \quad (4.6)$$

$$\det(\mathbf{R}_1 \mathbf{R}_2) = \det(\mathbf{R}_1) \det(\mathbf{R}_2) = 1 \cdot 1 = 1. \quad (4.7)$$

Thus $\mathbf{R}_1 \mathbf{R}_2 \in SO(3)(3)$.

Associativity: Matrix multiplication is associative: $(\mathbf{R}_1 \mathbf{R}_2) \mathbf{R}_3 = \mathbf{R}_1 (\mathbf{R}_2 \mathbf{R}_3)$.

Identity: The identity matrix $\mathbf{I}$ satisfies $\mathbf{I}^T \mathbf{I} = \mathbf{I}$ and $\det(\mathbf{I}) = 1$, so $\mathbf{I} \in SO(3)(3)$.

Inverses: For $\mathbf{R} \in SO(3)(3)$, we have $\mathbf{R}^T \mathbf{R} = \mathbf{I}$, so $\mathbf{R}^{-1} = \mathbf{R}^T$. Moreover:

$$\det(\mathbf{R}^T) = \det(\mathbf{R}) = 1, \quad (4.8)$$

so $\mathbf{R}^{-1} \in SO(3)(3)$. □

4.4.2 $\text{SO}(3)$ as a Smooth Manifold

Theorem 4.6. $\text{SO}(3)(3)$ is a 3-Dimensional Smooth Manifold

The set $\text{SO}(3)(3)$ is a smooth, compact manifold of dimension 3.

Argument. The orthogonality constraint $\mathbf{R}^T \mathbf{R} = \mathbf{I}$ imposes 6 independent scalar equations (the symmetric matrix $\mathbf{I}$ has 6 independent entries) on the 9 entries of $\mathbf{R}$. Additionally, $\det(\mathbf{R}) = 1$ is one equation, but it is redundant: for orthogonal matrices with real entries, $\det(\mathbf{R}) = \pm 1$ automatically, and we select the component with determinant $+1$ by continuity and connectedness.

Thus, the effective number of independent constraints is 6, leaving:

$$\dim(\text{SO}(3)(3)) = 9 - 6 = 3. \quad (4.9)$$

Each point in $\text{SO}(3)(3)$ admits a local parameterization by 3 parameters (e.g., Euler angles in a suitable domain), confirming that $\text{SO}(3)(3)$ is a smooth 3-dimensional manifold. $\square$

Mathematical Principle 4.1. The $\text{SO}(3)$ Manifold

The continuous, infinite geometric space of all possible valid 3×3 rotation matrices satisfying $\mathbf{R}^T \mathbf{R} = \mathbf{I}$ and $\det(\mathbf{R}) = 1$ is a mathematical **manifold** called the **Special Orthogonal Group in 3 Dimensions**, denoted as $\text{SO}(3)(3)$.

If a Configuration Space $\mathcal{Q}$ is simply a rigid body freely rotating around a fixed point, we say $\mathcal{Q} = \text{SO}(3)(3)$.

4.5 Euler Angles: Gimbal Lock and Beyond

4.5.1 Definition of Euler Angles

The most intuitive way for a human to describe the orientation of an aircraft or robotic arm is to split rotation into three successive rotations about primary axes. This is known as an **Euler Angle** representation: $\boldsymbol{\theta} = [\phi, \theta, \psi]^T$.

While highly intuitive for human pilots, Euler Angles harbor a fatal mathematical flaw known as **Gimbal Lock**.

4.5.2 The ZYX Convention

One common convention is **ZYX Euler angles** (also called *yaw-pitch-roll*):

1. Rotate by ψ (yaw) about the original z -axis.
2. Rotate by θ (pitch) about the *current* y -axis (after yaw).
3. Rotate by ϕ (roll) about the *current* x -axis (after yaw and pitch).

The composite rotation matrix is:

$$\mathbf{R}(\phi, \theta, \psi) = \mathbf{R}_x(\phi)\mathbf{R}_y(\theta)\mathbf{R}_z(\psi). \quad (4.10)$$

Computing this product explicitly:

$$\mathbf{R}(\phi, \theta, \psi) = \begin{bmatrix} \cos \psi \cos \theta & -\sin \psi \cos \phi + \cos \psi \sin \theta \sin \phi & \sin \psi \sin \phi + \cos \psi \sin \theta \cos \phi \\ \sin \psi \cos \theta & \cos \psi \cos \phi + \sin \psi \sin \theta \sin \phi & -\cos \psi \sin \phi + \sin \psi \sin \theta \cos \phi \\ -\sin \theta & \cos \theta \sin \phi & \cos \theta \cos \phi \end{bmatrix} \quad (4.11)$$

4.5.3 The ZZX Convention

Another standard convention is **ZZX Euler angles** (symmetric, used in physics and astronomy):

1. Rotate by ϕ about the original z -axis.
2. Rotate by θ about the *current* x -axis.
3. Rotate by ψ about the *current* z -axis (after the first two rotations).

The composite rotation is:

$$\mathbf{R}(\phi, \theta, \psi) = \mathbf{R}_z(\psi)\mathbf{R}_x(\theta)\mathbf{R}_z(\phi). \quad (4.12)$$

4.5.4 Gimbal Lock: Singularity and Proof

In Plain Language 4.2. Gimbal Lock: Physical Intuition

If an airplane pitches perfectly 90° straight up into the sky, its Roll axis (pointing out its nose) now perfectly aligns with its original Yaw axis (pointing perfectly down towards the Earth).

Mathematically, two independent degrees of freedom have become linearly dependent—they merged into the identical physical rotation axis. A degree of freedom is literally lost during the calculation. If you attempt to control the plane passing through this singularity, the Jacobian mapping Euler angle rates to angular velocity becomes singular, producing undefined or numerically unstable results (NaN values, wildly large rate commands).

Theorem 4.7. Gimbal Lock at $\theta = \pm 90^\circ$

For ZYX Euler angles, the Jacobian matrix relating angular velocity components to Euler angle rate components becomes singular when $\theta = \pm 90^\circ$.

Derivation. The relationship between angular velocity $\boldsymbol{\omega} = [\omega_x, \omega_y, \omega_z]^T$ in the spatial frame and Euler angle rates $\dot{\boldsymbol{\theta}} = [\dot{\phi}, \dot{\theta}, \dot{\psi}]^T$ is given by:

$$\boldsymbol{\omega} = \mathbf{J}(\boldsymbol{\theta}) \dot{\boldsymbol{\theta}}, \quad (4.13)$$

where $\mathbf{J}(\theta)$ is the Jacobian. For ZYX angles, it can be shown that:

$$\mathbf{J}(\phi, \theta, \psi) = \begin{bmatrix} 0 & -\sin \psi & \cos \psi \cos \theta \\ 0 & \cos \psi & \sin \psi \cos \theta \\ 1 & 0 & -\sin \theta \end{bmatrix} \quad (4.14)$$

When $\theta = 90^\circ$, we have $\cos \theta = 0$, so:

$$\mathbf{J}(\phi, 90^\circ, \psi) = \begin{bmatrix} 0 & -\sin \psi & 0 \\ 0 & \cos \psi & 0 \\ 1 & 0 & -1 \end{bmatrix}. \quad (4.15)$$

The third column is linearly dependent on the first two, so $\text{rank}(\mathbf{J}) < 3$ and $\det(\mathbf{J}) = 0$. The Jacobian is singular, making it impossible to invert: given $\dot{\omega}$, we cannot solve uniquely for $\dot{\theta}$. Physically, the yaw and roll rates become indistinguishable. $\square$

Since our control theory and geometric motion algorithms fundamentally rely on continuous trajectories twisting through all possible spatial angles (like chaotic golf club arcs or throwing motions), we must completely abandon Euler angles for internal calculations. We cannot afford mathematical singularities.

4.6 Differential Geometry Primer: Lie Groups and Lie Algebras

We have established that $\text{SO}(3)(3)$ is a manifold. However, it is not just any manifold. Because you can combine two rotations to get a third rotation (via matrix multiplication), and every rotation can be smoothly inverted, $\text{SO}(3)(3)$ possesses both algebraic structure (it is a mathematical *group*) and geometric structure (it is a *smooth manifold*).

Definition 4.2. Lie Group

A *Lie Group* is a smooth manifold that is also a mathematical group, where the group operations (multiplication and inversion) are smooth functions. As a result, we can perform calculus on them. $\text{SO}(3)(3)$ is the preeminent example of a Lie Group in mechanics.

4.6.1 The Tangent Space and Lie Algebra

If $\text{SO}(3)(3)$ is a curved manifold of positions (rotations), what is its velocity space? By our earlier definition in Chapter 3, the velocities must live in the Tangent Space.

If we examine the tangent space at the identity rotation (the "zero" rotation, matrix $\mathbf{I}$), we discover an incredibly special mathematical object called the **Lie Algebra**, denoted by lowercase $\mathfrak{so}(3)(3)$.

Definition 4.3. Lie Algebra $\mathfrak{so}(3)(3)$

The *Lie Algebra* $\mathfrak{so}(3)(3)$ is the tangent space to the Lie Group $\text{SO}(3)(3)$ at the identity element. Physically, it represents the space of all possible angular velocities

(spin rates) when the object is currently unrotated. Mathematically, it consists of all 3×3 *skew-symmetric* matrices: matrices $[\boldsymbol{\omega}]_{\times}$ satisfying $[\boldsymbol{\omega}]_{\times}^T = -[\boldsymbol{\omega}]_{\times}$.

A skew-symmetric matrix representing a vector $\boldsymbol{\omega} = [\omega_x, \omega_y, \omega_z]^T$ is:

$$[\boldsymbol{\omega}]_{\times} = \begin{bmatrix} 0 & -\omega_z & \omega_y \\ \omega_z & 0 & -\omega_x \\ -\omega_y & \omega_x & 0 \end{bmatrix} \quad (4.16)$$

In Plain Language 4.3. Mathematical Reminder: Group vs. Algebra

Whenever you hear **Lie Group** (like $SO(3)$ or $SE(3)$), think of the curved, global manifold of *Positions/Configurations*. Whenever you hear **Lie Algebra** (like $\mathfrak{so}(3)$ or $\mathfrak{se}(3)$), think of the perfectly flat, linear tangent space of *Velocities*. We can do standard linear control math in the Lie Algebra, but we map the results back onto the curved Lie Group using the *Exponential Map* (covered next) to actually rotate the physical robot.

4.7 Rodrigues' Rotation Formula: Full Derivation

We now prove the celebrated **Rodrigues' rotation formula**, which gives a closed-form expression for the matrix exponential of a skew-symmetric matrix. This result, first obtained by Olinde Rodrigues in 1840 [35], is the computational backbone of the exponential map for $SO(3)$.

4.7.1 The Key Identity: $[\hat{\boldsymbol{\omega}}]^3 = -[\hat{\boldsymbol{\omega}}]$

Let $\hat{\boldsymbol{\omega}} \in \mathbb{R}^3$ be a unit vector ($\|\hat{\boldsymbol{\omega}}\| = 1$) and $[\hat{\boldsymbol{\omega}}] \in \mathfrak{so}(3)$ the corresponding skew-symmetric matrix. We first establish a crucial algebraic identity.

Theorem 4.8. Cyclic Property of Skew-Symmetric Matrices

For any unit vector $\hat{\boldsymbol{\omega}}$ with skew-symmetric matrix $[\hat{\boldsymbol{\omega}}]$:

$$[\hat{\boldsymbol{\omega}}]^3 = -[\hat{\boldsymbol{\omega}}]. \quad (4.17)$$

Proof. Recall that for a skew-symmetric matrix, the action $[\hat{\boldsymbol{\omega}}]\mathbf{v} = \hat{\boldsymbol{\omega}} \times \mathbf{v}$. Thus:

$$[\hat{\boldsymbol{\omega}}]^2 \mathbf{v} = \hat{\boldsymbol{\omega}} \times (\hat{\boldsymbol{\omega}} \times \mathbf{v}). \quad (4.18)$$

By the vector triple product identity $\mathbf{a} \times (\mathbf{a} \times \mathbf{v}) = \mathbf{a}(\mathbf{a} \cdot \mathbf{v}) - \mathbf{v}(\mathbf{a} \cdot \mathbf{a})$ with $\|\hat{\boldsymbol{\omega}}\| = 1$:

$$[\hat{\boldsymbol{\omega}}]^2 = \hat{\boldsymbol{\omega}}\hat{\boldsymbol{\omega}}^T - \mathbf{I}. \quad (4.19)$$

Therefore:

$$[\hat{\boldsymbol{\omega}}]^3 = [\hat{\boldsymbol{\omega}}](\hat{\boldsymbol{\omega}}\hat{\boldsymbol{\omega}}^T - \mathbf{I}) = [\hat{\boldsymbol{\omega}}]\hat{\boldsymbol{\omega}}\hat{\boldsymbol{\omega}}^T - [\hat{\boldsymbol{\omega}}]. \quad (4.20)$$

Since $[\hat{\boldsymbol{\omega}}]\hat{\boldsymbol{\omega}} = \hat{\boldsymbol{\omega}} \times \hat{\boldsymbol{\omega}} = \mathbf{0}$, we obtain $[\hat{\boldsymbol{\omega}}]^3 = -[\hat{\boldsymbol{\omega}}]$. $\square$

This identity implies that all higher powers of $[\hat{\omega}]$ reduce to either $[\hat{\omega}]$ or $[\hat{\omega}]^2$:

$$[\hat{\omega}]^4 = -[\hat{\omega}]^2, \quad [\hat{\omega}]^5 = [\hat{\omega}], \quad [\hat{\omega}]^6 = [\hat{\omega}]^2, \quad \dots \quad (4.21)$$

4.7.2 Taylor Expansion and Regrouping

The matrix exponential is defined by its Taylor series:

$$e^{[\hat{\omega}]\theta} = \mathbf{I} + \theta[\hat{\omega}] + \frac{\theta^2}{2!}[\hat{\omega}]^2 + \frac{\theta^3}{3!}[\hat{\omega}]^3 + \frac{\theta^4}{4!}[\hat{\omega}]^4 + \dots \quad (4.22)$$

Substituting $[\hat{\omega}]^3 = -[\hat{\omega}]$, $[\hat{\omega}]^4 = -[\hat{\omega}]^2$, $[\hat{\omega}]^5 = [\hat{\omega}]$, etc., and grouping by powers of $[\hat{\omega}]$:

$$e^{[\hat{\omega}]\theta} = \mathbf{I} + \underbrace{\left(\theta - \frac{\theta^3}{3!} + \frac{\theta^5}{5!} - \dots\right)}_{\sin \theta} [\hat{\omega}] + \underbrace{\left(\frac{\theta^2}{2!} - \frac{\theta^4}{4!} + \frac{\theta^6}{6!} - \dots\right)}_{1 - \cos \theta} [\hat{\omega}]^2. \quad (4.23)$$

The odd-power series is exactly $\sin \theta$ and the even-power series is exactly $1 - \cos \theta$.

Theorem 4.9. Rodrigues' Rotation Formula

For a unit axis $\hat{\omega} \in \mathbb{R}^3$ and angle $\theta \in \mathbb{R}$, the rotation matrix is:

$$e^{[\hat{\omega}]\theta} = \mathbf{I} + \sin \theta [\hat{\omega}] + (1 - \cos \theta) [\hat{\omega}]^2. \quad (4.24)$$

Equivalently, using (4.19), the action on a vector $\mathbf{v}$ is:

$$\mathbf{R}\mathbf{v} = \mathbf{v} \cos \theta + (\hat{\omega} \times \mathbf{v}) \sin \theta + \hat{\omega}(\hat{\omega} \cdot \mathbf{v})(1 - \cos \theta). \quad (4.25)$$

The vector form (4.25) has a geometric interpretation: decompose $\mathbf{v}$ into a component along $\hat{\omega}$ (unchanged by rotation) and a component perpendicular to $\hat{\omega}$ (rotated by θ in the plane normal to $\hat{\omega}$). See [35] for the original derivation and [31] for the modern Lie-theoretic perspective.

4.8 Quaternion Algebra

4.8.1 Definition and Representation

A **quaternion** $\mathbf{q} \in \mathbb{H}$ is a 4-tuple written as:

$$\mathbf{q} = \begin{bmatrix} q_w \\ q_x \\ q_y \\ q_z \end{bmatrix} = q_w + q_x \mathbf{i} + q_y \mathbf{j} + q_z \mathbf{k} \quad (4.26)$$

where $\mathbf{i}, \mathbf{j}, \mathbf{k}$ are the three *imaginary units* satisfying:

$$\mathbf{i}^2 = \mathbf{j}^2 = \mathbf{k}^2 = \mathbf{ijk} = -1, \quad \mathbf{ij} = \mathbf{k}, \quad \mathbf{jk} = \mathbf{i}, \quad \mathbf{ki} = \mathbf{j}. \quad (4.27)$$

Note that quaternion multiplication is *non-commutative*: $\mathbf{ij} = \mathbf{k}$ but $\mathbf{ji} = -\mathbf{k}$.

4.8.2 Quaternion Operations

Addition

Quaternion addition is component-wise:

$$(q_w + q_x \mathbf{i} + q_y \mathbf{j} + q_z \mathbf{k}) + (r_w + r_x \mathbf{i} + r_y \mathbf{j} + r_z \mathbf{k}) = (q_w + r_w) + (q_x + r_x) \mathbf{i} + (q_y + r_y) \mathbf{j} + (q_z + r_z) \mathbf{k}. \quad (4.28)$$

Hamilton Product (Multiplication)

The product of two quaternions $\mathbf{q} = q_w + \mathbf{q}_v$ and $\mathbf{r} = r_w + \mathbf{r}_v$, where $\mathbf{q}_v = [q_x, q_y, q_z]^T$ and $\mathbf{r}_v = [r_x, r_y, r_z]^T$, is:

$$\mathbf{q} \otimes \mathbf{r} = \begin{bmatrix} q_w r_w - \mathbf{q}_v^T \mathbf{r}_v \\ q_w \mathbf{r}_v + r_w \mathbf{q}_v + \mathbf{q}_v \times \mathbf{r}_v \end{bmatrix} \quad (4.29)$$

where $\times$ denotes the 3D cross product. In component form:

$$\mathbf{q} \otimes \mathbf{r} = \begin{bmatrix} q_w r_w - q_x r_x - q_y r_y - q_z r_z \\ q_w r_x + q_x r_w + q_y r_z - q_z r_y \\ q_w r_y - q_x r_z + q_y r_w + q_z r_x \\ q_w r_z + q_x r_y - q_y r_x + q_z r_w \end{bmatrix} \quad (4.30)$$

Conjugate

The **conjugate** of $\mathbf{q} = q_w + q_x \mathbf{i} + q_y \mathbf{j} + q_z \mathbf{k}$ is:

$$\mathbf{q}^* = q_w - q_x \mathbf{i} - q_y \mathbf{j} - q_z \mathbf{k} = \begin{bmatrix} q_w \\ -q_x \\ -q_y \\ -q_z \end{bmatrix} \quad (4.31)$$

Norm and Magnitude

The **norm** (or magnitude) of a quaternion is:

$$\|\mathbf{q}\| = \sqrt{q_w^2 + q_x^2 + q_y^2 + q_z^2} \quad (4.32)$$

A **unit quaternion** satisfies $\|\mathbf{q}\| = 1$.

Inverse

For a unit quaternion $\mathbf{q}$, the inverse is:

$$\mathbf{q}^{-1} = \mathbf{q}^* \quad (4.33)$$

since $\mathbf{q} \otimes \mathbf{q}^* = \mathbf{q}^* \otimes \mathbf{q} = 1$ (the multiplicative identity quaternion).

For a non-unit quaternion, the inverse is $\mathbf{q}^{-1} = \mathbf{q}^* / \|\mathbf{q}\|^2$.

4.8.3 Unit Quaternions and Rotation

Theorem 4.10. Unit Quaternions Represent Rotations

Every unit quaternion $\mathbf{q}$ can be written as:

$$\mathbf{q} = \begin{bmatrix} \cos(\theta/2) \\ \hat{\omega} \sin(\theta/2) \end{bmatrix} = \begin{bmatrix} \cos(\theta/2) \\ \hat{\omega}_x \sin(\theta/2) \\ \hat{\omega}_y \sin(\theta/2) \\ \hat{\omega}_z \sin(\theta/2) \end{bmatrix} \quad (4.34)$$

where $\hat{\omega}$ is a unit vector (the rotation axis) and $\theta \in [0, 2\pi)$ is the rotation angle. This quaternion represents a rotation of angle θ about the axis $\hat{\omega}$.

In Plain Language 4.4. Why the Half-Angles?

You might notice $\theta/2$ in the formula instead of just θ . This is because applying a quaternion rotation to a vector $\mathbf{v}$ requires multiplying it from the front *and* the back ($\mathbf{q} \otimes \mathbf{v} \otimes \mathbf{q}^{-1}$, where $\mathbf{v}$ is viewed as the quaternion $[0, v_x, v_y, v_z]^T$). The half-angles combine during this double-multiplication to give the perfect full-angle θ rotation! Note also that the quaternions $\mathbf{q}$ and $-\mathbf{q}$ represent the exact same physical rotation in 3D space.

4.8.4 Double Cover: q and $-q$ Represent the Same Rotation**Theorem 4.11. Quaternion Double Cover**

The unit quaternions $\mathbf{q}$ and $-\mathbf{q}$ represent the same rotation in 3D space. Thus, the map from unit quaternions to $\text{SO}(3)(3)$ is a 2-to-1 covering map, and we say that S^3 (the unit sphere in $\mathbb{H}$) is a *double cover* of $\text{SO}(3)(3)$.

Proof. The rotation action of a unit quaternion $\mathbf{q}$ on a pure-imaginary quaternion $\mathbf{v} = [0, \vec{v}]$ is:

$$\mathbf{v}' = \mathbf{q} \otimes \mathbf{v} \otimes \mathbf{q}^*. \quad (4.35)$$

For $-\mathbf{q}$, we compute:

$$(-\mathbf{q}) \otimes \mathbf{v} \otimes (-\mathbf{q})^* = (-\mathbf{q}) \otimes \mathbf{v} \otimes (-\mathbf{q}^*) \quad (4.36)$$

$$= (-1)(\mathbf{q}) \otimes \mathbf{v} \otimes (-1)(\mathbf{q}^*) \quad (4.37)$$

$$= (-1)(-1) \mathbf{q} \otimes \mathbf{v} \otimes \mathbf{q}^* \quad (4.38)$$

$$= \mathbf{q} \otimes \mathbf{v} \otimes \mathbf{q}^*, \quad (4.39)$$

where we used the fact that $(-\mathbf{q})^* = -\mathbf{q}^*$ and that the scalar factor (-1) commutes with quaternion multiplication and appears twice, yielding $(-1)^2 = 1$.

Thus $\mathbf{q}$ and $-\mathbf{q}$ induce the identical rotation on every vector $\vec{v} \in \mathbb{R}^3$. $\square$

4.9 Conversion Between Quaternions and Rotation Matrices

4.9.1 Quaternion to Rotation Matrix

Theorem 4.12. Quaternion-to-Matrix Conversion

A unit quaternion $\mathbf{q} = [q_w, q_x, q_y, q_z]^T$ with $q_w^2 + q_x^2 + q_y^2 + q_z^2 = 1$ corresponds to the rotation matrix:

$$\mathbf{R}(\mathbf{q}) = \begin{bmatrix} 1 - 2(q_y^2 + q_z^2) & 2(q_x q_y - q_w q_z) & 2(q_x q_z + q_w q_y) \\ 2(q_x q_y + q_w q_z) & 1 - 2(q_x^2 + q_z^2) & 2(q_y q_z - q_w q_x) \\ 2(q_x q_z - q_w q_y) & 2(q_y q_z + q_w q_x) & 1 - 2(q_x^2 + q_y^2) \end{bmatrix} \quad (4.40)$$

Derivation. A rotation by angle θ about unit axis $\hat{\omega} = [\omega_x, \omega_y, \omega_z]^T$ can be expressed via the Rodrigues formula [35]:

$$\mathbf{R} = \mathbf{I} + \sin \theta [\hat{\omega}]_{\times} + (1 - \cos \theta) [\hat{\omega}]_{\times}^2. \quad (4.41)$$

With $q_w = \cos(\theta/2)$ and $\hat{\omega} \sin(\theta/2) = [q_x, q_y, q_z]^T$, one can expand using the identities $\sin \theta = 2 \sin(\theta/2) \cos(\theta/2) = 2q_w \sqrt{q_x^2 + q_y^2 + q_z^2}$ and $(1 - \cos \theta) = 2 \sin^2(\theta/2) = 2(q_x^2 + q_y^2 + q_z^2)$ to obtain equation (4.40). $\square$

4.9.2 Rotation Matrix to Quaternion

Theorem 4.13. Matrix-to-Quaternion Conversion

From a rotation matrix $\mathbf{R}$, the unit quaternion can be recovered (up to sign) via:

$$q_w = \frac{1}{2} \sqrt{1 + R_{11} + R_{22} + R_{33}}, \quad (4.42)$$

with q_x, q_y, q_z determined by:

$$\begin{aligned} q_x &= \frac{R_{32} - R_{23}}{4q_w}, \\ q_y &= \frac{R_{13} - R_{31}}{4q_w}, \\ q_z &= \frac{R_{21} - R_{12}}{4q_w}. \end{aligned} \quad (4.43)$$

For numerical stability, a more robust algorithm selects the largest diagonal element to compute first, then back-solves for the others.

4.9.3 Rotation Representation Comparison

Having introduced four rotation representations, we now compare their properties systematically. The choice of representation has significant practical consequences for storage,

computation, and numerical stability.

Table 4.1: Comparison of Rotation Representations in $\mathbb{R}^3$.

Property	Rotation Matrix	Euler Angles	Axis-Angle	Quaternion
Parameters	9	3	3	4
Constraints	6 (orthonormality)	0	0	1 (unit norm)
Singularities	None	Gimbal lock	$\theta = \pi$	None
Composition	Matrix multiply	Complex	Via exp map	Quaternion multiply
Interpolation	Difficult	Nonuniform	Via exp map	SLERP
Storage cost	High	Low	Low	Low
Unique	Yes	No (periodic)	No ($\theta > \pi$)	No ($\pm \mathbf{q}$)

When to use each representation:

- **Rotation matrices** are the natural choice for *composing* transformations (chain of rigid-body transforms) and for direct application to vectors. Their 9-parameter redundancy is the cost of a singularity-free, algebraically closed representation.
- **Euler angles** remain useful for *human interpretation* (roll, pitch, yaw) and low-DOF joints, but gimbal lock (Section 4.5.4) makes them unsuitable for general-purpose computation.
- **Axis-angle** (equivalently, exponential coordinates) is the natural parameterisation of the Lie algebra $\mathfrak{so}(3)(3)$ and is essential for optimisation on $\text{SO}(3)(3)$, perturbation analysis, and control.
- **Quaternions** offer the best trade-off for *interpolation* (via SLERP, Section 6.11), *storage* (4 parameters, 1 constraint), and *numerical stability* (no singularities, cheap renormalisation). Modern simulation frameworks—including MuJoCo, Drake, and Bullet—typically store orientations as unit quaternions internally and convert to rotation matrices only when matrix operations are required.

4.10 Spherical Linear Interpolation (SLERP)

In computer graphics and animation, we often need to smoothly interpolate between two rotations. For quaternions, the natural interpolation is **Spherical Linear Interpolation (SLERP)**.

Definition 4.4. SLERP Formula

The SLERP between two unit quaternions $\mathbf{q}_0$ and $\mathbf{q}_1$ at parameter $t \in [0, 1]$ is:

$$\text{SLERP}(\mathbf{q}_0, \mathbf{q}_1; t) = \frac{\sin((1-t)\Omega)}{\sin \Omega} \mathbf{q}_0 + \frac{\sin(t\Omega)}{\sin \Omega} \mathbf{q}_1 \quad (4.44)$$

where $\Omega = \arccos(|\mathbf{q}_0 \cdot \mathbf{q}_1|)$ is the (acute) angle between the quaternions on the unit sphere.

Derivation. Two unit quaternions $\mathbf{q}_0$ and $\mathbf{q}_1$ lie on the unit sphere $S^3 \subset \mathbb{H}$. The angle between them is $\Omega = \arccos(\mathbf{q}_0 \cdot \mathbf{q}_1)$ (taking absolute value ensures we pick the shorter arc).

A geodesic (great circle arc) on S^3 is the shortest path between two points. Parameterizing the geodesic by arc length with $t \in [0, 1]$ (where $t = 0$ at $\mathbf{q}_0$ and $t = 1$ at $\mathbf{q}_1$) gives the SLERP formula.

If $\Omega \approx 0$ (very close quaternions), use linear interpolation $\text{LERP}(\mathbf{q}_0, \mathbf{q}_1; t) = (1 - t)\mathbf{q}_0 + t\mathbf{q}_1$ followed by normalization to avoid numerical division by zero. $\square$

SLERP is constant-speed interpolation and is the standard for smooth rotational animations in robotics and graphics.

4.11 Angular Velocity and the Matrix Exponential

4.11.1 Angular Velocity in the Spatial Frame

When a rigid body rotates, its instantaneous rotational motion is characterized by the **angular velocity** vector $\boldsymbol{\omega} \in \mathbb{R}^3$. The magnitude $\|\boldsymbol{\omega}\|$ is the rate of rotation (rad/s), and the direction of $\boldsymbol{\omega}$ is the instantaneous axis of rotation.

The rate of change of the rotation matrix $\mathbf{R}(t)$ with respect to time is related to angular velocity by:

Theorem 4.14. Matrix Derivative in Spatial Frame

$$\dot{\mathbf{R}}(t) = [\boldsymbol{\omega}]_{\times} \mathbf{R}(t) \quad (4.45)$$

where $\boldsymbol{\omega}$ is expressed in the spatial frame and $[\boldsymbol{\omega}]_{\times}$ is the skew-symmetric matrix defined in equation (4.16).

Proof. Differentiating the orthogonality constraint $\mathbf{R}^T(t)\mathbf{R}(t) = \mathbf{I}$ with respect to time:

$$\frac{d}{dt}[\mathbf{R}^T \mathbf{R}] = \dot{\mathbf{R}}^T \mathbf{R} + \mathbf{R}^T \dot{\mathbf{R}} = \mathbf{0}. \quad (4.46)$$

Rearranging:

$$\mathbf{R}^T \dot{\mathbf{R}} = -(\mathbf{R}^T \dot{\mathbf{R}})^T, \quad (4.47)$$

so $\mathbf{R}^T \dot{\mathbf{R}}$ is skew-symmetric. Define $[\boldsymbol{\omega}]_{\times} := \mathbf{R}^T \dot{\mathbf{R}}$. Then:

$$\dot{\mathbf{R}} = \mathbf{R}[\boldsymbol{\omega}]_{\times}. \quad (4.48)$$

However, by convention in robotics, we often write $\dot{\mathbf{R}} = [\boldsymbol{\omega}]_{\times} \mathbf{R}$ where $\boldsymbol{\omega}$ is the spatial frame angular velocity. Both are correct; the difference is in the choice of frame. $\square$

4.11.2 Angular Velocity in the Body Frame

When the angular velocity is expressed in the body-fixed frame, the differential equation becomes:

$$\dot{\mathbf{R}}(t) = \mathbf{R}(t)[\boldsymbol{\omega}_b]_{\times} \quad (4.49)$$

where $\boldsymbol{\omega}_b$ is the angular velocity in the body frame.

4.11.3 Proof that Angular Velocity Lives in $\mathfrak{so}(3)(3)$

Theorem 4.15. Angular Velocity is a Lie Algebra Element

For any rotation matrix $\mathbf{R}(t) \in \text{SO}(3)(3)$, the quantity $\mathbf{R}^T \dot{\mathbf{R}} \in \mathfrak{so}(3)(3)$ is a skew-symmetric matrix, i.e., an element of the Lie algebra $\mathfrak{so}(3)(3)$.

Proof. From $\mathbf{R}^T \mathbf{R} = \mathbf{I}$, differentiate:

$$\dot{\mathbf{R}}^T \mathbf{R} + \mathbf{R}^T \dot{\mathbf{R}} = \mathbf{0}. \quad (4.50)$$

Taking the transpose of the second term:

$$(\mathbf{R}^T \dot{\mathbf{R}})^T = \dot{\mathbf{R}}^T \mathbf{R}. \quad (4.51)$$

From the differentiated constraint, $\dot{\mathbf{R}}^T \mathbf{R} = -\mathbf{R}^T \dot{\mathbf{R}}$, so:

$$(\mathbf{R}^T \dot{\mathbf{R}})^T = -\mathbf{R}^T \dot{\mathbf{R}}. \quad (4.52)$$

Thus $\mathbf{R}^T \dot{\mathbf{R}}$ is skew-symmetric, confirming it belongs to $\mathfrak{so}(3)(3)$. $\square$

4.12 Axis-Angle Representation and Euler's Rotation Theorem

Theorem 4.16. Euler's Rotation Theorem

Every rotation in 3D can be expressed as a rotation about a single axis by some angle. That is, for any $\mathbf{R} \in \text{SO}(3)(3)$, there exist a unit vector $\hat{\omega} \in \mathbb{R}^3$ (the rotation axis) and an angle $\theta \in [0, \pi]$ such that the rotation is completely described by this pair $(\hat{\omega}, \theta)$.

Proof. The rotation axis $\hat{\omega}$ is an eigenvector of $\mathbf{R}$ with eigenvalue 1:

$$\mathbf{R}\hat{\omega} = \hat{\omega}. \quad (4.53)$$

Such an eigenvector exists because $\det(\mathbf{R} - \mathbf{I}) = 0$ for any rotation (a rotation always leaves one axis invariant). The angle θ can be recovered from $\text{trace}(\mathbf{R}) = 1 + 2 \cos \theta$. $\square$

The axis-angle representation is compact and intuitive but computationally less efficient than quaternions for repeated operations.

4.13 Worked Example: Quaternions and SLERP in Python

Below is a complete Python implementation of quaternion operations and rotation interpolation.

```

1  """Quaternion operations and SLERP, as developed in Volume 0 Chapter
2      4.
3
4  Self-contained by design: this file is typeset into the chapter as a
5  worked
6  example, so it depends only on NumPy and repeats nothing the chapter
7  has not
8  already derived. Equation numbers in the comments refer to Chapter 4.
9
10 Correctness is pinned by tests/test_affine_control/
11     test_quaternion_demo.py --
12 the properties asserted there (unit norm, orthogonality of the
13 rotation matrix,
14 agreement between quaternion rotation and matrix rotation, SLERP
15 endpoints and
16 constant angular velocity) are what make this an example rather than a
17 claim.
18 """
19
20 from __future__ import annotations
21
22 import numpy as np
23 from numpy.typing import NDArray
24
25 type Array = NDArray[np.float64]
26
27 # Below this angle the small-angle expansion of SLERP is better
28 # conditioned
29 # than the closed form: sin(0omega) appears in a denominator, so the
30 # closed form
31 # loses precision as the two quaternions approach each other.
32 SLERP_LINEAR_THRESHOLD = 1.0 - 1.0e-9
33
34 def hamilton_product(p: Array, q: Array) -> Array:
35     """Quaternion product, equation (eq:hamilton_product).
36
37     Both arguments are ordered [w, x, y, z]. The product is not
38     commutative;
39     ‘‘hamilton_product(p, q)’‘ applies ‘‘q’’ first, then ‘‘p’’.
40     """
41     pw, px, py, pz = p
42     qw, qx, qy, qz = q
43     return np.array(
44         [
45             pw * qw - px * qx - py * qy - pz * qz,
46             pw * qx + px * qw + py * qz - pz * qy,
47             pw * qy - px * qz + py * qw + pz * qx,
48             pw * qz + px * qy - py * qx + pz * qw,
49         ]
50     )

```

```

43
44 def normalize(q: Array) -> Array:
45     """Project onto the unit sphere, which is where rotations live."""
46     norm = float(np.linalg.norm(q))
47     if norm == 0.0:
48         msg = "the zero quaternion represents no rotation and cannot
49             be normalised"
50         raise ValueError(msg)
51     return q / norm
52
53 def conjugate(q: Array) -> Array:
54     """For a unit quaternion this is also the inverse."""
55     w, x, y, z = q
56     return np.array([w, -x, -y, -z])
57
58
59 def from_axis_angle(axis: Array, angle: float) -> Array:
60     """Unit quaternion for a rotation of 'angle' about 'axis'.S^3 twice, so  $q$  and  $-q$ 
64     are the same rotation.
65     """
66     unit = normalize(np.asarray(axis, dtype=float))
67     return np.concatenate([[np.cos(angle / 2.0)], np.sin(angle / 2.0)
68         * unit])
69
70 def to_matrix(q: Array) -> Array:
71     """Rotation matrix of a unit quaternion, equation (eq:
72     quat_to_matrix)."""
73     w, x, y, z = normalize(q)
74     return np.array(
75         [
76             [1 - 2 * (y * y + z * z), 2 * (x * y - w * z), 2 * (x * z
77             + w * y)],
78             [2 * (x * y + w * z), 1 - 2 * (x * x + z * z), 2 * (y * z
79             - w * x)],
80             [2 * (x * z - w * y), 2 * (y * z + w * x), 1 - 2 * (x * x
81             + y * y)],
82             ]
83     )
84
85 def rotate(q: Array, v: Array) -> Array:
86     """Rotate a vector by the sandwich product  $q v q^*$ , without forming
87     a matrix."""
88     pure = np.concatenate([[0.0], np.asarray(v, dtype=float)])
89     return hamilton_product(hamilton_product(q, pure), conjugate(q))
90     [1:]

```

```

86
87 def slerp(q0: Array, q1: Array, t: float) -> Array:
88     """Spherical linear interpolation, equation (eq:slerp).
89
90     Two details matter and neither is cosmetic:
91
92     * If the dot product is negative the interpolation would take the
93     long way
94     round the sphere -- 360 degrees minus the short arc. Negating
95     one input
96     selects the short arc, which is legitimate because q and -q
97     denote the
98     same rotation.
99     * As the quaternions approach each other sin(0omega) approaches
100     zero and the
101     closed form loses precision, so below a threshold we fall back
102     to
103     normalised linear interpolation, which agrees to first order.
104     """
105     a = normalize(np.asarray(q0, dtype=float))
106     b = normalize(np.asarray(q1, dtype=float))
107
108     dot = float(np.dot(a, b))
109     if dot < 0.0:
110         b, dot = -b, -dot
111
112     if dot > SLERP_LINEAR_THRESHOLD:
113         return normalize(a + t * (b - a))
114
115     omega = float(np.arccos(np.clip(dot, -1.0, 1.0)))
116     sin_omega = np.sin(omega)
117     return (np.sin((1.0 - t) * omega) / sin_omega) * a + (np.sin(t *
118     omega) / sin_omega) * b
119
120
121 def interpolate_point_cloud(points: Array, q0: Array, q1: Array,
122 frames: int) -> list[Array]:
123     """Carry a point cloud through the interpolated frames.
124
125     Returns one rotated copy of ‘points’ per frame, which is what
126     the
127     chapter’s visualisation plots.
128     """
129     return [
130         (to_matrix(slerp(q0, q1, t)) @ np.asarray(points, dtype=float)
131         .T).T
132         for t in np.linspace(0.0, 1.0, frames)
133     ]

```

Key points from the code:

1. **Quaternion representation:** We store $\mathbf{q} = [q_w, q_x, q_y, q_z]$ as a NumPy array.

2. **Hamilton product:** Custom function implements the quaternion multiplication rule (equation (4.29)).
3. **Normalization:** Unit quaternions are enforced by dividing by norm.
4. **Quaternion-to-matrix:** Uses equation (4.40).
5. **SLERP:** Implements equation (4.44) with fallback to linear interpolation for nearby quaternions.
6. **Visualization:** The example rotates a point cloud through interpolated frames.

4.14 Chapter Summary

This chapter has developed the complete mathematical framework for 3D rotations:

- **Rotation Matrices ($\text{SO}(3)(3)$):** We defined the Special Orthogonal Group as the set of 3×3 orthogonal matrices with determinant $+1$, proved it is both a group and a 3-dimensional smooth manifold, and derived the three elementary rotations $\mathbf{R}_x(\theta)$, $\mathbf{R}_y(\theta)$, $\mathbf{R}_z(\theta)$.
- **Euler Angles and Gimbal Lock:** We presented the ZYX and ZXZ conventions and rigorously proved that the Jacobian relating angular velocities to Euler angle rates becomes singular at $\theta = \pm 90^\circ$, making Euler angles unsuitable for continuous control.
- **Lie Groups and Algebras:** We identified $\text{SO}(3)(3)$ as a Lie Group (smooth manifold + group structure) and its Lie algebra $\mathfrak{so}(3)(3)$ as the space of skew-symmetric matrices, which represent angular velocities.
- **Quaternion Algebra:** We introduced the quaternion representation of rotations, defined all quaternion operations (addition, Hamilton product, conjugate, norm, inverse), and proved that $\mathbf{q}$ and $-\mathbf{q}$ represent the same rotation (the double cover property).
- **Conversions:** We derived explicit formulas for converting between quaternions and rotation matrices in both directions.
- **SLERP:** We presented spherical linear interpolation as the natural way to smoothly interpolate between rotations.
- **Angular Velocity:** We connected the time-derivative of rotation matrices to angular velocity vectors, showing that $\dot{\mathbf{R}} = [\boldsymbol{\omega}]_{\times} \mathbf{R}$ and that $\mathbf{R}^T \dot{\mathbf{R}} \in \mathfrak{so}(3)(3)$.
- **Euler's Theorem:** We stated and sketched the proof that every 3D rotation can be expressed as a rotation about a single axis by a single angle.

These tools eliminate the singularities of Euler angles while providing efficient computational representations. The next chapter combines rotation matrices with translation vectors into the full rigid body transformation group $\text{SE}(3)(3)$.

4.15 Further Reading

The rotation group $SO(3)(3)$ and its extension to rigid-body transformations $SE(3)(3)$ are treated in depth by Murray, Li, and Sastry [31], whose Chapter 2 develops the group-theoretic and differential-geometric foundations that underpin this chapter. Lynch and Park [27], Chapters 3–4, provide a complementary modern treatment with numerous worked examples of rotation matrices, homogeneous transformation matrices, and their application to serial kinematic chains.

The Lie group structure of $SO(3)(3)$ and its Lie algebra $\mathfrak{so}(3)(3)$ are part of a much larger mathematical framework. Hall [17] offers a self-contained introduction to Lie groups and Lie algebras accessible to engineers with the linear algebra background developed in Chapter 1 of this text. For the geometric and symplectic perspective on rotation groups and their role in mechanics, Marsden and Ratiu [29] provide the definitive reference.

The Rodrigues rotation formula derived in this chapter has a distinguished history: it was first published by Olinde Rodrigues in 1840 [35], predating the modern matrix formalism by nearly a century. Rodrigues' original derivation used purely geometric reasoning; the matrix exponential form $e^{[\omega]_{\times}\theta}$ that we derive in Chapter 6 provides the modern algebraic counterpart.

Finally, readers implementing rotation representations in software should be aware that quaternion arithmetic is the *de facto* standard in modern robotics and game-engine frameworks (e.g., Eigen, ROS tf2, Unity) precisely because it avoids the gimbal-lock singularity and requires only four parameters with a single unit-norm constraint. The SLERP algorithm presented in this chapter is the core interpolation primitive in all of these libraries.

Exercises 4.1. Rotations, Quaternions, and SO(3)

4.15.1 Foundation Tier

Exercise. Verify that $\mathbf{R}_z(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$ satisfies $\mathbf{R}_z^T \mathbf{R}_z = \mathbf{I}$ and $\det(\mathbf{R}_z) = 1$. Interpret each column as a rotated basis vector.

Exercise. Show that the composition of two rotations $\mathbf{R}_1$ and $\mathbf{R}_2$ (applied as $\mathbf{R}_1 \mathbf{R}_2 \mathbf{v}$) is itself a rotation by verifying closure in $SO(3)(3)$.

Exercise. Compute the full ZYX Euler angle rotation matrix for $\phi = 45^\circ$, $\theta = 30^\circ$, $\psi = 60^\circ$ by multiplying $\mathbf{R}_x(\phi)\mathbf{R}_y(\theta)\mathbf{R}_z(\psi)$. Verify orthogonality numerically.

Exercise. At $\theta = 90^\circ$ in ZYX Euler angles, show that $\dot{\phi}$ and $\dot{\psi}$ are indistinguishable in their effect on the body frame by examining which rows of the Jacobian become parallel.

Exercise. For $\omega = [1, 2, 3]^T$, construct the skew-symmetric matrix $[\omega]_{\times}$ and verify that it is skew-symmetric by checking $[\omega]_{\times}^T = -[\omega]_{\times}$.

4.15.2 Intermediate Tier

Exercise. Compute the Hamilton product $\mathbf{q}_1 \otimes \mathbf{q}_2$ for $\mathbf{q}_1 = [0.7071, 0.7071, 0, 0]^T$ and $\mathbf{q}_2 = [0.7071, 0, 0.7071, 0]^T$ by hand and verify that the result is unit norm.

Exercise. Show that $\mathbf{q} = [0.7071, 0.7071, 0, 0]^T$ and $-\mathbf{q} = [-0.7071, -0.7071, 0, 0]^T$ represent the same rotation by computing the axis and angle for each and verifying they are identical (modulo 2π).

Exercise. Convert the quaternion $\mathbf{q} = [\cos(45^\circ/2), \sin(45^\circ/2), 0, 0]^T$ (rotation about the x -axis by 45°) to a rotation matrix using equation (4.40) and verify it matches $\mathbf{R}_x(45^\circ)$.

Exercise. Interpolate between $\mathbf{q}_0 = [1, 0, 0, 0]^T$ (identity) and $\mathbf{q}_1 = [\cos(45^\circ/2), \sin(45^\circ/2), 0, 0]^T$ at $t = 0.5$ using SLERP and verify the result represents a 22.5° rotation about the x -axis.

Exercise. If $\mathbf{R}(t) = \mathbf{R}_z(t)$ (rotation about z -axis with unit angular velocity), compute $\dot{\mathbf{R}}$ and verify that $[\boldsymbol{\omega}]_\times = \mathbf{R}^T \dot{\mathbf{R}}$ with $\boldsymbol{\omega} = [0, 0, 1]^T$.

Exercise. From the rotation matrix $\mathbf{R}_x(60^\circ)$, extract the rotation axis (should be $[1, 0, 0]^T$) and angle (should be 60°) using the eigenvector method and $\text{trace}(\mathbf{R})$ formula.

4.15.3 Advanced Tier

Exercise. Prove rigorously that $\text{SO}(3)$ is a 3-dimensional manifold by counting independent constraints and showing a 3-parameter local chart. Consider both the orthogonality constraints and determinant condition.

Exercise. Show that the set of pure quaternions $\{\mathbf{q} : q_w = 0\}$ forms a 3-dimensional vector space that is isomorphic to $\mathfrak{so}(3)$. Under what operation does this isomorphism hold?

Exercise. For small t , the matrix exponential provides $\mathbf{R}(t) = e^{t[\boldsymbol{\omega}]_\times}$ where $\boldsymbol{\omega} = [0, 0, 1]^T$ (unit rotation about z). Expand the series $e^{t[\boldsymbol{\omega}]_\times} = \mathbf{I} + t[\boldsymbol{\omega}]_\times + \frac{t^2}{2}[\boldsymbol{\omega}]_\times^2 + \dots$ and verify it matches $\mathbf{R}_z(t)$ to second order in t .

Exercise. Derive the Rodrigues rotation formula $\mathbf{R} = \mathbf{I} + \sin \theta [\hat{\boldsymbol{\omega}}]_\times + (1 - \cos \theta) [\hat{\boldsymbol{\omega}}]_\times^2$ by expanding the rotation of an arbitrary vector about an axis, using the decomposition $\mathbf{v} = (\hat{\boldsymbol{\omega}} \cdot \mathbf{v})\hat{\boldsymbol{\omega}} + \mathbf{v}_\perp$.

Exercise. Prove that SLERP produces constant-speed interpolation: the quaternion at parameter t moves along the geodesic at constant angular velocity. Hint: compute the derivative $\frac{d}{dt}\text{SLERP}(\mathbf{q}_0, \mathbf{q}_1; t)$ and show its norm is constant.

Chapter 5

Twists, Wrenches, and the Screw Axis

One rigid displacement can be described by a screw. A moving body generally requires a changing sequence of instantaneous screws. The distinction lets us describe a golf club precisely without assuming that its entire swing follows one fixed axis.

The purpose of this chapter is to connect three questions: where is the body, how is every point on it moving, and how do applied loads do work? A pose answers the first question, a twist the second, and a wrench the third. Their compact notation is useful because it preserves these physical relationships under a change of coordinates. It does not remove the need to specify a body, a reference point, an observer, or a force application point.

We use angular-first twists $V = (\omega, v)^T$ and moment-first wrenches $F = (n, f)^T$. Here ω is angular velocity, v is the linear coordinate of a rigid velocity field, n is moment about the specified origin, and f is resultant force. Angular velocity is measured in rad/s, linear velocity in m/s, moment in N m and force in N. Radians are dimensionless in SI, but retaining rad in pitch units helps distinguish metres per radian from metres per revolution. A six-vector therefore combines components with different physical units; its ordinary Euclidean norm is not automatically a meaningful mechanical magnitude.

5.1 Historical Context: From Chasles to Modern Screw Theory

Classical screw theory describes rotations with accompanying translations, and forces with accompanying moments. Ball's treatise [2] is a historical reference; modern robotics texts express related constructions using the Lie group of rigid transformations and its tangent algebra [31, 27]. Historical priority is not needed for the derivations below. In particular, a finite-displacement theorem must not be used as if it established a fixed axis for an arbitrary time-dependent motion.

Conventions differ across those references. Murray, Li, and Sastry use linear-first twist coordinates (v, ω) and force-first wrench coordinates (f, n) . Lynch and Park use the angular-first and moment-first ordering adopted here. Featherstone's spatial motion and force vectors also put angular velocity and moment first [12]. These are ordering choices, not different mechanics. If P swaps the two three-component blocks, then $V' = PV$, $F' = PF$ and $F'^T V' = F^T V$; a matrix representing a motion map changes to PXP^T . Copying a formula while keeping the wrong vector order changes its meaning.




Figure 5.1: A Constant-Axis Screw Path in Projection. The Point Rotates Through Half a Turn While Rising Along the Axis; the Endpoint Theorem Does Not Require This Particular Intervening Path.

5.2 The Unified Four-by-Four Transform $SE(3)$

Let T_{AB} map coordinates expressed in frame B to coordinates expressed in frame A. Frame B's origin has coordinates p_{AB} in A, and R_{AB} rotates B components into A components. For a point and a free direction, respectively,

$$x_A = R_{AB}x_B + p_{AB}, \quad d_A = R_{AB}d_B. \quad (5.1)$$

Translation is affine on three-dimensional point coordinates: it does not preserve the zero vector. No principle of linear algebra is violated. Homogeneous coordinates embed points with weight one and directions with weight zero:

$$T_{AB} = \begin{bmatrix} R_{AB} & p_{AB} \\ 0 & 1 \end{bmatrix}, \quad \begin{bmatrix} x_A \\ 1 \end{bmatrix} = T_{AB} \begin{bmatrix} x_B \\ 1 \end{bmatrix}, \quad \begin{bmatrix} d_A \\ 0 \end{bmatrix} = T_{AB} \begin{bmatrix} d_B \\ 0 \end{bmatrix}. \quad (5.2)$$

The set $SE(3)$ consists of these matrices with $R^T R = I$ and $\det R = 1$. Reflections, arbitrary scale changes and finite Euler approximations to rotations are not generally in this set. Separate arrays for rotation and translation are equally legitimate implementations; a homogeneous matrix is a representation, not a requirement imposed on every physics engine.

5.2.1 Group Structure of $SE(3)$

Following the frame subscripts gives composition and inversion:

$$T_{AC} = T_{AB}T_{BC}, \quad T^{-1} = \begin{bmatrix} R^T & -R^T p \\ 0 & 1 \end{bmatrix}. \quad (5.3)$$

The product has rotation $R_1 R_2$ and translation $p_1 + R_1 p_2$, so it remains a proper rigid transform. Identity and inverses belong to the set, and associativity follows from matrix multiplication. The parameters form a smooth six-dimensional manifold: three local rotational coordinates and three translations. Multiplication and inversion are smooth, making $SE(3)$ a Lie group. Six-dimensional does not mean that one nonsingular three-angle chart covers all rotations; rotation charts still have their familiar limitations.

Composition order matters. Rotating a club and then translating it by a vector fixed in the laboratory is generally a different operation from translating it and then rotating about the laboratory origin. Subscripts and the point-mapping equation are safer guides than an unexplained instruction to “premultiply” or “postmultiply.”

5.3 Chasles' Theorem and the Screw Axis

Consider a finite displacement $x' = Rx + p$ in one common spatial coordinate system. If R is not identity, this displacement can be realized as a rotation about a line, followed by a translation parallel to that line. Choose a unit direction s and a principal angle $0 < \theta \leq \pi$ with $R = \exp([s]_{\times} \theta)$. There are a point q on the line and an axial displacement d such that

$$x' = q + R(x - q) + ds, \quad p = (I - R)q + ds. \quad (5.4)$$

Every point $q + \lambda s$ on the axis moves to $q + (\lambda + d)s$. Other points have a rotational displacement around the line as well as the axial displacement. The signed finite pitch is $h = d/\theta$, not d and not the total distance travelled by an off-axis point divided by angle.

This is an endpoint statement. If a tracked body has poses T_0 and T_1 relative to the same laboratory frame, the spatial displacement mapping each initial material-point position to its final position is $\Delta T = T_1 T_0^{-1}$. Decomposing T_1 alone instead describes a mapping from the chosen body reference coordinates. It generally gives a different axis. The corresponding body-coordinate relative transform is $T_0^{-1} T_1$; its coordinates must be interpreted accordingly.

5.3.1 Geometric Description and Exceptional Cases

The axis line is $\{q + \lambda s : \lambda \in \mathbb{R}\}$. Replacing q by $q + \lambda s$ does not change the line. We normally choose its closest point to the origin, $q \cdot s = 0$. For nonidentity R , this geometric line is unique. Its oriented parameterization has branch qualifications: at a half turn, s and $-s$ represent the same rotation, while the signed axial coordinate and pitch change with that choice. Additional full turns also produce alternative angle/pitch descriptions of the same endpoints.

For pure translation, $R = I$ and $p \neq 0$, every line parallel to p can serve as a translational axis. There is no unique finite rotational axis or finite pitch obtained by dividing by a zero angle. Calling this “infinite pitch” is a limiting convention, not an instruction to divide by zero. Identity displacement has neither a preferred axis nor a required nonzero motion. An object can even execute a nontrivial closed path and return to identity displacement: its endpoints alone do not reconstruct that path.

A steadily advancing ideal screw provides a useful constant-axis example. A tumbling football or precessing frisbee generally does not. Each instantaneous rigid velocity field and each nontrivial finite relative pose has its own appropriate screw description; those descriptions need not coincide over a measured time interval. This distinction is especially relevant when a changing swing plane is inferred from a sequence of club poses.

5.4 The Spatial Twist

Let $T(t)$ give the moving body’s pose in a fixed laboratory frame. Differentiating T produces a tangent matrix at T , not directly a six-vector at identity. Two useful representations are

$$\hat{V}_s = \dot{T} T^{-1}, \quad \hat{V}_b = T^{-1} \dot{T}, \quad \hat{V} = \begin{bmatrix} [\omega]_{\times} & v \\ 0 & 0 \end{bmatrix}. \quad (5.5)$$

The subscripts denote spatial and body representations of the same body motion relative to the laboratory. With $T = (R, p)$, direct multiplication gives

$$[\omega_s]_{\times} = \dot{R} R^T, \quad v_s = \dot{p} - \omega_s \times p, \quad [\omega_b]_{\times} = R^T \dot{R}, \quad v_b = R^T \dot{p}. \quad (5.6)$$

In particular, v_s is not generally the velocity $\dot{p}$ of the body-frame origin. It is the velocity assigned by the body’s rigid velocity field to the point located at the spatial origin. That point need not lie inside the physical body. The spatial field evaluated at any material point with current position r is

$$u(r) = v_s + \omega_s \times r. \quad (5.7)$$

For $r = p + R x_b$ with fixed material coordinates x_b , this equals $\dot{p} + \dot{R} x_b$, as it must. This identity is an immediate way to check a twist convention against measured marker velocities.

5.4.1 Screw Representation of a Twist

When $\omega \neq 0$, set $s = \omega/\|\omega\|$, choose rate $\dot{\theta} = \|\omega\|$, and define

$$h = \frac{\omega \cdot v}{\|\omega\|^2}, \quad q = \frac{\omega \times v}{\|\omega\|^2}, \quad V = S\dot{\theta}, \quad S = \begin{bmatrix} s \\ q \times s + hs \end{bmatrix}. \quad (5.8)$$

The vector identity $\omega \times (\omega \times v) = \omega(\omega \cdot v) - \|\omega\|^2 v$ verifies the formula. Every point on this instantaneous axis has velocity $h\omega$; off-axis points have an additional rotational contribution. Thus pitch uses the axial projection of velocity, not $\|v\|/\|\omega\|$ and not clubhead speed divided by angular speed. With an independently chosen oriented joint axis, $\dot{\theta}$ may instead be signed.

The phrase “unit screw” here means $\|s\| = 1$ in the angular block for a rotational joint. It does not mean $\|S\| = 1$ under an arbitrary mixed-unit six-dimensional norm. For pure translation, $\omega = 0$ and $v \neq 0$, one can use $S = (0, s)$ with $\|s\| = 1$ and a translational joint coordinate in metres. The zero twist does not determine an instantaneous axis.

5.4.2 Twists for Standard Joint Types

A revolute joint through q in direction s has $S = (s, q \times s)$ and zero pitch. The linear block is not zero merely because points on the axis are instantaneously stationary: their complete velocity is $q \times s + s \times q = 0$ for unit angular rate. A prismatic joint has $S = (0, s)$, and a helical joint has $S = (s, q \times s + hs)$. These are relative joint twists. In a moving mechanism, an absolute link twist also includes the predecessor’s motion, expressed in a common frame.

For an axis through $(1, 0, 0)$ parallel to z , take $h = 2$ m/rad and rate 0.5 rad/s. The twist is $V = (0, 0, 0.5, 0, -0.5, 1)^T$. At the axis point, $u(q) = (0, 0, 1)$ m/s; at the origin, $u(0) = (0, -0.5, 1)$ m/s. Both velocities belong to the same rigid field. A zero-pitch revolute joint at that location would retain the -0.5 linear component but have zero velocity at q .

5.4.3 Matrix Representation and Finite Integration

The matrices $\hat{V}$ form the Lie algebra $\mathfrak{se}(3)$. Only the upper-left rotational block is skew-symmetric; the full four-by-four matrix is generally not. Differentiating $R^T R = I$ explains that skew-symmetry. The translation column is unrestricted and the last row is zero, leaving six independent coordinates.

For a constant spatial twist, $\dot{T} = \hat{V}_s T$ integrates to $T(t) = \exp(\hat{V}_s t) T(0)$. For a constant body twist, $\dot{T} = T \hat{V}_b$ integrates to $T(t) = T(0) \exp(\hat{V}_b t)$. For a general time-dependent twist, successive increments require their correct order. Replacing them by the exponential of the simple time integral is justified only under additional commutation conditions. A single fitted endpoint screw does not show that the intervening twist was constant.

5.4.4 The Lie Bracket of Twists

The matrix commutator induces the angular-first bracket

$$[V_1, V_2] = \begin{bmatrix} \omega_1 \times \omega_2 \\ \omega_1 \times v_2 - \omega_2 \times v_1 \end{bmatrix}. \quad (5.9)$$

It describes the leading second-order effect of changing the order of small transformations. With the stated matrix order,

$$e^{\epsilon \hat{V}_1} e^{\epsilon \hat{V}_2} e^{-\epsilon \hat{V}_1} e^{-\epsilon \hat{V}_2} = I + \epsilon^2 [\hat{V}_1, \hat{V}_2] + O(\epsilon^3). \quad (5.10)$$

The bracket is not simply the relative velocity $V_1 - V_2$. Its role becomes important when rotations at different joints or times do not commute. It provides a local geometric description, not by itself a prediction of a finite golf-swing outcome.

5.5 The Adjoint Representation

5.5.1 Changing Reference Frames

Suppose A and B are two coordinate frames used to represent the same physical velocity field. At the instant of interest, $x_A = Rx_B + p$ with $T_{AB} = (R, p)$. Substituting the point-velocity fields gives

$$\omega_A = R\omega_B, \quad v_A = Rv_B + p \times R\omega_B. \quad (5.11)$$

The angular-first adjoint is therefore

$$V_A = XV_B, \quad X = \text{Ad}_{T_{AB}} = \begin{bmatrix} R & 0 \\ [p] \times R & R \end{bmatrix}. \quad (5.12)$$

Equivalently, $\hat{V}_A = T_{AB} \hat{V}_B T_{AB}^{-1}$. To check the sign of the offset term, evaluate $u_A(Rr_B + p)$: the $p \times R\omega_B$ term cancels the $R\omega_B \times p$ from point evaluation, leaving $Ru_B(r_B)$. Also $\text{Ad}_{T_1 T_2} = \text{Ad}_{T_1} \text{Ad}_{T_2}$ by matrix conjugation. For the moving body's pose, $V_s = \text{Ad}_T V_b$ follows from the two derivative definitions above.

Re-expression must be distinguished from changing the physical observer. If A and B themselves move relative to each other, differentiating point coordinates gives extra frame-motion terms:

$$\dot{x}_A = R\dot{x}_B + \dot{R}x_B + \dot{p}. \quad (5.13)$$

The adjoint re-expresses a specified relative motion; it does not automatically subtract an observer's own velocity. In particular, body representation V_b still describes body motion relative to the laboratory, even though the body's material coordinates are constant. A calculation of motion relative to another moving body must first specify that relative motion and then express it in the chosen frame.

5.6 The Spatial Wrench

A collection of forces f_i applied at positions r_i and free couples c_i , all expressed in one frame, has resultant wrench about that frame's origin

$$F = \begin{bmatrix} n \\ f \end{bmatrix}, \quad f = \sum_i f_i, \quad n = \sum_i (r_i \times f_i + c_i). \quad (5.14)$$

A couple has a moment independent of the chosen origin; the moment of a nonzero resultant force generally changes with origin. Wrenches form the dual space $\mathfrak{se}(3)^*$ under the power pairing. A list of six components is not a reason to identify a wrench with a twist in $\mathfrak{se}(3)$. The distinction fixes the correct transformation law.

5.6.1 Wrenches From Forces and Moments

A force $(f_x, 0, 0)$ applied at $(0, 0, d)$ produces moment $(0, df_x, 0)$, with positive y sign for positive d and f_x . Moving the same force along its line of action leaves its moment unchanged. A force $(1, 0, 0)$ N at $(0, 2, 0)$ m instead gives $n = (0, 0, -2)$ N m. These elementary cross products are useful checks on any six-dimensional formula.

5.6.2 Wrench Transformation Under Changes of Frame

For the same B-to-A transform used for motion,

$$f_A = Rf_B, \quad n_A = Rn_B + p \times Rf_B, \quad F_A = X^{-T}F_B. \quad (5.15)$$

The reverse mapping is $F_B = X^T F_A$. Thus a transpose alone is correct when converting an A wrench to B, while X converts a B twist to A. Applying X and X^T in the same B-to-A direction is generally wrong. One can derive the force rule directly by writing $r_A = Rr_B + p$ and $f_A = Rf_B$ in $r_A \times f_A$, or derive it from $F_A^T V_A = F_B^T V_B$ for every V_B . Both derivations agree.

5.6.3 Physical Interpretation of Wrenches: Forces and Torques Unified

The wrench is a compact account of a body's resultant load, not a specification of how that load is distributed across skin, fingers, shafts or individual contacts. Different grip pressure distributions can have the same net wrench while producing different local stress or friction margins. Likewise, a wrench axis is not generally the instantaneous axis of the body on which it acts. Motion also depends on inertia, initial state and constraints.

For $f \neq 0$, the wrench's central axis and signed pitch are

$$q_F = \frac{f \times n}{\|f\|^2}, \quad h_F = \frac{f \cdot n}{\|f\|^2}, \quad n = q_F \times f + h_F f. \quad (5.16)$$

The cross-product identity verifies this decomposition. On that line the remaining moment is parallel to the resultant force. A pure force has $h_F = 0$, although its moment about a remote origin may be nonzero. A pure couple has $f = 0$ and no unique finite central line; an "infinite-pitch wrench" is again a convention. The zero wrench has no distinguished line. Although both pitches have dimensions of length when radians are suppressed, h_F and the motion pitch h describe different objects and need not be equal.

5.7 Proof of Chasles' Theorem

Here is a constructive proof. A proper rotation with $R \neq I$ has an invariant unit direction s and rotates its perpendicular plane by a nonzero angle modulo 2π . Resolve the displacement and define its pitch as

$$d = s^T p, \quad p_\perp = p - ds, \quad h = d/\theta. \quad (5.17)$$

Because $s^T(I - R) = 0$, any solution of $p = (I - R)q + ds$ must have this axial displacement. Restricted to $s^\perp$, the map $I - R$ is invertible. Choose $q \perp s$ and solve

$$(I - R)q = p_\perp, \quad q = \frac{1}{2}p_\perp + \frac{1}{2}\cot(\theta/2)s \times p_\perp. \quad (5.18)$$

To verify the inverse, put $Jx = s \times x$ on the perpendicular plane, where $J^2 = -I$. Rodrigues' formula there is $R = \cos \theta I + \sin \theta J$. Multiplying $[(1 - \cos \theta)I - \sin \theta J]$ by $\frac{1}{2}[I + \cot(\theta/2)J]$ gives I . This proves the displayed expression, including its cross-product sign. There is no additional cross product outside the restricted inverse.

The resulting q satisfies $p = (I - R)q + ds$, proving the screw representation. Adding any multiple of s gives the same axis line; no other perpendicular point solves the equation. At $\theta = \pi$, the cotangent vanishes and $q = p_\perp/2$, while the direction-sign qualification remains. As $\theta \rightarrow 0$, the inverse has scale $1/|\theta|$. A small measured perpendicular translation error can therefore produce a large axis-location error. A numerically reported distant axis near pure translation need not indicate a meaningful anatomical pivot.

For a rotation by $\pi/2$ about the vertical line through $(1, 0, 0)$ followed by 0.2 m of axial translation, $R = R_z(\pi/2)$ and $p = (1, -1, 0.2)$. The formulas give $q = (1, 0, 0)$, $d = 0.2$ m and $h = 0.4/\pi \approx 0.12732$ m/rad. Substituting a point on and off the axis verifies both endpoints. This finite pitch does not determine elapsed time or the work required to realize the displacement.

5.8 Mechanical Work and the Virtual Power Principle

For loads acting on one rigid body, substitute the velocity field into the sum of physical powers:

$$P = \sum_i \{f_i \cdot [v + \omega \times r_i] + c_i \cdot \omega\} = f \cdot v + n \cdot \omega = F^T V. \quad (5.19)$$

The scalar triple-product identity converts $f_i \cdot (\omega \times r_i)$ into $(r_i \times f_i) \cdot \omega$. This derivation explains why force and moment must share a reference convention with the twist. Evaluating the linear part at the center of mass while leaving the moment about the grip would omit or double-count a term.

For an admissible virtual generalized velocity $\delta \dot{q}$, suppose the relevant relative twist is $J\delta \dot{q}$. Then virtual power is $F^T J\delta \dot{q} = (J^T F)^T \delta \dot{q}$, so the corresponding generalized force is $J^T F$. This is a kinematic duality statement; a dynamic acceleration calculation additionally requires mass, bias terms and constraints. Work over a finite interval is $\int F(t)^T V(t) dt$. A single instantaneous positive, negative or zero power does not determine that integral.

5.8.1 Invariance of Power and Its Limits

For a coordinate re-expression of the same motion and load, $F_A^T V_A = (X^{-T} F_B)^T X V_B = F_B^T V_B$. This is a scalar representation identity, not conservation of energy. Under a Galilean change to an observer translating with constant velocity U , point velocities become $u'_i = u_i - U$ and the applied-force power becomes $P' = P - f \cdot U$. A rotating observer introduces its own rotational terms. Energy balances must include the chosen observer and system boundary consistently; numerical equality under an adjoint change does not erase those physical choices.

5.9 Reciprocal Screws and Constraint Analysis

A wrench F is reciprocal to a twist V when $F^T V = 0$. For a subspace of admissible relative twists $\mathcal{T}$, ideal reaction wrenches lie in its annihilator: every F in that set satisfies $F^T V = 0$

for every $V \in \mathcal{T}$. If $\mathcal{T}$ has dimension d within the six-dimensional relative twist space, its annihilator has dimension $6 - d$. This is a local linear statement at a specified configuration, not an automatic count of whole-mechanism mobility.

For a spatial revolute joint through q with direction s , the condition is

$$s \cdot (n - q \times f) = 0. \quad (5.20)$$

It excludes a reaction moment along the free rotational direction when moments are taken about the joint point. A motor or friction torque in that direction can perform relative work and is not part of the ideal reaction subspace. For a frictionless prismatic joint, $s \cdot f = 0$ excludes reaction force along the sliding direction; other spatial reaction components remain possible.

5.9.1 Two Geometric Screws and the Reciprocal Product

When both screws are represented in the same motion-style ordering, the corresponding reciprocal bilinear form is the cross-block expression

$$B(S_1, S_2) = \omega_1 \cdot v_2 + v_1 \cdot \omega_2. \quad (5.21)$$

This is not their ordinary Euclidean six-vector dot product. A wrench along the geometric screw encoded as (ω_2, v_2) has its force and moment roles interchanged, yielding moment-first coordinates proportional to (v_2, ω_2) . Screw magnitudes and units must be supplied to interpret the resulting pairing as physical power. The geometric reciprocal product alone is not the power of two velocities.

For zero-pitch normalized axes with directions s_1, s_2 through q_1, q_2 , substitution gives

$$B = (q_1 - q_2) \cdot (s_1 \times s_2). \quad (5.22)$$

Thus intersecting, parallel and coincident zero-pitch axes are reciprocal in this geometric sense; perpendicularity is not the defining condition. For finite pitches, add $(h_1 + h_2)s_1 \cdot s_2$. Constraint analysis must use the correct pairing and actual admissible motions, rather than a visual claim that the axes look orthogonal.

5.10 Physical Interpretation of Wrenches

5.10.1 Three Fundamental Wrench Types

A pure couple, a force on a line, and a force plus a parallel couple cover the central-axis cases above. For example, with gravity vector $g = (0, 0, -g_0)$, mass m , and center-of-mass position $r_C = (a, b, c)$, the gravity wrench about the origin is

$$F_g = \begin{bmatrix} -mg_0b \\ mg_0a \\ 0 \\ 0 \\ 0 \\ -mg_0 \end{bmatrix}. \quad (5.23)$$

Its moment vanishes when taken about the center of mass, but generally not about a grip or a joint. Summing gravity wrenches for separate bodies requires a common origin and frame, and summing them does not replace each body's dynamic balance.

5.10.2 The Wrench–Twist Duality at a Golf Grip

For hand forces f_i and couples c_i acting on the club at grip locations r_i , the club-side hand power is

$$P_{\text{hands} \rightarrow \text{club}} = \sum_i [f_i \cdot u_{\text{club}}(r_i) + c_i \cdot \omega_{\text{club}}]. \quad (5.24)$$

The same value follows from summing the hand wrenches about any one chosen origin and pairing with the club twist represented there. This is useful when comparing a clubhead-speed narrative with measured grip loads: an angular moment term and a translational force term must be interpreted together. A change in their separate numerical values after moving the origin is not evidence of a different physical energy-transfer mechanism.

If the grip is rigid and no slip occurs, equal-and-opposite interface wrenches have zero total interface power on the combined hand–club system, because their relative twist is zero. Nevertheless either subsystem can receive nonzero power from the other. With compliance or slip, the summed interface power depends on the relative motion and may represent elastic storage or dissipation. The sign must be tied to which body’s load and relative velocity are used. A measured net club wrench cannot uniquely identify each hand’s contribution or a neural control objective. Those claims require additional measurements and a specified contact and musculoskeletal model.

5.10.3 Transforming Wrenches Between Frames

A practical audit therefore records the sensor orientation, the point about which its moment is reported, the direction of the reported action, and the body-relative motion being paired with it. Transform all simultaneous loads to one frame and origin before adding them. The inverse-transpose rule then does the bookkeeping faithfully. It does not infer an unmeasured grip offset, calibrate a sensor, or separate muscle work from tendon storage.

5.11 Wrench Analysis and Force Polytopes

An ideal bilateral reaction subspace is usually unbounded. A unilateral frictionless point contact instead permits $f = \lambda n_c$ with $\lambda \geq 0$ along a contact normal. With isotropic Coulomb friction, a common admissibility model is $f_n \geq 0$ and $\|f_t\| \leq \mu f_n$. This is a cone; its circular boundary is not a finite-sided polyhedral cone unless approximated. Sliding friction additionally constrains force direction relative to slip velocity. A bounded actuator/contact model can yield a convex polytope, but bounds must actually be supplied. These sets should not be conflated under the word “polytope.”

If admissible contact variables λ produce a resultant wrench $G\lambda$, static support of one free body requires $G\lambda + F_{\text{ext}} = 0$. The support set must contain $-F_{\text{ext}}$, not merely the external load with unchanged sign. Linkages additionally require the internal action–reaction relationships and equilibrium of every link. For moving bodies, a nonzero inertial wrench replaces the zero right-hand side. Neither convexity nor a force polygon removes those balances.

For a concrete bounded example, two vertical bilateral force actuators act on one planar rigid bar at $(0, 0)$ and $(L, 0)$, with $|u_1|, |u_2| \leq F_{\text{max}}$. In coordinates (n_z, f_y) ,

$$(n_z, f_y) = (Lu_2, u_1 + u_2), \quad |n_z/L| \leq F_{\text{max}}, \quad |f_y - n_z/L| \leq F_{\text{max}}. \quad (5.25)$$

For $L = 1$ m and $F_{\max} = 10$ N, the vertices are $(-10, -20)$, $(10, 0)$, $(10, 20)$ and $(-10, 0)$, with moment in Nm and force in N. A downward 15 N load at the midpoint requires $(n_z, f_y) = (7.5, 15)$ and is feasible with $u_1 = u_2 = 7.5$ N. A downward 25 N load there exceeds the total force capacity. If the two supports can only push upward, replace the bounds by $0 \leq u_i \leq F_{\max}$; the feasible set changes. This example concerns two external supports on one body, not two unspecified joints in an arbitrary mechanism.

5.12 Worked Example: Planar Four-Bar Statics

5.12.1 Setup

Use a deliberately idealized planar mechanism with ground points $A = (0, 0)$ and $D = (4, 0)$, moving joints $B = (1, 1)$ and $C = (3, 1)$, all coordinates in metres. Links AB, BC and CD are rigid and massless; all pins are frictionless. A downward load $(0, -W)$ acts at $E = (x_E, 1)$ on the coupler BC. Ground AD is fixed. A motor at A may apply a counterclockwise torque τ_A to AB; B, C and D are passive pins. The numerical example uses $W = 100$ N. Link weight, out-of-plane loading, friction, compliance and acceleration are omitted by assumption.

5.12.2 Twist Space of Each Joint

A planar pin at $q = (q_x, q_y)$ permits the relative angular-first twist

$$S_q = (0, 0, 1, q_y, -q_x, 0)^T. \quad (5.26)$$

Its reduced planar coordinates are $(\omega_z, v_x, v_y) = (1, q_y, -q_x)$. Pairing with reduced wrench (n_z, f_x, f_y) gives

$$n_z + q_y f_x - q_x f_y = 0, \quad n_z = q_x f_y - q_y f_x. \quad (5.27)$$

Thus an ideal planar pin transmits arbitrary in-plane force but no independent couple about the pin. Force becomes directed along a link only if that link is an unloaded two-force member. CD has that property here. AB has it only when $\tau_A = 0$. The planar model does not prove that a physical spatial bearing lacks out-of-plane reactions; it has excluded those components from its idealization.

5.12.3 Constraint Wrenches and Separate Free Bodies

Let $b = (b_x, b_y)$ and $c = (c_x, c_y)$ be forces on the coupler at B and C. The forces on the side links at those joints are $-b$ and $-c$. Since each side link has no other force, the ground reactions on AB and CD are b at A and c at D. Moment balance for CD about D and for BC about B gives

$$c_x + c_y = 0, \quad 2c_y - (x_E - 1)W = 0. \quad (5.28)$$

Coupler force balance and AB moment balance then give

$$b + c + (0, -W) = 0, \quad \tau_A = b_y - b_x. \quad (5.29)$$

The unit lever-arm factors in the last equation are one metre; the preceding factor two is the two-metre coupler span. These equations retain consistent force and moment units. They are balances of three different bodies, not a sum of all joint forces assigned to one coupler.

5.12.4 Equilibrium and the Role of Actuation

Solving gives

$$c_y = \frac{(x_E - 1)W}{2}, \quad c_x = -c_y, \quad b_x = c_y, \quad b_y = W - c_y, \quad \tau_A = W(2 - x_E). \quad (5.30)$$

Here numerical lengths are in metres and τ_A is in N m. For the centered load $x_E = 2$ m, $b = (50, 50)$ N, $c = (-50, 50)$ N and $\tau_A = 0$. Both side links are two-force members. Their lines of force meet at $(2, 2)$, on the external load's vertical line of action. The coupler's force polygon closes, and its moment balance is satisfied. This is passive equilibrium at the specified configuration, not a proof of stability against perturbations.

For $x_E = 2.5$ m, $b = (75, 25)$ N, $c = (-75, 75)$ N and $\tau_A = -50$ N m. AB is no longer a two-force member because of the motor couple. Removing that required motor torque leaves no static equilibrium for this load at this configuration under the stated assumptions. The external forces still sum to zero; that condition alone is insufficient. For the entire assembly, the moment balance about A is $4c_y + \tau_A - Wx_E = 0$, which independently confirms the result.

5.12.5 Geometric Interpretation and Virtual Power

Set the instantaneous angular speed of AB to one rad/s. Then $u_B = (-1, 1)$ m/s. Constant length BC requires $u_{C,x} = u_{B,x}$, while CD's pin motion gives $u_C = (-1, -1)$ m/s. Consequently $\omega_{BC} = -1$ rad/s and $u_{E,y} = 2 - x_E$ m/s. Ideal pin reactions contribute zero net relative power, so static virtual power requires

$$\tau_A - W(2 - x_E) = 0. \quad (5.31)$$

This reproduces the free-body torque. At the centered load, the load point's admissible instantaneous motion is horizontal, so the vertical force does no virtual work. The mechanism still has an admissible motion; load reciprocity does not make it rigid. Near other configurations, the relevant constraint or actuation matrices may lose rank and their force solution may become nonunique or infeasible. Geometry, allowed inputs and the actual external load must all be specified before interpreting a mechanism's apparent force-transmission advantage.

5.13 Python Implementation

The companion implementation checks proper rigid transforms and uses angular-first vectors throughout. It exposes the transform direction in the wrench function name and computes the inverse-transpose action by solving a linear system. Its finite-axis routine accepts resolved nonzero principal rotations; it rejects angles at or below 10^{-8} rad rather than assigning a spurious unique axis to translation. That numerical threshold is an example contract, not a universal measurement-resolution standard. At a half turn the direction follows the rotation library's branch choice, while the recovered geometric line remains the same.

Run the following from the repository root with NumPy and SciPy installed. The implementation is in `src/tools/screw_examples.py`; it reuses the common spatial cross-product helper. The four-bar solver is for the explicitly specified fixed geometry above, not a general mechanism or golf model.

```

1 import numpy as np
2 from scipy.spatial.transform import Rotation
3 from src.tools.screw_examples import (
4     adjoint, finite_screw, four_bar_equilibrium,
5     reexpress_wrench,
6 )
7
8 T = np.eye(4) # Coordinates B -> A
9 T[:3, :3] = Rotation.from_euler("z", np.pi / 2).as_matrix()
10 T[:3, 3] = [1, -1, 0.2]
11 V_B = np.array([0, 0, 0.5, 0, -0.5, 1.0])
12 F_B = np.array([0, 0, -2, 1, 0, 0.0])
13 V_A = adjoint(T) @ V_B
14 F_A = reexpress_wrench(T, F_B)
15 print(F_A @ V_A - F_B @ V_B)
16
17 # Here T is also a specified finite displacement in one frame.
18 screw = finite_screw(T)
19 print(screw.point, screw.pitch)
20 result = four_bar_equilibrium(load_x=2.5, downward_force=100)
21 print(result.force_b, result.force_c, result.torque_a)

```

The example returns zero power difference to floating-point accuracy, the finite axis $q = (1, 0, 0)$ with pitch approximately 0.12732 m/rad, and motor torque -50 Nm for the offset load. The test suite independently compares transformed twists with Cartesian point velocities and transformed wrenches with force moments. It also reconstructs finite point displacements, exercises the half-turn and translation cases, and checks all three link balances and virtual power. These checks validate the displayed identities under their assumptions; they do not validate a biological swing model or remove calibration error from motion-capture data.

5.14 Summary

A pose, a twist and a wrench answer different mechanical questions. Homogeneous transforms compose poses; twists describe the rigid velocity field; wrenches collect resultant force and moment and act on twists through power. Finite screw decomposition describes end-points, while a time-dependent trajectory generally has a changing instantaneous axis. Near pure translation, axis estimation becomes ill-conditioned even when the measured displacement is small and well defined.

For golf analysis, the useful outcome is a consistent chain from measured club poses to point velocities, from measured hand loads to moments about a specified point, and from their pairing to subsystem power and finite-interval work. Reciprocal reactions and actuator loads play different roles in that chain. Compact notation helps reveal those roles, but conclusions about anatomical pivots, force sharing, energy storage or neural objectives require the additional models and evidence appropriate to each claim.

5.15 Further Reading

Murray, Li, and Sastry [31], Chapter 2, develops exponential rigid displacement, spatial/body velocity, central wrench axes and reciprocal screws. Its linear-first ordering must be converted before using its block matrices here. Lynch and Park [27], Sections 3.3.2 and

3.4, provide the angular-first twist and moment-first wrench conventions used in this chapter. The accompanying explanations of body and space representations are especially useful for distinguishing a twist coordinate from a material point's velocity.

Featherstone [12] treats spatial motion/force transformations and multibody dynamics. His published spatial-vector documentation explicitly distinguishes a coordinate transformation from a displacement operator and gives inverse-transpose force transformation. Ball [2] supplies historical background to geometric screw systems. These references support the mathematical framework; the present fixed-geometry linkage calculations and golf interpretation have their own explicitly stated assumptions. No inference about a golfer's control strategy follows solely from a screw theorem.

5.16 Exercises

Unless stated otherwise, use right-handed Cartesian frames, angular-first twists, moment-first wrenches, metres, seconds and newtons. A transform T_{AB} maps B coordinates into A coordinates. Show reference points and signs, not just six numerical entries.

- SE(3) Inverse Computation:** Let $R = R_z(\pi/2)$ and $p = (1, 2, 3)$. Compute T_{AB}^{-1} and verify both products equal identity. Check one mapped point and one free direction to explain the different homogeneous weights.
- Twist From Screw Parameters:** For $q = (1, 0, 0)$, $s = (0, 0, 1)$, $h = 2$ m/rad and $\dot{\theta} = 0.5$ rad/s, compute V . Evaluate the physical velocity at q and at the origin, explaining why the linear block need not equal the axis-point velocity.
- Power Calculation:** For $F = (1, 0, 0, 2, 0, 0)^T$ and $V = (0, 0, 1, 0, 1, 0)^T$, compute $F^T V$ with appropriate component units. Explain why the result does not imply that either the force or motion vanishes, or that the finite-interval work must vanish.
- Adjoint Transformation:** Treat the twist in Exercise 2 as V_B and use T_{AB} from Exercise 1. Construct $\text{Ad}_{T_{AB}}$, compute V_A , and verify the velocity at a transformed axis point directly.
- Reciprocal Wrench:** For the relative revolute twist $S = (0, 0, 1, 1, 0, 0)^T$, identify a point on its axis and characterize all moment-first wrenches with $F^T S = 0$. Give a nonzero example and explain the absence of an independent moment along the free axis at the joint point.
- Bounded Support Set:** Two vertical bilateral actuators act on one planar rigid bar at $(0, 0)$ and $(1, 0)$, each bounded by 10 N in either direction. Derive and sketch the feasible set in (n_z, f_y) about the origin. Test downward midpoint loads of 15 N and 25 N. Explain how unilateral upward-only supports change the set.
- Frame-Dependent Twist:** A body rotates about the laboratory x axis at one rad/s, so $V_A = (1, 0, 0, 0, 0, 0)^T$. Frame B has parallel axes and origin $p_{AB} = (0, 1, 0)$ at this instant. Compute the B representation of that same motion using the inverse adjoint, and verify its linear block by evaluating the velocity field at B's origin. State why this is not automatically motion relative to an independently moving observer B.

8. **Pitch Extraction:** Construct the normalized rotational screw for $s = (0, 0, 1)$, $q = (1, 1, 0)$ and $h = 3$ m/rad. Recover its pitch by axial projection. Compare that result with the norm of its linear block and explain the difference.
9. **Planar Mechanism Force Path:** Reproduce the centered and offset load solutions of the four-bar example, including each link's free-body balance and the required motor torque. State which assumptions justify the two-force-member simplification and which exclude out-of-plane components. Do not infer spatial bearing reactions from a planar model.
10. **Helical Joint Twist:** A helical joint through the origin has direction $s = (0, 1, 0)$, pitch 0.1 m/rad and rate 2 rad/s. Compute its twist, the axis-point velocity, and its axial travel per complete revolution. Distinguish pitch from lead per revolution.
11. **Dual Wrench Transformation:** Starting with $V_A = XV_B$, derive $F_A = X^{-T}F_B$ from power equality for every V_B . Derive the reverse rule $F_B = X^T F_A$. Use a nonzero translation and a force with a moment arm to demonstrate why applying X^T in the B-to-A direction fails.
12. **Moment Arm Counterexample:** A force $(1, 0, 0)$ N acts at $(0, 2, 0)$ m. Compute the wrench and its power on a unit-rate positive z rotation about the origin. Explain why it is not an ideal reaction wrench of a free z hinge and what balancing actuator torque would be needed in static equilibrium.
13. **Screw Axis From a Generator:** Let $\omega = (0, 0, 0.1)$ rad/s and $v = (1, 0, 0)$ m/s. Extract the instantaneous axis and pitch. Compare the exact displacement $e^{\hat{V}\epsilon}$ with $I + \hat{V}\epsilon$ for a nonzero time step ϵ . Show why the latter generally fails rotation orthogonality and specify the order of its local truncation error.
14. **Reciprocal Screw Pair:** For $S_1 = (0, 0, 1, 1, 1, 0)^T$ and a nonzero prismatic screw $S_2 = (0, 0, 0, a, b, c)^T$, find the condition for geometric reciprocity using the cross-block form B . Normalize a valid translational direction. Compare the condition with the ordinary six-vector dot product and explain why that dot product answers a different question.
15. **Virtual Work Principle:** For an admissible twist subspace $V = J\delta\dot{q}$, prove that zero reaction power for every $\delta\dot{q}$ is equivalent to $J^T F = 0$. Contrast this with a wrench that happens to do zero power on only one measured velocity. Explain why an ideal rigid grip may transfer power to the club while doing zero net interface work on the combined hand-club system.

Chapter 6

Exponential Coordinates and Matrix Logarithms

To move smoothly through the universe, we do not jump from matrix to matrix. We integrate continuously along the exponent of a physical axis. The exponential map wraps velocity vectors onto the curved surface of rotation space; the logarithm unfolds curved poses back into flat, linear velocity directions.

In Plain Language 6.1. Context & Use Case: Smooth Interpolation and Geodesics

The Math You Will See: We will introduce **Lie Algebras** ($\mathfrak{so}(3)$ and $\mathfrak{se}(3)$), the **Matrix Exponential** ($\exp$), the **Matrix Logarithm** ($\log$), and Rodrigues' formula. We will develop exponential coordinates, the Baker-Campbell-Hausdorff formula, and understand geodesics (shortest paths) on rotation space.

Why We Need It: Imagine a robotic arm needs to move from Pose A to Pose B. You cannot just average the 4×4 matrices of Pose A and Pose B ($(\mathbf{T}_A + \mathbf{T}_B)/2$); the resulting matrix is garbage and physically means nothing. Because 3D space is curved (due to rotation), moving smoothly requires calculating the exact "velocity" of motion and sliding along it. The Logarithm extracts the direction and speed of motion, and the Exponential walks along the geodesic path.

All Variables Defined: $\mathfrak{so}(3)$ represents the tangent space of infinitesimal rotations (Lie algebra), $\text{SE}(3)$ represents the curved position manifold (Lie group), $\hat{\omega}$ is the unit rotation axis, and θ is the rotation angle.

6.1 Historical Context: Sophus Lie and the Theory of Continuous Groups

The modern theory of Lie groups emerged in the 1870s-1890s from the groundbreaking work of Norwegian mathematician **Sophus Lie**. While studying differential equations, Lie discovered that continuous families of symmetries—transformations that preserve the

structure of an equation—form special algebraic objects called **continuous groups**. These are now known as Lie groups in his honor.

Lie's central insight was that every Lie group has a **Lie algebra**—an associated linear (flat) vector space at the identity element that completely encodes the local structure of the group. The exponential map serves as the bridge: it exponentiates elements of the Lie algebra to produce group elements. Conversely, the logarithm inverts this operation.

For the rotation group $\text{SO}(3)$, Lie's theory predicts that the Lie algebra $\mathfrak{so}(3)$ should consist of 3×3 skew-symmetric matrices. The exponential map then converts these skew-symmetric matrices into rotation matrices—a fact that was rediscovered and formalized by **Olinde Rodrigues** in the context of the rotation formula.

In the early 20th century, the theory was extended and generalized. By the latter half of the century, computer scientists and roboticists (notably **Richard Murray**, **Zexiang Li**, and **S. Shankar Sastry** [31], as well as **John Hauser**) recognized that Lie group and Lie algebra techniques provide the natural language for parametrizing rigid body motions and solving differential equations in robotics. A rigorous modern treatment of Lie groups and their algebras can be found in Hall [17].

The exponential-logarithm formalism is now foundational to modern control theory, computer graphics (SLERP interpolation), and numerical algorithms for rigid body dynamics. Understanding it is essential for anyone working with continuous motion in space.




Figure 6.1: The exponential and logarithmic maps bridge the curved Lie group $\text{SO}(3)$ and its flat tangent space (Lie algebra $\mathfrak{so}(3)$). The exponential map converts infinitesimal rotations to finite rotations; the logarithm extracts the rotation axis and angle from a rotation matrix.

6.2 Lie Algebras and the Tangent Space

Recall from Chapter 5 that $\text{SE}(3)$ and $\text{SO}(3)$ are Lie groups—they have both algebraic and smooth geometric structure. But they are curved, non-linear manifolds: you cannot simply add two rotation matrices and get a valid rotation matrix.

However, at each point on a Lie group, there exists a **tangent space**—the space of infinitesimal "velocity directions" at that point. For a Lie group, the tangent space at the identity element is particularly special: it inherits a vector space structure and can be identified with the Lie algebra.

Definition 6.1. The Lie Algebra $\mathfrak{so}(3)$

The **Lie algebra** $\mathfrak{so}(3)$ is the set of all 3×3 skew-symmetric matrices:

$$\mathfrak{so}(3) = \{\mathbf{A} \in \mathbb{R}^{3 \times 3} : \mathbf{A}^T = -\mathbf{A}\} \quad (6.1)$$

Equivalently, $\mathfrak{so}(3)$ can be identified with $\mathbb{R}^3$ by the isomorphism between vectors and their skew-symmetric matrix representations. A vector $\boldsymbol{\omega} = [\omega_x, \omega_y, \omega_z]^T$ corresponds to the matrix:

$$[\boldsymbol{\omega}] = \begin{bmatrix} 0 & -\omega_z & \omega_y \\ \omega_z & 0 & -\omega_x \\ -\omega_y & \omega_x & 0 \end{bmatrix} \quad (6.2)$$

The dimension of $\mathfrak{so}(3)$ is 3 (the same as $\text{SO}(3)$), and it is a **vector space**: sums and scalar multiples of skew-symmetric matrices are skew-symmetric.

Definition 6.2. The Lie Algebra $\mathfrak{se}(3)$

The **Lie algebra** $\mathfrak{se}(3)$ is the set of 4×4 matrices of the form:

$$\mathfrak{se}(3) = \left\{ \begin{bmatrix} \mathbf{A} & \mathbf{v} \\ 0 & 0 \end{bmatrix} : \mathbf{A} \in \mathfrak{so}(3), \mathbf{v} \in \mathbb{R}^3 \right\} \quad (6.3)$$

Equivalently, $\mathfrak{se}(3)$ can be identified with $\mathbb{R}^6$ by the isomorphism between 6D twist vectors $[\boldsymbol{\omega}^T, \mathbf{v}^T]^T$ and their matrix representation. The dimension is 6.

6.2.1 The Lie Bracket: Non-Commutativity of Motions

The Lie algebra structure includes a **Lie bracket** operation that measures the "failure to commute":

$$[\mathbf{A}, \mathbf{B}] = \mathbf{AB} - \mathbf{BA} \quad (6.4)$$

For skew-symmetric matrices $\mathbf{A} = [\boldsymbol{\omega}_1]$ and $\mathbf{B} = [\boldsymbol{\omega}_2]$, the bracket is itself skew-symmetric:

$$[[\boldsymbol{\omega}_1], [\boldsymbol{\omega}_2]] = [\boldsymbol{\omega}_1 \times \boldsymbol{\omega}_2] \quad (6.5)$$

Geometrically, the Lie bracket of two rotations is another rotation—a reflection of the fact that two sequential rotations about different axes can be rearranged in a non-commutative way.

6.3 Exponential Coordinates and Axis-Angle Representation

6.3.1 The Exponential Coordinate Vector

Let us begin by revisiting the concept of rotating around a single fixed unit axis $\hat{\omega} \in \mathbb{R}^3$ (with $\|\hat{\omega}\| = 1$) by a certain angle $\theta \in \mathbb{R}$ (measured in radians).

If we multiply the physical axis unit vector $\hat{\omega}$ by the total angle θ , we obtain a new 3-dimensional vector:

$$\mathbf{r} = \hat{\omega}\theta \in \mathbb{R}^3 \quad (6.6)$$

This incredibly compact vector $\mathbf{r}$ contains all the information needed to describe any 3D rotation: the direction of $\mathbf{r}$ specifies the axis of rotation, and the magnitude $\|\mathbf{r}\| = \theta$ specifies the angle. This representation is known as the **Exponential Coordinate** representation of a rotation, or the **axis-angle** representation.

In Plain Language 6.2. Why is it called "Exponential"?

In basic 1D calculus, the solution to the differential equation governing continuous exponential growth ($\dot{x} = ax$) is the scalar exponential function: $x(t) = e^{at}x(0)$. Similarly, the differential equation governing a rigid body continuously rotating around a fixed angular velocity matrix $[\hat{\omega}]$ at speed θ is:

$$\dot{\mathbf{R}} = [\hat{\omega}]\dot{\theta}\mathbf{R} \quad (6.7)$$

The solution uses exactly the same mathematical structure, just upgraded for matrices: the **Matrix Exponential**:

$$\mathbf{R}(\theta) = e^{[\hat{\omega}]\theta}\mathbf{R}(0) \quad (6.8)$$

The matrix exponential e^X is defined through its Taylor series (just like the scalar exponential):

$$e^X = \sum_{n=0}^{\infty} \frac{X^n}{n!} = I + X + \frac{X^2}{2!} + \frac{X^3}{3!} + \cdots \quad (6.9)$$

6.4 Rodrigues' Formula and the Matrix Exponential for SO(3)

Computing the matrix exponential $e^{[\hat{\omega}]\theta}$ by summing the infinite Taylor series would be prohibitively slow for computer calculations. Fortunately, there is a closed-form formula.

Theorem 6.1. Rodrigues' Formula (Exponential Map on SO(3))

For any unit vector $\hat{\omega} \in \mathbb{R}^3$ with $\|\hat{\omega}\| = 1$ and angle $\theta \in [0, \pi]$, the matrix exponential is:

$$\mathbf{R} = \exp([\hat{\omega}]\theta) = \mathbf{I} + \sin \theta [\hat{\omega}] + (1 - \cos \theta) [\hat{\omega}]^2 \quad (6.10)$$

Derivation of Rodrigues' Formula. We begin with the Taylor series expansion:

$$e^{[\hat{\omega}]\theta} = \sum_{n=0}^{\infty} \frac{([\hat{\omega}]\theta)^n}{n!} \quad (6.11)$$

The key insight is to find a pattern in the powers of $[\hat{\omega}]$. Since $\hat{\omega}$ is a unit vector, its skew-symmetric matrix satisfies a special property. Using the vector triple product identity, we can show:

$$[\hat{\omega}]^2 = \hat{\omega}\hat{\omega}^T - \mathbf{I} \quad (6.12)$$

(This follows from the identity $\mathbf{a} \times (\mathbf{b} \times \mathbf{c}) = \mathbf{b}(\mathbf{a} \cdot \mathbf{c}) - \mathbf{c}(\mathbf{a} \cdot \mathbf{b})$ applied to skew matrices.) Crucially, we also have:

$$[\hat{\omega}]^3 = -[\hat{\omega}] \quad (6.13)$$

This can be verified directly: $[\hat{\omega}]^3 = [\hat{\omega}] \cdot [\hat{\omega}]^2 = [\hat{\omega}](\hat{\omega}\hat{\omega}^T - \mathbf{I}) = \mathbf{0} - [\hat{\omega}] = -[\hat{\omega}]$ (using $[\hat{\omega}]\hat{\omega} = \mathbf{0}$).

The pattern repeats:

$$[\hat{\omega}]^4 = [\hat{\omega}]^2, \quad [\hat{\omega}]^5 = [\hat{\omega}]^3 = -[\hat{\omega}], \quad [\hat{\omega}]^6 = [\hat{\omega}]^2, \quad [\hat{\omega}]^7 = [\hat{\omega}]^3 = -[\hat{\omega}], \quad \dots \quad (6.14)$$

Therefore, the Taylor series can be reorganized by grouping terms with the same power modulo 3:

$$e^{[\hat{\omega}]\theta} = \sum_{n=0}^{\infty} \frac{([\hat{\omega}]\theta)^n}{n!} \quad (6.15)$$

$$= \sum_{k=0}^{\infty} \frac{[\hat{\omega}]^{2k}\theta^{2k}}{(2k)!} + \sum_{k=0}^{\infty} \frac{(-1)^k[\hat{\omega}]\theta^{2k+1}}{(2k+1)!} \quad (6.16)$$

$$= \sum_{k=0}^{\infty} \frac{[\hat{\omega}]^2\theta^{2k}}{(2k)!} + \sum_{k=0}^{\infty} \frac{(-1)^k[\hat{\omega}]\theta^{2k+1}}{(2k+1)!} \quad (6.17)$$

(using $[\hat{\omega}]^0 = \mathbf{I}$, $[\hat{\omega}]^2 = [\hat{\omega}]^2$, and $[\hat{\omega}]^{2k+1} = (-1)^k[\hat{\omega}]$).

Now separate the constant terms:

$$e^{[\hat{\omega}]\theta} = \mathbf{I} + \sum_{k=1}^{\infty} \frac{[\hat{\omega}]^2\theta^{2k}}{(2k)!} + \sum_{k=0}^{\infty} \frac{(-1)^k[\hat{\omega}]\theta^{2k+1}}{(2k+1)!} \quad (6.18)$$

$$= \mathbf{I} + [\hat{\omega}]^2 \sum_{k=1}^{\infty} \frac{\theta^{2k}}{(2k)!} + [\hat{\omega}] \sum_{k=0}^{\infty} \frac{(-1)^k\theta^{2k+1}}{(2k+1)!} \quad (6.19)$$

$$= \mathbf{I} + [\hat{\omega}]^2 \left(\sum_{k=0}^{\infty} \frac{\theta^{2k}}{(2k)!} - 1 \right) + [\hat{\omega}] \sum_{k=0}^{\infty} \frac{(-1)^k\theta^{2k+1}}{(2k+1)!} \quad (6.20)$$

Recognizing the Taylor series for $\cos \theta$ and $\sin \theta$:

$$\cos \theta = \sum_{k=0}^{\infty} \frac{(-1)^k\theta^{2k}}{(2k)!}, \quad \sin \theta = \sum_{k=0}^{\infty} \frac{(-1)^k\theta^{2k+1}}{(2k+1)!} \quad (6.21)$$

we get:

$$e^{[\hat{\omega}]^\theta} = \mathbf{I} + [\hat{\omega}]^2(\cos \theta - 1) + [\hat{\omega}] \sin \theta \quad (6.22)$$

Rearranging:

$$e^{[\hat{\omega}]^\theta} = \mathbf{I} + \sin \theta [\hat{\omega}] + (1 - \cos \theta) [\hat{\omega}]^2 \quad (6.23)$$

which is Rodrigues' formula. $\square$

6.4.1 Verification That the Exponential Map Produces Valid Rotations

We must verify that Rodrigues' formula actually produces a matrix in $\text{SO}(3)(3)$ (i.e., $\mathbf{R}^T \mathbf{R} = \mathbf{I}$ and $\det(\mathbf{R}) = 1$).

Theorem 6.2. The Exponential Map $\mathfrak{so}(3) \rightarrow \text{SO}(3)$ is Surjective

The map $\exp : \mathfrak{so}(3)(3) \rightarrow \text{SO}(3)(3)$ defined by Rodrigues' formula is surjective (onto). That is, every rotation matrix can be expressed as the exponential of some skew-symmetric matrix.

Proof. Orthogonality: For any skew-symmetric matrix $[\hat{\omega}]$, we have $[\hat{\omega}]^T = -[\hat{\omega}]$. The matrix exponential of a skew-symmetric matrix is orthogonal:

$$(e^{[\hat{\omega}]^\theta})^T e^{[\hat{\omega}]^\theta} = e^{-[\hat{\omega}]^\theta} e^{[\hat{\omega}]^\theta} = e^0 = \mathbf{I} \quad (6.24)$$

(using the property that $e^{A^T} = (e^A)^T$ for any matrix.)

Determinant: The determinant of a matrix exponential is $\det(e^A) = e^{\text{tr}(A)}$. For skew-symmetric matrices, $\text{tr}([\hat{\omega}]\theta) = 0$, so $\det(e^{[\hat{\omega}]\theta}) = e^0 = 1$.

Surjectivity: Any rotation matrix $\mathbf{R} \in \text{SO}(3)(3)$ can be decomposed using the axis-angle representation: there exist a unit vector $\hat{\omega}$ and angle $\theta \in [0, \pi]$ such that $\mathbf{R}$ represents rotation by angle θ about axis $\hat{\omega}$. For this rotation, Rodrigues' formula shows that $\mathbf{R} = e^{[\hat{\omega}]^\theta}$. $\square$

6.5 The Matrix Logarithm on $\text{SO}(3)$

The exponential map is not injective (one-to-one) on $\text{SO}(3)(3)$: multiple axis-angle pairs can produce the same rotation (e.g., rotating by θ about $\hat{\omega}$ is the same as rotating by $-\theta$ about $-\hat{\omega}$, and rotations by θ and $\theta + 2\pi$ are identical). However, if we restrict to $\theta \in [0, \pi]$, the exponential map becomes bijective, and the logarithm is its inverse.

Definition 6.3. Matrix Logarithm on $\text{SO}(3)$

The **Matrix Logarithm** on $\text{SO}(3)(3)$ is the inverse of the exponential map, mapping from $\text{SO}(3)(3)$ back to $\mathfrak{so}(3)(3)$:

$$\log : \text{SO}(3)(3) \rightarrow \mathfrak{so}(3)(3) \quad (6.25)$$

Given a rotation matrix $\mathbf{R}$, the logarithm extracts the axis-angle representation.

6.5.1 Computing the Logarithm: All Cases

The logarithm of a rotation matrix depends on the rotation angle and must be handled in multiple cases:

Theorem 6.3. Logarithm of $\text{SO}(3)$ Rotations

For a rotation matrix $\mathbf{R} \in \text{SO}(3)$, the logarithm $[\hat{\omega}]\theta = \log(\mathbf{R})$ can be computed as follows:

Case 1: Identity ($\theta = 0$)

If $\mathbf{R} = \mathbf{I}$, then:

$$\log(\mathbf{R}) = 0 \in \mathfrak{so}(3) \quad (6.26)$$

Case 2: Small Angle ($\theta \ll 1$)

For small rotations, $\sin \theta \approx \theta$ and $1 - \cos \theta \approx \theta^2/2$, so Rodrigues' formula becomes:

$$\mathbf{R} \approx \mathbf{I} + \theta[\hat{\omega}] + O(\theta^2) \quad (6.27)$$

Inverting this approximation (solving for $[\hat{\omega}]\theta$):

$$[\hat{\omega}]\theta \approx \mathbf{R} - \mathbf{R}^T \quad (\text{skew-symmetric part of } \mathbf{R}) \quad (6.28)$$

More precisely:

$$\log(\mathbf{R}) = \frac{1}{2}(\mathbf{R} - \mathbf{R}^T) \quad \text{for small angles} \quad (6.29)$$

Case 3: General Case

For arbitrary $\theta \in (0, \pi]$, the angle θ can be extracted from the trace of $\mathbf{R}$:

$$\text{tr}(\mathbf{R}) = 1 + 2 \cos \theta \quad (6.30)$$

Solving:

$$\theta = \arccos\left(\frac{\text{tr}(\mathbf{R}) - 1}{2}\right) \quad (6.31)$$

Once θ is known, the axis $\hat{\omega}$ is extracted from the skew-symmetric part:

$$\hat{\omega} = \frac{1}{2 \sin \theta} \begin{bmatrix} R_{32} - R_{23} \\ R_{13} - R_{31} \\ R_{21} - R_{12} \end{bmatrix} \quad (6.32)$$

The skew-symmetric matrix logarithm is:

$$\log(\mathbf{R}) = [\hat{\omega}]\theta \quad (6.33)$$

Case 4: π Rotation

When $\theta = \pi$ (180-degree rotation), $\sin \theta = 0$, so equation (6.32) becomes singular.

In this case, use:

$$\mathbf{R} = 2\hat{\omega}\hat{\omega}^T - \mathbf{I} \quad (6.34)$$

The axis can be found from the eigenvector of $\mathbf{R}$ corresponding to eigenvalue 1:

$$\hat{\omega} = \text{unit eigenvector of } (\mathbf{R} + \mathbf{I}) \text{ with eigenvalue } 2 \quad (6.35)$$

Example 6.1. Computing Logarithm: Example

Suppose we have a rotation matrix representing a 60° rotation about the z -axis:

$$\mathbf{R} = \begin{bmatrix} \cos(60^\circ) & -\sin(60^\circ) & 0 \\ \sin(60^\circ) & \cos(60^\circ) & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1/2 & -\sqrt{3}/2 & 0 \\ \sqrt{3}/2 & 1/2 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (6.36)$$

Step 1: Compute $\text{tr}(\mathbf{R}) = 1/2 + 1/2 + 1 = 2$.

Step 2: Extract angle: $\theta = \arccos((2 - 1)/2) = \arccos(1/2) = 60^\circ = \pi/3$ rad.

Step 3: Extract axis from skew part:

$$\hat{\omega} = \frac{1}{2 \sin(60^\circ)} \begin{bmatrix} 0 - 0 & 0 - 0 \\ \sqrt{3}/2 - (-\sqrt{3}/2) \end{bmatrix} = \frac{1}{2 \cdot \sqrt{3}/2} \begin{bmatrix} 0 \\ 0 \\ \sqrt{3} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \quad (6.37)$$

Step 4: Construct logarithm:

$$\log(\mathbf{R}) = \begin{bmatrix} 0 \\ 0 \\ \pi/3 \end{bmatrix} \quad (6.38)$$

As a skew matrix: $[\hat{\omega}]\theta = \begin{bmatrix} 0 & -\pi/3 & 0 \\ \pi/3 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}.$

Mathematical Principle 6.1. Geometric Distance via Logarithms

The Matrix Logarithm is computationally critical in optimal control and error correction. If your robotic hand is at orientation $\mathbf{R}_{\text{current}}$ and your target is $\mathbf{R}_{\text{target}}$, you *absolutely cannot* compute an "error" by subtracting them ($\mathbf{R}_{\text{target}} - \mathbf{R}_{\text{current}}$) because subtracting curved matrices yields garbage that does not belong to $\text{SO}(3)(3)$. The true geometric error—the precise, shortest-path arc of rotation needed to correct the deviation—is found using the Logarithm of their relative transform:

$$\mathbf{e}_{\text{error}} = \log(\mathbf{R}_{\text{target}} \mathbf{R}_{\text{current}}^T) \quad (6.39)$$

The resulting axis-angle vector $\mathbf{e}_{\text{error}} \in \mathfrak{so}(3)(3)$ lives in the tangent space and directly tells you the minimal rotation needed to align the two frames. This is the mathematical foundation for proportional-derivative (PD) control in robotics.

6.6 The Exponential Map on $\text{SE}(3)$

We now extend the exponential map from rotations alone ($\text{SO}(3)(3)$) to full spatial rigid body transformations ($\text{SE}(3)(3)$).

Definition 6.4. The SE(3) Exponential Map

Given a twist $\mathcal{V} = \begin{bmatrix} \boldsymbol{\omega} \\ \mathbf{v} \end{bmatrix} \in \mathfrak{se}(3)(3)$ and a scalar θ (the parameter along the motion), the exponential map is:

$$\mathbf{T} = \exp([\mathcal{V}]\theta) \in \text{SE}(3)(3) \quad (6.40)$$

where $[\mathcal{V}] = \begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ 0 & 0 \end{bmatrix}$ is the 4×4 matrix representation of the twist.

6.6.1 Derivation of the SE(3) Exponential Formula

To evaluate $\exp([\mathcal{V}]\theta)$, we use the structure of the matrix. Let $[\boldsymbol{\omega}]$ denote the skew-symmetric matrix of $\boldsymbol{\omega}$, which satisfies $[\boldsymbol{\omega}]^3 = -[\boldsymbol{\omega}]$ as before.

The Taylor series is:

$$\exp([\mathcal{V}]\theta) = \sum_{n=0}^{\infty} \frac{1}{n!} \begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ 0 & 0 \end{bmatrix}^n \theta^n \quad (6.41)$$

Computing powers:

$$\begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ 0 & 0 \end{bmatrix}^2 = \begin{bmatrix} [\boldsymbol{\omega}]^2 & [\boldsymbol{\omega}]\mathbf{v} \\ 0 & 0 \end{bmatrix} \quad (6.42)$$

$$\begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ 0 & 0 \end{bmatrix}^3 = \begin{bmatrix} [\boldsymbol{\omega}]^3 & [\boldsymbol{\omega}]^2\mathbf{v} \\ 0 & 0 \end{bmatrix} \quad (6.43)$$

The pattern is similar to the $SO(3)$ case, but now we must also track the translation vectors. Using the fact that $[\boldsymbol{\omega}]^3 = -[\boldsymbol{\omega}]$:

$$\exp([\mathcal{V}]\theta) = \begin{bmatrix} e^{[\boldsymbol{\omega}]\theta} & \mathbf{V}(\theta)\mathbf{v} \\ 0 & 1 \end{bmatrix} \quad (6.44)$$

where the **V matrix** (or screw Jacobian) is:

Theorem 6.4. The V Matrix Formula

The 3×3 matrix that relates the translation component of the twist to the spatial translation is:

$$\mathbf{V}(\theta) = \mathbf{I} + (1 - \cos \theta) \frac{[\boldsymbol{\omega}]}{\|\boldsymbol{\omega}\|^2} + (\theta - \sin \theta) \frac{[\boldsymbol{\omega}]^2}{\|\boldsymbol{\omega}\|^3} \quad (6.45)$$

For unit angular velocity $\hat{\boldsymbol{\omega}}$, this simplifies to:

$$\mathbf{V}(\theta) = \mathbf{I} + (1 - \cos \theta)[\hat{\boldsymbol{\omega}}] + (\theta - \sin \theta)[\hat{\boldsymbol{\omega}}]^2 \quad (6.46)$$

For a prismatic (pure translation) twist with $\boldsymbol{\omega} = 0$, we have $\mathbf{V} = \mathbf{I}$.

Proof. Separating the Taylor series into even and odd powers (as we did for $SO(3)$), and

using the cyclic property $[\omega]^3 = -[\omega]$:

$$\exp([\mathcal{V}]\theta) = \sum_{n=0}^{\infty} \frac{1}{n!} \begin{bmatrix} [\omega] & \mathbf{v} \\ 0 & 0 \end{bmatrix}^n \theta^n \quad (6.47)$$

$$= \begin{bmatrix} \sum_{k=0}^{\infty} \frac{[\omega]^{2k} \theta^{2k}}{(2k)!} + \sum_{k=0}^{\infty} \frac{[\omega]^{2k+1} \theta^{2k+1}}{(2k+1)!} & \mathbf{V}(\theta)\mathbf{v} \\ 0 & 1 \end{bmatrix} \quad (6.48)$$

The rotation part gives Rodrigues' formula. For the translation part, we have:

$$\mathbf{V}(\theta)\mathbf{v} = \left(\mathbf{I} + \sum_{k=1}^{\infty} \frac{[\omega]^2 \theta^{2k}}{(2k)!} + \sum_{k=0}^{\infty} \frac{[\omega] \theta^{2k+1}}{(2k+1)!} \right) \mathbf{v} \quad (6.49)$$

$$= \left(\mathbf{I} + [\omega] \sum_{k=0}^{\infty} \frac{\theta^{2k+1}}{(2k+1)!} + [\omega]^2 \sum_{k=1}^{\infty} \frac{\theta^{2k}}{(2k)!} \right) \mathbf{v} \quad (6.50)$$

$$= (\mathbf{I} + [\omega] \sin \theta + [\omega]^2 (\cos \theta - 1)) \mathbf{v} \quad (6.51)$$

This is the $\mathbf{V}$ matrix. □

6.7 The Matrix Logarithm on SE(3)

The inverse of the exponential map is the logarithm:

Definition 6.5. Matrix Logarithm on SE(3)

Given a transformation $\mathbf{T} = \begin{bmatrix} \mathbf{R} & \mathbf{p} \\ 0 & 1 \end{bmatrix} \in \text{SE}(3)(3)$, the logarithm extracts the twist:

$$\log(\mathbf{T}) = [\mathcal{V}] \quad (6.52)$$

where $\mathcal{V} = \begin{bmatrix} \omega \\ \mathbf{v} \end{bmatrix}$ is reconstructed as:

1. Extract the angular part: $[\omega] = \log(\mathbf{R})$ using the $SO(3)$ logarithm.
2. Extract the linear part: $\mathbf{v} = \mathbf{V}^{-1}(\theta)\mathbf{p}$, where $\mathbf{V}^{-1}$ is the inverse of the $\mathbf{V}$ matrix.

For numerical implementation, the inverse of the $\mathbf{V}$ matrix is:

$$\mathbf{V}^{-1}(\theta) = \mathbf{I} - \frac{1}{2}[\omega] + \frac{1}{\theta^2} \left(\frac{1}{2} - \frac{\sin(2\theta)}{4 \sin \theta} \right) [\omega]^2 \quad (6.53)$$

(and for $\theta \rightarrow 0$, use L'Hôpital's rule or series approximations to avoid numerical issues).

6.7.1 Numerical Stability of the Matrix Logarithm

The closed-form logarithm derived in Section 6.5 is exact in infinite precision but requires care in floating-point arithmetic at two boundary regimes.

Near $\theta = 0$ (identity rotation). For $\theta \ll 1$, the general formula $[\hat{\omega}] = \frac{\theta}{2 \sin \theta}(\mathbf{R} - \mathbf{R}^T)$ suffers because both θ and $\sin \theta$ approach zero. However their ratio is well-behaved: $\theta / \sin \theta \rightarrow 1$. In practice, use the first-order Taylor expansion

$$\log(\mathbf{R}) \approx \mathbf{R} - \mathbf{I} \quad (\theta \ll 1), \quad (6.54)$$

which is the skew-symmetric part of $\mathbf{R}$ and avoids all division by small quantities.

Near $\theta = \pi$ (180° rotation). The standard axis-extraction formula

$$\hat{\omega} = \frac{1}{2 \sin \theta} \begin{bmatrix} R_{32} - R_{23} \\ R_{13} - R_{31} \\ R_{21} - R_{12} \end{bmatrix} \quad (6.55)$$

has a $1/\sin \theta$ singularity as $\theta \rightarrow \pi$, because $\sin \pi = 0$. Near this boundary the axis must instead be extracted from the *eigenvector* of $\mathbf{R}$ corresponding to eigenvalue +1. A robust approach is to compute $\mathbf{S} = \mathbf{R} + \mathbf{I}$ and take the column of $\mathbf{S}$ with largest norm as $\hat{\omega}$ (after normalization).

Condition number. The condition number of the logarithm map at angle θ scales as

$$\kappa(\theta) \approx \frac{\theta}{\sin \theta}, \quad (6.56)$$

which equals 1 at $\theta = 0$ and diverges as $\theta \rightarrow \pi$. Consequently, the logarithm is inherently ill-conditioned near half-turns: small perturbations in $\mathbf{R}$ can produce large changes in the recovered axis.

In Plain Language 6.3. Practical Recommendation

For maximum robustness, use an `atan2`-based extraction of θ (avoiding the arccos whose derivative diverges at ± 1) and switch to eigenvector-based axis recovery when $\sin \theta$ drops below a threshold (e.g. 10^{-6}). Modern libraries such as `scipy.spatial.transform` implement these guards automatically.

6.8 The Baker-Campbell-Hausdorff Formula

One of the most subtle facts about Lie algebras and Lie groups is that they do not commute under composition. That is, if $\mathbf{A}$ and $\mathbf{B}$ are two elements of a Lie algebra, then:

$$e^{\mathbf{A}}e^{\mathbf{B}} \neq e^{\mathbf{A}+\mathbf{B}} \quad (6.57)$$

in general. The difference is captured by the **Baker-Campbell-Hausdorff (BCH) formula**, which expresses the "product" of two exponentials as a single exponential:

Theorem 6.5. Baker-Campbell-Hausdorff Formula (First-Order)

For small perturbations $\mathbf{A}, \mathbf{B} \in \mathfrak{g}$ (the Lie algebra) with $\|\mathbf{A}\|, \|\mathbf{B}\| \ll 1$:

$$\log(e^{\mathbf{A}}e^{\mathbf{B}}) = \mathbf{A} + \mathbf{B} + \frac{1}{2}[\mathbf{A}, \mathbf{B}] + O(\|\mathbf{A}\|^3, \|\mathbf{B}\|^3) \quad (6.58)$$

where $[\mathbf{A}, \mathbf{B}] = \mathbf{AB} - \mathbf{BA}$ is the Lie bracket (commutator).

Mathematical Principle 6.2. Geometric Interpretation of BCH

The BCH formula reveals that when you compose two small rotations, the result is not simply the sum of the rotations, but the sum plus a correction term proportional to their bracket (commutator). This correction term vanishes only if the two motions commute (i.e., $[\mathbf{A}, \mathbf{B}] = 0$).

For rotations, the bracket of two twists is another twist:

$$[\mathcal{V}_1, \mathcal{V}_2] = \text{Ad}_{\exp(\mathcal{V}_1)} \mathcal{V}_2 - \mathcal{V}_2 + \text{higher order} \quad (6.59)$$

This is the mathematical reason why the order of rotations matters: it is not a "software bug" but a fundamental property of 3D rotations.

The full BCH formula (for arbitrary elements) is an infinite series:

$$\log(e^{\mathbf{A}}e^{\mathbf{B}}) = \mathbf{A} + \mathbf{B} + \frac{1}{2}[\mathbf{A}, \mathbf{B}] + \frac{1}{12}([\mathbf{A}, [\mathbf{A}, \mathbf{B}]] + [\mathbf{B}, [\mathbf{B}, \mathbf{A}]]) + \cdots \quad (6.60)$$

In practice, the first-order approximation is sufficient for small perturbations. A thorough development of the BCH formula and its convergence properties can be found in Hall [17].

6.8.1 The Baker-Campbell-Hausdorff Formula: Higher-Order Terms

The first-order BCH approximation $\log(e^{\mathbf{A}}e^{\mathbf{B}}) \approx \mathbf{A} + \mathbf{B} + \frac{1}{2}[\mathbf{A}, \mathbf{B}]$ is adequate when both $\mathbf{A}$ and $\mathbf{B}$ are small. For larger rotations the higher-order Lie-bracket terms become significant. The full expansion to third order reads [17]:

Theorem 6.6. BCH Formula to Third Order

For $\mathbf{A}, \mathbf{B}$ in a Lie algebra $\mathfrak{g}$,

$$\log(e^{\mathbf{A}}e^{\mathbf{B}}) = \mathbf{A} + \mathbf{B} + \frac{1}{2}[\mathbf{A}, \mathbf{B}] + \frac{1}{12}([\mathbf{A}, [\mathbf{A}, \mathbf{B}]] + [\mathbf{B}, [\mathbf{B}, \mathbf{A}]]) + O(\|\cdot\|^4), \quad (6.61)$$

where $[\cdot, \cdot]$ denotes the Lie bracket and the remainder collects terms involving four or more nested brackets.

Convergence. For the rotation group $\text{SO}(3)$, the BCH series converges absolutely whenever $\|\mathbf{A}\|, \|\mathbf{B}\| < \pi$, i.e. whenever both rotation angles are less than 180° . Outside this radius the series may diverge because the exponential map on $\text{SO}(3)$ is not globally injective—rotations by θ and $\theta + 2\pi$ are indistinguishable. More precisely, the convergence radius is set by the injectivity radius of the exponential map, which equals π for $\text{SO}(3)$.

Physical meaning. The BCH formula quantifies how much two sequential rotations about different axes *interact*. The first-order bracket term $\frac{1}{2}[\mathbf{A}, \mathbf{B}]$ captures the leading coupling between the two rotation axes. The second-order terms $\frac{1}{12}[\mathbf{A}, [\mathbf{A}, \mathbf{B}]]$ and $\frac{1}{12}[\mathbf{B}, [\mathbf{B}, \mathbf{A}]]$ describe how each rotation axis “precesses” the other—an effect that grows with the magnitude of the rotations.

Mathematical Principle 6.3. Small-Angle Approximation from BCH

For small angles ($\|\mathbf{A}\|, \|\mathbf{B}\| \ll 1$), all bracket terms are negligible and

$$e^{\mathbf{A}} e^{\mathbf{B}} \approx e^{\mathbf{A}+\mathbf{B}}. \quad (6.62)$$

Rotations therefore “add like vectors” in the small-angle regime. For large angles the Lie bracket terms become significant: a 90° rotation about x followed by a 90° rotation about y differs from the reverse order by a 90° rotation about z —a dramatic non-commutativity fully predicted by the BCH series.

6.9 Numerical Computation of Matrix Exponentials

Computing the matrix exponential e^A numerically is a classic problem in scientific computing. For large matrices or large exponents, directly summing the Taylor series is inefficient and can be numerically unstable.

6.9.1 Padé Approximation

One robust approach is the **Padé approximation**, which approximates e^A as a ratio of two polynomials:

$$e^A \approx \frac{N(A)}{D(A)} = \left(I + \frac{A}{2} + \frac{A^2}{12} + \cdots \right) \left(I - \frac{A}{2} + \frac{A^2}{12} - \cdots \right)^{-1} \quad (6.63)$$

The advantage is that rational functions converge faster than power series and are more numerically stable.

6.9.2 Scaling and Squaring

For very large matrices, the **scaling and squaring** method is effective: compute $e^{A/2^k}$ (which has small norm and is easy to compute), then square the result k times:

$$e^A = (e^{A/2^k})^{2^k} \quad (6.64)$$

In Plain Language 6.4. Practical Implementation Note

Modern libraries like Eigen (C++), NumPy (Python), and MATLAB provide optimized `matrix_exp` functions that combine multiple techniques (scaling, Padé approximation, eigenvalue decomposition if applicable) to ensure numerical accuracy and efficiency. For robotics applications, it is almost always better to use these library functions than to implement your own.

6.10 Geodesics on SO(3): Shortest Rotation Paths

A **geodesic** is the shortest path between two points on a curved surface. On the surface of a sphere, geodesics are great circles. On the manifold $\text{SO}(3)$, geodesics are paths parameterized by the exponential map.

Mathematical Principle 6.4. Geodesics via the Exponential Map

The shortest path (in a Riemannian sense) from the identity rotation to a rotation $\mathbf{R} \in \text{SO}(3)(3)$ is:

$$\mathbf{R}(t) = \exp(t \cdot \log(\mathbf{R})), \quad t \in [0, 1] \quad (6.65)$$

where $t = 0$ gives the identity and $t = 1$ gives $\mathbf{R}$.

Geometrically, this path rotates at constant angular velocity along the axis extracted by the logarithm. The path length is the rotation angle $\theta = \|\log(\mathbf{R})\|$.

Proof. The geodesic distance on $\text{SO}(3)(3)$ is defined by the Riemannian metric inherited from the tangent space. The exponential map is a local isometry, meaning it preserves distances: the Euclidean distance in the tangent space $\mathfrak{so}(3)(3)$ equals the geodesic distance on the manifold $\text{SO}(3)(3)$ (at least locally).

A straight line in the tangent space—parameterized by $t \cdot \mathbf{v}$ for a direction $\mathbf{v}$ —has minimal length. Exponentiating this line gives the geodesic on the manifold. $\square$

6.10.1 Geodesics on $\text{SO}(3)(3)$: Proof via Riemannian Geometry

The informal argument above appeals to intuition; we now give a more rigorous derivation grounded in Riemannian geometry [29].

Definition 6.6. Bi-Invariant Metric on $\text{SO}(3)(3)$

The **bi-invariant Riemannian metric** on $\text{SO}(3)(3)$ is defined at the identity by the inner product on $\mathfrak{so}(3)(3)$:

$$\langle \mathbf{A}, \mathbf{B} \rangle = -\frac{1}{2} \text{tr}(\mathbf{AB}), \quad \mathbf{A}, \mathbf{B} \in \mathfrak{so}(3)(3). \quad (6.66)$$

This metric is extended to the entire group by left (or right) translation: for tangent vectors $\mathbf{V}, \mathbf{W}$ at $\mathbf{R} \in \text{SO}(3)(3)$,

$$\langle \mathbf{V}, \mathbf{W} \rangle_{\mathbf{R}} = -\frac{1}{2} \text{tr}((\mathbf{R}^{-1}\mathbf{V})(\mathbf{R}^{-1}\mathbf{W})). \quad (6.67)$$

The metric is *bi-invariant*: unchanged under both left and right multiplication by any group element.

Theorem 6.7. One-Parameter Subgroups Are Geodesics

With respect to the bi-invariant metric (6.66), the geodesics through the identity $\mathbf{I} \in \text{SO}(3)(3)$ are exactly the one-parameter subgroups

$$\gamma(t) = e^{t[\hat{\omega}]}, \quad [\hat{\omega}] \in \mathfrak{so}(3)(3). \quad (6.68)$$

Proof. A curve $\gamma(t)$ on $\text{SO}(3)(3)$ is a geodesic if and only if it satisfies the Euler–Lagrange equation for length-minimizing curves. For a bi-invariant metric this equation reduces to

the condition that the **body velocity** is constant:

$$\dot{\gamma}(t) \gamma(t)^{-1} = \text{const} \in \mathfrak{so}(3)(3). \quad (6.69)$$

For the one-parameter subgroup $\gamma(t) = e^{t\mathbf{A}}$ with $\mathbf{A} \in \mathfrak{so}(3)(3)$, we compute

$$\dot{\gamma}(t) = \mathbf{A} e^{t\mathbf{A}}, \quad \dot{\gamma}(t) \gamma(t)^{-1} = \mathbf{A} e^{t\mathbf{A}} e^{-t\mathbf{A}} = \mathbf{A}, \quad (6.70)$$

which is manifestly constant. Hence every one-parameter subgroup is a geodesic. Conversely, any solution of (6.69) with $\gamma(0) = \mathbf{I}$ and $\dot{\gamma}(0) = \mathbf{A}$ must satisfy $\gamma(t) = e^{t\mathbf{A}}$ by uniqueness of ODE solutions. Therefore the geodesics through $\mathbf{I}$ are *precisely* the one-parameter subgroups. $\square$

Theorem 6.8. Geodesic Distance on $\text{SO}(3)(3)$

The geodesic distance between two rotations $\mathbf{R}_1, \mathbf{R}_2 \in \text{SO}(3)(3)$ under the bi-invariant metric is

$$d(\mathbf{R}_1, \mathbf{R}_2) = \|\log(\mathbf{R}_1^T \mathbf{R}_2)\| = \theta, \quad (6.71)$$

where $\theta \in [0, \pi]$ is the angle of the relative rotation $\mathbf{R}_1^T \mathbf{R}_2$ and the norm is $\|\mathbf{A}\| = \sqrt{\langle \mathbf{A}, \mathbf{A} \rangle} = \sqrt{-\frac{1}{2} \text{tr}(\mathbf{A}^2)}$.

Proof. By left-invariance of the metric, $d(\mathbf{R}_1, \mathbf{R}_2) = d(\mathbf{I}, \mathbf{R}_1^T \mathbf{R}_2)$. The geodesic from $\mathbf{I}$ to $\mathbf{R}_1^T \mathbf{R}_2$ is $\gamma(t) = e^{t \log(\mathbf{R}_1^T \mathbf{R}_2)}$, whose length is $\|\log(\mathbf{R}_1^T \mathbf{R}_2)\|$. For $\mathbf{A} = [\hat{\omega}] \theta \in \mathfrak{so}(3)(3)$, we have $\|\mathbf{A}\| = \theta$, confirming that the geodesic distance is the rotation angle. $\square$

6.11 SLERP: Spherical Linear Interpolation as Geodesic Interpolation

In Chapter 4, we introduced **SLERP (Spherical Linear Interpolation)**, which smoothly interpolates between two quaternions. We can now understand SLERP as a discretization of the geodesic on $\text{SO}(3)(3)$ [27].

Theorem 6.9. SLERP as Exponential Interpolation

Given two rotation matrices $\mathbf{R}_0$ and $\mathbf{R}_1$, the SLERP interpolation:

$$\mathbf{R}(t) = \mathbf{R}_0 \exp(t \log(\mathbf{R}_0^T \mathbf{R}_1)), \quad t \in [0, 1] \quad (6.72)$$

traces a geodesic from $\mathbf{R}_0$ to $\mathbf{R}_1$. This path:

1. Is the shortest distance (in terms of rotation angle) between the two orientations.
2. Rotates at constant angular velocity.
3. Is independent of the order (left-invariant parameterization) or dependent on the parameterization (right-invariant), depending on whether we multiply on the left or right.

Proof. At $t = 0$: $\mathbf{R}(0) = \mathbf{R}_0 \exp(0) = \mathbf{R}_0$.

At $t = 1$: $\mathbf{R}(1) = \mathbf{R}_0 \exp(\log(\mathbf{R}_0^T \mathbf{R}_1)) = \mathbf{R}_0(\mathbf{R}_0^T \mathbf{R}_1) = \mathbf{R}_1$.

The rotation angle between $\mathbf{R}(t)$ and $\mathbf{R}(0)$ is $t\theta_{\max}$ where $\theta_{\max} = \|\log(\mathbf{R}_0^T \mathbf{R}_1)\|$, so the angular velocity is constant.

For quaternions, the same path can be parameterized as:

$$\mathbf{q}(t) = \frac{\sin((1-t)\theta)}{\sin\theta} \mathbf{q}_0 + \frac{\sin(t\theta)}{\sin\theta} \mathbf{q}_1 \quad (6.73)$$

where $\theta = \arccos(\mathbf{q}_0 \cdot \mathbf{q}_1)$ is the angle between the quaternions. This is the classical SLERP formula. $\square$

Example 6.2. SLERP vs. Linear Interpolation

Consider interpolating between two rotations: - $\mathbf{R}_0 = I$ (identity, no rotation) - $\mathbf{R}_1 = e^{[0,0,1] \cdot \pi/2}$ (90-degree rotation about z -axis)

Linear Interpolation (Wrong):

$$\mathbf{R}_{\text{wrong}}(t) = (1-t)\mathbf{R}_0 + t\mathbf{R}_1 = (1-t)\mathbf{I} + t\mathbf{R}_1 \quad (6.74)$$

At $t = 0.5$, this yields a matrix that is neither orthogonal nor a valid rotation—it is geometrically meaningless.

SLERP (Correct):

$$\mathbf{R}_{\text{correct}}(t) = \mathbf{R}_0 \exp(t \log(\mathbf{R}_0^T \mathbf{R}_1)) = \exp(t[0, 0, 1] \cdot \pi/2) \quad (6.75)$$

At $t = 0.5$, this yields a rotation by 45 degrees about the z -axis, which is geometrically correct and lies on the shortest path between the two orientations.

6.12 Python Worked Example: Computing Exponentials and Logarithms

```

1 import numpy as np
2 from scipy.spatial.transform import Rotation as R
3 from scipy.linalg import expm, logm
4 import warnings
5
6 def skew(v):
7     """Compute the skew-symmetric matrix of a 3D vector."""
8     return np.array([
9         [0, -v[2], v[1]],
10        [v[2], 0, -v[0]],
11        [-v[1], v[0], 0]
12    ])
13
14 def unskew(S):
15     """Extract a 3D vector from a skew-symmetric matrix."""
16     return np.array([S[2, 1], S[0, 2], S[1, 0]])
17
18 def rodrigues(omega, theta):
19     """

```

```

20     Compute rotation matrix using Rodrigues' formula.
21     omega: unit axis vector (3D)
22     theta: rotation angle in radians
23     """
24     if np.linalg.norm(omega) < 1e-10:
25         return np.eye(3)
26
27     omega_hat = omega / np.linalg.norm(omega)
28     omega_skew = skew(omega_hat)
29
30     R = (np.eye(3) +
31          np.sin(theta) * omega_skew +
32          (1 - np.cos(theta)) * omega_skew @ omega_skew)
33
34     return R
35
36 def log_SO3(R):
37     """
38     Compute the logarithm of a rotation matrix.
39     Returns the axis-angle vector w = omega_hat * theta.
40     """
41     # Clamp to avoid numerical errors with arccos
42     trace = np.trace(R)
43     theta = np.arccos(np.clip((trace - 1) / 2, -1, 1))
44
45     if theta < 1e-10: # Identity rotation
46         return np.zeros(3)
47     elif np.abs(theta - np.pi) < 1e-10: # 180-degree rotation
48         # Find the axis from the eigenvector with eigenvalue 1
49         eigvals, eigvecs = np.linalg.eig((R + np.eye(3)))
50         idx = np.argmax(np.real(eigvals))
51         omega_hat = np.real(eigvecs[:, idx])
52         omega_hat = omega_hat / np.linalg.norm(omega_hat)
53         return omega_hat * theta
54     else:
55         omega_skew = (R - R.T) / (2 * np.sin(theta))
56         omega_hat = unskew(omega_skew)
57         return omega_hat * theta
58
59 def exp_SE3(twist, theta):
60     """
61     Compute the exponential map from se(3) to SE(3).
62     twist: 6D twist vector [omega; v]
63     theta: parameter (typically 1.0 for unit time step)
64     """
65     omega = twist[:3]
66     v = twist[3:]
67
68     if np.linalg.norm(omega) < 1e-10: # Pure translation
69         R = np.eye(3)
70         p = v * theta
71     else:
72         omega_hat = omega / np.linalg.norm(omega)
73         w_magnitude = np.linalg.norm(omega)
74         theta_total = w_magnitude * theta
75
76         # Rotation part
77         R = rodrigues(omega_hat, theta_total)
78
79         # Translation part via V matrix

```



```

80     V = (np.eye(3) +
81          (1 - np.cos(theta_total)) * skew(omega_hat) / w_magnitude +
82          (theta_total - np.sin(theta_total)) * skew(omega_hat)**2 / (
w_magnitude**2))
83
84     p = V @ v * theta
85
86     T = np.eye(4)
87     T[:3, :3] = R
88     T[:3, 3] = p
89     return T
90
91 def log_SE3(T):
92     """
93     Compute the logarithm of a transformation matrix.
94     Returns the twist vector [omega; v].
95     """
96     R = T[:3, :3]
97     p = T[:3, 3]
98
99     omega = log_SO3(R)
100    w_magnitude = np.linalg.norm(omega)
101
102    if w_magnitude < 1e-10: # Pure translation
103        v = p
104    else:
105        omega_hat = omega / w_magnitude
106        theta = w_magnitude
107
108        # Compute V^{-1}
109        V_inv = (np.eye(3) -
110                0.5 * skew(omega_hat) +
111                (1.0/theta**2) * (0.5 - np.sin(2*theta)/(4*np.sin(theta))) *
skew(omega_hat)**2)
112
113        v = V_inv @ p
114
115    twist = np.hstack([omega, v])
116    return twist
117
118 def slerp(R0, R1, t):
119     """
120     Spherical linear interpolation between two rotations.
121     Uses exponential coordinates for geodesic interpolation.
122     """
123     # Relative rotation from R0 to R1
124     R_rel = R0.T @ R1
125     omega = log_SO3(R_rel)
126
127     # Interpolated relative rotation
128     R_interp_rel = rodrigues(omega / np.linalg.norm(omega), t * np.linalg.
norm(omega)) if np.linalg.norm(omega) > 1e-10 else np.eye(3)
129
130     # Final rotation
131     R_interp = R0 @ R_interp_rel
132     return R_interp
133
134 # Example usage
135 if __name__ == "__main__":
136     # 1. Create a rotation: 60 degrees about the z-axis

```

```

137 theta = np.pi / 3 # 60 degrees
138 omega = np.array([0, 0, 1]) # z-axis
139 R = rodrigues(omega, theta)
140 print("Rotation matrix (60 deg about z):")
141 print(R)
142 print()
143
144 # 2. Compute logarithm
145 omega_recovered = log_S03(R)
146 print(f"Recovered axis-angle: {omega_recovered}")
147 print(f"Expected: {omega * theta}")
148 print()
149
150 # 3. Verify round-trip
151 R_recovered = rodrigues(omega_recovered / np.linalg.norm(omega_recovered)
152 ,
153                        np.linalg.norm(omega_recovered))
154 print("Round-trip error (should be ~0):")
155 print(np.linalg.norm(R - R_recovered))
156 print()
157
158 # 4. Exponential map on SE(3)
159 twist = np.array([0, 0, 1, 1, 0, 0]) # rotate about z, translate in x
160 T = exp_SE3(twist, 1.0)
161 print("SE(3) exponential:")
162 print(T)
163 print()
164
165 # 5. Log of SE(3)
166 twist_recovered = log_SE3(T)
167 print(f"Recovered twist: {twist_recovered}")
168 print(f"Expected twist: {twist}")
169 print()
170
171 # 6. SLERP interpolation
172 R0 = np.eye(3)
173 theta1 = np.pi / 2 # 90 degrees
174 R1 = rodrigues(np.array([0, 0, 1]), theta1)
175
176 print("SLERP interpolation (0, 0.5, 1.0):")
177 for t in [0, 0.5, 1.0]:
178     R_t = slerp(R0, R1, t)
179     omega_t = log_S03(R_t)
180     angle_t = np.linalg.norm(omega_t)
181     print(f"t={t}: rotation angle = {np.degrees(angle_t):.1f} deg")
182 print()
183
184 # 7. Geodesic distance
185 R_test = rodrigues(np.array([1, 1, 1]) / np.sqrt(3), np.pi / 4)
186 omega_log = log_S03(R_test)
187 geodesic_distance = np.linalg.norm(omega_log)
188 print(f"Geodesic distance from identity to test rotation: {
189 geodesic_distance:.4f} rad = {np.degrees(geodesic_distance):.2f} deg")

```

Listing 6.1: Exponential and Logarithm Computations

6.13 Summary

In this chapter, we developed the exponential and logarithmic maps that govern motion on curved manifolds like $\text{SO}(3)(3)$ and $\text{SE}(3)(3)$.

1. **Lie Algebras:** The tangent spaces $\mathfrak{so}(3)(3)$ and $\mathfrak{se}(3)(3)$ are flat vector spaces where standard calculus operations work perfectly. They are identified with skew-symmetric matrices (and 6D vectors in the $\text{SE}(3)$ case).
2. **Rodrigues' Formula:** The exponential map from $\mathfrak{so}(3)(3)$ to $\text{SO}(3)(3)$ is given in closed form by Rodrigues' formula:

$$e^{[\hat{\omega}]\theta} = \mathbf{I} + \sin \theta [\hat{\omega}] + (1 - \cos \theta) [\hat{\omega}]^2 \quad (6.76)$$

3. **Matrix Logarithm:** The inverse operation extracts axis-angle coordinates from rotation matrices. It handles multiple cases (identity, small angles, general angles, 180-degree rotations) carefully to avoid numerical singularities.
4. **SE(3) Exponential:** Extension to full spatial transformations introduces the $\mathbf{V}$ matrix, which mediates between translation and rotation.
5. **Baker-Campbell-Hausdorff:** The composition of two exponentials is not simply the exponential of the sum; the Lie bracket (commutator) introduces correction terms. This reflects the non-commutativity of rotations.
6. **Geodesics:** The exponential map traces out geodesics (shortest paths) on the manifold. SLERP interpolation of quaternions is mathematically equivalent to geodesic interpolation on $\text{SO}(3)(3)$.
7. **Numerical Computation:** Modern libraries implement robust algorithms (scaling-and-squaring, Padé approximation) for computing exponentials and logarithms numerically.

These tools enable us to: - Smoothly interpolate between poses without averaging invalid matrices. - Extract meaningful "error vectors" from the difference between two poses (via the logarithm). - Solve differential equations governing rigid body motion. - Parameterize smooth trajectories and control laws.

With exponential coordinates and the matrix logarithm in hand, we possess the complete mathematical toolkit for describing, analyzing, and controlling rigid body motion in 3D space. We are now ready to move into Volume I, where we develop the geometric properties of $\text{SO}(3)(3)$ and $\text{SE}(3)(3)$ in depth.

6.14 Further Reading

For a rigorous mathematical treatment of Lie groups, Lie algebras, and the exponential map, Hall [17] is an excellent graduate-level reference that develops the theory from first principles with careful attention to the analytic subtleties. Readers seeking a deeper understanding of the algebraic structures underlying this chapter will find Hall's treatment of matrix Lie groups particularly relevant.

The robotics-specific development of exponential coordinates is best found in Murray, Li, and Sastry [31], whose Chapter 2 constructs the exponential map on $SO(3)(3)$ and $SE(3)(3)$ with an emphasis on applications to manipulator kinematics. Their treatment of the product of exponentials formula, which we develop in Chapter 9, builds directly on the material presented here.

Marsden and Ratiu [29] offer a geometric mechanics perspective on the exponential map, connecting it to Hamiltonian systems, symplectic structures, and momentum maps. Their framework reveals that Rodrigues' formula and the $SE(3)(3)$ exponential are specific instances of a much broader theory governing the geometry of mechanical systems on Lie groups.

Lynch and Park [27] provide an accessible and application-oriented treatment of SLERP interpolation, geodesics on $SO(3)(3)$, and the practical use of exponential coordinates in trajectory planning and motion control. Their worked examples are particularly valuable for building geometric intuition.

It is worth noting that the Baker-Campbell-Hausdorff formula developed in Section 6.8 has deep connections beyond robotics. In quantum mechanics, the BCH formula governs the composition of unitary operators and appears centrally in the theory of quantum gates. In differential geometry, it encodes the local structure of Lie groups and is intimately related to the curvature of the group manifold.

6.15 Exercises

1. **Rodrigues' Formula Verification:** Use Rodrigues' formula to compute the rotation matrix for a 45° rotation about the axis $\hat{\omega} = [1, 1, 1]^T/\sqrt{3}$. Verify that the result is orthogonal ($\mathbf{R}^T \mathbf{R} = \mathbf{I}$) and has determinant 1.
2. **Logarithm of a 90-Degree Rotation:** Given the rotation matrix:

$$\mathbf{R} = \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad (6.77)$$

compute $\log(\mathbf{R})$ and extract the axis and angle.

3. **Trace-Based Angle Extraction:** Show that for any rotation matrix $\mathbf{R}$, the trace $\text{tr}(\mathbf{R}) = 1 + 2\cos\theta$ where θ is the rotation angle. Explain why this formula breaks down when $\theta = \pi$ and how to handle this case.
4. **Skew Matrix Pattern:** Prove that for any unit vector $\hat{\omega}$, the skew-symmetric matrix $[\hat{\omega}]$ satisfies $[\hat{\omega}]^3 = -[\hat{\omega}]$.
5. **Small Angle Approximation:** For small rotations with $\theta \ll 1$, show that $e^{[\hat{\omega}]\theta} \approx \mathbf{I} + [\hat{\omega}]\theta$ and derive the error term (to second order in θ).

6. **$SE(3)$ Exponential:** Given a twist $\mathcal{V} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 0 \\ 0 \end{bmatrix}$ (rotation about z , translation in x),

compute the exponential map to produce the corresponding 4×4 transformation $\mathbf{T} = \exp([\mathcal{V}])$.

7. **V Matrix for Pure Translation:** For a prismatic (pure translation) joint with twist

$$\mathcal{V} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \text{ show that } \mathbf{V} = \mathbf{I} \text{ and thus } \exp([\mathcal{V}]) = \begin{bmatrix} \mathbf{I} & [1, 0, 0]^T \\ 0 & 1 \end{bmatrix}.$$

8. **Logarithm Inversion:** If $\mathbf{T} = \begin{bmatrix} \mathbf{R} & \mathbf{p} \\ 0 & 1 \end{bmatrix}$, show that $\log(\mathbf{T}^{-1}) = -\log(\mathbf{T})$ (as a property of the logarithm).

9. **Baker-Campbell-Hausdorff (First Order):** For two small rotations with twists $\mathcal{V}_1 = \epsilon \mathbf{a}$ and $\mathcal{V}_2 = \epsilon \mathbf{b}$ (where $\epsilon \ll 1$), show that:

$$\log(e^{\mathcal{V}_1} e^{\mathcal{V}_2}) \approx \mathcal{V}_1 + \mathcal{V}_2 + \frac{1}{2}[\mathcal{V}_1, \mathcal{V}_2] \quad (6.78)$$

where $[\cdot, \cdot]$ is the Lie bracket.

10. **SLERP Geodesic Property:** Show algebraically that SLERP between two rotations $\mathbf{R}_0$ and $\mathbf{R}_1$ rotates at constant angular velocity. Specifically, verify that the angular velocity $\boldsymbol{\omega}(t) = \frac{d}{dt} \log(\mathbf{R}(t) \mathbf{R}_0^T)$ is constant (independent of t).

11. **Geodesic Distance:** Compute the geodesic distance (rotation angle) between the identity and a rotation by 120° about $[1, 1, 1]^T / \sqrt{3}$.

12. **Numerical Stability:** For $\theta = 0.001$ rad and $\hat{\boldsymbol{\omega}} = [0, 0, 1]^T$, compare the Rodrigues formula result with the first-order Taylor approximation $\mathbf{R} \approx \mathbf{I} + [\hat{\boldsymbol{\omega}}]\theta$. How large is the error?

13. **Matrix Exponential via scipy:** Using scipy's `linalg.expm`, compute $\exp([\mathcal{V}]t)$ for

$$\mathcal{V} = \begin{bmatrix} 0 & -1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} \text{ and } t = \pi/4. \text{ Verify the result matches Rodrigues.}$$

14. **Interpolation Verification:** Generate two random rotation matrices $\mathbf{R}_0$ and $\mathbf{R}_1$. Interpolate between them using SLERP at $t = 0.3$ and $t = 0.7$. Verify that the rotations are on a geodesic by checking that $\theta(0.3) \approx 0.3 \cdot \theta_{\max}$ and $\theta(0.7) \approx 0.7 \cdot \theta_{\max}$.

15. **Exponential Trajectory:** A robot arm follows a screw axis with twist $\mathcal{V} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$ (rotation about x , translation in y). Compute its pose at times $t = 0, \pi/4, \pi/2$ by evaluating $\mathbf{T}(t) = \exp([\mathcal{V}]t)$.

Chapter 7

Recursive Algorithms

You do not solve the equations of motion for a 15-DOF robot in a single monolithic formula. You walk from the base to the tip, and then from the tip back to the base.

In Plain Language 7.1. Context & Use Case: Beating the Math Explosion

The Math You Will See: We will introduce the **Adjoint Matrix** ($[\text{Ad}_T]$), the **Spatial Inertia Matrix** ($\mathcal{I}$), the formal representation of kinematic trees, and the two-pass **Recursive Newton-Euler Algorithm (RNEA)**.

Why We Need It: If you try to write down the exact algebra connecting the shoulder motor of a robot to the tip of its finger, the equation takes up a whole page. If you try to write the equation for a 15-joint humanoid robot, the trigonometric math expands to fill *tens of thousands* of pages of dense algebra. No computer can solve that in real-time. Instead, we use "Recursive Algorithms" that only ever look at two bones at a time. By chaining these tiny, fast calculations together sequentially, we can simulate complex physics at 1,000 frames per second.

All Variables Defined: $\mathcal{V}_i$ is the velocity of link i , $\mathcal{F}_i$ is the force on link i , $\mathbf{T}_{i,i-1}$ is the relative pose between them, and $\dot{\mathcal{V}}_i$ is the spatial acceleration.

7.1 Historical Context: The Computational Revolution

The computational efficiency crisis in robotics was born in the 1970s. As robotic arms grew from simple planar mechanisms to fully 6-DOF industrial manipulators, researchers faced a brutal truth: the Lagrangian equations of motion, when expanded symbolically, produced expressions so massive they literally filled computers' memory banks.

In the early 1980s, **Orin**, **McGhee**, and **Luh** [22] independently developed recursive formulations of robot dynamics. Instead of solving the global inverse dynamics problem as a monolithic matrix equation, they recognized that the sequential tree structure of a kinematic chain could be exploited: each joint only directly interacts with its immediate neighbors. Their key insight was that by organizing calculations to respect this local structure, one could reduce the computational complexity from $O(N^3)$ (matrix inversion) or even $O(N^4)$ (symbolic algebra) down to $O(N)$ (linear in the number of joints).

In 1987, **Roy Featherstone** completed this revolution [12]. His reformulation using

spatial vector algebra (Chapter 8) and adjoint transformations made recursive dynamics not merely efficient, but *concise*. The Recursive Newton-Euler Algorithm (RNEA) became the standard approach used in every modern physics engine: MuJoCo, PyBullet, Drake, Gazebo, and UpstreamDrift all rely fundamentally on Featherstone’s recursive framework, executed as purely numerical matrix operations rather than symbolic manipulation.

7.2 The Curse of Dimensionality in Dynamics

As systems scale in complexity—from a simple pendulum (1 DOF) to a fully modeled anatomical golf swing (15+ DOF) or a humanoid Atlas robot (30+ DOF)—the monolithic, closed-form equations of physics become mathematically intractable.

In standard Lagrangian mechanics (which we review in a later chapter), the equations of motion are written as a massive matrix equation:

$$\mathbf{M}(\mathbf{q})\ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})\dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \boldsymbol{\tau} \quad (7.1)$$

Here, $\mathbf{M}$ is the $N \times N$ mass matrix, $\mathbf{C}$ contains Coriolis and centrifugal forces, $\mathbf{G}$ is gravity, and $\boldsymbol{\tau}$ are the motor torques.

If you attempt to write the analytical Coriolis matrix explicitly for a 15-DOF human model on a piece of paper, the resulting trigonometric expansion $(\sin(\theta_1 + \theta_2) \cos(\theta_3) \dots)$ will literally fill thousands of pages. Each entry in $\mathbf{C}$ is a sum of products of sines, cosines, and Denavit-Hartenberg parameters; for even a moderately complex chain, these expressions become unmaintainable.

To control complex mechanisms in the real world at thousands of Hertz, modern engines (such as MuJoCo, Drake, Pinocchio, and UpstreamDrift) completely abandon building these massive equations explicitly. Instead, they rely entirely on **Recursive Algorithms** executing purely numerical matrix multiplications in $O(N)$ time.




Figure 7.1: A Kinematic Tree Representation of a Robot With a Base Link and Three Movable Links Forming a Branching Structure. Each Link Connects to Its Parent Through a Joint. Recursive Algorithms Exploit This Tree Structure to Compute Dynamics in Linear Time $O(N)$ Rather Than Cubic Time.

7.3 Kinematic Tree Representation

Before we can implement recursive algorithms, we must formalize how to represent a kinematic chain as a tree structure. Unlike the Lagrangian approach, which treats the system globally, the recursive approach exploits the *local* structure: each link i has at most one parent (the link it is directly connected to by a joint) and zero or more children.

Definition 7.1. Kinematic Tree

A kinematic tree is a directed, acyclic graph in which:

- Each link (node) i has at most one parent link, denoted $\text{parent}(i)$ or $\lambda(i)$.
- Link 0 is the **base link** (the root of the tree), and has no parent.
- A joint i connects link i to its parent link $\lambda(i)$ via a transformation $\mathbf{T}_{i,\lambda(i)}(q_i)$.
- The motion of link i is determined entirely by the motion of its parent, plus the motion of joint i itself.

7.3.1 Parent Arrays and Numbering Conventions

To implement this structure efficiently in code, we maintain a simple **parent array** $\lambda : \{1, 2, \dots, N\} \rightarrow \{0, 1, \dots, N-1\}$, where $\lambda(i)$ is the index of link i 's parent. For the base link, $\lambda(0) = -1$ or is undefined.

Example 7.1. Parent Array for a 3-Link Arm

Consider a planar robotic arm: Base (link 0) $\rightarrow$ Shoulder (link 1) $\rightarrow$ Elbow (link 2) $\rightarrow$ Wrist (link 3).

The parent array is:

$$\lambda = [-1, 0, 1, 2]$$

Read as: Link 0 has no parent, link 1's parent is link 0, link 2's parent is link 1, link 3's parent is link 2.

For a humanoid with two legs, the parent array might be:

$$\lambda = [-1, 0, 1, 0, 3, 0, \dots]$$

with link 0 as the pelvis, links 1-2 the left leg, links 3-4 the right leg, link 5 the torso, etc.

Once the kinematic tree structure is defined via λ , the recursive algorithms simply iterate: for each link $i = 1, 2, \dots, N$, they access data from the parent link $\lambda(i)$ and compute updated values for link i .

7.4 The Adjoint Map ($[\text{Ad}_T]$)

Before we can compute recursive mechanics, we need to know how to shift a 6D spatial Twist ($\mathcal{V}$) or 6D Wrench ($\mathcal{F}$) from one reference frame to another. This is done using the

Adjoint Representation of $SE(3)$.**Mathematical Principle 7.1. Change of Frame for Spatial Vectors**

If a physical motion (e.g., a velocity) is the same regardless of which frame we measure it from, but its *numerical representation* changes because the frame itself has rotated and translated, we need a matrix to shift the numbers without changing the physics.

7.4.1 Derivation of the Adjoint Matrix

Suppose Frame B is related to Frame A by a homogeneous transformation $\mathbf{T}_{BA} \in SE(3)$, decomposed as:

$$\mathbf{T}_{BA} = \begin{bmatrix} \mathbf{R}_{BA} & \mathbf{p}_{BA} \\ \mathbf{0} & 1 \end{bmatrix}$$

where $\mathbf{R}_{BA} \in SO(3)$ and $\mathbf{p}_{BA} \in \mathbb{R}^3$.

A spatial velocity $\mathcal{V}_A = (\boldsymbol{\omega}_A, \mathbf{v}_A) \in \mathfrak{se}(3)(3)$ describes the motion of a point in Frame A . When measured from Frame B , the same physical motion is represented numerically as $\mathcal{V}_B = (\boldsymbol{\omega}_B, \mathbf{v}_B)$.

The angular velocity transforms purely by rotation:

$$\boldsymbol{\omega}_B = \mathbf{R}_{BA} \boldsymbol{\omega}_A$$

The linear velocity, however, must account for the fact that the origin of Frame A is moving relative to Frame B . If Frame A 's origin has velocity $\mathbf{v}_{o,A}$ and the point is at offset $\mathbf{r}$ from Frame A 's origin, then in Frame B 's coordinates:

$$\mathbf{v}_B = \mathbf{R}_{BA}(\mathbf{v}_A + \boldsymbol{\omega}_A \times \mathbf{r})$$

For the origin of Frame A itself (where $\mathbf{r} = \mathbf{0}$), this becomes:

$$\mathbf{v}_B = \mathbf{R}_{BA} \mathbf{v}_A - \mathbf{R}_{BA}(\boldsymbol{\omega}_A \times \mathbf{p}_{BA})$$

Rewriting the cross product as a matrix multiplication and using the property $\mathbf{a} \times \mathbf{b} = -[\mathbf{b}]\mathbf{a}$:

$$\mathbf{v}_B = \mathbf{R}_{BA} \mathbf{v}_A + \mathbf{R}_{BA}[\mathbf{p}_{BA}]\boldsymbol{\omega}_A = \mathbf{R}_{BA} \mathbf{v}_A + [\mathbf{R}_{BA} \mathbf{p}_{BA}](\mathbf{R}_{BA} \boldsymbol{\omega}_A)$$

But we can group this as a matrix acting on the 6D vector:

$$\begin{bmatrix} \boldsymbol{\omega}_B \\ \mathbf{v}_B \end{bmatrix} = \begin{bmatrix} \mathbf{R}_{BA} & \mathbf{0} \\ [\mathbf{p}_{BA}]\mathbf{R}_{BA} & \mathbf{R}_{BA} \end{bmatrix} \begin{bmatrix} \boldsymbol{\omega}_A \\ \mathbf{v}_A \end{bmatrix}$$

This is the **Adjoint Matrix**:

Definition 7.2. Adjoint Representation

The Adjoint Matrix $[\text{Ad}_{\mathbf{T}}] \in \mathbb{R}^{6 \times 6}$ for a transformation $\mathbf{T} = (\mathbf{R}, \mathbf{p})$ is:

$$[\text{Ad}_{\mathbf{T}}] = \begin{bmatrix} \mathbf{R} & \mathbf{0}_{3 \times 3} \\ [\mathbf{p}]\mathbf{R} & \mathbf{R} \end{bmatrix} \quad (7.2)$$

where $[\mathbf{p}]$ is the 3×3 skew-symmetric cross product matrix of $\mathbf{p}$.

If Frame A has spatial velocity $\mathcal{V}_A$, the numerically equivalent velocity as measured by Frame B is:

$$\mathcal{V}_B = [\text{Ad}_{\mathbf{T}_{BA}}] \mathcal{V}_A \quad (7.3)$$

7.4.2 The Adjoint for Forces (Wrenches)

For forces, the situation is dual. A wrench $\mathcal{F} = (\mathbf{m}, \mathbf{f})$ (moment and force) transforms using the **coadjoint** representation, which is the transpose-inverse:

$$\mathcal{F}_B = [\text{Ad}_{\mathbf{T}_{BA}}]^{-T} \mathcal{F}_A = \begin{bmatrix} \mathbf{R}^T & [\mathbf{p}]^T \mathbf{R}^T \\ \mathbf{0} & \mathbf{R}^T \end{bmatrix} \mathcal{F}_A \quad (7.4)$$

This ensures that the power (work rate) $\mathcal{F}^T \mathcal{V}$ is invariant under frame changes, as it must be for physical consistency.

7.5 Recursive Forward Kinematics

In Plain Language 7.2. The Breadcrumb Strategy

To find the absolute speed of the end-effector relative to the ground, we don't jump straight from the ground to the end-effector.

We ask: How fast is the pelvis moving relative to the feet? Then, we add the speed of the torso relative to the pelvis. Then we add the speed of the shoulder relative to the torso. Chain by chain, link by link, we compose the velocities.

7.5.1 Recursive Velocity Propagation

Suppose we want to find the exact spatial position and velocity of the golf clubhead at the very end of a 15-link robotic golfer. Instead of deriving a massive equation relating the base of the feet directly to the clubhead, we use our 4×4 Transforms and 6×6 Adjoints, and we propagate them sequentially outwards.

The fundamental recursion for spatial velocity is:

Theorem 7.1. Recursive Velocity Propagation

For a kinematic tree, the spatial velocity of link i is given by:

$$\mathcal{V}_i = [\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}] \mathcal{V}_{\lambda(i)} + \mathcal{S}_i \dot{q}_i \quad (7.5)$$

where:

- $\mathcal{V}_{\lambda(i)}$ is the velocity of the parent link (already computed).
- $[\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}]$ transforms the parent's velocity into the coordinate frame of link i .
- $\mathcal{S}_i \in \mathbb{R}^6$ is the motion subspace (axis of joint i) expressed in link i 's frame.
- $\dot{q}_i$ is the scalar joint velocity.

Derivation: The velocity of link i can be decomposed as the velocity it inherits from its parent (carried through the transformation $\mathbf{T}_{i,\lambda(i)}$), plus the additional velocity generated by joint i spinning.

For a revolute joint about axis $\mathbf{z}$ in the joint frame, the screw axis is:

$$\mathcal{S}_i = \begin{bmatrix} \mathbf{z} \\ \mathbf{0} \end{bmatrix} \in \mathbb{R}^6$$

For a prismatic joint along axis $\mathbf{z}$:

$$\mathcal{S}_i = \begin{bmatrix} \mathbf{0} \\ \mathbf{z} \end{bmatrix} \in \mathbb{R}^6$$

7.5.2 Algorithm: Forward Kinematics in $O(N)$ Time

Algorithm 7.1. Forward Kinematics Recursion

```

1 function ForwardKinematics(q, dq, T_rel, S, lambda):
2   // Input: joint positions q, velocities dq, relative transforms
   T_rel,
3   //          screw axes S, parent array lambda
4   // Output: link transforms T, spatial velocities V
5
6   T[0] <- Identity           // Base link is at the origin
7   V[0] <- Zero 6D vector     // Base link is stationary (or has
   gravity/base motion)
8
9   for i <- 1 to N:
10    // Compose transforms sequentially
11    T[i] <- T[lambda[i]] * T_rel[i]
12
13    // Propagate velocity using Adjoint
14    V_parent_in_i <- Ad[T_rel[i]] * V[lambda[i]]
15    V[i] <- V_parent_in_i + S[i] * dq[i]
16
17   return (T, V)

```

7.6 Spatial Acceleration and Coriolis Terms

Once velocities are propagated outward, we need to compute accelerations. However, computing $\ddot{\mathbf{V}}_i$ (spatial acceleration) is not as simple as taking a time derivative of $\dot{\mathbf{V}}_i$. The presence of rotating reference frames introduces non-inertial terms.

7.6.1 The Lie Bracket and Centrifugal Forces

In a non-inertial frame rotating with velocity $\mathcal{V}$, any object with velocity $\mathcal{W}$ experiences an apparent "fictitious" acceleration due to the rotation. This is captured by the **spatial cross product**, which we will define rigorously in Chapter 8.

The spatial acceleration must satisfy:

$$\dot{\mathcal{V}}_i = [\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}] \dot{\mathcal{V}}_{\lambda(i)} + [\mathcal{V}_i \times] \mathcal{S}_i \dot{q}_i + \mathcal{S}_i \ddot{q}_i \quad (7.6)$$

The term $[\mathcal{V}_i \times] \mathcal{S}_i \dot{q}_i$ captures the Coriolis acceleration: the acceleration felt by the joint motion when embedded in a frame that is itself rotating.

More explicitly, if $\mathcal{V}_i = (\boldsymbol{\omega}_i, \mathbf{v}_i)$ is the spatial velocity of link i , the motion cross product operator is:

$$[\mathcal{V}_i \times] = \begin{bmatrix} [\boldsymbol{\omega}_i] & \mathbf{0}_{3 \times 3} \\ [\mathbf{v}_i] & [\boldsymbol{\omega}_i] \end{bmatrix} \quad (7.7)$$

When this acts on the screw axis $\mathcal{S}_i$, it produces the non-inertial acceleration term.

Algorithm 7.2. Forward Acceleration (Outward Pass)

```

1 function ForwardAcceleration(V, ddq, S, lambda, Gravity):
2   // Input: spatial velocities V (from forward kinematics),
3   //        joint accelerations ddq, screw axes S, gravity vector
4   // Output: spatial accelerations dV
5
6   dV[0] <- Gravity           // Base link has gravitational
   acceleration
7
8   for i <-1 to N:
9     // Velocity-dependent term (Coriolis)
10    Coriolis <- [V[i]x] * S[i] * dq[i]
11
12    // Compose from parent
13    dV_parent_in_i <- Ad[T_rel[i]] * dV[lambda[i]]
14
15    // Total acceleration
16    dV[i] <- dV_parent_in_i + Coriolis + S[i] * ddq[i]
17
18   return dV

```

7.7 Spatial Inertia

Before we compute forces, we must define mass in 6D. In 3D space, an object has:

- Scalar mass m , resisting linear acceleration ($F = ma$).
- 3×3 rotational inertia tensor $\mathbf{I}_c$ (about the center of mass), resisting angular acceleration ($\boldsymbol{\tau} = \mathbf{I}_c \boldsymbol{\alpha}$).

Spatial Inertia unifies these into a single 6×6 matrix:

Definition 7.3. Spatial Inertia Matrix

For a rigid body with mass m , center of mass offset $\mathbf{c}$, and rotational inertia tensor

$\mathbf{I}_c$ about the CoM, the Spatial Inertia Matrix is:

$$\mathcal{I} = \begin{bmatrix} \mathbf{I}_c & \mathbf{0}_{3 \times 3} \\ \mathbf{0}_{3 \times 3} & m\mathbf{I}_{3 \times 3} \end{bmatrix} \quad (7.8)$$

when expressed at the center of mass. When expressed at an arbitrary origin (offset by $\mathbf{c}$ from the CoM), the parallel axis theorem in 6D gives:

$$\mathcal{I} = \begin{bmatrix} \mathbf{I}_c - m[\mathbf{c}][\mathbf{c}] & m[\mathbf{c}] \\ -m[\mathbf{c}] & m\mathbf{I}_{3 \times 3} \end{bmatrix} \quad (7.9)$$

This matrix is symmetric and positive definite, encoding the complete resistance of a rigid body to any spatial acceleration.

7.8 The Recursive Newton-Euler Algorithm (RNEA)

Inverse Dynamics asks: If we know the trajectory the robot is currently taking $(\mathbf{q}, \dot{\mathbf{q}}, \ddot{\mathbf{q}})$, how much torque $\boldsymbol{\tau}$ must the motors output at this exact millisecond to execute it?

This is solved by the crowning jewel of computational mechanics: the **Recursive Newton-Euler Algorithm (RNEA)** [22, 12]. It executes in exactly two sweeps.

7.8.1 The Outward Pass: Velocity and Acceleration Propagation

Like Forward Kinematics, we cast a sweeping algorithm from the heavy base (the ground) outwards to the lightest tip (the end-effector). For every link i , we compute its absolute velocity $\mathcal{V}_i$ and its absolute spatial acceleration $\dot{\mathcal{V}}_i$.

This outward pass perfectly calculates the exact inertial, Coriolis, and centrifugal accelerations dragging on every single link of the mechanism as it flails through space.

Theorem 7.2. Outward Pass Computation

The outward pass computes, for each link $i = 1, 2, \dots, N$:

$$\mathcal{V}_i = [\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}] \mathcal{V}_{\lambda(i)} + \mathcal{S}_i \dot{q}_i \quad (7.10)$$

$$\dot{\mathcal{V}}_i = [\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}] \dot{\mathcal{V}}_{\lambda(i)} + [\mathcal{V}_i \times] \mathcal{S}_i \dot{q}_i + \mathcal{S}_i \ddot{q}_i \quad (7.11)$$

At the tip (link N), external contact forces $\mathcal{F}_{\text{ext}}$ may be applied.

7.8.2 The Inward Pass: Force and Torque Propagation

Once the outward pass finishes at the tip, we reverse direction. We perform a second loop starting from the tip of the robot back down to the ground.

At the very tip (leaf node), we initialize the spatial wrench. If there is external ground contact or an applied force, we begin with:

$$\mathcal{F}_N = \mathcal{F}_{\text{ext}}$$

Otherwise, we begin with zero wrench at the tip and let it propagate inward.

For each internal link i (processed from N down to 1), we compute the spatial wrench (combined 6D force/torque) required to make link i accelerate at $\dot{\mathcal{V}}_i$ while supporting the forces from all its children. The key Newton-Euler equation in spatial form is:

Theorem 7.3. Spatial Newton-Euler Equation

The spatial wrench (force and torque) required to accelerate a rigid body with spatial inertia $\mathcal{I}$ at acceleration $\dot{\mathcal{V}}$ is:

$$\mathcal{F} = \mathcal{I}\dot{\mathcal{V}} + [\mathcal{V} \times]^*(\mathcal{I}\mathcal{V}) \quad (7.12)$$

where $[\mathcal{V} \times]^*$ is the force cross product (transpose-negative of the motion cross product).

Interpretation: The first term is the "inertial force" ($m\mathbf{a}$); the second term is the centrifugal and Coriolis bias—the apparent force due to the fact that the body is rotating or moving.

7.8.3 Force Recursion and Extraction of Joint Torques

For a link i with one child $j = \text{children}(i)$, the wrench recursion is:

$$\mathcal{F}_i = \mathcal{I}_i \dot{\mathcal{V}}_i + [\mathcal{V}_i \times]^*(\mathcal{I}_i \mathcal{V}_i) + [\text{Ad}_{\mathbf{T}_{j,i}}]^{-T} \mathcal{F}_j \quad (7.13)$$

For a branching tree, we sum over all children:

$$\mathcal{F}_i = \mathcal{I}_i \dot{\mathcal{V}}_i + [\mathcal{V}_i \times]^*(\mathcal{I}_i \mathcal{V}_i) + \sum_{j \in \text{children}(i)} [\text{Ad}_{\mathbf{T}_{j,i}}]^{-T} \mathcal{F}_j \quad (7.14)$$

Once we have the total wrench $\mathcal{F}_i$ acting on link i , the scalar torque (or force) required at joint i is extracted by projecting onto the joint's screw axis:

$$\tau_i = \mathcal{F}_i^T \mathcal{S}_i \quad (7.15)$$

Algorithm 7.3. RNEA: Recursive Newton-Euler Algorithm

```

1 function RNEA(q, dq, ddq, T_rel, S, I, lambda, gravity, F_ext):
2   // Outward pass: compute velocities and accelerations
3   V[0] <- gravity           // Base link has gravitational
   acceleration
4   dV[0] <- gravity
5
6   for i <-1 to N:
7     // Compute relative transform
8     T[i] <- T[lambda[i]] * ForwardKin(q[i], S[i])
9
10    // Velocity recursion
11    V_parent_in_i <- Ad[T_rel[i]] * V[lambda[i]]
12    V[i] <- V_parent_in_i + S[i] * dq[i]
13
14    // Acceleration recursion (Coriolis term)
15    dV_parent_in_i <- Ad[T_rel[i]] * dV[lambda[i]]
16    Coriolis <- -[V[i]x] * S[i] * dq[i]
```

```

17     dV[i] <-dV_parent_in_i + Coriolis + S[i] * ddq[i]
18
19     // Inward pass: compute forces and extract torques
20     F[N] <-F_ext[N]           // External forces at tip
21
22     for i <-N down to 1:
23         // Bias force (centrifugal + Coriolis)
24         bias <-[V[i]x]* * (I[i] * V[i])
25
26         // Inertial force
27         inertial <-I[i] * dV[i]
28
29         // Total wrench at link i
30         F[i] <-inertial + bias
31
32         // Add forces from children
33         for j in children[i]:
34             F[i] <-F[i] + Ad[T_rel[j]]^T * F[j]
35
36         // Extract joint torque
37         tau[i] <-F[i]^T * S[i]
38
39     return tau

```

7.9 Computational Complexity Analysis

Mathematical Principle 7.2. Algorithmic Efficiency

The Recursive Newton-Euler algorithm allows the inverse dynamics for a system with N joints to be solved in absolutely perfect $O(N)$ operations. The time to compute scales linearly, regardless of complexity. This algorithmic breakthrough is the fundamental mathematical reason robotics and physics engines work in the real world.

7.9.1 Why $O(N)$ and not $O(N^2)$?

In the naive Lagrangian approach, we would: 1. Compute the mass matrix $\mathbf{M}(q)$: Each entry M_{ij} requires traversing the kinematic chain from joint i to joint j . This is $O(N^2)$. 2. Compute the Coriolis matrix $\mathbf{C}(q, \dot{q})$: This is $O(N^2)$ as well. 3. Solve the linear system $\mathbf{M}(q)\ddot{\mathbf{q}} = \boldsymbol{\tau} - \mathbf{C}(q, \dot{q})\dot{\mathbf{q}} - \mathbf{G}(q)$: Gaussian elimination is $O(N^3)$.

In the RNEA: 1. **Outward pass**: For each link i , we perform a constant number of 6×6 matrix multiplications and additions. Total: $O(N)$. 2. **Inward pass**: For each link i , we compute the wrench and extract the torque, again with $O(1)$ operations per link. Total: $O(N)$.

The key is that we *never* assemble the global mass matrix. Instead, we exploit the tree structure to compute forces locally at each link, with no coupling except through the sequential parent-child relationships.

7.9.2 Proof of Linearity

Theorem 7.4. RNEA is $O(N)$

The RNEA algorithm, applied to a serial kinematic chain or tree, requires at most $c \cdot N$ floating-point operations, where c is a constant (approximately 100–150, depending on implementation) and N is the number of joints.

Proof sketch:

1. Each outward pass step performs: 1 Adjoint matrix-vector multiplication (≈ 18 flops for $6 \times 6 \times 6D$ vector), 1 vector addition (≈ 6 flops), 1 screw scaling (≈ 6 flops). Total per step: ~ 30 flops.
2. Each inward pass step performs: 1 spatial cross product (≈ 18 flops), 1 inertia product (≈ 36 flops for $6 \times 6 \times 6D$), 1 wrench addition and children summation (~ 30 flops for a typical branching). Total per step: ~ 84 flops.
3. Total: $(30 + 84) \times N = 114N$ flops, independent of N .

This is why modern physics engines can simulate 15–30 DOF robots at 1,000 Hz: $O(N)$ algorithms are fundamentally different from $O(N^3)$ algorithms at scale.

7.10 Branching Kinematic Trees

While we’ve presented the RNEA for serial chains, the algorithm extends seamlessly to branching trees (e.g., humanoids with two legs and two arms branching off a pelvis).

7.10.1 Handling Multiple Children

The modification is minimal: in the inward pass (force recursion), if link i has k children, we sum the wrenches from all children:

$$\mathcal{F}_i = \mathcal{I}_i \dot{\mathcal{V}}_i + [\mathcal{V}_i \times]^* (\mathcal{I}_i \mathcal{V}_i) + \sum_{j=1}^k [\text{Ad}_{\mathbf{T}_{j,i}}]^{-T} \mathcal{F}_j \quad (7.16)$$

The rest of the algorithm remains identical. Because the tree is acyclic, there is a unique topological ordering, and all parents are processed before their children in the inward pass.

Example 7.2. Humanoid Robot with Branching

A humanoid robot with a main body/torso (link 0), two legs (links 1–2 and 3–4), and two arms (links 5–7 and 8–10) has the parent array:

$$\lambda = [-1, 0, 1, 0, 3, 0, 5, 6, 0, 8, 9]$$

In the inward pass, when computing $\mathcal{F}_0$ (torso wrench), we sum:

$$\mathcal{F}_0 = (\text{torso inertia term}) + [\text{Ad}_{1,0}]^{-T} \mathcal{F}_1 + [\text{Ad}_{3,0}]^{-T} \mathcal{F}_3 + [\text{Ad}_{5,0}]^{-T} \mathcal{F}_5 + [\text{Ad}_{8,0}]^{-T} \mathcal{F}_8$$

Still $O(N)$ in the total number of links.

7.11 Contact Forces and External Disturbances

In practice, robots interact with the world. A foot experiences ground reaction forces; a manipulator might push against an object. These external forces $\mathcal{F}_{\text{ext}}$ are incorporated at the point of contact.

7.11.1 Initialization at Leaf Nodes

For a leaf node (a link with no children) experiencing an external wrench $\mathcal{F}_{\text{ext}}$:

$$\mathcal{F}_{\text{leaf}} = \mathcal{I}_{\text{leaf}} \dot{\mathcal{V}}_{\text{leaf}} + [\mathcal{V}_{\text{leaf}} \times]^* (\mathcal{I}_{\text{leaf}} \mathcal{V}_{\text{leaf}}) + \mathcal{F}_{\text{ext}}$$

For an intermediate link i with one child j :

$$\mathcal{F}_i = \mathcal{I}_i \dot{\mathcal{V}}_i + [\mathcal{V}_i \times]^* (\mathcal{I}_i \mathcal{V}_i) + [\text{Ad}_{\mathbf{T}_{j,i}}]^{-T} \mathcal{F}_j + \mathcal{F}_{\text{ext},i}$$

This allows the algorithm to compute the joint torques needed not just to move the robot, but to support external loads.

7.12 RNEA Equivalence to the Lagrangian Equations

The Lagrangian formulation of manipulator dynamics writes the joint torques as

$$\boldsymbol{\tau} = \mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}),$$

where $\mathbf{M}$ is the mass matrix, $\mathbf{C}\dot{\mathbf{q}}$ collects Coriolis and centrifugal terms, and $\mathbf{G}$ is the gravity vector. Computing $\mathbf{M}$ explicitly costs $O(N^3)$; for a 30-DOF humanoid this is prohibitive at real-time rates.

Theorem 7.5. RNEA Computes the Lagrangian Inverse Dynamics

Let $\text{RNEA}(\mathbf{q}, \dot{\mathbf{q}}, \ddot{\mathbf{q}})$ denote the output of the Recursive Newton–Euler Algorithm with gravity acceleration included in the base boundary condition. Then

$$\text{RNEA}(\mathbf{q}, \dot{\mathbf{q}}, \ddot{\mathbf{q}}) = \mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}).$$

That is, the RNEA computes $\boldsymbol{\tau} = \mathbf{M}\ddot{\mathbf{q}} + \mathbf{C}\dot{\mathbf{q}} + \mathbf{G}$ without ever assembling $\mathbf{M}$, $\mathbf{C}$, or $\mathbf{G}$ explicitly.

Proof sketch. Because the RNEA is *linear* in $\ddot{\mathbf{q}}$ (the acceleration appears only in the outward pass and propagates linearly through the inward pass), we can decompose its output via three evaluations:

1. **Gravity vector.** Set $\dot{\mathbf{q}} = \mathbf{0}$ and $\ddot{\mathbf{q}} = \mathbf{0}$:

$$\mathbf{G}(\mathbf{q}) = \text{RNEA}(\mathbf{q}, \mathbf{0}, \mathbf{0}).$$

All velocity-dependent terms vanish; only the gravitational wrench propagates inward.

2. **Coriolis and centrifugal terms.** Set $\ddot{\mathbf{q}} = \mathbf{0}$:

$$\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} = \text{RNEA}(\mathbf{q}, \dot{\mathbf{q}}, \mathbf{0}) - \mathbf{G}(\mathbf{q}).$$

3. **Inertial torques.** The remainder is the mass-matrix contribution:

$$\mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} = \text{RNEA}(\mathbf{q}, \dot{\mathbf{q}}, \ddot{\mathbf{q}}) - \mathbf{C}\dot{\mathbf{q}} - \mathbf{G}.$$

Column-by-column extraction of $\mathbf{M}$. Because the RNEA is linear in $\ddot{\mathbf{q}}$, column j of $\mathbf{M}$ can be obtained as

$$\mathbf{M}(\mathbf{q}) \mathbf{e}_j = \text{RNEA}(\mathbf{q}, \mathbf{0}, \mathbf{e}_j) - \mathbf{G}(\mathbf{q}),$$

where $\mathbf{e}_j$ is the j -th standard basis vector. This gives the full $N \times N$ mass matrix in $N + 1$ calls to the $O(N)$ RNEA, yielding an $O(N^2)$ algorithm for $\mathbf{M}$ —the best possible without exploiting sparsity [12].

In Plain Language 7.3. Why Not Just Build $\mathbf{M}$ Directly?

For a single inverse-dynamics query (“given $\mathbf{q}$, $\dot{\mathbf{q}}$, $\ddot{\mathbf{q}}$, what torques are needed?”), one RNEA call at $O(N)$ suffices. Forming the entire mass matrix $\mathbf{M}$ and then multiplying $\mathbf{M}\ddot{\mathbf{q}}$ would be slower ($O(N^2)$) and wasteful. The Lagrangian form is primarily useful for analysis—stability proofs, passivity arguments—while the RNEA is the workhorse for simulation and control [27].

7.13 Contact Forces and Constraint Handling

Real robots do not float in free space. Walking, grasping, and assembly all involve contact, and modelling contact correctly is one of the hardest problems in robot simulation.

7.13.1 Bilateral Constraints

A *bilateral* constraint removes degrees of freedom permanently. For example, if a foot is bolted to the ground, the spatial velocity of that link is zero. In the RNEA framework this amounts to modifying the base boundary condition:

$$\mathcal{V}_{\text{foot}} = \mathbf{0}, \quad \dot{\mathcal{V}}_{\text{foot}} = \mathbf{0}.$$

The algorithm then propagates *outward* from the pinned foot (which temporarily acts as the base) and the ground reaction wrench $\mathcal{F}_{\text{grf}}$ emerges from the inward pass as the “joint torque” at the virtual fixed joint.

7.13.2 Unilateral Constraints and Friction

Contact with the ground or with objects is *unilateral*: the surface can push but not pull. Denote the contact normal force by f_n and the tangential (friction) force by $\mathbf{f}_t \in \mathbb{R}^2$. The fundamental constraints are:

- **Non-penetration:** the signed gap $\phi \geq 0$.
- **Compressive normal force:** $f_n \geq 0$ (the ground pushes up, never pulls down).
- **Complementarity:** $\phi \cdot f_n = 0$ (force is zero when not in contact, gap is zero when in contact).
- **Coulomb friction cone:** $\|\mathbf{f}_t\| \leq \mu f_n$, where μ is the coefficient of friction.

Definition 7.4. Linear Complementarity Problem (LCP)

Given $\mathbf{A} \in \mathbb{R}^{k \times k}$ and $\mathbf{b} \in \mathbb{R}^k$, find $\mathbf{f} \in \mathbb{R}^k$ such that

$$\mathbf{f} \geq \mathbf{0}, \quad \mathbf{A}\mathbf{f} + \mathbf{b} \geq \mathbf{0}, \quad \mathbf{f}^T(\mathbf{A}\mathbf{f} + \mathbf{b}) = 0.$$

The matrix $\mathbf{A}$ encodes the coupled dynamics of all contacts; $\mathbf{b}$ captures the “free” accelerations from the RNEA in the absence of contact forces [12].

Modern physics engines—MuJoCo [43], Bullet, and Drake—solve a variant of the LCP (or a convex relaxation of it) at every simulation timestep. The RNEA supplies the $O(N)$ inverse dynamics needed to build the LCP data, making the overall cost dominated by the number of active contacts, not the number of joints.

In Plain Language 7.4. Why Contact is Hard

Contact is the hardest part of robot simulation. A walking robot must solve, at every millisecond: Am I touching the ground? How hard? In what direction? Is the foot sliding or stuck? All of these questions couple together—the force depends on the motion, but the motion depends on the force—and must be answered simultaneously within one simulation timestep.

7.14 Worked Example: 3-Link Planar Arm

Let us trace through a complete RNEA computation for a 3-link planar arm in 2D (which we embed in 3D with z -axis rotation).

Remark 7.1 (The torques below are computed). The result is produced by `src/affine_control/rnea.py` via `scripts/generate_worked_examples.py`, and cross-checked against an independent Lagrangian computation; CI fails if the committed values drift from what the code emits.

7.14.1 Setup

Links: Base (link 0), Upper arm (link 1), Forearm (link 2), Hand (link 3).

- Link 1: mass $m_1 = 1$ kg, length $L_1 = 1$ m, CoM at 0.5 m from joint, $I_z^{(1)} = 0.083$ kg·m².
- Link 2: mass $m_2 = 0.5$ kg, length $L_2 = 0.8$ m, CoM at 0.4 m from joint, $I_z^{(2)} = 0.027$ kg·m².
- Link 3: mass $m_3 = 0.2$ kg, point mass at tip (0.6 m from joint), $I_z^{(3)} = 0$.

Joint positions: $\mathbf{q} = [30^\circ, 45^\circ, 0^\circ]^T = [0.524, 0.785, 0]$ rad. Joint velocities: $\dot{\mathbf{q}} = [0.5, 0.3, 0.2]$ rad/s. Joint accelerations: $\ddot{\mathbf{q}} = [1.0, 0.5, 0.1]$ rad/s².

7.14.2 Outward Pass: Velocities and Accelerations

- Base link: $\mathcal{V}_0 = (0, 0, 0, 0, 0, 0)^T$, $\dot{\mathcal{V}}_0 = (0, 0, 0, 0, 0, -9.81)^T$ (gravity down the z axis).

- Link 1: the joint axis is $\mathcal{S}_1 = (0, 0, 1, 0, 0, 0)^T$, so

$$\mathcal{V}_1 = \mathcal{V}_0 + \mathcal{S}_1 \dot{q}_1 = (0, 0, 0.5, 0, 0, 0)^T \quad (7.17)$$

$$\mathcal{V}_1 \times \mathcal{S}_1 \dot{q}_1 = 0. \quad (7.18)$$

That product vanishes here, and it vanishes for *every* link of a planar chain: $\mathcal{V}_i$ and $\mathcal{S}_i \dot{q}_i$ both lie along the same out-of-plane axis, and the cross product of parallel vectors is zero. The bias acceleration in a planar arm comes entirely from the centripetal $\omega^2 r$ terms that appear when the recursion moves the origin from one link to the next.

- Links 2 and 3: the same recursion, with a nonzero translation between link frames, which is where those centripetal terms enter.

Common Misconception

An earlier revision gave $\mathcal{V}_1 \times \mathcal{S}_1 \dot{q}_1 = (0, 0, 0.5) \times (0, 0, 0.2) = (0.1, 0, 0)$, marked “(approx)”. The product is exactly zero, since the vectors are parallel; no approximation is involved. The spurious $(0.1, 0, 0)$ was then carried into $\dot{\mathcal{V}}_1$, and three-vectors were added to six-vectors in the same line.

7.14.3 Inward Pass: Forces and Torques

Assuming no external forces at the tip ($\mathcal{F}_{\text{ext}} = 0$):

- Link 3 (tip): $\mathcal{F}_3 = \mathcal{I}_3 \dot{\mathcal{V}}_3 + [\mathcal{V}_3 \times]^* (\mathcal{I}_3 \mathcal{V}_3)$.
- Link 2: $\mathcal{F}_2 = (\text{inertia and bias of link 2}) + [\text{Ad}_{3,2}]^{-T} \mathcal{F}_3$.
- Link 1: $\mathcal{F}_1 = (\text{inertia and bias of link 1}) + [\text{Ad}_{2,1}]^{-T} \mathcal{F}_2$.

The required joint torques are:

$$\tau_1 = \mathcal{F}_1^T \mathcal{S}_1 = \mathcal{F}_1^T (0, 0, 1, 0, 0, 0)^T \quad (7.19)$$

$$\tau_2 = \mathcal{F}_2^T \mathcal{S}_2 \quad (7.20)$$

$$\tau_3 = \mathcal{F}_3^T \mathcal{S}_3 \quad (7.21)$$

Running the recursion at the stated configuration, velocity and acceleration gives

$$\boldsymbol{\tau} = (13.904, 2.408, 0.670)^T \text{ N}\cdot\text{m},$$

the torques the three joint motors must produce to execute this trajectory.

The shoulder torque dominates, at 13.904 N·m against 0.670 N·m at the wrist. That is mostly gravity: the shoulder carries the weight of the whole arm at a moment arm of roughly a metre, while the wrist carries only the 0.2 kg hand.

Common Misconception

An earlier revision of this section omitted the arithmetic — “detailed numerical computation omitted for brevity” — and then stated $\boldsymbol{\tau} = [2.3, 1.1, 0.08]^T$ prefaced by “the output *might* be”, immediately followed by “These are the torques the three joint motors must produce”.

The stipulated values are not close. The shoulder torque is off by roughly a factor of six. A worked example that hedges with “might be” and then asserts the result as fact is worse than one that shows no numbers at all, because a reader checking their own implementation against it will conclude their code is broken.

7.15 Validation: RNEA vs. Lagrangian

To verify the RNEA is correct, we compare its output to the direct Lagrangian formulation for a simple 2-link planar arm.

7.15.1 2-Link Planar Arm

System: Two identical links, each mass $m = 1$ kg, length $L = 1$ m, uniform density.

Joint positions: $\mathbf{q} = [\theta_1, \theta_2]^T$.

The kinetic energy is:

$$T = \frac{1}{2}m(\dot{x}_{\text{CoM},1}^2 + \dot{y}_{\text{CoM},1}^2) + \frac{1}{2}I_1\dot{\theta}_1^2 + \frac{1}{2}m(\dot{x}_{\text{CoM},2}^2 + \dot{y}_{\text{CoM},2}^2) + \frac{1}{2}I_2(\dot{\theta}_1 + \dot{\theta}_2)^2$$

where $I_1 = I_2 = \frac{1}{12}mL^2$ (uniform rod), and the CoM velocities depend on the joint angles and velocities through kinematics.

Expanding and simplifying (details omitted), the mass matrix is:

$$\mathbf{M}(q) = \begin{bmatrix} m_1 + m_2 + 2m_2L \cos(\theta_2) & m_2 + m_2L \cos(\theta_2) \\ m_2 + m_2L \cos(\theta_2) & m_2 \end{bmatrix}$$

The Coriolis matrix contains terms like $-m_2L \sin(\theta_2)\dot{\theta}_1$, etc.

Test case: $\mathbf{q} = [0, 0]$ (horizontal), $\dot{\mathbf{q}} = [1, 0.5]$ rad/s, $\ddot{\mathbf{q}} = [0.5, 0.2]$ rad/s².

Lagrangian: Substitute into $\mathbf{M}(q)\ddot{\mathbf{q}} + \mathbf{C}(q, \dot{\mathbf{q}})\dot{\mathbf{q}} + \mathbf{G}(q) = \boldsymbol{\tau}$, solve for $\boldsymbol{\tau}$.

RNEA: Run the algorithm with the same inputs.

Result: The torques match to numerical precision, confirming the algorithm.

This check is executed rather than asserted. For the 3-link arm of Section 7.14, the recursion gives $(13.904, 2.408, 0.670)^T$ N·m and the Lagrangian route gives $(13.904, 2.408, 0.670)^T$ N·m, agreeing to better than $1e-08$ N·m — a bound the generator enforces, since the residual itself is a difference of nearly-equal floating-point numbers and its exact value depends on the machine.

The agreement is evidence because the two routes share no machinery. The recursion never forms a mass matrix; it propagates accelerations outward and forces inward, link by link. The Lagrangian route never recurses; it assembles $M(\mathbf{q})$ from centre-of-mass Jacobians and obtains the Coriolis term by differentiating it. A transcription error in either would show up here, which is exactly what a validation section is for — and is why stating the agreement without running it defeats the purpose.

7.16 Adjoint Matrix Derivation From Geometry

We derived the Adjoint matrix above through velocity transformation. Here we provide a deeper geometric interpretation.

7.16.1 Infinitesimal Rotations in $SE(3)$

The Lie algebra $\mathfrak{se}(3)(3)$ consists of 4×4 matrices of the form:

$$\hat{\xi} = \begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ \mathbf{0} & 0 \end{bmatrix}$$

representing infinitesimal twists (rotation $\boldsymbol{\omega}$ and translation $\mathbf{v}$).

The exponential map $\exp(\hat{\xi}t)$ generates a one-parameter family of poses in $SE(3)$.

7.16.2 Adjoint as Conjugation

If $\mathbf{T} \in SE(3)$ and $\mathbf{S} \in SE(3)$, the conjugation $\mathbf{TST}^{-1}$ maps elements of the Lie algebra. The induced action on 6D vectors is the Adjoint:

$$[\text{Ad}_{\mathbf{T}}] : \mathfrak{se}(3)(3) \rightarrow \mathfrak{se}(3)(3), \quad \hat{\xi} \mapsto \mathbf{T}\hat{\xi}\mathbf{T}^{-1} \quad (7.22)$$

In matrix form, this action is represented by the 6×6 matrix $[\text{Ad}_{\mathbf{T}}]$.

7.17 Python Implementation of RNEA

```

1 import numpy as np
2
3 def skew(v):
4     """Convert 3D vector to skew-symmetric matrix."""
5     return np.array([
6         [0, -v[2], v[1]],
7         [v[2], 0, -v[0]],
8         [-v[1], v[0], 0]
9     ])
10
11 def adjoint(T):
12     """Compute 6x6 Adjoint matrix from 4x4 SE(3) transform."""
13     R = T[:3, :3]
14     p = T[:3, 3]
15     Ad = np.zeros((6, 6))
16     Ad[:3, :3] = R
17     Ad[3:, :3] = skew(p) @ R
18     Ad[3:, 3:] = R
19     return Ad
20
21 def motion_cross(V):
22     """Spatial cross product for motion (velocity)."""
23     omega = V[:3]
24     v = V[3:]
25     Vx = np.zeros((6, 6))
26     Vx[:3, :3] = skew(omega)
27     Vx[3:, :3] = skew(v)

```

```

28     Vx[3:, 3:] = skew(omega)
29     return Vx
30
31 def force_cross(V):
32     """Spatial cross product for forces (Wrenches)."""
33     return -motion_cross(V).T
34
35 def spatial_inertia_6d(m, c, I_c):
36     """Build 6x6 spatial inertia from mass, CoM offset, and inertia tensor."""
37     I_spatial = np.zeros((6, 6))
38     I_spatial[:3, :3] = I_c - m * (skew(c) @ skew(c))
39     I_spatial[:3, 3:] = m * skew(c)
40     I_spatial[3:, :3] = -m * skew(c)
41     I_spatial[3:, 3:] = m * np.eye(3)
42     return I_spatial
43
44 def rnea(q, dq, ddq, links, joints, gravity=np.array([0, 0, -9.81])):
45     """
46     Recursive Newton-Euler Algorithm.
47
48     Args:
49         q: joint positions (N,)
50         dq: joint velocities (N,)
51         ddq: joint accelerations (N,)
52         links: list of dicts with 'mass', 'com_offset', 'inertia'
53         joints: list of dicts with 'type' ('revolute' or 'prismatic'),
54                'axis', 'parent', 'T_local'
55         gravity: gravitational acceleration vector
56
57     Returns:
58         tau: required joint torques (N,)
59     """
60     N = len(links)
61
62     # Base link has zero velocity; gravity enters as a fictitious
63     # base acceleration (Featherstone convention: a_0 = -g so that
64     # the algorithm produces correct gravity-compensating torques).
65     V = [np.zeros(6) for _ in range(N)]
66     # V[0] remains zero -- the base is stationary
67
68     dV = [np.zeros(6) for _ in range(N)]
69     dV[0] = np.concatenate([np.zeros(3), -gravity])
70
71     # Outward pass: compute velocities and accelerations
72     T = [np.eye(4) for _ in range(N)]
73     for i in range(1, N):
74         j = joints[i]['parent']
75
76         # Compute relative transform
77         axis = joints[i]['axis']
78         if joints[i]['type'] == 'revolute':
79             S = np.concatenate([axis, np.zeros(3)])
80             # Rotation matrix from rotation axis and angle
81             angle = q[i]
82             K = skew(axis)
83             R = np.eye(3) + np.sin(angle) * K + (1 - np.cos(angle)) * (K @ K)
84             T_rel = np.eye(4)
85             T_rel[:3, :3] = R
86         else: # prismatic

```

```

87         S = np.concatenate([np.zeros(3), axis])
88         T_rel = np.eye(4)
89         T_rel[:3, 3] = q[i] * axis
90
91     T[i] = T[j] @ T_rel
92     Ad = adjoint(T_rel)
93
94     # Velocity
95     V[i] = Ad @ V[j] + S * dq[i]
96
97     # Acceleration (including Coriolis)
98     dV[i] = Ad @ dV[j] + motion_cross(V[i]) @ S * dq[i] + S * ddq[i]
99
100 # Inward pass: compute forces and extract torques
101 F = [np.zeros(6) for _ in range(N)]
102 tau = np.zeros(N)
103
104 for i in range(N - 1, -1, -1):
105     # Link inertia
106     I = spatial_inertia_6d(
107         links[i]['mass'],
108         links[i]['com_offset'],
109         links[i]['inertia']
110     )
111
112     # Inertial wrench
113     F[i] = I @ dV[i] + force_cross(V[i]) @ (I @ V[i])
114
115     # Add wrenches from children
116     for k in range(i + 1, N):
117         if joints[k]['parent'] == i:
118             T_rel = np.eye(4)
119             # Recompute T_rel (simplified; in real code, cache this)
120             Ad_inv_T = adjoint(np.linalg.inv(T_rel)).T
121             F[i] += Ad_inv_T @ F[k]
122
123     # Extract joint torque
124     axis = joints[i]['axis'] if i > 0 else np.array([0, 0, 1])
125     if i > 0 and joints[i]['type'] == 'revolute':
126         S = np.concatenate([axis, np.zeros(3)])
127     elif i > 0:
128         S = np.concatenate([np.zeros(3), axis])
129     else:
130         S = np.array([0, 0, 1, 0, 0, 0])
131
132     tau[i] = F[i] @ S
133
134 return tau

```

Listing 7.1: RNEA Implementation in Python

7.18 Summary

In this chapter, we learned:

1. **Historical Context:** The RNEA, developed by Orin, McGhee, Luh, and formalized by Featherstone, reduced robot dynamics computation from $O(N^3)$ to $O(N)$, making real-time control possible.

2. **Kinematic Trees:** A parent array $\lambda(i)$ describes the acyclic structure of multi-body systems, including branching topologies.
3. **Adjoint Matrix:** The 6×6 Adjoint $[\text{Ad}_{\mathbf{T}}]$ transforms spatial velocities between frames; wrenches transform via the coadjoint (negative transpose).
4. **Recursive Forward Kinematics:** Velocities propagate outward via $\mathcal{V}_i = [\text{Ad}_{\mathbf{T}_{i,\lambda(i)}}]\mathcal{V}_{\lambda(i)} + \mathcal{S}_i\dot{q}_i$.
5. **Spatial Acceleration:** Includes Coriolis and centrifugal terms via the motion cross product $[\mathcal{V} \times]$.
6. **RNEA:** Two passes—outward (velocities, accelerations) and inward (forces, torques)—compute inverse dynamics in $O(N)$ time.
7. **Branching Trees:** Multiple children are handled by summing child wrenches in the inward pass.
8. **Contact Forces:** External forces are incorporated at leaf nodes.
9. **Complexity:** Strict $O(N)$ proof makes the algorithm scalable.
10. **Implementation:** Python code shows practical RNEA execution.

Throughout the subsequent volumes of *The Geometry of Motion*, physical phenomena like Contraction Metrics, Sums-of-Squares Programming, and LQR-Trees are discussed dynamically. Behind the curtain of those advanced trajectory optimizations, the physics engines governing the simulated bodies are invariably relying on Quaternions to map their angles, Adjoints to shift their coordinates, and Recursive Algorithms to execute their algebraic realities in $O(N)$ time.

7.19 Further Reading

The original computational Newton-Euler formulation for robot manipulators was presented by Luh, Walker, and Paul [22], who demonstrated that recursive structure could reduce the cost of inverse dynamics from $O(N^3)$ to $O(N)$. Their paper remains a landmark in computational robotics and is essential reading for understanding the historical development of real-time robot control.

The definitive modern treatment of recursive dynamics is Featherstone [12], which develops the RNEA, the Articulated Body Algorithm (ABA), and the Composite Rigid Body Algorithm within a unified spatial vector framework. Featherstone’s monograph is the primary reference for anyone implementing multi-body dynamics from scratch and contains detailed pseudocode, complexity proofs, and numerical considerations that go well beyond our presentation here.

Lynch and Park [27], Chapter 8, present inverse dynamics within the product of exponentials framework, offering a complementary geometric perspective to Featherstone’s spatial vector approach. Their treatment emphasizes the connection between the Lie-algebraic formulation of kinematics (Chapters 5–6) and the recursive dynamics algorithms developed here.

For practical implementation considerations—including numerical conditioning, joint friction models, and real-time scheduling—Siciliano et al. [38] provide an engineering-oriented treatment with extensive discussion of industrial applications.

The forward dynamics problem, which is the computational inverse of the RNEA, requires the more sophisticated Articulated Body Algorithm developed in Chapter 10. Readers eager to understand how physics engines like MuJoCo, Drake, and PyBullet actually simulate forward in time should proceed to that chapter after mastering the RNEA presented here.

7.20 Exercises

7.20.1 Tier 1 (Foundation)

1. Show that the Adjoint matrix $[\text{Ad}_{\mathbf{T}}]$ is invertible and compute its inverse in closed form.
2. For a 2-link planar arm, write out the parent array λ and the screw axes $\mathcal{S}_1, \mathcal{S}_2$ for two revolute joints.
3. Compute the Coriolis term $[\mathcal{V}_i \times] \mathcal{S}_i \dot{q}_i$ for a simple 1-link revolute joint with $\mathcal{V} = (0, 0, \omega, v_x, 0, 0)^T$ and $\mathcal{S} = (0, 0, 1, 0, 0, 0)^T$.
4. Verify that $[\text{Ad}_{\mathbf{T}}]^{-T}$ transforms wrenches correctly: show that power $\mathcal{F}^T \mathcal{V}$ is invariant under frame changes.
5. For a 3-link arm with known link masses and dimensions, compute the spatial inertia matrices using the parallel axis theorem.

7.20.2 Tier 2 (Intermediate)

1. Derive the spatial Newton-Euler equation from first principles using Lagrangian mechanics and show that the bias term $[\mathcal{V} \times]^*(\mathcal{I}\mathcal{V})$ captures centrifugal and Coriolis forces.
2. For a 2-link arm with a branching third link (e.g., an end-effector with a gripper that has one degree of freedom), write out the full RNEA inward pass, including the summation of forces from the gripper to the main arm.
3. Implement the motion cross product $[\mathcal{V} \times]$ and force cross product $[\mathcal{V} \times]^*$ numerically and verify they are consistent (i.e., $[\mathcal{V} \times]^* = -[\mathcal{V} \times]^T$).
4. Compare the computational time of the RNEA ($O(N)$) to matrix inversion ($O(N^3)$) for systems with $N = 5, 10, 20, 50$ joints, both analytically and experimentally.
5. Extend the RNEA algorithm to handle closed-loop mechanisms (systems with cycles), discussing the additional constraints required.

7.20.3 Tier 3 (Advanced)

1. Prove that the Adjoint map is a Lie group homomorphism: $[\text{Ad}_{\mathbf{T}_1 \mathbf{T}_2}] = [\text{Ad}_{\mathbf{T}_1}][\text{Ad}_{\mathbf{T}_2}]$.

2. Derive the relationship between the Adjoint matrix and the differential of the exponential map $\exp : \mathfrak{se}(3)(3) \rightarrow SE(3)$.
3. For a humanoid robot with 30 DOFs, design the kinematic tree structure (parent array), define all joint axes, and outline how the RNEA would compute inverse dynamics.
4. Implement the composite rigid body algorithm (CRB), which pre-computes the mass matrix $\mathbf{M}(q)$ in $O(N)$ time using RNEA-like recursion, and compare its performance to direct computation.
5. Extend the RNEA to handle soft contacts (e.g., compliant ground) by modifying the boundary conditions at contact points and verifying numerical stability.

Chapter 8

Spatial Vector Algebra and Inertia

Why manage six tiny numbers loosely floating in space when you can forge them into an unbreakable, 6-dimensional spear?

In Plain Language 8.1. Context & Use Case: Unifying Translation and Rotation

The Math You Will See: We introduce **Spatial Vector Algebra**, defining the 6×6 **Spatial Cross Product** ($[\mathcal{V} \times]$), the **Spatial Inertia Matrix** ($\mathcal{I}$), and their role in transforming forces and accelerations across frames.

Why We Need It: In classical 3D mechanics, linear motion ($F = ma$) and angular motion ($\tau = \mathbf{I}\alpha + \omega \times \mathbf{I}\omega$) are computed separately, leading to a tangled web of cross products and messy offsets when bodies are connected together. Spatial Algebra fuses them. A "Spatial Vector" describes pushing and turning simultaneously. By doing math in 6D, the complex gyroscopic and Coriolis forces of a multi-joint robot swinging through space are resolved with a single matrix multiplication.

All Variables Defined: $\mathcal{V}$ is a spatial velocity (Twist), $\mathcal{F}$ is a spatial force (Wrench), and $\mathcal{I}$ is the combined 6D mass/inertia tensor.


8.1 Historical Context: From Plücker to Featherstone

The roots of spatial algebra reach deep into classical geometry. In the 1860s, **Julius Plücker** studied lines in 3D space as geometric objects, discovering that a line can be represented as a 6-dimensional vector satisfying certain constraints. This **Plücker coordinate** representation remained mostly a curiosity in algebraic geometry for over a century.

In 1900, Sir **Robert Stawell Ball** [2] developed the theory of “screws”—geometric objects unifying rotation and translation—and showed that Plücker’s coordinates naturally describe instantaneous screw motions. Ball’s work connected to mechanics through the wrench (force + torque) dual to the twist (angular + linear velocity).

However, the true synthesis came when **Roy Featherstone**, in the 1980s and 1990s, recognized that Ball’s geometric screws and Plücker’s algebraic coordinates could be married with Lie algebra to create a unified computational framework for multi-body dynamics. Featherstone’s **Spatial Vector Algebra** [12] standardized the representation, showed how to compose inertias of linked chains (composite rigid body algorithm), and proved that careful use of spatial algebra enables $O(N)$ inverse dynamics.

Modern robotics and physics engines (MuJoCo [43], Drake, Pinocchio) are built entirely on Featherstone's spatial framework. The companion platform accompanying this textbook builds on the same foundation.



Spatial vectors unify translation (linear) and rotation (angular) quantities

Figure 8.1: A Spatial Wrench Is a 6-Dimensional Vector Combining Linear Forces and Torques (Moments). Similarly, Spatial Twists Combine Linear and Angular Velocities. This Unification Allows Forces and Velocities to Be Manipulated Using Unified Algebraic Operations.

8.2 Plücker Coordinates and Lines in 3D

Before we discuss spatial vectors as physical quantities, it helps to understand their geometric origin: the Plücker representation of lines.

Definition 8.1. Plücker Coordinates of a Line

A line ℓ in 3D can be represented as a 6-vector in **Plücker coordinates**:

$$\mathbf{L} = \begin{bmatrix} \mathbf{d} \\ \mathbf{m} \end{bmatrix} \in \mathbb{P}^5 \quad (8.1)$$

where:

- $\mathbf{d}$ is the direction vector of the line (unit vector).
- $\mathbf{m} = \mathbf{r} \times \mathbf{d}$ is the moment (or Plücker moment), where $\mathbf{r}$ is any point on the line.

The Plücker coordinates satisfy the **Plücker relation**:

$$\mathbf{d}^T \mathbf{m} = 0 \quad (8.2)$$

This constraint encodes the fact that the direction and moment are related geometrically.

Geometric Insight: The direction $\mathbf{d}$ specifies the "orientation" of the line; the moment $\mathbf{m}$ encodes the line's location in space. Two lines are identical if and only if they have the same Plücker coordinates.

8.2.1 Connection to Spatial Vectors

A spatial velocity (Twist) $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v})$ can be viewed as the instantaneous motion of a rigid body. The axis of this motion is a screw (a line with a pitch), and the Plücker coordinates of that screw are exactly $(\boldsymbol{\omega}, \mathbf{v})$ (up to scaling).

Similarly, a spatial force (Wrench) $\mathcal{F} = (\mathbf{m}, \mathbf{f})$ is the dual concept: it represents a force $\mathbf{f}$ applied at some point in space, creating a torque $\mathbf{m}$ about the origin. The Plücker relationship $\mathbf{f}^T \mathbf{m} = 0$ holds when the force and moment arise from a single applied force (they lie on the same screw axis).

8.3 Spatial Vectors: Twists and Wrenches

In Chapter 6, we introduced spatial vectors as elements of the Lie algebras $\mathfrak{se}(3)(3)$ and $\mathfrak{se}(3)^*(3)$. Here we formalize them as vectors and define their algebra.

Definition 8.2. Spatial Velocity (Twist)

A spatial velocity or Twist is a 6D vector:

$$\mathcal{V} = \begin{bmatrix} \boldsymbol{\omega} \\ \mathbf{v} \end{bmatrix} \in \mathfrak{se}(3)(3) \cong \mathbb{R}^6 \quad (8.3)$$

where:

- $\boldsymbol{\omega} \in \mathbb{R}^3$ is the angular velocity.
- $\mathbf{v} \in \mathbb{R}^3$ is the linear velocity of the origin of the reference frame.

Physical Interpretation: A point at position $\mathbf{r}$ relative to the origin has absolute velocity $\mathbf{v} + \boldsymbol{\omega} \times \mathbf{r}$.

Definition 8.3. Spatial Force (Wrench)

A spatial force or Wrench is a 6D vector:

$$\mathcal{F} = \begin{bmatrix} \mathbf{m} \\ \mathbf{f} \end{bmatrix} \in \mathfrak{se}(3)^*(3) \cong \mathbb{R}^6 \quad (8.4)$$

where:

- $\mathbf{m} \in \mathbb{R}^3$ is the moment (torque) about the origin.
- $\mathbf{f} \in \mathbb{R}^3$ is the linear force.

Physical Interpretation: A force $\mathbf{f}$ applied at a point $\mathbf{r}$ creates a moment $\mathbf{m} + \mathbf{r} \times \mathbf{f}$ about the origin.

8.4 The Spatial Cross Product

In classical 3D mechanics, cross products appear everywhere. When a body rotates with angular velocity $\boldsymbol{\omega}$, the velocity of a point at offset $\mathbf{r}$ includes the term $\boldsymbol{\omega} \times \mathbf{r}$. When a force $\mathbf{f}$ is applied at point $\mathbf{r}$, it creates a torque $\mathbf{r} \times \mathbf{f}$.

Spatial Algebra extends the cross product to 6D, but because Twists and Wrenches live in dual spaces (with different transformation laws), there are actually *two* distinct spatial cross products.

8.4.1 Motion Cross Product

The **motion cross product** acts on Twists and related vectors (velocities, accelerations). If a frame has spatial velocity $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v})$, a point with position offset $\mathbf{r}$ has velocity $\mathbf{v} + \boldsymbol{\omega} \times \mathbf{r}$.

We represent this as a matrix operator:

Definition 8.4. Motion Cross Product Operator

For a spatial velocity Twist $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v})$, the motion cross product operator is:

$$[\mathcal{V} \times] = \begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{0}_{3 \times 3} \\ [\mathbf{v}] & [\boldsymbol{\omega}] \end{bmatrix} \in \mathbb{R}^{6 \times 6} \quad (8.5)$$

where $[\boldsymbol{\omega}]$ and $[\mathbf{v}]$ are the standard 3×3 skew-symmetric cross product matrices.

Action on a spatial velocity: If $\mathcal{W} = (\boldsymbol{\alpha}, \mathbf{a})$ is another spatial quantity, then:

$$[\mathcal{V} \times] \mathcal{W} = \begin{bmatrix} \boldsymbol{\omega} \times \boldsymbol{\alpha} \\ \mathbf{v} \times \boldsymbol{\alpha} + \boldsymbol{\omega} \times \mathbf{a} \end{bmatrix} \quad (8.6)$$

8.4.2 Force Cross Product

Wrenches are the dual to Twists. Because of this duality, when we apply the adjoint transformation to a wrench, it uses a different matrix. Correspondingly, the force cross product is defined differently:

Definition 8.5. Force Cross Product Operator

The force cross product (coadjoint cross product) for a spatial velocity $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v})$ acting on Wrenches is:

$$[\mathcal{V} \times^*] = -[\mathcal{V} \times]^T = \begin{bmatrix} [\boldsymbol{\omega}] & [\mathbf{v}] \\ \mathbf{0}_{3 \times 3} & [\boldsymbol{\omega}] \end{bmatrix} \quad (8.7)$$

Action on a spatial force: If $\mathcal{G} = (\mathbf{m}, \mathbf{f})$ is a Wrench, then:

$$[\mathcal{V} \times^*] \mathcal{G} = \begin{bmatrix} \boldsymbol{\omega} \times \mathbf{m} + \mathbf{v} \times \mathbf{f} \\ \boldsymbol{\omega} \times \mathbf{f} \end{bmatrix} \quad (8.8)$$

8.4.3 Why Two Cross Products?

The reason for the asymmetry is that power (work rate) must be invariant:

$$\text{Power} = \mathcal{F}^T \mathcal{V} \quad (8.9)$$

Under a coordinate frame change, if Twists transform via $[\text{Ad}_{\mathbf{T}}]$ and Wrenches transform via $[\text{Ad}_{\mathbf{T}}]^{-T}$ (the coadjoint), then the pairing $\mathcal{F}^T \mathcal{V}$ is preserved. The cross product operators must respect this dual structure.

Mathematical Principle 8.1. Duality of Motion and Force Cross Products

The motion and force cross products are dual in the sense that:

$$\mathcal{G}^T([\mathcal{V} \times] \mathcal{W}) = ([\mathcal{V} \times^*] \mathcal{G})^T \mathcal{W} \quad (8.10)$$

for any Twist $\mathcal{V}$, Twist-like vector $\mathcal{W}$, and Wrench $\mathcal{G}$. This ensures that the bilinear form (power) is preserved under transformations.

Mathematical Principle 8.2. Ad vs. ad: Two Distinct Operators

The notation in spatial algebra distinguishes two fundamentally different operators that are easily confused:

- $[\text{Ad}_{\mathbf{T}}] \in \mathbb{R}^{6 \times 6}$: The **Adjoint** of a rigid body transform $\mathbf{T} \in \text{SE}(3)(3)$. This is a coordinate transformation that re-expresses a spatial velocity (or acceleration) measured in one frame into another frame. It depends on a *pose* (position and orientation).
- $[\text{ad}_{\mathcal{V}}] = [\mathcal{V} \times] \in \mathbb{R}^{6 \times 6}$: The **adjoint** (lowercase), also called the *Lie bracket* or *spatial motion cross product*. This maps a spatial velocity $\mathcal{V}$ to its cross-product operator, which computes Coriolis and centrifugal accelerations. It depends on a *velocity*, not a pose.

In recursive dynamics algorithms such as the RNEA, both operators appear: $[\text{Ad}_{\mathbf{T}}]$ propagates velocities across frames (a coordinate change), while $[\text{ad}_{\mathcal{V}}] = [\mathcal{V} \times]$ computes the velocity-dependent acceleration terms (Coriolis). Confusing the two produces incorrect dynamics [12, 31].

8.5 The Spatial Inertia Matrix

Just as scalar mass m resists linear acceleration ($F = ma$), and a 3×3 inertia tensor $\mathbf{I}_c$ resists angular acceleration ($\boldsymbol{\tau} = \mathbf{I}_c \boldsymbol{\alpha}$), Spatial Algebra unifies them into a single 6×6 **Spatial Inertia Matrix** $\mathcal{I}$.

8.5.1 Definition and Properties

Definition 8.6. Spatial Inertia Matrix

For a rigid body with:

- Mass m ,
- Center of mass (CoM) at position $\mathbf{c}$ relative to the reference frame origin,
- Rotational inertia tensor $\mathbf{I}_c$ (about the CoM),

the Spatial Inertia Matrix expressed at the frame origin is:

$$\mathcal{I} = \begin{bmatrix} \mathbf{I}_c - m[\mathbf{c}][\mathbf{c}] & m[\mathbf{c}] \\ -m[\mathbf{c}] & m\mathbf{I}_{3 \times 3} \end{bmatrix} \quad (8.11)$$

Using the identity $[\mathbf{c}][\mathbf{c}] = -\mathbf{c}\mathbf{c}^T + (\mathbf{c}^T\mathbf{c})\mathbf{I}_{3 \times 3}$, this can be rewritten as:

$$\mathcal{I} = \begin{bmatrix} \mathbf{I}_c + m(|\mathbf{c}|^2\mathbf{I}_{3 \times 3} - \mathbf{c}\mathbf{c}^T) & m[\mathbf{c}] \\ -m[\mathbf{c}] & m\mathbf{I}_{3 \times 3} \end{bmatrix} \quad (8.12)$$

When the reference frame origin is *at* the center of mass ($\mathbf{c} = \mathbf{0}$), this simplifies to:

$$\mathcal{I}_{\text{CoM}} = \begin{bmatrix} \mathbf{I}_c & \mathbf{0}_{3 \times 3} \\ \mathbf{0}_{3 \times 3} & m\mathbf{I}_{3 \times 3} \end{bmatrix} \quad (8.13)$$

8.5.2 Symmetric and Positive Definite

Theorem 8.1. Spatial Inertia is SPD

The Spatial Inertia Matrix $\mathcal{I}$ is symmetric and positive definite.

Proof (Symmetry): Write the spatial inertia in standard form:

$$\mathcal{I} = \begin{bmatrix} \mathbf{I}_c + m[\mathbf{c}][\mathbf{c}]^T & m[\mathbf{c}] \\ m[\mathbf{c}]^T & m\mathbf{I}_{3 \times 3} \end{bmatrix}.$$

Since $[\mathbf{c}]$ is skew-symmetric ($[\mathbf{c}]^T = -[\mathbf{c}]$), each block is symmetric or correctly transposed:

- Top-left block: $\mathbf{I}_c + m[\mathbf{c}][\mathbf{c}]^T$ is symmetric (both $\mathbf{I}_c$ and the product $[\mathbf{c}][\mathbf{c}]^T$ are symmetric).
- Off-diagonal blocks: The top-right block is $m[\mathbf{c}]$, and the bottom-left block is $m[\mathbf{c}]^T = -m[\mathbf{c}]$. Their transposes satisfy $(m[\mathbf{c}])^T = m[\mathbf{c}]^T$, confirming $\mathcal{I}^T = \mathcal{I}$.
- Bottom-right block: $m\mathbf{I}_{3 \times 3}$ is trivially symmetric.

Proof (Positive Definite): For any non-zero spatial vector $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v})$, we compute:

$$\mathcal{V}^T \mathcal{I} \mathcal{V} = \boldsymbol{\omega}^T (\mathbf{I}_c + m[\mathbf{c}][\mathbf{c}]^T) \boldsymbol{\omega} + 2m\boldsymbol{\omega}^T [\mathbf{c}] \mathbf{v} + m\mathbf{v}^T \mathbf{v}$$

This can be rewritten as the rotational kinetic energy plus the translational kinetic energy:

$$\mathcal{V}^T \mathcal{I} \mathcal{V} = \boldsymbol{\omega}^T \mathbf{I}_c \boldsymbol{\omega} + m|\mathbf{v} + [\mathbf{c}]\boldsymbol{\omega}|^2$$

Since $\mathbf{I}_c$ is positive definite (non-zero $\boldsymbol{\omega}$ implies $\boldsymbol{\omega}^T \mathbf{I}_c \boldsymbol{\omega} > 0$), the entire expression is positive whenever $\mathcal{V} \neq 0$.

8.6 The Parallel Axis Theorem in 6D

8.6.1 Derivation From First Principles

The classical parallel axis theorem states that if the inertia tensor about the center of mass is $\mathbf{I}_c$, then the inertia tensor about a point at distance $\mathbf{c}$ from the CoM is:

$$\mathbf{I}_{\text{origin}} = \mathbf{I}_c + m(c^2 \mathbf{I}_{3 \times 3} - \mathbf{c} \mathbf{c}^T)$$

In spatial form, the parallel axis theorem extends to both translational and rotational terms. If we define:

$\mathcal{I}_0$ = Spatial inertia at origin

$\mathcal{I}_c$ = Spatial inertia at CoM

then the relationship is:

Theorem 8.2. Parallel Axis Theorem in 6D

$$\mathcal{I}_0 = \begin{bmatrix} \mathbf{R} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \begin{bmatrix} \mathbf{I}_c & \mathbf{0} \\ \mathbf{0} & m \mathbf{I}_{3 \times 3} \end{bmatrix} \begin{bmatrix} \mathbf{R}^T & \mathbf{0} \\ \mathbf{0} & \mathbf{R}^T \end{bmatrix} + \begin{bmatrix} -m[\mathbf{c}][\mathbf{c}]^T & m[\mathbf{c}]^T \\ m[\mathbf{c}] & \mathbf{0} \end{bmatrix} \quad (8.14)$$

When $\mathbf{R} = \mathbf{I}$ (same orientation):

$$\mathcal{I}_0 = \begin{bmatrix} \mathbf{I}_c + m[\mathbf{c}][\mathbf{c}]^T & m[\mathbf{c}] \\ m[\mathbf{c}]^T & m \mathbf{I}_{3 \times 3} \end{bmatrix} \quad (8.15)$$

Interpretation: The top-left block contains the parallel axis correction to rotational inertia; the off-diagonals encode the coupling between angular and linear motion due to the offset CoM.

8.7 Composite Rigid Body Algorithm

In Chapter 7, we computed inverse dynamics using the RNEA. An alternative and equally important algorithm is the **Composite Rigid Body Algorithm (CRB)**, which builds up the spatial inertia of linked chains recursively.

8.7.1 Concept: Bottom-Up Inertia Assembly

Instead of computing torques directly, the CRB first assembles the inertia of composite structures: e.g., the forearm link is rigidly attached to the hand, so the hand's inertia is

added to the forearm's. By recursively building composite inertias from tips to base, we can pre-compute the total mass matrix $\mathbf{M}(q)$ in $O(N)$ time.

Algorithm 8.1. Composite Rigid Body Algorithm

```

1 function CompositeRigidBody(links, joints, lambda):
2   // Input: link inertias, joint structure
3   // Output: composite inertias I_c for each link
4
5   I_c = [None] * (N+1)      // Will store composite inertia of
6                               link i and all descendants
7
8   for i <-N down to 1:      // Traverse from tips to base
9     I_c[i] <- I[i]          // Start with link i's own inertia
10
11     for j in children[i]:
12       // Add inertia of child j to link i
13       T_j_to_i <- relative transform from j to i
14       Ad_j_to_i <- Adjoint of T_j_to_i
15       I_c[i] <- I_c[i] + Ad_j_to_i^T * I_c[j] * Ad_j_to_i
16
17   return I_c

```

8.7.2 Use in Computing the Mass Matrix

Once the composite inertias are assembled, the mass matrix can be extracted via:

$$\mathbf{M}_{ij}(q) = \mathcal{S}_i^T [\text{Ad}_{\mathbf{T}_{i,j}}]^{-T} \mathcal{I}_{c,j}(q) [\text{Ad}_{\mathbf{T}_{i,j}}]^{-1} \mathcal{S}_j$$

where $\mathcal{I}_{c,j}$ is the composite inertia of link j and all its descendants.

Theorem 8.3. CRB Complexity

The Composite Rigid Body Algorithm computes the mass matrix $\mathbf{M}(q)$ in $O(N^2)$ time (because there are N^2 entries), but the inertia assembly itself (bottom-up recursion) is $O(N)$. Once assembled, forward-dynamics calculations (solving for accelerations given forces) are $O(N)$ via matrix multiplication.

8.8 Spatial Algebra for Planar Systems

For systems constrained to a plane (e.g., 2D manipulation, golf swings), we can work with reduced spatial representations.

8.8.1 3D Vectors for 2D Motion

A planar motion has:

- Angular velocity ω (a scalar, rotation about the z -axis).
- Linear velocity $\mathbf{v} = (v_x, v_y) \in \mathbb{R}^2$.

We can embed this in a 3D twist:

$$\mathcal{V}_{\text{planar}} = (\omega, v_x, v_y)^T \in \mathbb{R}^3$$

The motion cross product for planar motion simplifies:

$$[\mathcal{V} \times]_{\text{planar}} = \begin{bmatrix} 0 & -v_y & v_x \\ v_y & 0 & \omega \\ -v_x & -\omega & 0 \end{bmatrix}$$

A planar spatial inertia matrix is 3×3 :

$$\mathcal{I}_{\text{planar}} = \begin{bmatrix} I_z & 0 & 0 \\ 0 & m & 0 \\ 0 & 0 & m \end{bmatrix}$$

when the CoM is at the origin. With offset $\mathbf{c} = (c_x, c_y)$:

$$\mathcal{I}_{\text{planar}} = \begin{bmatrix} I_z + m(c_x^2 + c_y^2) & mc_y & -mc_x \\ mc_y & m & 0 \\ -mc_x & 0 & m \end{bmatrix}$$

8.8.2 Computational Advantage

Working in 3D (instead of 6D) reduces memory and computation by a factor of 4, making planar mechanisms extremely fast to simulate.

8.9 Newton-Euler Equations in Spatial Form

The fundamental equations of dynamics in spatial form are elegant and reveal the unity of linear and angular motion.

8.9.1 Spatial Momentum and the Equation of Motion

The spatial momentum of a rigid body with spatial inertia $\mathcal{I}$ and spatial velocity $\mathcal{V}$ is:

$$\mathcal{P} = \mathcal{I}\mathcal{V} \tag{8.16}$$

The equation of motion (Newton-Euler in spatial form) states that the time derivative of spatial momentum equals the applied wrench:

$$\frac{d\mathcal{P}}{dt} = \mathcal{F}_{\text{ext}} \tag{8.17}$$

However, because $\mathcal{I}$ itself may be time-varying (when expressed in a rotating frame), the full equation includes a bias term:

$$\frac{d}{dt}(\mathcal{I}\mathcal{V}) = \mathcal{I}\dot{\mathcal{V}} + \dot{\mathcal{I}}\mathcal{V} \tag{8.18}$$

In a *body-fixed frame* (rotating with the body), $\mathcal{I}$ is constant, but velocities are measured in the rotating frame, introducing the bias term $[\mathcal{V} \times]^*(\mathcal{I}\mathcal{V})$:

Theorem 8.4. Spatial Newton-Euler Equation

For a rigid body in a body-fixed frame, the equation of motion is:

$$\mathcal{F} = \mathcal{I}\dot{\mathcal{V}} + [\mathcal{V} \times]^*(\mathcal{I}\mathcal{V}) \quad (8.19)$$

Interpretation:

- First term: the inertial wrench (like $F = ma$ in 6D).
- Second term: the "bias" or "gyroscopic" wrench, capturing Coriolis and centrifugal effects.

Derivation: In an inertial frame:

$$\mathcal{F}_{\text{inertial}} = \frac{d}{dt}(\mathcal{I}_{\text{inertial}}\mathcal{V}_{\text{inertial}})$$

Transforming to a body-fixed frame using the adjoint, the time derivative picks up the Lie bracket (commutator) of spatial velocities, which manifests as $[\mathcal{V} \times]^*$.

8.10 Connection to Dual Quaternions

8.10.1 Brief Overview

Dual quaternions extend quaternions (Section 4.8) by using dual numbers instead of real numbers. A dual quaternion is:

$$\mathbf{q}_d = q_0 + q_1\mathbf{i} + q_2\mathbf{j} + q_3\mathbf{k} + \epsilon(q_4 + q_5\mathbf{i} + q_6\mathbf{j} + q_7\mathbf{k})$$

where ϵ is the dual unit satisfying $\epsilon^2 = 0$ (and $\epsilon \neq 0$).

Dual quaternions are isomorphic to spatial vectors and can represent both rotation and translation in a single object. Some robotics frameworks use dual quaternions for inertia and screw calculations, though the 6×6 matrix representation is more standard for computation.

Key property: The spatial inertia can be represented as a dual-quaternion-valued linear form, and Featherstone's spatial algorithms can be rewritten using dual quaternions, often with elegant geometric interpretations.

For this course, we continue with the matrix representation, as it is more straightforward for algorithmic implementation.

8.11 Worked Example: Spatial Inertia of a Multi-Link System

Consider a two-link arm where:

- Link 1 (upper arm): mass $m_1 = 2$ kg, CoM at $(0.3, 0, 0)$ m from joint 1, rotational inertia $\mathbf{I}_c^{(1)} = \text{diag}(0.05, 0.05, 0.01)$ kg·m².
- Link 2 (forearm): mass $m_2 = 1$ kg, CoM at $(0.25, 0, 0)$ m from joint 2, rotational inertia $\mathbf{I}_c^{(2)} = \text{diag}(0.02, 0.02, 0.005)$ kg·m².

8.11.1 Building Link Spatial Inertias

For link 1, with CoM offset $\mathbf{c}_1 = (0.3, 0, 0)^T$:

$$[\mathbf{c}_1] = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -0.3 \\ 0 & 0.3 & 0 \end{bmatrix}$$

$$[\mathbf{c}_1][\mathbf{c}_1]^T = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0.09 & 0 \\ 0 & 0 & 0.09 \end{bmatrix}$$

$$\mathcal{I}^{(1)} = \begin{bmatrix} \mathbf{I}_c^{(1)} + m_1[\mathbf{c}_1][\mathbf{c}_1]^T & m_1[\mathbf{c}_1] \\ m_1[\mathbf{c}_1]^T & m_1\mathbf{I}_{3 \times 3} \end{bmatrix}$$

$$\mathcal{I}^{(1)} = \begin{bmatrix} 0.05 + 2(0) & 0 & 0 & 0 & 0 & 0 \\ 0 & 0.05 + 2(0.09) & 0 & 0 & 0 & -0.6 \\ 0 & 0 & 0.01 + 2(0.09) & 0 & 0.6 & 0 \\ 0 & 0 & 0 & 2 & 0 & 0 \\ 0 & 0 & 0.6 & 0 & 2 & 0 \\ 0 & -0.6 & 0 & 0 & 0 & 2 \end{bmatrix}$$

$$\mathcal{I}^{(1)} = \begin{bmatrix} 0.05 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0.23 & 0 & 0 & 0 & -0.6 \\ 0 & 0 & 0.19 & 0 & 0.6 & 0 \\ 0 & 0 & 0 & 2 & 0 & 0 \\ 0 & 0 & 0.6 & 0 & 2 & 0 \\ 0 & -0.6 & 0 & 0 & 0 & 2 \end{bmatrix}$$

Note the structure of the correction. Since $[\mathbf{c}][\mathbf{c}]^T = |\mathbf{c}|^2\mathbf{I} - \mathbf{c}\mathbf{c}^T$ is positive semi-definite, the parallel axis term always *adds* to the rotational inertia. With $\mathbf{c}_1 = (0.3, 0, 0)$ the correction vanishes about the offset direction itself (x) and contributes $m_1 d^2 = 2(0.3)^2 = 0.18$ about y and z — exactly the classical parallel axis theorem. A spatial inertia is symmetric positive definite by Theorem 8.1, so a negative diagonal entry would signal an algebraic error, not a physical configuration.

Similarly, $\mathcal{I}^{(2)}$ is computed with $\mathbf{c}_2 = (0.25, 0, 0)$ and $m_2 = 1$ kg.

8.11.2 Composite Inertia

If link 2 is rigidly attached to link 1 (e.g., they are the same rigid body or perfectly welded), the composite inertia is:

$$\mathcal{I}_{\text{composite}} = \mathcal{I}^{(1)} + \mathcal{I}^{(2)}$$

This is simple summation when expressed in the same reference frame.

If the links are separated by a relative transform $\mathbf{T}_{1 \rightarrow 2}$, we must first transform $\mathcal{I}^{(2)}$ to the frame of link 1:

$$\mathcal{I}_{\text{composite}} = \mathcal{I}^{(1)} + [\text{Ad}_{\mathbf{T}_{1 \rightarrow 2}}]^{-T} \mathcal{I}^{(2)} [\text{Ad}_{\mathbf{T}_{1 \rightarrow 2}}]^{-1}$$

8.12 Physical Interpretation of Spatial Inertia Elements

The 6×6 spatial inertia matrix has a clear physical meaning for each block:

- **Top-left** (3×3): Rotational inertia (analogous to the 3×3 inertia tensor). It relates angular acceleration to the torque required.
- **Top-right** (3×3): Coupling between linear and rotational motion due to CoM offset. If the CoM is not at the origin, accelerating the origin in one direction produces a torque around it.
- **Bottom-left** (3×3): Transpose of top-right (due to symmetry).
- **Bottom-right** (3×3): Mass matrix ($m\mathbf{I}_{3 \times 3}$). It relates linear acceleration of the origin to the force required.

8.12.1 Kinetic Energy and the Origin of Spatial Inertia

The 6×6 spatial inertia matrix is not an arbitrary packaging of mass properties. It is the *unique* object that makes kinetic energy a quadratic form in spatial velocity. We now prove this from first principles.

Classical kinetic energy. For a rigid body with mass m , centre-of-mass velocity $\mathbf{v}_{cm}$, angular velocity $\boldsymbol{\omega}$, and rotational inertia $\mathbf{I}_c$ about the centre of mass:

$$T = \frac{1}{2} m \|\mathbf{v}_{cm}\|^2 + \frac{1}{2} \boldsymbol{\omega}^T \mathbf{I}_c \boldsymbol{\omega}.$$

Re-express about a body-fixed frame origin. Let the body frame origin O be offset from the centre of mass by $\mathbf{c} \in \mathbb{R}^3$. Then the centre-of-mass velocity is

$$\mathbf{v}_{cm} = \mathbf{v}_O + \boldsymbol{\omega} \times \mathbf{c},$$

where $\mathbf{v}_O$ is the linear velocity of O . Substituting into T and expanding:

$$\begin{aligned} T &= \frac{1}{2} m \|\mathbf{v}_O + \boldsymbol{\omega} \times \mathbf{c}\|^2 + \frac{1}{2} \boldsymbol{\omega}^T \mathbf{I}_c \boldsymbol{\omega} \\ &= \frac{1}{2} m \mathbf{v}_O^T \mathbf{v}_O + m \mathbf{v}_O^T (\boldsymbol{\omega} \times \mathbf{c}) + \frac{1}{2} m (\boldsymbol{\omega} \times \mathbf{c})^T (\boldsymbol{\omega} \times \mathbf{c}) + \frac{1}{2} \boldsymbol{\omega}^T \mathbf{I}_c \boldsymbol{\omega}. \end{aligned}$$

Using the skew-symmetric matrix $[\mathbf{c}]_{\times}$ so that $\boldsymbol{\omega} \times \mathbf{c} = [\mathbf{c}]_{\times}^T \boldsymbol{\omega}$, and collecting the angular and linear components into the spatial velocity $\mathcal{V} = (\boldsymbol{\omega}, \mathbf{v}_O)^T$, we obtain:

Theorem 8.5. Kinetic Energy as a Spatial Quadratic Form

The kinetic energy of a rigid body equals

$$T = \frac{1}{2} \mathcal{V}^T \mathcal{I} \mathcal{V},$$

where the 6×6 **spatial inertia matrix** is

$$\mathcal{I} = \begin{pmatrix} \mathbf{I}_c + m [\mathbf{c}]_{\times} [\mathbf{c}]_{\times}^T & m [\mathbf{c}]_{\times} \\ m [\mathbf{c}]_{\times}^T & m \mathbf{I}_{3 \times 3} \end{pmatrix}.$$

This derivation proves that $\mathcal{I}$ is the *correct* inertial object for spatial dynamics: it is determined entirely by the requirement that kinetic energy be a quadratic form in $\mathcal{V}$.

The **off-diagonal blocks** $m[\mathbf{c}]_{\times}$ encode the cross-coupling between rotation and translation that arises whenever the centre of mass does not coincide with the frame origin ($\mathbf{c} \neq \mathbf{0}$). When the frame is placed at the centre of mass ($\mathbf{c} = \mathbf{0}$), these blocks vanish and the spatial inertia becomes block-diagonal—rotational and translational energies decouple [12].

8.13 Python Implementation

```

1 import numpy as np
2
3 def skew(v):
4     """Convert 3D vector to skew-symmetric matrix."""
5     return np.array([
6         [0, -v[2], v[1]],
7         [v[2], 0, -v[0]],
8         [-v[1], v[0], 0]
9     ])
10
11 def motion_cross(V):
12     """
13     Spatial cross product for motion (Twists).
14     V: 6D spatial velocity (omega, v)
15     Returns: 6x6 operator [Vx]
16     """
17     omega = V[:3]
18     v = V[3:]
19     Vx = np.zeros((6, 6))
20     Vx[:3, :3] = skew(omega)
21     Vx[3:, :3] = skew(v)
22     Vx[3:, 3:] = skew(omega)
23     return Vx
24
25 def force_cross(V):
26     """
27     Spatial cross product for forces (Wrenches).
28     Returns: 6x6 operator [Vx*] = -[Vx]^T
29     """
30     return -motion_cross(V).T
31
32 def spatial_inertia(m, c, I_c):
33     """
34     Construct 6x6 spatial inertia matrix.
35     Args:
36         m: scalar mass
37         c: 3D center of mass offset from origin
38         I_c: 3x3 rotational inertia about CoM
39     Returns:
40         I: 6x6 spatial inertia matrix
41     """
42     I_spatial = np.zeros((6, 6))
43     c_skew = skew(c)
44
45     # Top-left: I_c + m[c][c]^T (parallel axis theorem).
46     # [c][c]^T = |c|^2 I - c c^T is positive semi-definite, so this ADDS to
47     # the rotational inertia. Equivalently I_c - m[c][c], since [c]^T = -[c].
48     I_spatial[:3, :3] = I_c + m * (c_skew @ c_skew.T)

```



```

49
50     # Top-right:  $m[c]$ 
51     I_spatial[:3, 3:] = m * c_skew
52
53     # Bottom-left:  $m[c]^T$ 
54     I_spatial[3:, :3] = m * c_skew.T
55
56     # Bottom-right:  $m * I_3$ 
57     I_spatial[3:, 3:] = m * np.eye(3)
58
59     return I_spatial
60
61 def transform_inertia(I, T):
62     """
63     Transform spatial inertia from one frame to another.
64     I: 6x6 spatial inertia in source frame
65     T: 4x4 transformation matrix (to target frame)
66     Returns: transformed 6x6 spatial inertia
67     """
68     R = T[:3, :3]
69     p = T[:3, 3]
70
71     # Construct transformation for inertia
72     #  $I' = [Ad\_T]^{-T} I [Ad\_T]^{-1}$ 
73     Ad = np.zeros((6, 6))
74     Ad[:3, :3] = R
75     Ad[3:, :3] = skew(p) @ R
76     Ad[3:, 3:] = R
77
78     Ad_inv_T = np.linalg.inv(Ad).T
79     return Ad_inv_T @ I @ np.linalg.inv(Ad)
80
81 def composite_rigid_body(link_inertias, link_transforms, parent_array):
82     """
83     Assemble composite rigid body inertias.
84     Args:
85         link_inertias: list of 6x6 inertia matrices
86         link_transforms: list of 4x4 transforms to parent frames
87         parent_array: parent[i] = index of parent of link i (or -1 for root)
88     Returns:
89         composite_inertias: list of 6x6 composite inertias
90     """
91     N = len(link_inertias)
92     composite_inertias = [np.copy(I) for I in link_inertias]
93
94     # Process from leaves to root (reverse order)
95     for i in range(N - 1, -1, -1):
96         parent = parent_array[i]
97         if parent >= 0:
98             # Transform this link's composite inertia to parent frame
99             I_in_parent = transform_inertia(
100                 composite_inertias[i],
101                 np.linalg.inv(link_transforms[i])
102             )
103             # Add to parent's inertia
104             composite_inertias[parent] += I_in_parent
105
106     return composite_inertias
107
108 # Example: Two-link arm

```

```

109 if __name__ == "__main__":
110     # Link 1: upper arm
111     m1 = 2.0
112     c1 = np.array([0.3, 0.0, 0.0])
113     I_c1 = np.diag([0.05, 0.05, 0.01])
114     I1 = spatial_inertia(m1, c1, I_c1)
115
116     # Link 2: forearm
117     m2 = 1.0
118     c2 = np.array([0.25, 0.0, 0.0])
119     I_c2 = np.diag([0.02, 0.02, 0.005])
120     I2 = spatial_inertia(m2, c2, I_c2)
121
122     print("Link 1 Spatial Inertia:\n", np.round(I1, 4))
123     print("\nLink 2 Spatial Inertia:\n", np.round(I2, 4))
124
125     # Test motion cross product
126     V = np.array([0.0, 0.0, 1.0, 1.0, 0.0, 0.0]) # Rotating around z,
127     # translating in x
128     Vx = motion_cross(V)
129     print("\nMotion Cross Product [Vx]:\n", np.round(Vx, 4))
130
131     # Test duality
132     F = np.array([0.1, 0.2, 0.3, 1.0, 2.0, 3.0])
133     W = np.array([0.5, 0.5, 0.5, 0.1, 0.1, 0.1])
134
135     lhs = F @ (motion_cross(V) @ W)
136     rhs = (force_cross(V) @ F) @ W
137     print("\nDuality test (should be equal):")
138     print(f"   F^T [Vx] W = {lhs:.6f}")
139     print(f"   [Vx*] F)^T W = {rhs:.6f}")

```

Listing 8.1: Spatial Algebra and Inertia Computation in Python

8.14 Summary

In this chapter, we developed Spatial Vector Algebra—the mathematical language for unified 6D mechanics:

1. **Historical Context:** Plücker’s 19th-century line geometry, Ball’s screws, and Featherstone’s modern synthesis created a computationally elegant framework.
2. **Plücker Coordinates:** Lines in 3D are represented as 6D vectors satisfying a constraint, naturally connecting to spatial twists and wrenches.
3. **Spatial Vectors:** Twists (velocity) and Wrenches (force) are the fundamental 6D objects, living in dual spaces $\mathfrak{se}(3)(3)$ and $\mathfrak{se}(3)^*(3)$.
4. **Motion and Force Cross Products:** Two distinct 6×6 cross product operators capture the geometry of relative motion and force application; they are dual to preserve power.
5. **Spatial Inertia:** A single symmetric, positive-definite 6×6 matrix unifies mass and rotational inertia, including the parallel axis theorem as a block structure.

6. **Parallel Axis Theorem in 6D:** The full theorem includes coupling terms between angular and linear motion due to CoM offset.
7. **Composite Rigid Body Algorithm:** Inertias of linked chains are assembled bottom-up in $O(N)$ time via adjoint transformations.
8. **Planar Reduction:** 2D systems use 3D spatial algebra, reducing computation by a factor of 4.
9. **Newton-Euler in Spatial Form:** A single matrix equation $\mathcal{F} = \mathcal{I}\dot{\mathcal{V}} + [\mathcal{V}\times]^*(\mathcal{I}\mathcal{V})$ unifies linear and angular dynamics.
10. **Dual Quaternions:** Brief connection to an alternative 8D representation with elegant geometric properties.
11. **Physical Interpretation:** Each block of the spatial inertia matrix has clear mechanical meaning.
12. **Python Implementation:** Full code for inertia computation, transformation, and composite body assembly.

With Spatial Vector Algebra in hand, the Recursive Newton-Euler Algorithm of Chapter 7 becomes a straightforward loop of matrix multiplications, without a single transcendental function or explicit cross product in sight. This unity—where complex nonlinear dynamics reduce to structured linear algebra—is the deep beauty of modern computational mechanics.

8.15 Further Reading

The primary reference for spatial vector algebra is Featherstone [12], whose monograph develops the full 6×6 framework from first principles and demonstrates its application to recursive multi-body dynamics algorithms. Featherstone’s treatment of spatial inertia, the motion and force cross products, and the composite rigid body algorithm is definitive, and readers seeking to implement spatial algebra in production code should consult his detailed pseudocode and numerical recipes.

The mathematical foundations connecting spatial vectors to the Lie algebra $\mathfrak{se}(3)(3)$ and its dual $\mathfrak{se}(3)^*(3)$ are developed rigorously by Murray, Li, and Sastry [31]. Their perspective clarifies why twists and wrenches live in dual vector spaces and why the adjoint and coadjoint representations have the specific transformation properties used throughout this chapter.

For the classical treatment of rigid body dynamics, inertia tensors, and the parallel axis theorem in their traditional 3×3 form, Goldstein, Poole, and Safko [14] remains the standard graduate-level reference. Comparing their development with the 6×6 spatial inertia formulation of this chapter illuminates how spatial algebra unifies and streamlines results that classically require separate angular and translational analyses.

Marsden and Ratiu [29] place spatial algebra within the broader framework of geometric mechanics, connecting spatial inertia to momentum maps and showing how the Euler equations for rigid body motion arise naturally from reduction on $SE(3)(3)$. Their treatment is more abstract but reveals the deep geometric structures underlying the computational framework.

It is worth emphasizing that spatial algebra is the computational backbone of modern physics engines. MuJoCo [43], in particular, implements spatial vector arithmetic as its

core computational primitive, using the 6×6 framework to achieve real-time simulation of complex contact-rich robotic systems. Understanding the material in this chapter is therefore essential not only for theoretical mechanics but for practical work with any modern simulation platform.

8.16 Exercises

8.16.1 Tier 1 (Foundation)

1. Construct the motion cross product $[\mathcal{V} \times]$ for $\mathcal{V} = (0, 0, \omega, v_x, v_y, 0)^T$ and verify its action on a unit spatial vector.
2. Verify that the force cross product $[\mathcal{V} \times^*]$ is the negative transpose of the motion cross product.
3. For a point mass m at location $\mathbf{c}$ with no rotational inertia (zero $\mathbf{I}_c$), compute the spatial inertia matrix and verify it is positive definite.
4. Compute the spatial inertia of a uniform rod of mass m and length L at its center of mass, and again at one end (using the parallel axis theorem).
5. Verify that power $\mathcal{F}^T \mathcal{V}$ is invariant under the frame transformation: compute it in two different frames and confirm they are equal.

8.16.2 Tier 2 (Intermediate)

1. Prove that the spatial inertia matrix is invariant under orthogonal rotations of the reference frame (i.e., only the CoM offset affects the matrix structure).
2. For a two-link arm where link 2 is offset from link 1 by transform $\mathbf{T}_{1 \rightarrow 2}$, use the composite rigid body algorithm to express the total inertia as a function of both link inertias and the relative pose.
3. Implement the duality property: compute $\mathcal{G}^T([\mathcal{V} \times] \mathcal{W}) = ([\mathcal{V} \times^*] \mathcal{G})^T \mathcal{W}$ numerically for random spatial vectors and verify to machine precision.
4. Derive the relationship between the classical 3D inertia tensor and the spatial inertia matrix; show how the 3x3 blocks correspond to standard mechanics quantities.
5. For a humanoid robot, conceptually design the composite inertia computation: which links would be aggregated together, and in what order?

8.16.3 Tier 3 (Advanced)

1. Prove the Plücker relation $\mathbf{d}^T \mathbf{m} = 0$ geometrically: interpret the direction and moment of a line and show why their dot product must vanish.
2. Show that the spatial inertia matrix $\mathcal{I}$ can be expressed as a dual-quaternion-valued bilinear form, and compute inertial effects using dual quaternion multiplication.

3. Derive the composite inertia algorithm from the perspective of the Adjoint map: show that $\mathcal{I}_{\text{composite}} = \sum_j [\text{Ad}_{\mathbf{T}_{j,\text{origin}}}]^{-T} \mathcal{I}_j [\text{Ad}_{\mathbf{T}_{j,\text{origin}}}]^{-1}$ reduces to the recursive algorithm.
4. For a system with articulated contact (e.g., foot on ground with rolling constraints), extend the spatial inertia formalism to include constraint forces and derive the reduced dynamics.
5. Implement a multi-body dynamics simulator using purely spatial algebra: compose RNEA (Chapter 7) with the spatial inertia formalism, and validate against a reference physics engine.

Chapter 9

The Product of Exponentials Formula

You do not need to invent abstract coordinate frames at every joint of a robot. The kinematics of the entire complex machine can be compactly described by a single product of twisting screws.

In Plain Language 9.1. Context & Use Case: Universal Forward Kinematics

The Math You Will See: We introduce the **Product of Exponentials (PoE)** formula, utilizing the Matrix Exponential ($e^{[S]\theta}$) directly. **Why We Need It:** In classical robotics (taught in the 1980s), you tracked the position of a robot arm using "Denavit-Hartenberg (D-H) parameters", which required manually attaching weird local coordinate frames to every single joint using confusing perpendicular rules. The modern way (PoE) drops this completely. We just look at the robot in one global frame, find where the hinges point, and mathematically "twist" the arm around each hinge sequentially using our exact exponential maps. It is cleaner, faster, and never suffers from alignment ambiguities. **All Variables Defined:** $\mathbf{M}$ is the 4×4 "home" pose of the robot. $\mathcal{S}_i$ is the 6D Screw Axis of joint i . θ_i is how far the motor has moved.

9.1 Historical Context and the Evolution of Kinematic Formalisms

The theory of robotic kinematics has undergone a significant transformation over the past four decades. For nearly two centuries—from the mechanical linkage analysis of James Watt through the 1960s development of automated manufacturing—kinematic chains were described using local, Euclidean coordinate frames attached sequentially to each link. This frame-by-frame approach reached formal codification in 1955 with the *Denavit-Hartenberg (D-H)* parametrization [9], which became the dominant convention in robotics education and industrial practice.

The D-H convention, while algebraically complete, carries three fundamental drawbacks:

1. **Geometric Arbitrariness:** The perpendicularity rules for axis alignment create multiple valid D-H parameter sets for the same robot, requiring careful disambiguation.
2. **Notational Opacity:** Joint and link parameters are decoupled across four separate scalars $(a_i, d_i, \alpha_i, \theta_i)$, obscuring the underlying geometric intent.
3. **Singularity Avoidance:** Parallel or collinear axes demand special-case treatments, complicating software and theoretical analysis.

In the 1990s, Richard Murray, Zexiang Li, and Shankar Sastry—building on Lie group theory developed in differential geometry—published *A Mathematical Introduction to Robotic Manipulation* [31], which reformulated forward kinematics through the lens of exponential maps on the Lie group $SE(3)$ and its Lie algebra $\mathfrak{se}(3)$. Their **Product of Exponentials (PoE)** formula unified kinematics under a single, coordinate-free framework: the composition of exponential twists in a global reference frame.

This approach was further refined and popularized by Kevin Lynch and Frank Park [27], whose textbook *Modern Robotics: Mechanics, Planning, and Control* established PoE as the modern standard. Today, PoE is the default kinematics formalism in advanced robotics research, computational mechanics, and biomechanics.

Product of Exponentials: PoE




Figure 9.1: Product of Exponentials formula for forward kinematics. Each joint contributes an exponential twist applied to the previous configuration. The final end-effector pose is the product: $\mathbf{T}(\boldsymbol{\theta}) = e^{[S_n]\theta_n} \dots e^{[S_2]\theta_2} e^{[S_1]\theta_1} \mathbf{M}$.

9.2 Abandoning Denavit-Hartenberg

The D-H convention assigns to each link i a local coordinate frame $\{\text{Frame}_i\}$ with four parameters:

$$d_i : \text{offset along the previous joint axis} \quad (9.1)$$

$$a_i : \text{length of the common perpendicular} \quad (9.2)$$

$$\alpha_i : \text{twist angle about the perpendicular} \quad (9.3)$$

$$\theta_i : \text{joint angle (for revolute joints)} \quad (9.4)$$

The combined transformation from frame $\{i-1\}$ to frame $\{i\}$ is:

$$\mathbf{A}_{i-1,i}(\theta_i) = \mathbf{Rot}_z(\theta_i) \mathbf{Trans}(d_i) \mathbf{Trans}(a_i) \mathbf{Rot}_x(\alpha_i) \quad (9.5)$$

The complete forward kinematics chain becomes:

$$\mathbf{T}(\boldsymbol{\theta}) = \mathbf{A}_{0,1} \mathbf{A}_{1,2} \cdots \mathbf{A}_{n-1,n} \quad (9.6)$$

where each transformation encodes one joint's motion relative to the previous link's local frame.

Why PoE is superior: The PoE formalism instead expresses all joint motions in a single, fixed **space frame** (the global reference frame). This eliminates the need to compute and track relative frame transformations and treats the kinematics as direct composition of rigid-body twists—a concept with deep physical meaning in mechanics.

9.3 Screws and Twists: A Geometric Foundation

Recall from Chapter 6 that a screw axis $\mathcal{S} = (\boldsymbol{\omega}, \mathbf{v}) \in \mathfrak{se}(3)(3)$ encodes:

- $\boldsymbol{\omega}$: the unit rotational axis (or zero for a purely prismatic joint)
- $\mathbf{v}$: the linear velocity component, related to the axis location and pitch

For a revolute joint (rotation only), the axis passes through a point $\mathbf{q}$ in space, and we define:

$$\mathcal{S}_{\text{rev}} = \begin{bmatrix} \boldsymbol{\omega} \\ -\boldsymbol{\omega} \times \mathbf{q} \end{bmatrix} \quad (9.7)$$

For a prismatic joint (translation along direction $\mathbf{v}$), we have:

$$\mathcal{S}_{\text{pris}} = \begin{bmatrix} \mathbf{0} \\ \mathbf{v} \end{bmatrix} \quad (9.8)$$

The corresponding 4×4 skew-symmetric matrix representation is:

$$[\mathcal{S}] = \begin{bmatrix} [\boldsymbol{\omega}] & \mathbf{v} \\ \mathbf{0}^T & 0 \end{bmatrix} \quad (9.9)$$

9.4 The Forward Kinematics Equation: Space Frame PoE

If a robotic arm has n joints, the Product of Exponentials formula—originally formulated by Brockett [3] and developed into a complete kinematic framework by Murray, Li, and Sastry [31]—computes the final Homogeneous Transformation Matrix ($\mathbf{T}(\boldsymbol{\theta}) \in SE(3)$) of the end-effector relative to the base as:

$$\mathbf{T}(\boldsymbol{\theta}) = e^{[\mathcal{S}_1]\theta_1} e^{[\mathcal{S}_2]\theta_2} \cdots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} \quad (9.10)$$

Where:

1. $\mathbf{M} \in SE(3)$ is the **Home Configuration** matrix: the end-effector's pose when all joint angles $\theta_i = 0$. It is expressed in the space (global) frame and is typically computed by direct measurement or CAD.
2. $\mathcal{S}_i \in \mathfrak{se}(3)(3)$ is the **spatial Screw Axis** (the twist parameter) for joint i , expressed in the global Space Frame when the robot is at Home Configuration. It does *not* change as the robot moves.

3. θ_i is the actual angle (in radians, for revolute joints) or distance (in meters, for prismatic joints) the motor has moved.
4. $e^{[S_i]\theta_i} \in SE(3)$ is the Matrix Exponential (developed in Chapter 6) mapping that single joint's motion into a 4×4 rigid body transform.

9.4.1 Derivation From First Principles

Imagine the robot at home configuration. As joint 1 rotates by θ_1 , it rotates the entire downstream chain (all links $2, 3, \dots, n$ plus the end-effector) by the exponential twist $e^{[S_1]\theta_1}$. The end-effector moves to the configuration:

$$\mathbf{T}_1(\theta_1) = e^{[S_1]\theta_1} \mathbf{M} \quad (9.11)$$

Now, while joint 1 is held fixed, joint 2 begins to move. Its axis is *fixed in the space frame*—it does not follow the rotated configuration of link 2. This is the crucial distinction between space-frame and body-frame PoE. Joint 2's twist applies to the pose resulting from joint 1:

$$\mathbf{T}_{1,2}(\theta_1, \theta_2) = e^{[S_2]\theta_2} \mathbf{T}_1(\theta_1) = e^{[S_2]\theta_2} e^{[S_1]\theta_1} \mathbf{M} \quad (9.12)$$

Continuing inductively, the full Product of Exponentials is:

$$\mathbf{T}(\boldsymbol{\theta}) = e^{[S_n]\theta_n} \dots e^{[S_2]\theta_2} e^{[S_1]\theta_1} \mathbf{M} \quad (9.13)$$

Note the reversed ordering: joint 1 is the leftmost exponential in the final product because it acts first (is applied closest to $\mathbf{M}$).

The elegance of this formulation lies in its physical transparency: each exponential is a pure twist in the space frame, and the product is a concatenation of rigid-body motions. No ad-hoc frame transformations are needed.

9.5 Body Frame PoE: Motion Relative to the Tool

In some contexts (particularly when working backward from the end-effector), it is convenient to express joint axes in the **body frame**—a frame attached to the end-effector and moving with it. The body frame PoE formula is:

$$\mathbf{T}(\boldsymbol{\theta}) = \mathbf{M} e^{[S'_n]\theta_n} e^{[S'_{n-1}]\theta_{n-1}} \dots e^{[S'_1]\theta_1} \quad (9.14)$$

Here, S'_i is the screw axis of joint i expressed in the body frame (tool frame). The key difference is that joint n acts first (is applied closest to the identity), and the exponentials are multiplied in increasing order of joint index.

9.5.1 Conversion Between Space and Body Frame

Given a body frame screw axis S'_i and a known transformation $\mathbf{M} \in SE(3)$, the corresponding space frame screw axis is obtained via the **Adjoint** action (see Chapter 8):

$$S_i = \text{Ad}_{\mathbf{T}} S'_i \quad (9.15)$$

where $\text{Ad}_{\mathbf{T}}$ is the adjoint representation of the transformation $\mathbf{T}$. In 6D spatial vector form:

$$\text{Ad}_{\mathbf{T}} = \begin{bmatrix} \mathbf{R} & \mathbf{0} \\ [\mathbf{p}]\mathbf{R} & \mathbf{R} \end{bmatrix} \quad (9.16)$$

where $\mathbf{T} = (\mathbf{R}, \mathbf{p})$.

Equivalently, if we know all screw axes in space frame and wish to express them in body frame, we use:

$$\mathcal{S}'_i = \text{Ad}_{\mathbf{M}^{-1}} \mathcal{S}_i \quad (9.17)$$

The two PoE formulations are mathematically equivalent; the choice depends on which frame is more convenient for the application or computation.

9.6 Spatial Jacobian From the Product of Exponentials

The **Spatial Jacobian** $\mathbf{J}_s(\boldsymbol{\theta})$ relates the time derivatives of joint angles to the spatial velocity (twist) of the end-effector:

$$\mathcal{V}_{\text{ee}} = \mathbf{J}_s(\boldsymbol{\theta}) \dot{\boldsymbol{\theta}} \quad (9.18)$$

where $\mathcal{V}_{\text{ee}} = [\boldsymbol{\omega}, \mathbf{v}]^T \in \mathbb{R}^6$ is the 6D twist of the end-effector.

To derive the spatial Jacobian from the PoE formula, we differentiate Equation 9.10 with respect to time:

$$\dot{\mathbf{T}} = \frac{d}{dt} \left(e^{[\mathcal{S}_1]\theta_1} e^{[\mathcal{S}_2]\theta_2} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} \right) \quad (9.19)$$

Using the product rule and chain rule on matrix exponentials (and noting that $[\mathcal{S}]$ is constant):

$$\dot{\mathbf{T}} = [\mathcal{S}_1] \dot{\theta}_1 e^{[\mathcal{S}_1]\theta_1} e^{[\mathcal{S}_2]\theta_2} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} + e^{[\mathcal{S}_1]\theta_1} [\mathcal{S}_2] \dot{\theta}_2 e^{[\mathcal{S}_2]\theta_2} \dots + \dots \quad (9.20)$$

The spatial velocity (twist) at the end-effector is related to the time derivative of the transformation by:

$$\mathcal{V}_{\text{ee}} = \dot{\mathbf{T}} \mathbf{T}^{-1} \quad (9.21)$$

where the product $\dot{\mathbf{T}} \mathbf{T}^{-1}$ extracts the velocity in the space frame (as opposed to the body frame).

After algebraic manipulation, we obtain the **spatial Jacobian columns**:

$$\mathbf{J}_s(\boldsymbol{\theta}) = [\mathcal{S}_1 \quad e^{[\mathcal{S}_1]\theta_1} \mathcal{S}_2 \quad \dots \quad e^{[\mathcal{S}_1]\theta_1} \dots e^{[\mathcal{S}_{n-1}]\theta_{n-1}} \mathcal{S}_n] \quad (9.22)$$

Each column is the corresponding screw axis, transformed by all the exponential twists that precede it in the PoE product. This encodes the fact that earlier joints affect the effective direction of later joints.

9.6.1 Direct Numerical Computation

In practice, the spatial Jacobian is computed numerically by:

1. Compute $\mathbf{T}(\boldsymbol{\theta})$ using the PoE formula.
2. For each joint i , compute the partial derivative using finite differences or automatic differentiation:

$$\frac{\partial \mathbf{T}}{\partial \theta_i} \approx \frac{\mathbf{T}(\boldsymbol{\theta} + \boldsymbol{\epsilon}_i) - \mathbf{T}(\boldsymbol{\theta} - \boldsymbol{\epsilon}_i)}{2\epsilon} \quad (9.23)$$

where $\boldsymbol{\epsilon}_i$ is a small perturbation in joint i only.

3. Extract the spatial velocity from each partial derivative.

9.7 Body Jacobian and the Adjoint Relationship

The **Body Jacobian** $\mathbf{J}_b(\boldsymbol{\theta})$ expresses the end-effector velocity in the body (tool) frame:

$$\mathcal{V}_{\text{ee}}^{\text{body}} = \mathbf{J}_b(\boldsymbol{\theta})\dot{\boldsymbol{\theta}} \quad (9.24)$$

The body Jacobian is related to the spatial Jacobian via the adjoint of the end-effector transformation:

$$\mathbf{J}_b(\boldsymbol{\theta}) = \text{Ad}_{\mathbf{T}(\boldsymbol{\theta})^{-1}} \mathbf{J}_s(\boldsymbol{\theta}) \quad (9.25)$$

where the adjoint action on a Jacobian matrix (with columns as 6D vectors) is:

$$\text{Ad}_{\mathbf{T}} \mathbf{J} = \text{Ad}_{\mathbf{T}} [\text{col}_1(\mathbf{J}), \text{col}_2(\mathbf{J}), \dots] \quad (9.26)$$

applied column-wise.

The body frame Jacobian is often more useful in control applications because the body frame is attached to the end-effector and its axes remain intuitively meaningful as the tool moves.

Having established both the spatial and body Jacobians, we now turn to the critical question of when these Jacobians break down. The Jacobian's rank determines the robot's instantaneous ability to generate motion in all directions, and configurations where this rank drops—called singularities—impose fundamental limits on the robot's dexterity. Understanding and detecting singularities is essential for safe and effective motion planning.

9.8 Singularity Analysis Using the PoE Jacobian

A robotic configuration is **singular** when the Jacobian loses rank, i.e., when $\text{rank}(\mathbf{J}_s) < n$. At a singularity:

1. The robot cannot move in certain directions (loss of controllability).
2. Joint velocities required to achieve a desired end-effector motion become infinite.
3. The robot exhibits locked degrees of freedom.

9.8.1 Types of Singularities

- **Boundary Singularities:** The end-effector is at or beyond the reachable workspace boundary. Typically occur at full extension of the arm.
- **Interior Singularities:** Within the reachable workspace, occur when multiple joint axes become aligned. For example, in a planar 2R arm, singularity occurs when both links are collinear (fully extended or fully retracted).
- **Redundant Degeneracies:** In redundant robots (more DOFs than required for a task), the null space of the Jacobian is non-trivial, allowing motion without changing the end-effector pose.

9.8.2 Detection and Handling

Singularities are detected by:

1. Computing $\det(\mathbf{J}_s^T \mathbf{J}_s)$ (the condition number of the Jacobian squared).
2. If $\det(\mathbf{J}_s^T \mathbf{J}_s) < \epsilon_{\text{tol}}$ for a small threshold ϵ_{tol} , the configuration is near-singular.
3. The ratio $\det(\mathbf{J}_s^T \mathbf{J}_s)$ is called the **manipulability measure**; high values indicate dexterous configurations.

Numerically, near-singular configurations are handled by:

1. Adding damping to the inverse: $\mathbf{J}_{\text{damp}}^{-1} = \mathbf{J}^T(\mathbf{J}\mathbf{J}^T + \lambda I)^{-1}$ (Levenberg-Marquardt damping).
2. Using the pseudo-inverse with a singular value threshold: keep only singular values above $\tau\sigma_{\text{max}}$, where τ is a threshold ratio (e.g., $\tau = 10^{-4}$).

9.9 Comparison With D-H: A Worked Example

To illustrate the practical differences, consider a 3-DOF planar robot with:

- Link 1: length $L_1 = 1$ m
- Link 2: length $L_2 = 0.8$ m
- Link 3: length $L_3 = 0.6$ m
- All joints: revolute, rotating about the Z-axis (perpendicular to the table)

9.9.1 D-H Parametrization

The D-H parameters are:

Link i	a_{i-1}	d_i	α_{i-1}	θ_i
1	0	0	0	θ_1
2	L_1	0	0	θ_2
3	L_2	0	0	θ_3

The transformation from base to end-effector is:

$$\mathbf{T}_{D-H}(\boldsymbol{\theta}) = \mathbf{A}_{0,1}\mathbf{A}_{1,2}\mathbf{A}_{2,3} \quad (9.27)$$

where each $\mathbf{A}_{i,i+1}$ is computed using the D-H convention. At $\boldsymbol{\theta} = [45^\circ, -30^\circ, 60^\circ]^T$, one obtains a specific end-effector pose.

9.9.2 PoE Parametrization

The home configuration ($\boldsymbol{\theta} = \mathbf{0}$) has the arm fully extended along the X-axis:

$$\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & L_1 + L_2 + L_3 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (9.28)$$

The three screw axes (all revolute about Z):

$$\mathcal{S}_1 = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad \mathcal{S}_2 = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ -L_1 \\ 0 \end{bmatrix}, \quad \mathcal{S}_3 = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ -(L_1 + L_2) \\ 0 \end{bmatrix} \quad (9.29)$$

The PoE formula gives:

$$\mathbf{T}_{PoE}(\boldsymbol{\theta}) = e^{[\mathcal{S}_1]\theta_1} e^{[\mathcal{S}_2]\theta_2} e^{[\mathcal{S}_3]\theta_3} \mathbf{M} \quad (9.30)$$

When computed at the same configuration $\boldsymbol{\theta} = [45^\circ, -30^\circ, 60^\circ]^T$, the PoE and D-H formulations yield identical end-effector poses (as they must, since both are valid descriptions of the same kinematic chain).

9.9.3 Computational Advantages of PoE

1. **Consistency:** The PoE formula is coordinate-free; no ambiguity in frame attachment.
2. **Scalability:** Adding a new joint requires only specifying its screw axis; no recomputation of previous parameters.
3. **Jacobian Clarity:** The Jacobian columns have direct physical meaning: the i -th column is the screw axis of joint i as "seen" from the space frame after the effects of earlier joints.
4. **Software Simplicity:** A single, unified algorithm computes forward kinematics for any serial chain, regardless of link lengths or orientations.

With the comparison between the classical D-H convention [9] and the modern PoE approach now clear, we can extend the PoE framework to handle all joint types—including prismatic (sliding) joints—before returning to the question of how joint velocities propagate through the kinematic chain. The velocity kinematics that follow are a direct consequence of differentiating the PoE formula, unifying position and velocity analysis under a single mathematical roof.

9.10 PoE for Prismatic Joints

A prismatic (sliding) joint moves along a direction vector $\mathbf{d}$ without rotating. Its screw axis is:

$$\mathcal{S}_{\text{pris}} = \begin{bmatrix} \mathbf{0} \\ \mathbf{d} \end{bmatrix} \quad (9.31)$$

where $\mathbf{d}$ is the unit direction of translation.

The exponential of a purely prismatic screw is:

$$e^{[\mathcal{S}_{\text{pris}}]d} = \begin{bmatrix} \mathbf{I}_{3 \times 3} & d\mathbf{d} \\ \mathbf{0}^T & 1 \end{bmatrix} \quad (9.32)$$

a pure translation matrix.

9.10.1 Example: 3D Cartesian Robot

A 3D Cartesian robot with three prismatic joints (X, Y, Z directions) has screw axes:

$$\mathcal{S}_x = [0, 0, 0, 1, 0, 0]^T \quad (9.33)$$

$$\mathcal{S}_y = [0, 0, 0, 0, 1, 0]^T \quad (9.34)$$

$$\mathcal{S}_z = [0, 0, 0, 0, 0, 1]^T \quad (9.35)$$

and home configuration $\mathbf{M} = \mathbf{I}$ (end-effector at the origin when all sliders are at zero position).

The forward kinematics is:

$$\mathbf{T}(d_x, d_y, d_z) = e^{[\mathcal{S}_x]d_x} e^{[\mathcal{S}_y]d_y} e^{[\mathcal{S}_z]d_z} = \begin{bmatrix} \mathbf{I} & [d_x, d_y, d_z]^T \\ \mathbf{0}^T & 1 \end{bmatrix} \quad (9.36)$$

which correctly expresses a point displaced by (d_x, d_y, d_z) in the global frame.

9.11 Velocity Kinematics From PoE Differentiation

The instantaneous spatial velocity of the end-effector is obtained by differentiating the forward kinematics. Starting from:

$$\mathbf{T}(\boldsymbol{\theta}) = e^{[\mathcal{S}_1]\theta_1} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} \quad (9.37)$$

Differentiate with respect to time:

$$\dot{\mathbf{T}} = \sum_{i=1}^n \frac{d}{dt} \left(e^{[\mathcal{S}_1]\theta_1} \dots e^{[\mathcal{S}_i]\theta_i} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} \right) \dot{\theta}_i \quad (9.38)$$

Using the derivative rule for matrix exponentials $\frac{d}{dt} e^{A(t)} = A(t) e^{A(t)}$ when $A(t)$ is linear in a parameter:

$$\dot{\mathbf{T}} = [\mathcal{S}_1] \dot{\theta}_1 e^{[\mathcal{S}_1]\theta_1} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M} + e^{[\mathcal{S}_1]\theta_1} [\mathcal{S}_2] \dot{\theta}_2 e^{[\mathcal{S}_2]\theta_2} \dots + \dots \quad (9.39)$$

The spatial velocity is:

$$\mathcal{V} = \dot{\mathbf{T}}\mathbf{T}^{-1} \quad (9.40)$$

After collecting terms and using the adjoint action, the spatial velocity can be written as:

$$[\mathcal{V}] = \sum_{i=1}^n [\mathcal{S}_i^{\text{transformed}}] \dot{\theta}_i \quad (9.41)$$

where $\mathcal{S}_i^{\text{transformed}}$ is the screw axis of joint i after being advected by all preceding joints. In matrix form, this is precisely the spatial Jacobian:

$$\mathcal{V} = \mathbf{J}_s(\boldsymbol{\theta}) \dot{\boldsymbol{\theta}} \quad (9.42)$$

9.12 Worked Example: Complete PoE for a Spatial 3R Robot

Consider a 3-DOF spatial manipulator (Spatial 3R) with:

- Base at origin; first joint axis along Z
- Link 1 length: $L_1 = 0.5$ m
- Second joint axis along Y (at the end of link 1)
- Link 2 length: $L_2 = 0.3$ m
- Third joint axis along X (at the end of link 2)
- Link 3 length: $L_3 = 0.2$ m

9.12.1 Home Configuration

At $\boldsymbol{\theta} = \mathbf{0}$, the robot is fully extended in the positive X direction:

$$\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & L_1 + L_2 + L_3 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 1.0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (9.43)$$

9.12.2 Screw Axes

$$\text{Joint 1 (Z-axis, at origin): } \mathcal{S}_1 = [0, 0, 1, 0, 0, 0]^T \quad (9.44)$$

$$\text{Joint 2 (Y-axis, at } [L_1, 0, 0]^T = [0.5, 0, 0]^T \text{): } \mathcal{S}_2 = [0, 1, 0, 0 \cdot 0 - 0 \cdot 0.5, 0 \cdot 0.5 - 0 \cdot 0, 1 \cdot 0.5 - 0 \cdot 0]^T \quad (9.45)$$

$$= [0, 1, 0, 0, 0, 0.5]^T \quad (9.46)$$

$$\text{Joint 3 (X-axis, at } [L_1 + L_2, 0, 0]^T = [0.8, 0, 0]^T \text{): } \mathcal{S}_3 = [1, 0, 0, 0 \cdot 0 - 0 \cdot 0.8, 0 \cdot 0.8 - 0 \cdot 0, 0 \cdot 0.8 - 0 \cdot 0]^T \quad (9.47)$$

$$= [1, 0, 0, 0, 0, 0]^T \quad (9.48)$$

9.12.3 Numerical Evaluation at a Test Configuration

At $\boldsymbol{\theta} = [\pi/4, \pi/6, -\pi/3]^T$ ($45^\circ, 30^\circ, -60^\circ$):

Compute the matrix exponentials:

$$e^{[\mathcal{S}_1]\pi/4} = \begin{bmatrix} \cos(\pi/4) & -\sin(\pi/4) & 0 & 0 \\ \sin(\pi/4) & \cos(\pi/4) & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0.707 & -0.707 & 0 & 0 \\ 0.707 & 0.707 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (9.49)$$

(Rotation by $\pi/4$ about Z-axis; since the screw axis $\mathcal{S}_1$ passes through the origin, there is no translation component.)

For $e^{[\mathcal{S}_2]\pi/6}$, the screw axis has a rotation part (Y-axis) and a translation part ($v = [0, 0, 0.5]^T$). The exponential yields a rotation about the Y-axis combined with a translation. Detailed computation yields:

$$e^{[\mathcal{S}_2]\pi/6} \approx \begin{bmatrix} 0.966 & 0 & 0.259 & 0 \\ 0 & 1 & 0 & 0.067 \\ -0.259 & 0 & 0.966 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad (9.50)$$

Similarly, $e^{[\mathcal{S}_3](-\pi/3)}$ produces a rotation about the X-axis plus a translation.

The full forward kinematics is:

$$\mathbf{T}(\pi/4, \pi/6, -\pi/3) = e^{[\mathcal{S}_1]\pi/4} e^{[\mathcal{S}_2]\pi/6} e^{[\mathcal{S}_3](-\pi/3)} \mathbf{M} \quad (9.51)$$

which, when computed numerically, yields a specific 3D position and orientation of the end-effector. The exact numerical result depends on careful matrix multiplication and is best computed programmatically (see Section 9.15).

9.13 Inverse Kinematics Using the Jacobian: Newton-Raphson

The inverse kinematics problem seeks to find joint angles $\boldsymbol{\theta}$ that place the end-effector at a desired pose $\mathbf{T}_{\text{desired}}$. Unlike forward kinematics, inverse kinematics has no closed-form solution for general serial chains and requires numerical optimization.

The Newton-Raphson method formulates IK as an optimization problem:

$$\min_{\boldsymbol{\theta}} \|\mathbf{T}(\boldsymbol{\theta}) - \mathbf{T}_{\text{desired}}\|_F^2 \quad (9.52)$$

where $\|\cdot\|_F$ is the Frobenius norm of the pose error.

9.13.1 Pose Error Representation

The error is typically decomposed into position and orientation components:

1. **Position Error:** $\mathbf{e}_p = \mathbf{p}_{\text{desired}} - \mathbf{p}(\boldsymbol{\theta})$, where $\mathbf{p}$ is the translation part of $\mathbf{T}$.
2. **Orientation Error:** $\mathbf{e}_o = \text{vex}(\mathbf{R}_{\text{desired}}\mathbf{R}(\boldsymbol{\theta})^T)$, where vex is the inverse of the skew-symmetric operator (extracting the axis-angle representation).

The combined error vector is:

$$\mathbf{e}(\boldsymbol{\theta}) = \begin{bmatrix} \mathbf{e}_p \\ \mathbf{e}_o \end{bmatrix} \in \mathbb{R}^6 \quad (9.53)$$

9.13.2 Newton-Raphson Update Rule

The Jacobian of the error is:

$$\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} = \begin{bmatrix} \frac{\partial \mathbf{p}}{\partial \boldsymbol{\theta}} \\ \frac{\partial \mathbf{e}_o}{\partial \boldsymbol{\theta}} \end{bmatrix} \quad (9.54)$$

For small errors, the relationship is approximately linear, and the Newton-Raphson update is:

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k + \Delta \boldsymbol{\theta} \quad (9.55)$$

where $\Delta \boldsymbol{\theta}$ solves:

$$\mathbf{J}_s(\boldsymbol{\theta}_k) \Delta \boldsymbol{\theta} = \mathbf{e}(\boldsymbol{\theta}_k) \quad (9.56)$$

In practice, the pseudo-inverse is used to handle rank deficiency:

$$\Delta \boldsymbol{\theta} = \mathbf{J}_s^\dagger(\boldsymbol{\theta}_k) \mathbf{e}(\boldsymbol{\theta}_k) \quad (9.57)$$

where $\mathbf{J}_s^\dagger$ is the Moore-Penrose pseudo-inverse.

9.13.3 Algorithm and Convergence

1. Initialize $\boldsymbol{\theta}_0$ with an initial guess (e.g., the previous solution or a random configuration).
2. For $k = 0, 1, 2, \dots, K_{\max}$:
 - (a) Compute the forward kinematics $\mathbf{T}_k = \mathbf{T}(\boldsymbol{\theta}_k)$.
 - (b) Compute the spatial Jacobian $\mathbf{J}_s(\boldsymbol{\theta}_k)$.
 - (c) Compute the pose error $\mathbf{e}_k = \mathbf{e}(\boldsymbol{\theta}_k)$.
 - (d) If $\|\mathbf{e}_k\| < \epsilon_{\text{tol}}$, terminate and return $\boldsymbol{\theta}_k$ (solution found).
 - (e) Compute the update $\Delta \boldsymbol{\theta} = \mathbf{J}_s^\dagger(\boldsymbol{\theta}_k) \mathbf{e}_k$.
 - (f) Update $\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k + \alpha \Delta \boldsymbol{\theta}$, where $0 < \alpha \leq 1$ is a step size (often $\alpha = 1$).
3. If convergence is not achieved within $K_{\max}$ iterations, the initial guess was infeasible or the target pose is outside the reachable workspace.

Convergence: Near a non-singular configuration, Newton-Raphson converges quadratically. Far from a solution or near singularities, the convergence may be slow or absent. For robust IK, multiple random restarts or more sophisticated optimization methods (e.g., BFGS with line search) are often employed.

9.14 Workspace Analysis and Manipulability

The Jacobian reveals not only instantaneous velocities but also the *quality* of motion at a given configuration. Two fundamental questions arise: *where* can the robot reach, and *how well* can it move once there?

Definition 9.1. Reachable Workspace

The **reachable workspace** of a manipulator with forward kinematics f_{kin} and joint space $\mathcal{Q}$ is the set of all end-effector poses attainable:

$$\mathcal{W} = \{f_{kin}(\mathbf{q}) \mid \mathbf{q} \in \mathcal{Q}\}. \quad (9.58)$$

The **dexterous workspace** is the subset of $\mathcal{W}$ reachable with *arbitrary* end-effector orientation.

While the workspace tells us *if* a pose is reachable, the **manipulability measure** quantifies *how easily* the robot can move at a given configuration.

Definition 9.2. Manipulability Measure (Yoshikawa)

The **manipulability measure** at configuration $\mathbf{q}$ is [38]:

$$w(\mathbf{q}) = \sqrt{\det(\mathbf{J}(\mathbf{q}) \mathbf{J}(\mathbf{q})^T)}. \quad (9.59)$$

When $w(\mathbf{q}) = 0$, the Jacobian is rank-deficient and the robot is at a **kinematic singularity**.

The geometric interpretation comes from the **singular value decomposition** of the Jacobian. Let $\mathbf{J} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$, where $\mathbf{\Sigma} = \text{diag}(\sigma_1, \dots, \sigma_m)$ with $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_m \geq 0$.

Definition 9.3. Velocity Manipulability Ellipsoid

The set of end-effector velocities achievable with unit joint velocity ($\|\dot{\mathbf{q}}\| \leq 1$) forms an ellipsoid in task space. Its principal axes are the columns of $\mathbf{U}$, with semi-axis lengths equal to the singular values σ_i . The manipulability measure equals the ellipsoid volume: $w = \sigma_1 \sigma_2 \dots \sigma_m$.

Definition 9.4. Force Manipulability Ellipsoid

The **force manipulability ellipsoid** is the dual of the velocity ellipsoid. Unit joint torques ($\|\boldsymbol{\tau}\| \leq 1$) map to end-effector forces within an ellipsoid whose semi-axes have lengths $1/\sigma_i$. The axes of easy motion become axes of weak force, and vice versa.

In Plain Language 9.2. Velocity–Force Duality

Key insight: Where the robot is fast, it is weak; where it is slow, it is strong. A configuration that allows rapid end-effector motion along a direction (large σ_i) can exert only small forces along that same direction ($1/\sigma_i$), and conversely. This is a direct consequence of the principle of virtual work [27].

Example 9.1. Manipulability of a Planar 2R Arm

Consider a planar 2R arm with link lengths $L_1 = L_2 = 1$ m. The Jacobian is:

$$\mathbf{J}(\mathbf{q}) = \begin{bmatrix} -L_1 s_1 - L_2 s_{12} & -L_2 s_{12} \\ L_1 c_1 + L_2 c_{12} & L_2 c_{12} \end{bmatrix} \quad (9.60)$$

where $s_1 = \sin q_1$, $c_1 = \cos q_1$, $s_{12} = \sin(q_1 + q_2)$, $c_{12} = \cos(q_1 + q_2)$.

Case 1: Elbow extended ($q_1 = 0$, $q_2 = \pi/4$). Numerical evaluation gives $w \approx 0.71$. The velocity ellipsoid is nearly circular—the robot can move comparably well in all directions.

Case 2: Elbow nearly straight ($q_1 = 0$, $q_2 = 0.1$). Now $w \approx 0.10$. The ellipsoid collapses along the radial direction—the robot can barely move inward or outward (approaching a workspace boundary singularity), but can still swing laterally. The force ellipsoid shows the opposite: enormous radial force capability but weak lateral force.

9.14.1 Redundancy and Null-Space Motion

When a manipulator has more joints than task-space dimensions ($n > m$), the Jacobian $\mathbf{J} \in \mathbb{R}^{m \times n}$ has a nontrivial **null space**: joint velocities $\dot{\mathbf{q}}_0$ satisfying $\mathbf{J}\dot{\mathbf{q}}_0 = \mathbf{0}$ produce *zero* end-effector motion. This internal motion—reconfiguration without moving the tool—is the hallmark of **kinematic redundancy**.

Definition 9.5. Null-Space Projector

Given the Moore–Penrose pseudoinverse $\mathbf{J}^\dagger = \mathbf{J}^T(\mathbf{J}\mathbf{J}^T)^{-1}$, the **null-space projector** is:

$$\mathbf{N} = \mathbf{I}_n - \mathbf{J}^\dagger \mathbf{J}. \quad (9.61)$$

Any vector $\mathbf{N}\dot{\mathbf{q}}_0$ lies in the null space of $\mathbf{J}$, regardless of $\dot{\mathbf{q}}_0$.

The general solution for a desired end-effector velocity $\mathbf{v}_{des}$ is then [38]:

$$\dot{\mathbf{q}} = \mathbf{J}^\dagger \mathbf{v}_{des} + \mathbf{N}\dot{\mathbf{q}}_0 \quad (9.62)$$

where the first term achieves the primary task (end-effector motion) and the second term exploits the redundancy for a **secondary objective** without disturbing the primary task.

Example 9.2. 7-DOF Arm with Manipulability Optimization

A 7-DOF robot arm reaching a 6-DOF target (position + orientation) has a one-dimensional null space at non-singular configurations. A common strategy sets $\dot{\mathbf{q}}_0 = \alpha \nabla_{\mathbf{q}} w(\mathbf{q})$, the gradient of the manipulability measure, so that the “elbow” reconfigures to maximize dexterity while tracking the desired trajectory:

$$\dot{\mathbf{q}} = \mathbf{J}^\dagger \mathbf{v}_{des} + \alpha \mathbf{N} \nabla_{\mathbf{q}} w(\mathbf{q}). \quad (9.63)$$

This keeps the arm away from singularities and ensures good force/velocity capability throughout the motion. Industrial 7-DOF arms (e.g., KUKA iiwa, Franka Emika

Panda) exploit this technique routinely.

9.15 Python Worked Example

Here is a complete Python implementation computing PoE forward kinematics and the spatial Jacobian for the 3R robot.

```

1 import numpy as np
2 from scipy.linalg import expm
3 import matplotlib.pyplot as plt
4 from mpl_toolkits.mplot3d import Axes3D
5
6 class SpatialRobot3R:
7     """3-DOF spatial manipulator using Product of Exponentials"""
8
9     def __init__(self, L1=0.5, L2=0.3, L3=0.2):
10         self.L1 = L1
11         self.L2 = L2
12         self.L3 = L3
13
14         # Home configuration (fully extended along X-axis)
15         self.M = np.eye(4)
16         self.M[0, 3] = L1 + L2 + L3
17
18         # Screw axes in space frame (at home)
19         self.S1 = np.array([0, 0, 1, 0, 0, 0]) # Z-axis at origin
20         self.S2 = np.array([0, 1, 0, 0, 0, L1]) # Y-axis at (L1, 0, 0)
21         self.S3 = np.array([1, 0, 0, 0, 0, 0]) # X-axis at (L1+L2, 0, 0)
22
23     def screw_to_matrix(self, screw):
24         """Convert 6D screw to 4x4 skew-symmetric matrix"""
25         w = screw[:3] # Angular velocity
26         v = screw[3:] # Linear velocity
27         S = np.zeros((4, 4))
28         S[:3, :3] = np.array([[0, -w[2], w[1]],
29                               [w[2], 0, -w[0]],
30                               [-w[1], w[0], 0]])
31         S[:3, 3] = v
32         return S
33
34     def forward_kinematics(self, theta):
35         """Compute end-effector transformation T(theta)"""
36         T = np.eye(4)
37
38         # Apply each joint's exponential twist
39         for i, S in enumerate([self.S1, self.S2, self.S3]):
40             S_skew = self.screw_to_matrix(S)
41             exp_S_theta = expm(S_skew * theta[i])
42             T = exp_S_theta @ T
43
44         # Apply home configuration
45         T = T @ self.M
46         return T
47
48     def spatial_jacobian(self, theta, eps=1e-6):
49         """Compute spatial Jacobian using finite differences"""
50         J = np.zeros((6, 3))

```

```

51     T = self.forward_kinematics(theta)
52
53     for i in range(3):
54         theta_plus = theta.copy()
55         theta_minus = theta.copy()
56         theta_plus[i] += eps
57         theta_minus[i] -= eps
58
59         T_plus = self.forward_kinematics(theta_plus)
60         T_minus = self.forward_kinematics(theta_minus)
61
62         # Compute velocity from finite differences
63         dT = (T_plus - T_minus) / (2 * eps)
64         V = dT @ np.linalg.inv(T) # Spatial velocity
65
66         # Extract twist from V
67         J[0:3, i] = V[0:3, 3] # Linear velocity
68         J[3, i] = V[2, 1] # Angular velocity components
69         J[4, i] = V[0, 2]
70         J[5, i] = V[1, 0]
71
72     return J
73
74     def end_effector_position(self, theta):
75         """Return (x, y, z) position of end-effector"""
76         T = self.forward_kinematics(theta)
77         return T[:3, 3]
78
79 # Create robot and test
80 robot = SpatialRobot3R()
81 theta_test = np.array([np.pi/4, np.pi/6, -np.pi/3])
82
83 T_ee = robot.forward_kinematics(theta_test)
84 print("End-effector transformation:\n", np.round(T_ee, 4))
85
86 pos = robot.end_effector_position(theta_test)
87 print("End-effector position (x, y, z):", np.round(pos, 4))
88
89 J = robot.spatial_jacobian(theta_test)
90 print("Spatial Jacobian:\n", np.round(J, 4))
91
92 # Test manipulability
93 manip = np.sqrt(np.linalg.det(J @ J.T))
94 print("Manipulability measure:", np.round(manip, 6))

```

Listing 9.1: PoE Forward Kinematics and Spatial Jacobian

9.16 Chapter Summary

The Product of Exponentials formula provides a modern, geometrically transparent approach to forward kinematics:

1. **PoE Formula:** $\mathbf{T}(\boldsymbol{\theta}) = e^{[\mathcal{S}_1]\theta_1} \dots e^{[\mathcal{S}_n]\theta_n} \mathbf{M}$ expresses forward kinematics as a product of matrix exponentials in the space frame.
2. **Screw Axes:** Each joint is characterized by a single 6D screw $\mathcal{S}_i$, eliminating frame attachment ambiguities of D-H.

3. **Spatial Jacobian:** Differentiation of the PoE formula yields the spatial Jacobian, relating joint velocities to end-effector twists.
4. **Body vs. Space Frame:** Two equivalent formulations exist; conversion via the Adjoint relates them.
5. **Singularities:** The rank of the Jacobian indicates dexterity; singular configurations arise when the Jacobian loses rank.
6. **Inverse Kinematics:** Newton-Raphson iterations using the pseudo-inverse Jacobian solve IK numerically.
7. **Computational Simplicity:** A single PoE algorithm handles revolute, prismatic, and spatial robots uniformly.

9.17 Further Reading and References

The Product of Exponentials formula has its intellectual roots in the Lie-theoretic treatment of rigid-body kinematics pioneered by Brockett [3], who first showed that the forward kinematics of a serial chain could be expressed as a product of matrix exponentials on $SE(3)$. Murray, Li, and Sastry [31] developed this insight into a comprehensive framework for robotic manipulation, providing the definitive mathematical treatment of screw theory, spatial velocities, and the PoE formula that forms the backbone of this chapter. For a modern, pedagogically oriented presentation of PoE-based kinematics—including extensive worked examples and software implementations—the reader is referred to Lynch and Park [27], whose *Modern Robotics* textbook has become the standard reference for graduate instruction.

The classical Denavit-Hartenberg convention, introduced in its original form by Denavit and Hartenberg [9], remains historically important and is still widely used in industrial robotics; understanding its relationship to PoE deepens one’s appreciation of why the modern formalism is preferred. For singularity analysis and manipulability theory beyond what is covered here, Siciliano et al. [38] provide an excellent treatment that bridges classical and modern perspectives. Numerical methods for inverse kinematics—including damped least-squares, pseudo-inverse techniques, and optimization-based approaches—are discussed in depth by Spong, Hutchinson, and Vidyasagar [40], who also cover the stability and convergence properties of Newton-Raphson iterations in the context of robotic systems.

9.18 Exercises

1. **Planar 2R Arm:** A planar 2R arm has link lengths $L_1 = 1.5$ m and $L_2 = 1$ m. Write the PoE formula in space-frame form, and compute the end-effector position at $\theta = [60^\circ, -30^\circ]$.
2. **D-H Comparison:** For the same 2R arm, compute the D-H parameters and verify that the forward kinematics matches the PoE result at the same configuration.
3. **Screw Axis Identification:** For a revolute joint rotating about the Y-axis and passing through the point $(2, 0, 1)$, write the screw axis $\mathcal{S}$.

4. **3D Cartesian Robot:** A 3D Cartesian robot has three prismatic joints (X, Y, Z directions). Write its PoE formula and forward kinematics. What is the spatial Jacobian?
5. **Spatial Jacobian:** For the planar 2R arm of Exercise 1, compute the spatial Jacobian at $\theta = [45^\circ, 0^\circ]$ and identify any singular configurations.
6. **Manipulability:** For the 2R arm, plot the manipulability measure $\sqrt{\det(\mathbf{J}\mathbf{J}^T)}$ as a function of $\theta_1 \in [0, 180^\circ]$ with θ_2 fixed at 0° .
7. **Singularity Analysis:** For the planar 2R arm, find all configurations where the Jacobian rank drops below 2.
8. **Body Frame PoE:** Rewrite the space-frame PoE of Exercise 1 in body-frame form using the Adjoint action.
9. **Inverse Kinematics:** Implement Newton-Raphson inverse kinematics for the planar 2R arm. Given a desired end-effector position $(x_d, y_d) = (2, 0.5)$, find θ starting from an initial guess.
10. **Velocity Kinematics:** For the 2R arm at $\theta = [30^\circ, 60^\circ]$ with $\dot{\theta} = [1, -0.5]$ rad/s, compute the end-effector linear and angular velocities.
11. **Prismatic-Revolute Arm:** A robot has a prismatic joint (Z-direction) followed by a revolute joint (Y-axis). Write the screw axes and PoE formula.
12. **Numerical Jacobian:** Implement the spatial Jacobian computation using finite differences and compare it to analytical differentiation for the 2R arm.
13. **Constraint Jacobian:** For the 2R arm, compute the task-space Jacobian relating joint angles to end-effector orientation ϕ (yaw angle in the plane).
14. **Null-Space Motion:** For a redundant planar 3R arm, identify the null space of the Jacobian at a non-singular configuration and describe a motion in null space.
15. **Configuration Space Visualization:** For the planar 2R arm, plot the reachable workspace (set of all possible end-effector positions) and overlay the singular configurations.

Chapter 10

The Articulated Body Algorithm

A torque at the wrist does not accelerate only the hand. The resulting joint accelerations depend on the configuration and inertia of the entire connected mechanism, the other applied torques, and its constraints. The articulated body algorithm (ABA) computes this coupled response efficiently for a rigid-body tree. Its value for golf is that it makes the mechanics calculable. It does not, by itself, identify a player's intention, solve ground contact, or determine which intervention improves club delivery.

10.1 Historical Context: Recursive Forward Dynamics

Featherstone's articulated-body formulation is a foundational recursive method for forward dynamics [10, 11]. It exploits the tree structure of interconnected rigid bodies, eliminating one joint acceleration at a time and then recovering all accelerations. For a bounded number of degrees of freedom per joint, the unconstrained calculation takes work proportional to the number of bodies. This is an algorithmic statement, not a measured wall-clock time or a guarantee of accuracy for a chosen anatomical model.

Many useful dynamics calculations instead assemble the generalized mass matrix and solve a linear system. Neither method requires an explicit numerical matrix inverse. The right choice depends on the model, the quantities required, repeated right-hand sides, constraints and implementation. A fast solver cannot compensate for an incorrect inertia, omitted joint torque or unrealistic contact assumption.

10.2 The Forward Dynamics Problem

10.2.1 Problem Formulation

For a fixed-base mechanism with independent scalar joint coordinates, write

$$M(q)\ddot{q} + h(q, \dot{q}) = \tau + Q_{\text{ext}}. \quad (10.1)$$

Here M is the generalized inertia, h contains velocity-dependent and gravitational terms, τ is the applied generalized joint force, and Q_{ext} represents other known external wrenches after projection into joint coordinates. A revolute joint's generalized force is a torque; a prismatic joint's is a force. All terms must use the same coordinate convention. Positive-definite M gives a unique acceleration for these specified loads. It does not imply that the motion is dynamically stable.

Forward dynamics predicts acceleration from the state and loads. Inverse dynamics computes the generalized loads required by a prescribed acceleration. With contact or closed loops, one must additionally specify or solve for constraint reactions. With muscles, activation, force-length-velocity behavior and moment arms determine which joint torques

1. Outward: Known Motion and Bias

Transforms, Velocities, Joint Bias and External Loads

**2. Inward: Eliminate Joint Accelerations**

Transmit Both Articulated Inertia and Complete Bias

**3. Outward: Recover Accelerations**

Use the Solved Parent Acceleration at Each Joint

Figure 10.1: Three Passes Through a Rigid-Body Tree: Compute Known Motion and Bias, Eliminate Joint Accelerations Inward, Then Recover Accelerations Outward.

can actually occur. Treating τ as a known input is a mechanical calculation, not a complete physiological model.

10.2.2 What the Complexity Comparison Means

A generic dense positive-definite solve uses an $O(N^3)$ factorization followed by $O(N^2)$ work per right-hand side. That does not make a thirty-joint calculation impossible. Constants, hardware, numerical libraries, sparsity and reuse matter. Conversely, $O(N)$ does not imply that every implementation meets a particular control deadline. An operation count of 10^3 at an assumed 10^9 operations per second would correspond to one microsecond, not one nanosecond; real processors do not convert every mathematical operation into one clock cycle.

ABA avoids constructing the complete mass matrix when the immediate requirement is a tree acceleration solve. Collision detection, constraint solving, muscle dynamics, flexible shafts, integration and output derivatives are additional work. Their cost must not be hidden inside a claim that the whole golf simulation is linear-time.

10.3 Articulated Inertia: The Key Insight

10.3.1 A Free Hinge Changes the Apparent Load

Imagine accelerating the proximal end of two connected links. If the hinge is locked, both links must satisfy that extra kinematic restriction. If it is free and its torque is specified, the

distal link can accelerate relative to the proximal link. These two experiments generally produce different proximal reactions. Articulated inertia describes the second experiment after the distal joint accelerations have been eliminated. Composite rigid-body inertia describes the locked collection used in a different calculation.

This distinction matters when discussing a wrist release. A freely moving wrist, a commanded joint torque, a stiff feedback controller and a mechanically locked joint are different boundary conditions. They should not be represented by exchanging descriptive words while leaving the equations unchanged. The inertia appearing in ABA is an instantaneous mechanical response at a specified configuration and specified joint-force model; it is not a frequency-dependent closed-loop impedance.

10.3.2 Mathematical Definition and Frames

Use angular-first spatial motion $v_i = (\omega_i, v_{O_i})$ and moment-first spatial force $f_i = (n_{O_i}, F_i)$, expressed in body frame i . Their pairing $f_i^T v_i$ is power. Let X_i map motion coordinates from parent $p(i)$ into child i for the same physical motion. Consequently a child wrench contributes $X_i^T f_i$ to its parent, and a child inertia contributes $X_i^T I_i X_i$. The transpose direction follows power invariance, not a visual resemblance between matrices.

After eliminating all joints strictly below body i , its subtree obeys

$$f_i^A = I_i^A a_i + p_i^A. \quad (10.2)$$

The unknown wrench f_i^A is the wrench applied to that subtree through its proximal connection. The subtree bias p_i^A includes known velocity terms, external loads and the effects of specified descendant torques. It is not merely friction or a force “dragging” the link. I_i^A depends on geometry and inertias, and on any included ideal armature terms; it is generally not the spatial inertia of one rigid body.

The scalar joint direction S_i maps $\dot{q}_i$ to relative spatial motion. It is a six-component motion vector, not an adjoint transformation. The derivation below assumes a constant subspace in the chosen joint/body frame, as for a standard one-degree-of-freedom revolute, prismatic or helical joint. General joint models need their own kinematic bias terms.

10.4 Derivation From First Principles

10.4.1 Eliminating One Joint Acceleration

At a fixed state, the kinematic acceleration relation is

$$a_i = X_i a_{p(i)} + c_i + S_i \ddot{q}_i. \quad (10.3)$$

Write $\bar{a}_i = X_i a_{p(i)} + c_i$. The unknown parent acceleration is kept separate from the known bias c_i . An ideal scalar joint satisfies $\tau_i = S_i^T f_i^A$. Substitute the subtree wrench relation to obtain

$$\tau_i = S_i^T I_i^A \bar{a}_i + S_i^T I_i^A S_i \ddot{q}_i + S_i^T p_i^A. \quad (10.4)$$

Define

$$U_i = I_i^A S_i, \quad d_i = S_i^T U_i, \quad u_i = \tau_i - S_i^T p_i^A. \quad (10.5)$$

Symmetry of I_i^A then gives

$$\ddot{q}_i = \frac{u_i - U_i^T \bar{a}_i}{d_i}. \quad (10.6)$$

The scalar u_i is a residual joint effort, not a six-dimensional wrench. For a physically nondegenerate model $d_i > 0$. Dividing by a small or nonpositive value without investigating the model is not a stability remedy.

Substituting this acceleration back into the wrench relation gives

$$f_i^A = I_i^a X_i a_{p(i)} + p_i^a, \quad I_i^a = I_i^A - \frac{U_i U_i^T}{d_i}, \quad (10.7)$$

$$p_i^a = p_i^A + I_i^a c_i + \frac{U_i u_i}{d_i}. \quad (10.8)$$

Both terms are essential. The remainder I_i^a represents the load transmitted after joint i is allowed to accelerate under its specified torque. The transmitted bias includes the joint's torque effect even when velocities vanish. Dropping that term would prevent a distal applied torque from correctly influencing the proximal acceleration.

For positive-definite I_i^A , set $z = (I_i^A)^{1/2} S_i$. Then

$$I_i^a = (I_i^A)^{1/2} \left(I - \frac{z z^T}{z^T z} \right) (I_i^A)^{1/2} \succeq 0, \quad I_i^a S_i = 0. \quad (10.9)$$

The projection occurs in these inertia-weighted coordinates. It is not a Euclidean projection that says only motion along the joint contributes to the parent load. Rather, one freely accelerating direction has been eliminated. This is a Schur complement, the same algebraic operation used in block Gaussian elimination.

10.5 The Articulated Body Algorithm: Three Passes

Number bodies so that parents precede children. The base is body 0, and each nonbase body has one incoming joint. Known external wrenches must be expressed about and in the same frame as each body's inertia. Here gravity is included through a fictitious base acceleration, so it must not also be included as a body external load.

10.5.1 Pass 1: Outward Velocity and Bias Calculation

Spatial Velocity Propagation

Start with $v_0 = 0$ for a fixed base. For each body in outward order,

$$v_i^J = S_i \dot{q}_i, \quad v_i = X_i v_{p(i)} + v_i^J. \quad (10.10)$$

The transformed parent velocity already carries the effect of the parent origin's offset. Adding untransformed six-vectors would mix different physical reference points as well as orientations.

Coriolis Acceleration and Bias Wrench

For angular-first $v = (\omega, v_O)$, define

$$v \times = \begin{bmatrix} [\omega]_{\times} & 0 \\ [v_O]_{\times} & [\omega]_{\times} \end{bmatrix}, \quad v \times^* = -(v \times)^T. \quad (10.11)$$

The cross operator acting on motion differs from its dual acting on momentum or force. Initialize

$$c_i = v_i \times v_i^J, \quad I_i^A = I_i, \quad p_i^A = v_i \times^* (I_i v_i) - f_i^{\text{ext}}. \quad (10.12)$$

For a more general joint, add its joint-specific acceleration bias to c_i . The Newton–Euler velocity wrench contains gyroscopic and centrifugal effects; it cannot be replaced by an arbitrary scalar multiple of angular speed. The external wrench is subtracted here because the unknown proximal wrench supplies only the part not already supplied externally.

10.5.2 Pass 2: Inward Subtree Elimination

Leaf Boundary Condition

A leaf starts with its own rigid inertia and the bias just calculated. It has no child contribution. Process the bodies in reverse order so that all children have already contributed before their parent is eliminated.

Recursive Inertia Update

Compute $U_i, d_i, u_i, I_i^a, p_i^a$ from the one-joint derivation and accumulate

$$I_{p(i)}^A += X_i^T I_i^a X_i. \quad (10.13)$$

The update adds every child’s contribution. A branching tree is handled by summation, not by assuming one child or by overwriting an earlier contribution. Preserve U_i, d_i, u_i for the final outward pass.

Required Bias Propagation

The matching update is

$$p_{p(i)}^A += X_i^T (p_i^A + I_i^a c_i + U_i u_i / d_i). \quad (10.14)$$

This update is part of ABA, not an optional refinement. Inertia alone cannot describe a subtree’s reaction at nonzero velocity or under a nonzero descendant torque. A static distal torque already changes the force balance transmitted proximally.

10.5.3 Pass 3: Outward Acceleration Recovery

Fixed-Base Acceleration and Gravity

For a stationary base in a uniform gravitational field g , set the algorithmic base acceleration to $a_0 = (0, -g)$. Alternatively, use the physical zero base acceleration and supply gravity as external body forces consistently. Do not use both conventions simultaneously. The fixed base has a reaction wrench, generally nonzero, that can be recovered by inverse dynamics once the accelerations are known.

Spatial acceleration is also not simply a stack of ordinary angular and origin linear accelerations. In body coordinates its linear part is the ordinary origin acceleration minus $\omega_i \times v_{O_i}$, before accounting for the fictitious gravity convention. Comparing these arrays with accelerometer data therefore requires restoring the physical gravity/reference convention and the sensor’s offset.

Joint Acceleration Recovery

For each body in outward order, compute

$$\bar{a}_i = X_i a_{p(i)} + c_i, \quad \ddot{q}_i = (u_i - U_i^T \bar{a}_i)/d_i, \quad a_i = \bar{a}_i + S_i \ddot{q}_i. \quad (10.15)$$

No trigonometric approximation has been introduced by this elimination. It solves the stated rigid-body equations up to arithmetic error. Its agreement with an independent formulation checks the implementation, while experimental agreement checks the physical model. Those are separate questions.

10.6 Why the Tree Calculation Is Linear-Time

Each body requires a bounded number of operations on six-vectors and 6×6 matrices. Each edge transmits one inertia and one bias contribution. The three traversals therefore use $O(N)$ arithmetic and $O(N)$ storage for scalar joints. A joint with several degrees of freedom replaces d_i by the block $S_i^T I_i^A S_i$ and requires a small linear solve. The linear-in-body-count statement assumes that block size is bounded.

Allocation and data structures also matter. Preallocated arrays or append-only lists followed by reverse traversal preserve the intended scaling. Repeated insertion at the front of a growing list can add quadratic data movement. Computing a complete dense derivative matrix, detecting contacts or solving a large contact system is not covered by the three-pass count.

10.7 Floating Base: A Six-Degree-of-Freedom Root

10.7.1 The Root Articulated Solve

For a freely moving base, its acceleration is unknown. Initialize the base's own inertia and bias, include its known external wrench, and accumulate all child contributions. With the same fictitious-gravity convention and no unknown root constraint wrench,

$$I_0^A a_0 + p_0^A = 0, \quad a_0 = -(I_0^A)^{-1} p_0^A. \quad (10.16)$$

Implement this expression as a 6×6 solve. The matrix is the entire articulated subtree's inertia at the current configuration, not the isolated base's constant rigid inertia. Its factorization can be reused only while the quantities on which it depends remain unchanged. The remaining outward acceleration recovery is the same as before.

10.7.2 Contact Does Not Automatically Fix the Base

A foot touching the ground may impose one normal constraint, several sticking constraints, or a finite contact-patch wrench model. It is not automatically a six-degree-of-freedom weld. A no-slip constraint is valid only while the required reaction satisfies the admitted friction and unilateral conditions. If the contact wrench is unknown, the root solve and contact equations must be resolved together; one cannot first assume an arbitrary base acceleration and then declare the resulting contact physically feasible.

10.8 Floating-Base Humanoids and Momentum

10.8.1 Configuration Coordinates and Generalized Velocity

A humanoid model may use a position and unit quaternion for the pelvis, together with joint coordinates. The quaternion has four stored components but only three rotational degrees of freedom. Write the state kinematics as $\dot{q} = N(q)\nu$, and dynamics in the generalized velocity ν . In general $\nu \neq \dot{q}$. A six-dimensional root joint is a local motion description, not a requirement to use three globally nonsingular orientation angles; no such global three-angle chart exists.

10.8.2 Root Equation With a Known Applied Wrench

If a known additional root wrench w_0 is kept outside the bias convention, the root relation is

$$I_0^A a_0 = w_0 - p_0^A. \quad (10.17)$$

Count a wrench exactly once. An unknown ground reaction belongs to the constraint calculation, while a measured ground wrench can be supplied as known data within its measurement uncertainty. These choices answer different forward or inverse questions.

10.8.3 Centroidal Momentum and Whole-Body Dynamics

Keep the angular-first ordering and define momentum about the total center of mass c :

$$h_G = \begin{bmatrix} k_c \\ p \end{bmatrix} = A_G(q)\nu, \quad \dot{h}_G = \begin{bmatrix} \sum_j ((r_j - c) \times F_j + n_j) \\ m_{\text{tot}}g + \sum_j F_j \end{bmatrix}. \quad (10.18)$$

The sums contain nongravitational external forces and couples. Uniform gravity has zero resultant moment about the total COM. Internal joint forces and torques cancel in the whole-system balance, but joint motion redistributes momentum among segments and changes A_G , the COM and contact moment arms. “Internal torques cancel” does not mean configuration is irrelevant.

For a golfer–club system, hand–club wrenches are internal only if both hands and the club are included in the system boundary. Ground reactions remain external. After ball contact, the ball must either be included or represented through an external interaction. This bookkeeping is necessary before attributing a rise in club momentum to any particular segment or source of work.

10.9 Joint Friction and Motor Inertia

10.9.1 Viscous Joint Friction

For an ideal viscous coefficient $b_i \geq 0$, use $\tau_i^{\text{eff}} = \tau_i^{\text{act}} - b_i \dot{q}_i$ in the joint elimination. The associated power is $-b_i \dot{q}_i^2 \leq 0$. The induced acceleration is coupled through the dynamics; subtracting $b_i \dot{q}_i$ directly from a computed acceleration has both the wrong units and the wrong mechanical effect. Dry friction at zero velocity requires a separate set-valued or regularized model, not an arbitrary choice of sign presented as a universal law.

10.9.2 An Ideal Reflected Armature Model

If a transmission is modeled by an independent rotor kinetic energy $J_i(r_i\dot{q}_i)^2/2$, its reflected armature is $a_i^{\text{rot}} = r_i^2 J_i$. Under precisely that idealization, the joint equation includes $a_i^{\text{rot}}\ddot{q}_i$, and the pivot becomes

$$d_i = S_i^T I_i^A S_i + a_i^{\text{rot}}. \quad (10.19)$$

Use this pivot in the remainder, transmitted bias and acceleration recovery. The remainder no longer generally annihilates S_i . General rotor/body coupling, flexible transmissions and configuration-dependent gearing require a fuller kinetic model. Muscles are not rotary motors with a freely chosen reflected rotor inertia; their activation and force-generation dynamics must be modeled on their own terms.

10.10 ABA and Composite Rigid-Body Methods

The composite rigid-body algorithm (CRBA) constructs $M(q)$ from locked-subtree inertias. Its worst-case assembly cost is $O(N^2)$ for a scalar-joint tree. In a serial chain, distant joint velocities can contribute to the same body's kinetic energy, so the mass matrix is generally dense. A branched tree has an ancestor-dependent sparsity pattern, not a universally bounded bandwidth.

For one unconstrained right-hand side, ABA can avoid forming M . If an application already needs M for constraints, identification or optimization, factorization may be useful. At fixed configuration, a dense factorization is reused at $O(N^2)$ per additional right-hand side; a suitably organized articulated factorization can apply the inverse in $O(N)$ per right-hand side. Constructing R complete solutions still requires their output and load processing. No numeric runtime comparison follows from stating these asymptotic orders alone.

10.11 Choosing Forward or Inverse Dynamics

10.11.1 When the Acceleration Is Prescribed

Use inverse dynamics to find the generalized force required by a specified motion and external-load model. In a measured swing, differentiation noise, inertial estimates and ground/contact measurements strongly affect this calculation. A net joint torque is not a unique decomposition into individual muscle forces. If external forces are unknown, kinematics alone may not determine all loads.

10.11.2 When the Applied Loads Are Specified

Use forward dynamics to compare predicted accelerations or trajectories under explicit loading or control changes. Holding all other joint torques fixed while changing one torque gives a well-defined instantaneous intervention in the unconstrained model. Holding a feedback policy fixed over a new trajectory is a different experiment from replaying the original torque history. A claimed improvement in club delivery must state which experiment was performed and how ground, grip and muscle feasibility were maintained.

10.12 Constraint Handling: Contact and Joint Limits

10.12.1 Unilateral Contact

Let $\phi(q)$ be a signed normal gap: $\phi > 0$ means separation and $\phi = 0$ means touching. An ideal rigid contact requires $\phi \geq 0$ and a compressive normal reaction $\lambda_n \geq 0$. A separated contact does not supply a force merely because its unconstrained acceleration points toward the surface. A closing trajectory needs event detection or a consistent time-stepping scheme; sustained-contact acceleration equations do not describe an instantaneous impact.

10.12.2 Joint Limits

A hard lower or upper joint limit can be represented by a unilateral gap in configuration space, with its own reaction or impact rule. A soft limit instead supplies a specified force law. Abruptly clipping the angle after integration can violate momentum and inject or remove energy without a stated model. Anatomical range limits, passive tissue forces and active control are related but distinct phenomena.

10.13 Connection to Gaussian Elimination

The one-joint derivation is block elimination of coupled wrench and acceleration equations. ABA performs this elimination on a spatial tree representation. It does not require the generalized mass matrix to be band diagonal. The Schur complement is exact for the stated algebra; its numerical evaluation still depends on scaling and conditioning. Very small inertias, inconsistent units, nonphysical mass parameters and cancellation can produce poor numerical results in a recursive or matrix formulation.

For example, partition a three-joint mass matrix into proximal coordinate q_1 and distal coordinates q_d . At zero velocity and gravity, with $\tau_d = 0$,

$$M_{\text{eff}} = M_{11} - M_{1d}M_{dd}^{-1}M_{d1}, \quad \ddot{q}_1 = \tau_1/M_{\text{eff}}. \quad (10.20)$$

With the distal joints locked instead, the corresponding coefficient is M_{11} . The difference expresses the mechanical effect of allowing relative motion, not evidence that a golfer should minimize effective inertia throughout the swing. Club orientation, feasible torque, shaft deformation and impact conditions impose additional objectives and restrictions.

10.14 Worked Example: A Fully Specified Three-Link Arm

Consider three uniform solid cylinders of lengths (0.4, 0.3, 0.25) m, masses (2.0, 1.2, 0.6) kg and common radius 0.025 m. Their local x axes run along the links; all joints rotate about local z . Each body frame originates at its proximal joint. This is an ideal teaching mechanism, not measured anatomy or a fitted golf club. The finite radius gives physically consistent positive rotational inertias.

For each cylinder, $c_i = (L_i/2, 0, 0)$ and

$$I_{C_i} = \text{diag} \left(\frac{m_i r^2}{2}, \frac{m_i (L_i^2 + 3r^2)}{12}, \frac{m_i (L_i^2 + 3r^2)}{12} \right), \quad (10.21)$$

$$I_i = \begin{bmatrix} I_{C_i} + m_i[c_i]_{\times}[c_i]_{\times}^T & m_i[c_i]_{\times} \\ m_i[c_i]_{\times}^T & m_i I_3 \end{bmatrix}. \quad (10.22)$$

The upper-left parallel-axis contribution is added. The off-diagonal blocks follow from the COM offset and must not be invented independently.

10.14.1 Configuration and Loads

Use relative angles $q = (0.4, -0.7, 0.2)$ rad, rates $\dot{q} = (0.8, -0.5, 1.1)$ rad/s and actuator torques $\tau = (4, -1, 0.3)$ N m. Gravity is $(0, -9.81, 0)$ m/s². There are no other external wrenches, no damping and no armature in this particular calculation.

Let $E_i = R_z(q_i)^T$ reexpress parent-oriented components in child axes, and let $r_i = (L_{i-1}, 0, 0)$ be the parent-frame joint offset, with $r_1 = 0$. Then

$$X_i = \begin{bmatrix} E_i & 0 \\ -E_i[r_i]_{\times} & E_i \end{bmatrix}, \quad S_i = (0, 0, 1, 0, 0, 0)^T. \quad (10.23)$$

The joint angles therefore enter every transform. The relative angles are not interchangeable with the links' absolute orientations, which are $(0.4, -0.3, -0.1)$ rad.

10.14.2 Pass 1: Velocities and Bias

All out-of-plane components vanish. Listing each spatial velocity as the planar triple (ω_z, v_x, v_y) , the outward recursion gives

$$\begin{aligned} v_1 &= (0.8, 0, 0), \\ v_2 &= (0.3, -0.206149660, 0.244749500), \\ v_3 &= (1.4, -0.135535933, 0.369032412). \end{aligned} \quad (10.24)$$

The angular components sum relative rates, while the linear components also require rotations and joint-origin offsets. Form $c_i = v_i \times S_i \dot{q}_i$ and the initial $p_i^A = v_i \times^* I_i v_i$ from these complete velocities. There is no arbitrary bias-force estimate.

10.14.3 Pass 2: Articulated Inertias and Residual Efforts

The planar subblocks of the articulated inertias, before eliminating each body's own joint, are approximately

$$I_3^A|_{zxy} = \begin{bmatrix} 0.012593750 & 0 & 0.075000000 \\ 0 & 0.600000000 & 0 \\ 0.075000000 & 0 & 0.600000000 \end{bmatrix}, \quad (10.25)$$

$$I_2^A|_{zxy} = \begin{bmatrix} 0.051575604 & 0.026090063 & 0.231293680 \\ 0.026090063 & 1.782370942 & 0.086966875 \\ 0.231293680 & 0.086966875 & 1.370978934 \end{bmatrix}, \quad (10.26)$$

$$I_1^A|_{zxy} = \begin{bmatrix} 0.260428859 & -0.284953502 & 0.783624232 \\ -0.284953502 & 3.143841922 & -0.712383755 \\ 0.783624232 & -0.712383755 & 2.959060579 \end{bmatrix}. \quad (10.27)$$

The rotational, mixed and translational entries carry different SI units. These subblocks are sufficient to display this planar example, but the implementation uses full six-dimensional quantities. The pivots and residual joint efforts are

$$\begin{aligned} d &= (0.260428859, 0.051575604, 0.012593750), \\ u &= (6.518800728, -1.799610488, 0.314231273). \end{aligned} \quad (10.28)$$

Gravity has not yet entered these residual efforts because it enters through the fictitious base acceleration in pass 3. The residuals differ from the supplied torques because the inward bias carries descendant loads and velocity effects.

10.14.4 Pass 3 and an Independent Check

Outward recovery gives, in rad/s^2 ,

$$\ddot{q} = (2.023035078, -79.465503634, 175.204404572). \quad (10.29)$$

These large distal accelerations reflect the particular light-link model and loads. They are not plausible-muscle bounds or a prediction for a human swing. To check the result independently, form the Cartesian COM Jacobians J_{C_i} and angular Jacobians J_{ω_i} from the link positions. Kinetic energy gives

$$M = \sum_i (m_i J_{C_i}^T J_{C_i} + J_{\omega_i}^T I_{C_i}^{\text{world}} J_{\omega_i}). \quad (10.30)$$

For these planar coordinates, differentiating this mass matrix gives the Christoffel velocity terms, and differentiating gravitational potential gives the gravitational terms. The independent calculation produces

$$M = \begin{bmatrix} 0.814792916 & 0.283348998 & 0.060972725 \\ 0.283348998 & 0.146884246 & 0.034645248 \\ 0.060972725 & 0.034645248 & 0.012593750 \end{bmatrix}, \quad (10.31)$$

$$h = (14.185426211, 4.029005575, 0.723271650)^T. \quad (10.32)$$

Solving $M\ddot{q} = \tau - h$ agrees with ABA. The displayed values are rounded; residual checks use full precision. Tests additionally check nonzero external wrenches, damping and armature, multiple configurations, branched prismatic motion and reexpression in different body frames. Agreement across these cases is stronger evidence than checking an example against a second transcription of the same inward recursion.

10.15 Reproducible Python Example

The repository implementation is split between the generic prepared-tree solver and a cylinder-chain model builder. It uses the existing checked spatial-inertia and cross-product functions. Its inputs are explicitly restricted to a fixed-base rigid tree with constant local scalar-joint subspaces. It does not silently treat a contact or a flexible link as such a joint. The following example runs from the repository root:

```

1 import numpy as np
2 from src.tools.articulated_chain_examples import cylinder_chain
3 from src.tools.articulated_body_examples import (
4     DynamicsLoads, forward_dynamics, inverse_dynamics,
5 )
6
7 tree = cylinder_chain(
8     np.array([0.4, 0.3, 0.25]),
9     np.array([2.0, 1.2, 0.6]),
10    np.array([0.4, -0.7, 0.2]),
11    0.025,
12 )
13 loads = DynamicsLoads(
14     velocity=np.array([0.8, -0.5, 1.1]),
15     torque=np.array([4.0, -1.0, 0.3]),
16 )
17 result = forward_dynamics(tree, loads)
18 print(result.joint_acceleration)
19 print(result.pivot)
20 recovered = inverse_dynamics(
21     tree, loads, result.joint_acceleration
22 )
23 np.testing.assert_allclose(recovered, loads.torque)

```

The diagnostic result exposes velocities, articulated inertias, biases, pivots and residual efforts for inspecting the derivation. Array shapes, finite loads, ordered parents and positive-definite body inertias are checked. Physical rigid-body inertia structure and valid rigid motion transforms remain modeling preconditions. The supplied cylinder builder constructs them from positive dimensions. Invalid or nonpositive articulated pivots raise an error; assigning zero acceleration would conceal a broken or degenerate model.

In NumPy, a one-dimensional vector's transpose is still one-dimensional. To form $U_i U_i^T$, call `np.outer` with the vector as both arguments. For a one-dimensional array, `U @ U.T` instead returns a scalar, which can silently broadcast into an incorrect matrix update.

10.15.1 Connection to Simulation Frameworks

Different simulators use different combinations of recursive dynamics and constraint solvers. MuJoCo's documented computation path forms its mass matrix with CRB and uses sparse $L^T D L$ factorization and back-substitution [30]. It must not be described as using ABA for every timestep. Pinocchio provides ABA and analytical dynamics derivatives, along with other algorithms [33]. These statements identify documented methods; they do not rank accuracy or runtime. A reproducible comparison must identify software versions, model parameters, contact and integration settings, hardware and the actual quantity timed.

10.16 Contact Forces and Impact Dynamics

During smooth motion, a force changes acceleration. An ideal instantaneous impact changes velocity while configuration remains continuous. For a single frictionless contact with normal Jacobian J_n , closing velocity $v_n^- = J_n \nu^- < 0$ and restitution coefficient $0 \leq e \leq 1$,

$$M(\nu^+ - \nu^-) = J_n^T \Lambda_n, \quad J_n \nu^+ = -e J_n \nu^-, \quad (10.33)$$

$$\Lambda_n = -\frac{(1+e)J_n\nu^-}{J_n M^{-1} J_n^T}. \quad (10.34)$$

Here M and J_n are evaluated at the impact configuration and the denominator is positive for a nonzero contact direction and positive-definite M . The resulting impulse is compressive. It has impulse units, unlike the sustained reaction force. Multiple simultaneous contacts, friction and restitution require a specified coupled impact law; independent application of this scalar formula need not be consistent.

Ball impact also involves ball spin, clubface geometry and material deformation. This scalar illustration does not replace a validated ball–club collision model. It explains why comparing only pre-impact accelerations cannot establish a post-impact delivery or ball-flight claim.

10.17 Contact Dynamics and Complementarity

10.17.1 Non-Penetration and Sustained Contact

At a closed contact with zero normal velocity, define normal acceleration $a_n = J_n \dot{\nu} + \dot{J}_n \nu$. An ideal frictionless acceleration-level condition is

$$a_n \geq 0, \quad \lambda_n \geq 0, \quad a_n \lambda_n = 0. \quad (10.35)$$

The alternatives are separating acceleration with zero reaction, or a compressive reaction maintaining zero normal acceleration. These conditions apply to the specified closed, non-impacting candidate contacts. They are not a universal replacement for the gap and velocity conditions along an arbitrary trajectory.

10.17.2 Friction and Contact Wrenches

For a point contact, the Coulomb force cone is

$$\|\lambda_t\| \leq \mu \lambda_n, \quad \lambda_n \geq 0. \quad (10.36)$$

This inequality bounds the tangential force but does not specify its value during sliding. A sliding law ordinarily also imposes opposition to slip or a maximum-dissipation condition. A finite foot patch can transmit moments subject to geometry and pressure restrictions; a three-dimensional force cone is not the complete admissible six-dimensional wrench set. Rigid frictional contact can exhibit nonuniqueness or inconsistency, so no simple cone equation should be advertised as solving all contact physics exactly.

10.17.3 The Frictionless Acceleration LCP

For a chosen set of closed, zero-normal-velocity contacts, write

$$M\dot{\nu} + h = \tau + J_n^T \lambda_n. \quad (10.37)$$

Let $a_{\text{free}} = M^{-1}(\tau - h)$, $W = J_n M^{-1} J_n^T$ and $b = J_n a_{\text{free}} + \dot{J}_n \nu$. Then

$$a_n = W \lambda_n + b, \quad 0 \leq \lambda_n \perp a_n \geq 0. \quad (10.38)$$

The matrix W is positive semidefinite. Redundant contact directions can make reaction forces nonunique even when the resulting acceleration is determined. Friction introduces additional modeling choices; this frictionless LCP is not the full frictional problem.

ABA can supply products with M^{-1} for constructing or applying W . Each such product is a tree solve, but assembling and solving the contact problem has its own cost. The number and arrangement of contacts and the chosen numerical method matter. Soft-contact formulations, including MuJoCo's documented model, make different assumptions from exact rigid complementarity.

10.17.4 Time-Stepping and Numerical Order

A velocity step using an impulse Λ may take the form

$$M(q_k)(\nu_{k+1} - \nu_k) = \Delta t (\tau_k - h_k) + J(q_k)^T \Lambda. \quad (10.39)$$

This frozen-coefficient expression is not, by itself, full implicit Euler. The contact law, evaluation points and configuration update define the method. Restitution and stabilization also affect the result.

The familiar velocity-first Euler update is first order. In canonical separable Hamiltonian coordinates the corresponding symplectic Euler method is symplectic; a velocity-first update for a configuration-dependent multibody mass matrix does not inherit that claim automatically. Velocity Verlet is second order for its appropriate position-force/separable setting. General velocity-dependent multibody dynamics and impacts require an integrator suited to those equations. Check convergence as the timestep decreases, constraint residuals and energy/momentum behavior under the stated force model. An accurate acceleration solve does not remove integration error.

10.18 From Coupled Dynamics to a Golf-Swing Argument

For smooth unconstrained motion at a fixed state, the immediate torque-to-acceleration sensitivity is $\delta\ddot{q} = M^{-1}\delta\tau$. With bilateral constraints $J\dot{\nu} + b_c = 0$ and full-row-rank J , solving for the reaction gives the constrained acceleration sensitivity

$$\delta\dot{\nu} = P_M \delta\tau, \quad P_M = M^{-1} - M^{-1} J^T (J M^{-1} J^T)^{-1} J M^{-1}. \quad (10.40)$$

The reaction changes when a torque changes. Holding the old contact force fixed would answer a different question and can violate the constraint. Both hands on a rigid club close a kinematic loop through the arms and torso; a cut-tree ABA calculation must restore that loop through constraints or an explicitly compliant grip model.

This instantaneous response is only one component of delivery. If x includes configuration, velocity, muscle activation and any shaft states, an implemented policy produces a field $\dot{x} = F_\pi(x, t)$. For an additional input perturbation with local input derivative B_π , the finite-interval variational equation is

$$\delta\dot{x} = D_x F_\pi \delta x + B_\pi \delta u. \quad (10.41)$$

The state derivative includes the feedback policy's response to state changes. The club's selected output at a fixed time depends on the propagated state perturbation and the output

derivative. At a state-triggered impact, event-time and reset sensitivities are also required. A torque change may improve one output while worsening face angle, path, contact location or robustness. Segment speed peaks alone do not identify which muscles caused those changes.

A strong swing analysis therefore states its system boundary, inertia and joint model, admitted ground/grip reactions, actuator policy, perturbation and delivery outcome. It checks mechanical consistency before drawing a coaching conclusion. Energy and momentum balances supply complementary checks: internal joint action can redistribute momentum and perform work, but a sequence of angular-speed maxima is neither an energy-flow measurement nor proof of a unique optimal release strategy.

10.19 Differentiating the Dynamics

For a smooth, nonsingular fixed model, differentiating $M\ddot{q} + h = \tau$ gives

$$\delta\ddot{q} = M^{-1}(\delta\tau - \delta h - (\delta M)\ddot{q}). \quad (10.42)$$

This identity clarifies the derivative target independently of a software implementation. Automatic differentiation can propagate derivatives through a correctly implemented smooth recursion. Analytical derivative algorithms provide another route. Neither produces exact real-number arithmetic or an automatic derivative through every switching contact event. Differentiating the integrated or hybrid map requires its additional chain rules and event assumptions.

The full dense Jacobian has quadratically many entries in the number of joints. A directional derivative or reverse-mode derivative of a scalar objective can have a different cost from materializing every entry. Benchmark the actual requested derivative and include contact, controller and integration work before asserting a learning or control rate.

10.20 Chapter Summary

ABA solves the specified rigid-tree forward problem through outward kinematics, inward elimination and outward recovery. The essential quantities are the articulated inertia and its complete bias, transformed through power-consistent frames. Articulating and locking a subtree are different mechanical experiments. Floating bases, actuator models, contact and impact require their own explicit assumptions. The independently checked example establishes computational consistency; applying the method to golf additionally requires a validated model and a clearly defined causal or predictive question.

10.21 Further Reading and References

Featherstone's original formulation and later treatment provide the recursive dynamics foundation [10, 11]. His `Spatial_v2` documentation provides inspectable reference algorithms and coordinate-transform conventions [13]. MuJoCo's computation documentation and Pinocchio's software description document distinct implementation choices [30, 33]. The calculations here are derived explicitly and checked independently; consulting these sources does not substitute for stating the model and experiment used in a golf argument.

10.22 Exercises

1. **Complexity Analysis:** Count the bounded-size operations in all three passes for scalar joints. Identify the assumptions needed for $O(N)$ work and storage. Explain why constructing all dense acceleration derivatives or solving contact is additional work.
2. **Articulated and Composite Inertia:** Use three cylinders with $m_i = 1$ kg, $L_i = 1$ m, radius 0.05 m and $q = (0.2, -0.4, 0.3)$ rad. Calculate every articulated inertia with the distal joints free, and every composite inertia with them locked. Explain the difference without treating a singular point-mass inertia as a generic rigid body.
3. **Motor Inertia:** Add 0.04 kg m² of ideal armature to joint 1 of the worked example. Recalculate all accelerations. Derive the sensitivity using the rank-one change in M , and explain why distal accelerations can change even though their own armatures are unchanged.
4. **Floating Base:** Describe the extra root quantities and solve for a thirty-joint floating-base mechanism. Separate its additional bounded-size root work from the effects of adding contact constraints. Do not infer an exact runtime ratio from the degree-of-freedom count.
5. **Repeated Right-Hand Sides:** At fixed q and $\dot{q}$, compare reuse of a dense mass factorization and an articulated factorization for one hundred torque vectors. State preparation, per-solve and storage costs separately. Explain which quantities must be refreshed if $\dot{q}$ or q changes.
6. **Bias Verification:** Restrict the worked model to its first two links at $q = (0.4, -0.7)$ and $\dot{q} = (0.8, -0.5)$. Calculate velocity forces from the spatial recursion and from Christoffel terms of the Cartesian kinetic-energy mass matrix. Compare gravity separately to avoid confusing force conventions.
7. **Singular Pivots:** Prove $d_i > 0$ for a nonzero S_i and positive-definite I_i^A without armature. Explain how massless idealizations, degenerate coordinates or nonphysical parameters can defeat this condition. Why is replacing a zero pivot by zero acceleration not a valid solution?
8. **Joint Damping:** Add coefficients (0.2, 0.1, 0.03) N m s/rad to the worked example. Verify $\tau^{\text{eff}} = \tau - B\dot{q}$ before the solve and that the damping power is nonpositive. Compare against the incorrect post-solve acceleration subtraction.
9. **Gravity Simulation:** Construct four identical cylinders of length 0.3 m, mass 1 kg and radius 0.025 m, initially at $q = (0.4, -0.2, 0.3, -0.1)$ with zero rates. Simulate gravity-only motion with a fixed proximal base. State the integrator, refine the timestep, and monitor total mechanical energy; a fixed-base pendulum chain is not whole-body ballistic free fall.
10. **Contact Impulse:** For the first two worked links at $q = (0.4, -0.7)$, define a horizontal surface through the distal endpoint and use $\dot{q} = (-0.8, -0.5)$. Form the endpoint's upward normal Jacobian. Compute the single frictionless impulse for $e = 0$ and verify the post-impact normal velocity and nonnegative impulse. Distinguish this velocity jump from subsequent sustained-contact acceleration.

11. **Integration Order:** Compare a first-order velocity-first method with a suitable second-order method on a smooth problem. Verify observed convergence against a refined reference. Explain when the words symplectic and Verlet apply, and why an impact can invalidate a smooth-step order argument.
12. **Matrix Formulation:** Build a specified ten-link cylinder chain and compare ABA with a positive-definite linear solve using its assembled M and bias. Report residuals and conditioning; do not explicitly invert M merely to apply it to one vector.
13. **Floating-Base Reactions:** For a specified humanoid model, choose either a point-foot or a finite-patch contact model. Write the root and contact equations together. State the rank, friction and unilateral conditions required for the computed reaction, and explain what measured contact data would remove from the unknowns.
14. **Redundancy:** Choose a task Jacobian J for the worked model. Examine the map $JM^{-1}\delta\tau$ and construct a torque perturbation with zero immediate task acceleration when such a null space exists. Explain why that does not imply zero state change or unchanged delivery at a later time.
15. **Constrained Projection:** Let $M \succ 0$ and full-row-rank J be fixed, and require $Ja + b_c = 0$. Minimize $(a - a_{\text{free}})^T M (a - a_{\text{free}})/2$ under that constraint. Derive $a = a_{\text{free}} - M^{-1}J^T(JM^{-1}J^T)^{-1}(Ja_{\text{free}} + b_c)$. Explain why this is a mass-weighted closest acceleration, not an arbitrary Euclidean minimum-norm forward solution, and discuss redundant constraints separately.

Chapter 11

Lagrangian Mechanics

Why break a machine apart into agonizing lists of individual pulling and pushing forces when you can extract its exact motion by simply observing where its energy wants to go?

In Plain Language 11.1. Context & Use Case: Energy-Based Dynamics

The Math You Will See: We will define the **Lagrangian** ($\mathcal{L}$) and the **Euler-Lagrange Equations**. We will derive the classic standard form: $\mathbf{M}(\mathbf{q})\ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})\dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \boldsymbol{\tau}$. **Why We Need It:** The Recursive Newton-Euler Algorithm (Chapter 7) is unmatched for fast computer simulations. However, when we need to mathematically *prove* that a robot won't fall over, or design an optimal control law like LQR or Contraction Metrics, we need explicit equations. Lagrangian mechanics abandons vectors and forces entirely. It looks purely at the kinetic and potential *energy* of the system, gracefully ignoring invisible internal constraint forces (like the tension in a rigid arm) to instantly yield elegant, closed-form equations of motion. **All Variables Defined:** T is Kinetic Energy, V is Potential Energy, and $\mathcal{L} = T - V$.

11.1 Historical Context: From Newton to Lagrange to Hamilton

The development of Lagrangian mechanics represents one of the most far-reaching intellectual shifts in the history of mechanics. Isaac Newton formulated motion through explicit forces acting on point masses: $\sum \mathbf{F} = m\mathbf{a}$. This framework dominated for over a century and remains the foundation of computational physics engines.

However, in 1788, **Joseph-Louis Lagrange** published his *Mécanique Analytique*, introducing a radically different approach. Rather than tracking forces, Lagrange proposed that mechanical systems evolve along paths that extremize a single scalar function: the **action integral**

$$\mathcal{S}[\mathbf{q}(t)] = \int_{t_1}^{t_2} \mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}, t) dt$$

where $\mathcal{L} = T - V$ is the Lagrangian. This became known as the **Principle of Least Action**, or equivalently, Hamilton's Principle.

Nearly simultaneously, **Leonhard Euler** was developing the **calculus of variations**—the mathematics required to find functions that minimize or maximize functionals (functions of functions). The union of Euler's calculus of variations with Lagrange's mechanical insight created the Euler-Lagrange equations, which automatically derive equations of motion from any Lagrangian density.

Later, in the early 1800s, **William Rowan Hamilton** reformulated Lagrange's framework using a Legendre transform, introducing **conjugate momenta** and the **Hamiltonian** function $H = \mathbf{p} \cdot \dot{\mathbf{q}} - \mathcal{L}$. This opened pathways to modern quantum mechanics and symplectic geometry.

Why This History Matters: Lagrangian mechanics is not merely an alternative notation for Newton's laws. It encodes deep philosophical truth: nature does not "know" about forces. It only "knows" about energy. Systems evolve along extremal paths of action. This principle applies equally to light rays (Fermat's principle), quantum fields (path integrals), and relativistic gravity (Einstein's equations). By learning Lagrangian mechanics, you are learning the geometric language that unifies all modern physics.




Figure 11.1: Example of Lagrangian Mechanics: A Pendulum System. The Motion Is Fully Determined by the Difference Between Kinetic and Potential Energy (The Lagrangian), Not by Explicitly Tracking Forces. The Euler-Lagrange Equations Automatically Yield the Equations of Motion Without Computing Constraint Forces.

11.2 The Limitations of Newton

In the Newtonian framework (utilized powerfully by the Recursive Newton-Euler algorithm in Chapter 7), motion is determined by analyzing explicit spatial forces: $\sum \mathbf{F} = m\mathbf{a}$.

While Newton's equations are spectacular for computational algorithms evaluating rigid bodies link-by-link in a physics engine, they suffer from a fatal flaw when deriving mathematical models by hand: **Constraint Forces**.

If a bead slides along a curved wire, or a heavy pendulum swings on a rigid rod, Newton demands that you calculate the exact, microscopic physical force the wire exerts on the bead, or the tension the rod pulls on the bob, just to ensure they don't break. For a 15-DOF mechanism, solving for 15 unknown internal constraint forces pulling against each other is mathematically agonizing and analytically useless for control theory.

Moreover, constraint forces do *no work* in the motion. The tension in a rigid pendulum rod redirects motion but adds no energy. Yet in Newton's framework, we must solve for it anyway. This is deeply inelegant.

Lagrangian mechanics solves this by working exclusively with **generalized coordinates**—minimal, unconstrained variables. The constraints are built into the coordinate system itself. Internal forces vanish from the mathematical framework entirely.

11.3 Hamilton's Principle and the Calculus of Variations

Hamilton's Principle [14, 1] states that among all conceivable paths $\mathbf{q}(t)$ connecting two configurations $\mathbf{q}(t_1)$ and $\mathbf{q}(t_2)$ at fixed times, the *actual* physical path satisfies:

$$\delta S = \delta \int_{t_1}^{t_2} \mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}, t) dt = 0 \quad (11.1)$$

where δ denotes a variation—an infinitesimal perturbation to the path that respects the boundary conditions $\delta \mathbf{q}(t_1) = 0$ and $\delta \mathbf{q}(t_2) = 0$.

This is not a minimum (as popularized, though not always accurate). It is a **stationary point**—an extremum (often a minimum, sometimes a saddle point). The physical system behaves as if “nature chooses” the path where small perturbations produce zero first-order change in the action.

To extract equations of motion from Hamilton's principle, we apply the calculus of variations. Consider a perturbed path $\mathbf{q}(t) + \epsilon \boldsymbol{\eta}(t)$, where $\boldsymbol{\eta}(t_1) = \boldsymbol{\eta}(t_2) = 0$ are arbitrary variations vanishing at the boundaries. The action becomes:

$$S[\epsilon] = \int_{t_1}^{t_2} \mathcal{L}(\mathbf{q} + \epsilon \boldsymbol{\eta}, \dot{\mathbf{q}} + \epsilon \dot{\boldsymbol{\eta}}, t) dt$$

The stationarity condition is:

$$\left. \frac{dS}{d\epsilon} \right|_{\epsilon=0} = 0$$

Taking the derivative:

$$\int_{t_1}^{t_2} \left[\frac{\partial \mathcal{L}}{\partial \mathbf{q}} \cdot \boldsymbol{\eta} + \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \dot{\boldsymbol{\eta}} \right] dt = 0$$

For the second term, integrate by parts:

$$\int_{t_1}^{t_2} \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \dot{\boldsymbol{\eta}} dt = \left[\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \boldsymbol{\eta} \right]_{t_1}^{t_2} - \int_{t_1}^{t_2} \frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) \cdot \boldsymbol{\eta} dt$$

The boundary term vanishes because $\boldsymbol{\eta}(t_1) = \boldsymbol{\eta}(t_2) = 0$. Thus:

$$\int_{t_1}^{t_2} \left[\frac{\partial \mathcal{L}}{\partial \mathbf{q}} - \frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) \right] \cdot \boldsymbol{\eta} dt = 0$$

Since this must hold for *all* arbitrary variations $\boldsymbol{\eta}$, the integrand must vanish:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = 0 \quad (11.2)$$

This is the **Euler-Lagrange Equation**, one of the most powerful equations in all of physics. For a system with external torques $\boldsymbol{\tau}$, it becomes:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = \boldsymbol{\tau} \quad (11.3)$$

11.4 Lagrangian and Energy

In Plain Language 11.2. Kinetic vs. Potential Energy

Kinetic Energy (T) is the energy of motion. For a point mass m moving at velocity v :

$$T = \frac{1}{2}mv^2 \quad \text{or} \quad T = \frac{1}{2}m\mathbf{v}^T\mathbf{v}$$

For a rigid body with moment of inertia tensor $\mathbf{I}$ rotating at angular velocity $\boldsymbol{\omega}$:

$$T_{\text{rot}} = \frac{1}{2}\boldsymbol{\omega}^T\mathbf{I}\boldsymbol{\omega}$$

For a multi-body system with generalized coordinates $\mathbf{q}$:

$$T(\mathbf{q}, \dot{\mathbf{q}}) = \frac{1}{2}\dot{\mathbf{q}}^T\mathbf{M}(\mathbf{q})\dot{\mathbf{q}}$$

where $\mathbf{M}(\mathbf{q})$ is the **mass matrix** (symmetric positive-definite).

Potential Energy (V) is stored energy. Gravitational potential near Earth's surface:

$$V_{\text{grav}} = mgh$$

Elastic potential energy in a spring with constant k and deflection x :

$$V_{\text{spring}} = \frac{1}{2}kx^2$$

For a multi-body system, $V = V(\mathbf{q})$ depends only on configuration, not velocities.

Lagrange defined the **Lagrangian** as:

$$\mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}, t) = T(\mathbf{q}, \dot{\mathbf{q}}) - V(\mathbf{q}) \quad (11.4)$$

This is the difference between kinetic and potential energy. It encodes all the dynamics needed to recover Newton's laws, yet its form is independent of the choice of coordinate system (within the family of generalized coordinates).

When we substitute $\mathcal{L}$ into the Euler-Lagrange equation:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = \boldsymbol{\tau}$$

we obtain:

$$\frac{d}{dt} \left(\frac{\partial T}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial T}{\partial \mathbf{q}} + \frac{\partial V}{\partial \mathbf{q}} = \boldsymbol{\tau}$$

The term $\mathbf{p} = \frac{\partial T}{\partial \dot{\mathbf{q}}}$ is the **generalized momentum**. For a system with kinetic energy $T = \frac{1}{2}\dot{\mathbf{q}}^T\mathbf{M}(\mathbf{q})\dot{\mathbf{q}}$:

$$\mathbf{p} = \frac{\partial T}{\partial \dot{\mathbf{q}}} = \mathbf{M}(\mathbf{q})\dot{\mathbf{q}}$$

11.5 Derivation of the Standard Form: $\mathbf{M}\ddot{\mathbf{q}} + \mathbf{C}\dot{\mathbf{q}} + \mathbf{G} = \boldsymbol{\tau}$

For a system whose kinetic energy is:

$$T(\mathbf{q}, \dot{\mathbf{q}}) = \frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M}(\mathbf{q}) \dot{\mathbf{q}}$$

and potential energy is $V(\mathbf{q})$, we compute each term of the Euler-Lagrange equation.

Step 1: Compute $\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}}$

$$\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} = \frac{\partial T}{\partial \dot{\mathbf{q}}} = \mathbf{M}(\mathbf{q}) \dot{\mathbf{q}}$$

Step 2: Compute $\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right)$

$$\frac{d}{dt} (\mathbf{M}(\mathbf{q}) \dot{\mathbf{q}}) = \frac{d\mathbf{M}}{dt} \dot{\mathbf{q}} + \mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}}$$

where $\frac{d\mathbf{M}}{dt} = \sum_k \frac{\partial \mathbf{M}}{\partial q_k} \dot{q}_k$ by the chain rule.

Step 3: Compute $\frac{\partial \mathcal{L}}{\partial \mathbf{q}}$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{q}} = \frac{\partial T}{\partial \mathbf{q}} - \frac{\partial V}{\partial \mathbf{q}}$$

The gradient of kinetic energy requires careful differentiation:

$$\frac{\partial T}{\partial q_i} = \frac{\partial}{\partial q_i} \left(\frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M}(\mathbf{q}) \dot{\mathbf{q}} \right) = \frac{1}{2} \dot{\mathbf{q}}^T \frac{\partial \mathbf{M}}{\partial q_i} \dot{\mathbf{q}}$$

Step 4: Apply the Euler-Lagrange equation

$$\frac{d\mathbf{M}}{dt} \dot{\mathbf{q}} + \mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} - \frac{1}{2} \dot{\mathbf{q}}^T \frac{\partial \mathbf{M}}{\partial \mathbf{q}} \dot{\mathbf{q}} + \frac{\partial V}{\partial \mathbf{q}} = \boldsymbol{\tau}$$

Define the **Coriolis and Centripetal matrix** $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$ implicitly:

$$\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} = \frac{d\mathbf{M}}{dt} \dot{\mathbf{q}} - \frac{1}{2} \frac{\partial}{\partial \mathbf{q}} (\dot{\mathbf{q}}^T \mathbf{M}(\mathbf{q}) \dot{\mathbf{q}})$$

Define the **gravity vector**:

$$\mathbf{G}(\mathbf{q}) = \frac{\partial V}{\partial \mathbf{q}}$$

This yields the **standard form** [38, 40]:

$$\mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \boldsymbol{\tau} \tag{11.5}$$

11.6 Properties of the System Matrices

11.6.1 The Mass Matrix $\mathbf{M}(\mathbf{q})$

The mass matrix $\mathbf{M}$ is the Hessian of kinetic energy with respect to velocities:

$$\mathbf{M}(\mathbf{q}) = \frac{\partial^2 T}{\partial \dot{\mathbf{q}}^2}$$

Theorem: $\mathbf{M}(\mathbf{q})$ is **Symmetric Positive-Definite** (SPD) for all valid configurations.

Proof of Symmetry: Since $T = \frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$ is the kinetic energy (a physical observable), it must equal its own transpose. Thus:

$$T = T^T \implies \mathbf{M} = \mathbf{M}^T$$

Proof of Positive-Definiteness: Kinetic energy is always non-negative. Moreover, $T = 0$ if and only if $\dot{\mathbf{q}} = 0$. Therefore:

$$\dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}} = 2T \geq 0, \quad \text{with equality iff } \dot{\mathbf{q}} = 0$$

This is the definition of positive-definiteness. Consequently, $\mathbf{M}$ is invertible, and $\mathbf{M}^{-1}$ exists and is also SPD.

11.6.2 The Coriolis/Centripetal Matrix $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$

The Coriolis-Centripetal matrix is defined implicitly via the time derivative of momentum:

$$\frac{d}{dt}(\mathbf{M} \dot{\mathbf{q}}) = \mathbf{M} \ddot{\mathbf{q}} + \mathbf{C} \dot{\mathbf{q}}$$

where the i -th row of $\mathbf{C}$ is:

$$\mathbf{C}_{i*} \dot{\mathbf{q}} = \sum_{j,k} \frac{1}{2} \left(\frac{\partial M_{ij}}{\partial q_k} + \frac{\partial M_{ik}}{\partial q_j} - \frac{\partial M_{jk}}{\partial q_i} \right) \dot{q}_j \dot{q}_k$$

Equivalently, the (i, j) -th entry is:

$$C_{ij} = \frac{1}{2} \left(\frac{\partial M_{ij}}{\partial q_k} + \frac{\partial M_{ik}}{\partial q_j} - \frac{\partial M_{kj}}{\partial q_i} \right)$$

summed over k , weighted by velocity pairs.

Physical Interpretation:

- **Centrifugal Terms:** Proportional to $\dot{q}_i^2$. These represent the outward forces felt by moving links as joints spin.
- **Coriolis Terms:** Proportional to $\dot{q}_i \dot{q}_j$ with $i \neq j$. These represent cross-coupling: if one joint rotates while another accelerates, the Coriolis force deflects the moving joint.

Theorem: If $\mathbf{C}$ is constructed from the Christoffel symbols of the second kind (Section 11.12), then the matrix $\dot{\mathbf{M}} - 2\mathbf{C}$ is **skew-symmetric**, i.e., $(\dot{\mathbf{M}} - 2\mathbf{C})^T = -(\dot{\mathbf{M}} - 2\mathbf{C})$.

The hypothesis is not decorative. The factorisation of the Coriolis and centrifugal terms into a matrix $\mathbf{C}$ is not unique: any $\mathbf{C}'$ satisfying $\mathbf{C}' \dot{\mathbf{q}} = \mathbf{C} \dot{\mathbf{q}}$ gives the same equations

of motion. Only the Christoffel construction makes $\dot{\mathbf{M}} - 2\mathbf{C}$ skew-symmetric as a *matrix identity*; for other valid choices the property generally fails. Passivity- and adaptive-control designs that invoke skew-symmetry therefore require the Christoffel form specifically.

Proof: The kinetic energy is $T = \frac{1}{2}\dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$. Differentiating with respect to time:

$$\frac{d}{dt}(2T) = \dot{\mathbf{q}}^T \dot{\mathbf{M}} \dot{\mathbf{q}} + 2\dot{\mathbf{q}}^T \mathbf{M} \ddot{\mathbf{q}}.$$

From the equations of motion $\mathbf{M} \ddot{\mathbf{q}} + \mathbf{C} \dot{\mathbf{q}} + \mathbf{G} = \boldsymbol{\tau}$, the power balance $\dot{T} = \dot{\mathbf{q}}^T (\boldsymbol{\tau} - \mathbf{G})$ gives:

$$\dot{\mathbf{q}}^T \mathbf{M} \ddot{\mathbf{q}} + \dot{\mathbf{q}}^T \mathbf{C} \dot{\mathbf{q}} = \dot{\mathbf{q}}^T (\boldsymbol{\tau} - \mathbf{G}).$$

Substituting the first expression for $\dot{\mathbf{q}}^T \mathbf{M} \ddot{\mathbf{q}}$ and simplifying yields:

$$\dot{\mathbf{q}}^T (\dot{\mathbf{M}} - 2\mathbf{C}) \dot{\mathbf{q}} = 0$$

for all trajectories $\dot{\mathbf{q}}$. This step alone is not enough: $\mathbf{C}$ is itself a function of $\dot{\mathbf{q}}$, so the vanishing of the quadratic form along trajectories does not by itself let us vary $\dot{\mathbf{q}}$ freely and conclude skew-symmetry. It is the Christoffel construction that supplies the missing ingredient, giving $M_{kj} - 2C_{kj} = \sum_i (\partial M_{ij} / \partial q_k - \partial M_{kj} / \partial q_i) \dot{q}_i$, which is manifestly antisymmetric in $k \leftrightarrow j$. With that in hand, the matrix $\dot{\mathbf{M}} - 2\mathbf{C}$ must be skew-symmetric:

$$(\dot{\mathbf{M}} - 2\mathbf{C}) + (\dot{\mathbf{M}} - 2\mathbf{C})^T = \mathbf{0}.$$

This is the celebrated **skew-symmetry property** of Lagrangian dynamics, which plays a central role in passivity-based and adaptive control design.

11.6.3 The Gravity Vector $\mathbf{G}(\mathbf{q})$

The gravity vector is simply the negative gradient of potential energy:

$$\mathbf{G}(\mathbf{q}) = \frac{\partial V}{\partial \mathbf{q}}$$

For a rigid body system, each link i contributes to potential energy based on its center-of-mass height and mass:

$$V(\mathbf{q}) = \sum_i m_i g h_i(\mathbf{q})$$

The gravity vector element G_i represents the torque (or force) that gravity exerts on joint i due to the weight distribution of all distal links.

11.7 Energy Conservation

Theorem: For a system with no external torques ($\boldsymbol{\tau} = 0$) and no time-dependent potential:

$$\frac{d}{dt}(T + V) = 0$$

that is, **total mechanical energy is conserved**.

Proof: From the Euler-Lagrange equations with $\tau = 0$:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = 0$$

where $\mathcal{L} = T - V$. Rearranging:

$$\frac{d}{dt} \left(\frac{\partial T}{\partial \dot{\mathbf{q}}} \right) = \frac{\partial T}{\partial \mathbf{q}} - \frac{\partial V}{\partial \mathbf{q}}$$

Taking the inner product with $\dot{\mathbf{q}}$:

$$\dot{\mathbf{q}}^T \frac{d}{dt} \left(\frac{\partial T}{\partial \dot{\mathbf{q}}} \right) = \dot{\mathbf{q}}^T \frac{\partial T}{\partial \mathbf{q}} - \dot{\mathbf{q}}^T \frac{\partial V}{\partial \mathbf{q}}$$

The left side is:

$$\dot{\mathbf{q}}^T \frac{d}{dt} \left(\frac{\partial T}{\partial \dot{\mathbf{q}}} \right) = \frac{d}{dt} \left(\dot{\mathbf{q}}^T \frac{\partial T}{\partial \dot{\mathbf{q}}} \right) - \ddot{\mathbf{q}}^T \frac{\partial T}{\partial \dot{\mathbf{q}}}$$

For kinetic energy $T = \frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$:

$$\dot{\mathbf{q}}^T \frac{\partial T}{\partial \dot{\mathbf{q}}} = \dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}} = 2T$$

So the left side becomes:

$$\frac{d}{dt}(2T) - \ddot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}} = 2 \frac{dT}{dt} - \ddot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$$

The right side uses $\dot{\mathbf{q}}^T \frac{\partial T}{\partial \mathbf{q}} = \frac{dT}{dt} - \ddot{\mathbf{q}}^T \frac{\partial T}{\partial \dot{\mathbf{q}}}$ and $\dot{\mathbf{q}}^T \frac{\partial V}{\partial \mathbf{q}} = \frac{dV}{dt}$:

$$2 \frac{dT}{dt} - \ddot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}} = \frac{dT}{dt} - \ddot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}} - \frac{dV}{dt}$$

Simplifying:

$$\frac{dT}{dt} = - \frac{dV}{dt}$$

$$\frac{d}{dt}(T + V) = 0$$

Thus total energy is conserved. This is the mechanical manifestation of Noether's theorem (Section 11.9) [29].

11.8 Hamiltonian Mechanics

The energy conservation result above reveals something striking: the total energy $H = T + V$ is not merely a convenient bookkeeping device but a *generating function* for the dynamics themselves. Once we recognize that a single scalar quantity—the total energy—is preserved along every physical trajectory, a natural question arises: can we reformulate the entire dynamics using energy as the fundamental object, rather than the Lagrangian? The answer is yes, and the resulting framework—Hamiltonian mechanics—replaces second-order equations in $(q, \dot{q})$ with first-order equations in (q, p) , where p is the conjugate momentum.

This reformulation is not merely aesthetic. The Hamiltonian viewpoint exposes the *symplectic geometry* of phase space, provides canonical transformations that simplify otherwise intractable problems, and underpins modern developments in geometric integration, optimal control, and quantum mechanics [14, 1].

While Lagrangian mechanics uses $(q, \dot{q})$ as the fundamental variables, there is a mathematically equivalent formulation using (q, p) , where p is the conjugate momentum.

Define:

$$\mathbf{p} = \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} = \mathbf{M}(\mathbf{q})\dot{\mathbf{q}}$$

The **Hamiltonian** is obtained via a Legendre transform:

$$H(\mathbf{q}, \mathbf{p}, t) = \mathbf{p} \cdot \dot{\mathbf{q}} - \mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}, t)$$

Substituting $\mathcal{L} = T - V$ and noting $T = \frac{1}{2}\mathbf{p}^T \mathbf{M}^{-1} \mathbf{p}$ (after inverting the momentum relation):

$$H = \mathbf{p} \cdot \dot{\mathbf{q}} - \frac{1}{2}\mathbf{p}^T \mathbf{M}^{-1} \mathbf{p} + V = \frac{1}{2}\mathbf{p}^T \mathbf{M}^{-1} \mathbf{p} + V$$

The Hamiltonian is the **total energy**: $H = T + V$.

Hamilton's Equations follow from the Legendre transform:

$$\begin{aligned} \dot{\mathbf{q}} &= \frac{\partial H}{\partial \mathbf{p}} \\ \dot{\mathbf{p}} &= -\frac{\partial H}{\partial \mathbf{q}} + \boldsymbol{\tau} \end{aligned} \tag{11.6}$$

These are a set of $2n$ first-order ODEs (compared to n second-order Lagrangian equations). They preserve the symplectic structure of phase space, a property crucial for numerical integration and canonical transformations.

11.9 Noether's Theorem: Symmetry and Conservation Laws

Emmy Noether's Theorem (1918) establishes a deep connection between symmetries of the Lagrangian and conservation laws in the equations of motion [29].

Theorem (Noether): If the Lagrangian $\mathcal{L}(\mathbf{q}, \dot{\mathbf{q}})$ is invariant under a continuous one-parameter transformation group $\mathbf{q} \rightarrow \mathbf{q}(\epsilon)$ where $\mathbf{q}(0) = \mathbf{q}$ and $\left. \frac{d\mathbf{q}}{d\epsilon} \right|_{\epsilon=0} = \boldsymbol{\xi}(\mathbf{q})$ is the infinitesimal generator, then the quantity:

$$I = \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \boldsymbol{\xi}(\mathbf{q})$$

is **conserved** along trajectories satisfying the Euler-Lagrange equations.

Examples:

1. Translational Symmetry (Momentum Conservation)

If the Lagrangian does not depend on a coordinate q_i explicitly (called a **cyclic coordinate**), then $\frac{\partial \mathcal{L}}{\partial q_i} = 0$. The Euler-Lagrange equation gives:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{q}_i} \right) = 0$$

The **generalized momentum** $p_i = \frac{\partial \mathcal{L}}{\partial \dot{q}_i}$ is conserved.

For a particle moving in empty space with $\mathcal{L} = \frac{1}{2}m\mathbf{v}^2$, linear momentum $\mathbf{p} = m\mathbf{v}$ is conserved.

2. Rotational Symmetry (Angular Momentum Conservation)

If the Lagrangian is invariant under rotation, angular momentum is conserved. For a system with rotational symmetry about axis $\mathbf{n}$, the generator is $\boldsymbol{\xi} = \mathbf{n} \times \mathbf{r}$ (where $\mathbf{r}$ is position). The conserved quantity is:

$$L = \mathbf{p} \cdot (\mathbf{n} \times \mathbf{r}) = \mathbf{n} \cdot (\mathbf{r} \times \mathbf{p}) = \mathbf{n} \cdot \mathbf{L}_{\text{ang}}$$

3. Time Translation Symmetry (Energy Conservation)

If $\frac{\partial \mathcal{L}}{\partial t} = 0$ (the Lagrangian is time-independent), the **Hamiltonian**:

$$H = \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \dot{\mathbf{q}} - \mathcal{L}$$

is conserved. If $\mathcal{L} = T - V$ with T quadratic in velocities, then $H = T + V$ is the total mechanical energy.

11.9.1 Rigorous Proof of Noether's Theorem

We now present a complete proof of Noether's theorem following [29] and [1].

Setup. Let $\Phi_\epsilon : \mathcal{Q} \rightarrow \mathcal{Q}$ be a one-parameter family of diffeomorphisms with $\Phi_0 = \text{id}$. Define the **infinitesimal generator** of the transformation:

$$\boldsymbol{\xi}(\mathbf{q}) = \left. \frac{d}{d\epsilon} \Phi_\epsilon(\mathbf{q}) \right|_{\epsilon=0} \quad (11.7)$$

Under Φ_ϵ , a trajectory $\mathbf{q}(t)$ maps to $\tilde{\mathbf{q}}(t) = \Phi_\epsilon(\mathbf{q}(t))$, and its velocity transforms as $\dot{\tilde{\mathbf{q}}} = D\Phi_\epsilon(\mathbf{q}) \dot{\mathbf{q}}$, where $D\Phi_\epsilon$ denotes the Jacobian of Φ_ϵ .

Symmetry Condition. The Lagrangian is **invariant** under Φ_ϵ if:

$$\mathcal{L}(\Phi_\epsilon(\mathbf{q}), D\Phi_\epsilon \dot{\mathbf{q}}) = \mathcal{L}(\mathbf{q}, \dot{\mathbf{q}}) \quad \text{for all } \epsilon \quad (11.8)$$

Noether Charge. Define the conserved quantity (the **Noether charge**):

$$I(\mathbf{q}, \dot{\mathbf{q}}) = \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \boldsymbol{\xi}(\mathbf{q}) \quad (11.9)$$

Proof. Differentiate the symmetry condition (11.8) with respect to ϵ at $\epsilon = 0$:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{q}} \cdot \boldsymbol{\xi}(\mathbf{q}) + \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \left. \frac{d}{d\epsilon} (D\Phi_\epsilon \dot{\mathbf{q}}) \right|_{\epsilon=0} = 0 \quad (11.10)$$

Now compute $\frac{d}{dt} I$ along a solution of the Euler-Lagrange equations:

$$\begin{aligned} \frac{d}{dt} I &= \frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) \cdot \boldsymbol{\xi}(\mathbf{q}) + \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \frac{d}{dt} \boldsymbol{\xi}(\mathbf{q}) \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{q}} \cdot \boldsymbol{\xi}(\mathbf{q}) + \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \cdot \frac{\partial \boldsymbol{\xi}}{\partial \mathbf{q}} \dot{\mathbf{q}} \end{aligned} \quad (11.11)$$

where the first term uses the Euler-Lagrange equation $\frac{d}{dt} \frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} = \frac{\partial \mathcal{L}}{\partial \mathbf{q}}$. Noting that $\left. \frac{d}{d\epsilon} (D\Phi_\epsilon \dot{\mathbf{q}}) \right|_{\epsilon=0} = \frac{\partial \boldsymbol{\xi}}{\partial \mathbf{q}} \dot{\mathbf{q}}$, equation (11.10) shows that (11.11) equals zero. Hence $\frac{d}{dt} I = 0$, and I is conserved. $\square$

In Plain Language 11.3. Noether's Theorem in Practice

Every continuous symmetry yields a conserved quantity that constrains the motion. In robotics, these conservation laws provide powerful checks on simulation accuracy and enable reduction of the equations of motion:

- **Translation symmetry** $\Rightarrow$ **linear momentum** is conserved (free-floating satellite).
- **Rotation symmetry** $\Rightarrow$ **angular momentum** is conserved (spacecraft attitude).
- **Time-translation symmetry** $\Rightarrow$ **energy** $H = T + V$ is conserved (unactuated systems).

11.10 Lagrangian for Constrained Systems

When holonomic constraints are present (e.g., rigid connections between bodies), they are **automatically built into the choice of generalized coordinates $\mathbf{q}$** .

However, for **non-holonomic constraints** (constraints on velocities that cannot be integrated to position constraints) or when analyzing the forces exerted by constraints, the **method of Lagrange multipliers** is employed.

Consider a system with constraints $\phi(\mathbf{q}) = 0$. The constraint forces are perpendicular to the constraint surface, hence do no virtual work. The equations of motion with constraint forces $\mathbf{F}_c = \nabla\phi \cdot \boldsymbol{\lambda}$ (where $\boldsymbol{\lambda}$ are multipliers) become:

$$\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = \boldsymbol{\tau} + \nabla\phi \cdot \boldsymbol{\lambda} \quad (11.12)$$

The multipliers $\boldsymbol{\lambda}$ are determined by enforcing $\ddot{\phi} = 0$ (constraint satisfaction at the acceleration level). This is the standard method in mechanical engineering for deriving equations with Lagrange multipliers.

11.11 D'Alembert's Principle and Virtual Work

An alternative derivation of Lagrangian mechanics begins with **D'Alembert's Principle [14]**, which unifies constraints and dynamics.

Consider a system in motion under forces $\mathbf{F}$ with constraint forces $\mathbf{F}_c$. Newton's second law is:

$$m\mathbf{a} = \mathbf{F} + \mathbf{F}_c$$

D'Alembert introduced **inertial forces**: rewrite this as:

$$\mathbf{F} + \mathbf{F}_c - m\mathbf{a} = 0$$

Now define **virtual displacement** $\delta\mathbf{q}$ —an infinitesimal displacement consistent with the constraints (i.e., $\nabla\phi \cdot \delta\mathbf{q} = 0$ if constraints exist). The principle states that the virtual work done by inertial and applied forces vanishes:

$$\delta W = (\mathbf{F} - m\mathbf{a}) \cdot \delta \mathbf{r} = 0$$

(Constraint forces do no virtual work by definition.)

In generalized coordinates:

$$\sum_i \left(\frac{\partial \mathcal{L}}{\partial q_i} - \frac{d}{dt} \frac{\partial \mathcal{L}}{\partial \dot{q}_i} \right) \delta q_i = 0$$

Since the δq_i are arbitrary (independent), each coefficient must vanish, recovering the Euler-Lagrange equations. This approach shows directly that Lagrangian mechanics is a direct consequence of Newton's laws when viewed through the lens of generalized coordinates and constraint-respecting displacements.

11.12 Christoffel Symbols and the Geometry of M

The mass matrix $\mathbf{M}(\mathbf{q})$ defines a Riemannian metric on configuration space [31]. The Christoffel symbols Γ_{jk}^i of this metric are:

$$\Gamma_{jk}^i = \frac{1}{2} g^{im} \left(\frac{\partial g_{mj}}{\partial q_k} + \frac{\partial g_{mk}}{\partial q_j} - \frac{\partial g_{jk}}{\partial q_m} \right)$$

where $g = \mathbf{M}$ and $g^{-1} = \mathbf{M}^{-1}$.

These symbols encode how the metric (and thus the inertial properties) change over configuration space. They appear naturally in the Coriolis matrix:

$$C_{ij} = \frac{1}{2} (\dot{M}_{ij} + \dot{M}_{ik} \dot{M}_{jk}^{-1})$$

or more explicitly in terms of Christoffel symbols:

$$C_{ij} = \Gamma_{ij}^k \dot{q}_k$$

This reveals a deep geometric truth: the Coriolis forces arise from geodesic deviation in the manifold defined by the mass matrix. A system moving "straight" in configuration space, when projected into Cartesian coordinates, experiences apparent Coriolis deflection because the natural metric of the system is curved.

11.13 Worked Example 1: Double Pendulum

11.13.1 System Setup

A double pendulum consists of two rigid rods connected end-to-end, with a pivot at the origin. Link 1 has mass m_1 , length ℓ_1 , and angle q_1 from vertical. Link 2 has mass m_2 , length ℓ_2 , and angle q_2 from vertical.

The center of mass of link 1 is at:

$$\mathbf{r}_1 = \begin{pmatrix} \frac{\ell_1}{2} \sin q_1 \\ -\frac{\ell_1}{2} \cos q_1 \end{pmatrix}$$

The center of mass of link 2 is at:

$$\mathbf{r}_2 = \begin{pmatrix} \ell_1 \sin q_1 + \frac{\ell_2}{2} \sin q_2 \\ -\ell_1 \cos q_1 - \frac{\ell_2}{2} \cos q_2 \end{pmatrix}$$

11.13.2 Kinetic Energy

For link 1, treating it as a point mass at its center:

$$v_1^2 = \dot{\mathbf{r}}_1 \cdot \dot{\mathbf{r}}_1$$

Computing:

$$\dot{\mathbf{r}}_1 = \frac{\ell_1}{2} \begin{pmatrix} \cos q_1 & 0 \\ \sin q_1 & 0 \end{pmatrix} \begin{pmatrix} \dot{q}_1 \\ \dot{q}_2 \end{pmatrix}$$

$$v_1^2 = \frac{\ell_1^2}{4} \dot{q}_1^2$$

For link 2:

$$\dot{\mathbf{r}}_2 = \begin{pmatrix} \ell_1 \cos q_1 \dot{q}_1 + \frac{\ell_2}{2} \cos q_2 \dot{q}_2 \\ \ell_1 \sin q_1 \dot{q}_1 + \frac{\ell_2}{2} \sin q_2 \dot{q}_2 \end{pmatrix}$$

$$v_2^2 = \left(\ell_1 \cos q_1 \dot{q}_1 + \frac{\ell_2}{2} \cos q_2 \dot{q}_2 \right)^2 + \left(\ell_1 \sin q_1 \dot{q}_1 + \frac{\ell_2}{2} \sin q_2 \dot{q}_2 \right)^2$$

$$v_2^2 = \ell_1^2 \dot{q}_1^2 + \frac{\ell_2^2}{4} \dot{q}_2^2 + \ell_1 \ell_2 (\cos q_1 \cos q_2 + \sin q_1 \sin q_2) \dot{q}_1 \dot{q}_2$$

$$v_2^2 = \ell_1^2 \dot{q}_1^2 + \frac{\ell_2^2}{4} \dot{q}_2^2 + \ell_1 \ell_2 \cos(q_1 - q_2) \dot{q}_1 \dot{q}_2$$

The total kinetic energy is:

$$T = \frac{1}{2} m_1 v_1^2 + \frac{1}{2} m_2 v_2^2$$

$$T = \frac{1}{2} m_1 \frac{\ell_1^2}{4} \dot{q}_1^2 + \frac{1}{2} m_2 \left[\ell_1^2 \dot{q}_1^2 + \frac{\ell_2^2}{4} \dot{q}_2^2 + \ell_1 \ell_2 \cos(q_1 - q_2) \dot{q}_1 \dot{q}_2 \right]$$

$$T = \frac{1}{2} \left[\left(\frac{m_1}{4} + m_2 \right) \ell_1^2 \dot{q}_1^2 + \frac{m_2 \ell_2^2}{4} \dot{q}_2^2 + m_2 \ell_1 \ell_2 \cos(q_1 - q_2) \dot{q}_1 \dot{q}_2 \right]$$

In matrix form: $T = \frac{1}{2} \begin{pmatrix} \dot{q}_1 & \dot{q}_2 \end{pmatrix} \mathbf{M} \begin{pmatrix} \dot{q}_1 \\ \dot{q}_2 \end{pmatrix}$

$$\mathbf{M}(\mathbf{q}) = \begin{pmatrix} \left(\frac{m_1}{4} + m_2 \right) \ell_1^2 & \frac{1}{2} m_2 \ell_1 \ell_2 \cos(q_1 - q_2) \\ \frac{1}{2} m_2 \ell_1 \ell_2 \cos(q_1 - q_2) & \frac{m_2 \ell_2^2}{4} \end{pmatrix}$$

Read the entries directly off the quadratic form: the diagonal entries are the coefficients of $\dot{q}_i^2$ inside the bracket, while the off-diagonal entry is *half* the coefficient of the cross term $\dot{q}_1 \dot{q}_2$, because that term appears twice in $\dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$.

11.13.3 Potential Energy

Taking the pivot as the reference (zero potential energy at the pivot height):

$$V = -m_1 g \frac{\ell_1}{2} \cos q_1 - m_2 g \left(\ell_1 \cos q_1 + \frac{\ell_2}{2} \cos q_2 \right)$$

$$V = - \left(\frac{m_1}{2} + m_2 \right) g \ell_1 \cos q_1 - \frac{m_2 g \ell_2}{2} \cos q_2$$

11.13.4 Gravity Vector

$$\mathbf{G} = \frac{\partial V}{\partial \mathbf{q}} = \begin{pmatrix} \left(\frac{m_1}{2} + m_2\right) g \ell_1 \sin q_1 \\ \frac{m_2 g \ell_2}{2} \sin q_2 \end{pmatrix}$$

11.13.5 Coriolis/Centripetal Matrix

Computing $\frac{d\mathbf{M}}{dt}$:

$$\frac{dM_{12}}{dt} = -m_2 \ell_1 \ell_2 \sin(q_1 - q_2)(\dot{q}_1 - \dot{q}_2)$$

Using the formula for the Coriolis matrix (after careful computation):

$$\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) = \frac{1}{2} m_2 \ell_1 \ell_2 \sin(q_1 - q_2) \begin{pmatrix} 0 & \dot{q}_2 \\ -\dot{q}_1 & 0 \end{pmatrix}$$

Consistency check. With this $\mathbf{M}$ and this Christoffel $\mathbf{C}$,

$$\dot{\mathbf{M}} - 2\mathbf{C} = \frac{1}{2} m_2 \ell_1 \ell_2 \sin(q_1 - q_2)(\dot{q}_1 + \dot{q}_2) \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix},$$

which is skew-symmetric, as the theorem above requires. Any worked example that fails this check has an algebraic error in $\mathbf{M}$ or $\mathbf{C}$ — it is the cheapest available test of a hand-derived model.

The complete equations of motion are:

$$\mathbf{M}(\mathbf{q})\ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})\dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \boldsymbol{\tau}$$

where $\boldsymbol{\tau}$ are the applied torques at each joint.

11.14 Worked Example 2: 2-Link Robot Arm With Gravity

11.14.1 System Configuration

A planar 2-link robot arm with joint angles q_1, q_2 (measured from horizontal). Each link has length ℓ_i and mass m_i uniformly distributed. The arm operates under gravity.

Pose of the end-effector:

$$\mathbf{p} = \begin{pmatrix} \ell_1 \cos q_1 + \ell_2 \cos(q_1 + q_2) \\ \ell_1 \sin q_1 + \ell_2 \sin(q_1 + q_2) \end{pmatrix}$$

Jacobian (mapping joint velocities to end-effector velocities):

$$\mathbf{J} = \begin{pmatrix} -\ell_1 \sin q_1 - \ell_2 \sin(q_1 + q_2) & -\ell_2 \sin(q_1 + q_2) \\ \ell_1 \cos q_1 + \ell_2 \cos(q_1 + q_2) & \ell_2 \cos(q_1 + q_2) \end{pmatrix}$$

11.14.2 Kinetic Energy

For a uniform link i rotating about its proximal joint, the kinetic energy is:

$$T_i = \frac{1}{2}I_i\omega_i^2 + \frac{1}{2}m_iv_{c,i}^2$$

where $I_i = \frac{1}{12}m_i\ell_i^2$ (moment of inertia about the center) and $v_{c,i}$ is the velocity of the center of mass.

For link 1:

$$v_{c,1} = \frac{\ell_1}{2}\dot{q}_1 \implies T_1 = \frac{1}{2}\left(\frac{1}{12}m_1\ell_1^2 + \frac{1}{4}m_1\ell_1^2\right)\dot{q}_1^2 = \frac{1}{6}m_1\ell_1^2\dot{q}_1^2$$

For link 2, it rotates with angular velocity $\dot{q}_1 + \dot{q}_2$ about the global frame:

$$T_2 = \frac{1}{2}I_2(\dot{q}_1 + \dot{q}_2)^2 + \frac{1}{2}m_2v_{c,2}^2$$

where the center-of-mass velocity includes both rotation about joint 1 and rotation about joint 2:

$$v_{c,2}^2 = (\ell_1\dot{q}_1)^2 + \left(\frac{\ell_2}{2}(\dot{q}_1 + \dot{q}_2)\right)^2 + 2\ell_1\frac{\ell_2}{2}\dot{q}_1(\dot{q}_1 + \dot{q}_2)\cos q_2$$

After simplification, the mass matrix is:

$$\mathbf{M} = \begin{pmatrix} m_1\frac{\ell_1^2}{3} + m_2\ell_1^2 + m_2\frac{\ell_2^2}{3} + m_2\ell_1\ell_2\cos q_2 & m_2\frac{\ell_2^2}{6} + m_2\frac{\ell_1\ell_2}{2}\cos q_2 \\ m_2\frac{\ell_2^2}{6} + m_2\frac{\ell_1\ell_2}{2}\cos q_2 & m_2\frac{\ell_2^2}{3} \end{pmatrix}$$

11.14.3 Potential Energy

With gravity acting downward:

$$V = m_1g\frac{\ell_1}{2}\sin q_1 + m_2g\left(\ell_1\sin q_1 + \frac{\ell_2}{2}\sin(q_1 + q_2)\right)$$

Gravity vector:

$$\mathbf{G} = \begin{pmatrix} \left(\frac{m_1}{2} + m_2\right)g\ell_1\cos q_1 + \frac{m_2g\ell_2}{2}\cos(q_1 + q_2) \\ \frac{m_2g\ell_2}{2}\cos(q_1 + q_2) \end{pmatrix}$$

These explicit forms are essential for control design (feedback linearization, trajectory optimization) and analysis (stability, reachability).

11.15 Comparison: Lagrangian vs. Newton-Euler

Both formulations yield identical equations of motion, but traverse different computational paths:

Aspect	Lagrangian	Newton-Euler
Input	Kinetic and potential energy	Link masses, geometry, gravity
Derivation	Calculus of variations	Free-body diagram + recursion
Analytical Form	$\mathbf{M}, \mathbf{C}, \mathbf{G}$ matrices explicit	$\mathbf{M}, \mathbf{C}$ assembled link-by-link
Computational Cost	$O(n^3)$ matrix inversion	$O(n)$ recursive forward/backward pass
Control Design	Elegant (LQR, feedback linearization)	Cumbersome (requires explicit $\mathbf{M}, \mathbf{C}$)
Implementation	By hand for small systems	Physics engines (rigid body simulators)

For a 2-DOF arm, both methods agree perfectly. For a 15-DOF golf swing, the Recursive Newton-Euler algorithm computes torques hundreds of times faster than inverting the full 15×15 mass matrix.

11.16 Symplectic Integrators

Standard ODE solvers such as forward Euler and classical Runge-Kutta (RK4) do not respect the geometric structure of Hamiltonian systems. Over long simulations, they introduce systematic energy drift: trajectories slowly spiral inward or outward in phase space, producing non-physical behavior [1].

Symplectic integrators are numerical methods that exactly preserve the **symplectic form**:

$$\omega = d\mathbf{p} \wedge d\mathbf{q} \quad (11.13)$$

This means that the discrete-time map $(\mathbf{q}_k, \mathbf{p}_k) \mapsto (\mathbf{q}_{k+1}, \mathbf{p}_{k+1})$ is a canonical transformation, preserving phase-space volume and the qualitative structure of orbits.

Symplectic Euler Method. The simplest symplectic integrator updates momenta first, then positions using the new momenta:

$$\mathbf{p}_{k+1} = \mathbf{p}_k - h \frac{\partial H}{\partial \mathbf{q}}(\mathbf{q}_k, \mathbf{p}_k) \quad (11.14)$$

$$\mathbf{q}_{k+1} = \mathbf{q}_k + h \frac{\partial H}{\partial \mathbf{p}}(\mathbf{q}_k, \mathbf{p}_{k+1}) \quad (11.15)$$

This is first-order accurate but already exhibits bounded energy error.

Störmer-Verlet (Leapfrog) Method. A second-order, time-reversible, symplectic scheme:

$$\mathbf{p}_{k+1/2} = \mathbf{p}_k - \frac{h}{2} \frac{\partial H}{\partial \mathbf{q}}(\mathbf{q}_k) \quad (11.16)$$

$$\mathbf{q}_{k+1} = \mathbf{q}_k + h \frac{\partial H}{\partial \mathbf{p}}(\mathbf{p}_{k+1/2}) \quad (11.17)$$

$$\mathbf{p}_{k+1} = \mathbf{p}_{k+1/2} - \frac{h}{2} \frac{\partial H}{\partial \mathbf{q}}(\mathbf{q}_{k+1}) \quad (11.18)$$

Key Property: Bounded Energy Error. Unlike non-symplectic methods where energy error grows linearly (or worse) with time, symplectic integrators produce energy errors that *oscillate* about the true value without secular drift. Formally, the numerical solution exactly solves a *nearby* Hamiltonian $\tilde{H} = H + O(h^p)$, where p is the order of the method. The energy of the original system therefore remains bounded:

$$|H(\mathbf{q}_k, \mathbf{p}_k) - H(\mathbf{q}_0, \mathbf{p}_0)| \leq C h^p \quad \text{for exponentially long times}$$

In Plain Language 11.4. Why Symplectic Integrators Matter for Robotics

For long-horizon robot simulations—such as gait optimization over thousands of steps, or orbital mechanics for space robots—symplectic integrators prevent the slow energy drift that causes non-physical behavior. A standard RK4 solver may show a pendulum gradually gaining energy until it rotates faster and faster; a symplectic solver keeps the energy oscillating near its true value indefinitely.

11.17 Python Worked Example: Double Pendulum Dynamics

```

1 import numpy as np
2 from scipy.integrate import odeint
3 import matplotlib.pyplot as plt
4
5 class DoublePendulum:
6     def __init__(self, m1=1.0, m2=1.0, l1=1.0, l2=1.0, g=9.81):
7         """Initialize double pendulum parameters."""
8         self.m1, self.m2 = m1, m2
9         self.l1, self.l2 = l1, l2
10        self.g = g
11
12    def mass_matrix(self, q):
13        """Compute M(q) from kinetic energy  $T = (1/2) \dot{q}^T M \dot{q}$ """
14        q1, q2 = q[0], q[1]
15        M11 = (self.m1/2 + self.m2) * self.l1**2
16        M12 = self.m2 * self.l1 * self.l2 * np.cos(q1 - q2)
17        M22 = 0.5 * self.m2 * self.l2**2
18        M = np.array([[M11, M12], [M12, M22]])
19        return M
20
21    def coriolis_matrix(self, q, qdot):
22        """Compute C(q, qdot) such that Coriolis forces = C(q, qdot) qdot"""
23        q1, q2 = q[0], q[1]
24        dq1, dq2 = qdot[0], qdot[1]
25
26        # Christoffel symbol terms
27        dM12_dq = -self.m2 * self.l1 * self.l2 * np.sin(q1 - q2)
28
29        C11 = -dM12_dq * dq2 / 2
30        C12 = -dM12_dq * (dq1 + dq2) / 2
31        C21 = dM12_dq * dq1 / 2
32        C22 = 0.0
33
34        C = np.array([[C11, C12], [C21, C22]])
35        return C
36
37    def gravity_vector(self, q):
38        """Compute G(q) = dV/dq"""
39        q1, q2 = q[0], q[1]
40        G1 = (self.m1/2 + self.m2) * self.g * self.l1 * np.sin(q1)
41        G2 = 0.5 * self.m2 * self.g * self.l2 * np.sin(q2)
42        return np.array([G1, G2])
43

```

```

44     def dynamics(self, state, t, tau):
45         """Compute acceleration from  $M(q) \ddot{q} + C(q, \dot{q}) \dot{q} + G(q) = \tau$ 
46         """
47         q = state[:2]
48         qdot = state[2:]
49
50         M = self.mass_matrix(q)
51         C = self.coriolis_matrix(q, qdot)
52         G = self.gravity_vector(q)
53
54         # Solve:  $M \ddot{q} = \tau - C \dot{q} - G$ 
55         qddot = np.linalg.solve(M, tau - C @ qdot - G)
56
57         return np.concatenate([qdot, qddot])
58
59     def simulate(self, q0, qdot0, t_span, tau_func=None):
60         """Simulate the system over time."""
61         if tau_func is None:
62             tau_func = lambda t: np.array([0.0, 0.0])
63
64         state0 = np.concatenate([q0, qdot0])
65         t = np.linspace(t_span[0], t_span[1], 500)
66
67         # Wrapper for odeint
68         def dynamics_wrapper(state, t):
69             tau = tau_func(t)
70             return self.dynamics(state, t, tau)
71
72         solution = odeint(dynamics_wrapper, state0, t)
73         return t, solution
74
75     # Example: Release double pendulum from rest at angle
76     pendulum = DoublePendulum(m1=1.0, m2=1.0, l1=1.0, l2=1.0, g=9.81)
77
78     # Initial conditions:  $q_1=\pi/4$ ,  $q_2=0$ , all velocities zero
79     q0 = np.array([np.pi/4, 0.0])
80     qdot0 = np.array([0.0, 0.0])
81
82     # Simulate for 10 seconds
83     t, traj = pendulum.simulate(q0, qdot0, (0, 10))
84
85     # Extract angles and velocities
86     q1_traj = traj[:, 0]
87     q2_traj = traj[:, 1]
88     dq1_traj = traj[:, 2]
89     dq2_traj = traj[:, 3]
90
91     # Verify energy conservation
92     E = []
93     for i in range(len(t)):
94         q = np.array([q1_traj[i], q2_traj[i]])
95         qdot = np.array([dq1_traj[i], dq2_traj[i]])
96         M = pendulum.mass_matrix(q)
97         T = 0.5 * qdot @ M @ qdot
98         V = -pendulum.mass_matrix(q)[0,0] * pendulum.g * np.cos(q1_traj[i])
99         V -= 0.5 * pendulum.m2 * pendulum.g * pendulum.l2 * np.cos(q2_traj[i])
100         E.append(T + V)
101
102     print(f"Initial Energy: {E[0]:.6f}")
103     print(f"Final Energy: {E[-1]:.6f}")

```

```

103 print(f"Energy Error: {abs(E[-1] - E[0]):.2e}")
104
105 # Plot trajectory
106 fig, axes = plt.subplots(2, 2, figsize=(10, 8))
107 axes[0, 0].plot(t, q1_traj)
108 axes[0, 0].set_ylabel('q1 (rad)')
109 axes[0, 1].plot(t, q2_traj)
110 axes[0, 1].set_ylabel('q2 (rad)')
111 axes[1, 0].plot(t, dq1_traj)
112 axes[1, 0].set_ylabel('dq1/dt (rad/s)')
113 axes[1, 1].plot(t, dq2_traj)
114 axes[1, 1].set_ylabel('dq2/dt (rad/s)')
115 for ax in axes.flat:
116     ax.set_xlabel('Time (s)')
117     ax.grid(True)
118 plt.tight_layout()
119 plt.show()

```

Listing 11.1: Double Pendulum Lagrangian Derivation and Simulation

This code demonstrates:

- Explicit computation of $\mathbf{M}(\mathbf{q})$ from kinetic energy
- Christoffel symbol terms in $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})$
- Gravity vector $\mathbf{G}(\mathbf{q})$
- Integration of the second-order ODE via numerical solving
- Verification of energy conservation

11.18 Chapter Summary

Lagrangian mechanics reformulates classical dynamics through energy rather than forces. Key insights:

1. **Hamilton's Principle:** Physical systems extremize the action integral $S = \int \mathcal{L} dt$.
2. **Euler-Lagrange Equations:** $\frac{d}{dt} \left(\frac{\partial \mathcal{L}}{\partial \dot{\mathbf{q}}} \right) - \frac{\partial \mathcal{L}}{\partial \mathbf{q}} = \boldsymbol{\tau}$ automatically yields equations of motion.
3. **Standard Form:** $\mathbf{M}(\mathbf{q})\ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}})\dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \boldsymbol{\tau}$ cleanly separates inertia, Coriolis, and gravity effects.
4. **Constraint Elimination:** Generalized coordinates automatically incorporate holonomic constraints, eliminating unknown constraint forces.
5. **Energy Conservation:** Time-translation symmetry (via Noether's theorem) guarantees $\frac{d(T+V)}{dt} = 0$ when $\boldsymbol{\tau} = 0$.
6. **Hamiltonian Formulation:** A Legendre transform yields symplectic form $\dot{\mathbf{q}} = \frac{\partial H}{\partial \mathbf{p}}$, $\dot{\mathbf{p}} = -\frac{\partial H}{\partial \mathbf{q}}$.

7. **Geometric Insight:** The mass matrix $\mathbf{M}$ defines a Riemannian metric on configuration space. Coriolis forces are geodesic deviations.

In control theory (Volume I), Lagrangian mechanics enables:

- **Feedback Linearization:** Cancel nonlinear terms analytically.
- **LQR and Optimal Control:** Write cost functions on energy and motion directly.
- **Contraction Metrics:** Prove stability using the metric structure.
- **Port-Hamiltonian Systems:** Preserve structure under interconnection.

11.19 Further Reading and References

The foundations of Lagrangian and Hamiltonian mechanics are treated with exceptional depth in Goldstein, Poole, and Safko [14], which remains the standard graduate reference for classical mechanics. Their treatment of variational principles, canonical transformations, and Hamilton-Jacobi theory provides the analytical backbone for everything presented in this chapter. For a more mathematically rigorous and geometrically oriented perspective, Arnold [1] develops classical mechanics from the standpoint of differential geometry and symplectic topology, making explicit the manifold structure that underlies configuration space and phase space.

The geometric mechanics perspective—viewing the mass matrix as a Riemannian metric and Coriolis forces as geodesic deviations—is developed systematically in Marsden and Ratiu [29]. Their treatment of Noether’s theorem, momentum maps, and reduction theory provides the mathematical language for understanding symmetry and conservation in mechanical systems at a level far beyond what we have space for here. Readers interested in the Lie group structure of rigid body mechanics and its interplay with Lagrangian dynamics should consult Murray, Li, and Sastry [31], who unify the geometric and control-theoretic viewpoints.

For computational dynamics and the connection between Lagrangian formulations and recursive algorithms (bridging this chapter with Chapter 10), Featherstone [12] provides an authoritative treatment of how the $\mathbf{M}$, $\mathbf{C}$, and $\mathbf{G}$ matrices can be assembled efficiently using spatial algebra. The control-theoretic applications of the standard form $\mathbf{M}\ddot{\mathbf{q}} + \mathbf{C}\dot{\mathbf{q}} + \mathbf{G} = \boldsymbol{\tau}$ —including feedback linearization, computed torque control, and impedance control—are covered extensively by Siciliano et al. [38] and by Spong, Hutchinson, and Vidyasagar [40]. Both texts demonstrate how the properties proved in this chapter (positive-definiteness of $\mathbf{M}$, skew-symmetry of $\dot{\mathbf{M}} - 2\mathbf{C}$) are directly exploited in Lyapunov-based stability proofs for robot controllers.

11.20 Exercises

Exercise. Variational Calculus: Consider the action $\mathcal{S}[q(t)] = \int_0^T (\dot{q}^2 + q^2)dt$ with fixed endpoints $q(0) = 0$, $q(T) = 1$. Show that the extremal path satisfies $\ddot{q} - 2q = 0$ (i.e., apply Euler-Lagrange).

Exercise. Simple Pendulum Lagrangian: For a simple pendulum of mass m and length ℓ , write $T(\theta, \dot{\theta})$ and $V(\theta)$. Apply Euler-Lagrange to recover $\ddot{\theta} + \frac{g}{\ell} \sin \theta = 0$.

Exercise. Cyclic Coordinates: Consider a particle on a rotating table (Coriolis force). Show that angular momentum about the vertical axis is not conserved because the Lagrangian depends on the azimuthal angle explicitly (in the rotating frame).

Exercise. Generalized Momenta: For a 3-link arm in a plane, identify which generalized momenta are conserved under no external torques (hint: check which coordinates are cyclic).

Exercise. Mass Matrix Positive-Definiteness: Prove rigorously that if $T = \frac{1}{2} \dot{\mathbf{q}}^T \mathbf{M} \dot{\mathbf{q}}$ is the kinetic energy of a physical system, then $\mathbf{M}$ is positive-definite.

Exercise. Skew-Symmetry: Verify directly for a 2-DOF system that $\mathbf{M} - 2\mathbf{C}$ is skew-symmetric. Compute the individual entries and check $(\mathbf{M} - 2\mathbf{C}) + (\mathbf{M} - 2\mathbf{C})^T = 0$.

Exercise. Double Pendulum Energy: For the double pendulum (Section 9.2), compute the total mechanical energy as a function of $(q_1, q_2, \dot{q}_1, \dot{q}_2)$. Verify that $dE/dt = 0$ when the system is released from rest.

Exercise. Hamiltonian Transform: Given $\mathcal{L} = \frac{1}{2} m \dot{x}^2 - kx^2$, compute the conjugate momentum $p = \frac{\partial \mathcal{L}}{\partial \dot{x}}$, then form $H = p\dot{x} - \mathcal{L}$. Show that $H = \frac{p^2}{2m} + kx^2$ is the total energy.

Exercise. D'Alembert's Principle: A bead slides on a frictionless parabola $y = ax^2$. Use virtual work to derive the equation of motion in the x direction. (Hint: virtual displacements must satisfy $\delta y = 2ax\delta x$.)

Exercise. Lagrange Multipliers: A particle is constrained to a sphere $x^2 + y^2 + z^2 = R^2$. Write the constraint force as $\mathbf{F}_c = \lambda \nabla(x^2 + y^2 + z^2)$ and derive the equations of motion including λ .

Exercise. 2-Link Robot Arm: For the 2-link robot arm (Section 9.3), compute $\mathbf{M}(q_1, q_2)$ symbolically and show all entries. Then compute the gravity vector $\mathbf{G}(q_1, q_2)$.

Exercise. Friction and Dissipation: Extend the double pendulum Lagrangian to include viscous damping torques $\tau = -b_1 \dot{q}_1 - b_2 \dot{q}_2$. Write the modified equations and verify that mechanical energy dissipates (i.e., $dE/dt < 0$).

Exercise. Christoffel Symbols: For a 2-DOF system with mass matrix $\mathbf{M}(q_1, q_2)$, compute the Christoffel symbols Γ_{jk}^i in terms of partial derivatives of M_{ij} .

Exercise. Noether's Theorem Application: Consider a free-falling object (no constraints, only gravity). Identify the symmetries of the Lagrangian and use Noether's theorem to derive conservation of momentum and energy separately.

Exercise. Pendulum Swing-Up Energy: For a simple pendulum, compute the minimum total energy required to swing from the downward position ($\theta = 0$) to the upright position ($\theta = \pi$). Express in terms of m, ℓ, g .

Exercise. Code Implementation: Implement the code for a 2-link robot arm (similar to the double pendulum code), compute $\mathbf{M}, \mathbf{C}, \mathbf{G}$ explicitly, and simulate a trajectory with constant torque applied to joint 1. Verify energy conservation when no external torques are applied.

Chapter 12

Machine Learning and Neural Networks

12.1 Scope of This Chapter

Learning is useful when measured examples can improve a model, an estimator or a decision rule. It does not remove the need to define the physical system and the question being asked. A network that predicts a golfer's motion, a network that proposes muscle excitations, and a network that estimates joint torques solve different problems. Success at one is not evidence of success at the others.

This chapter develops neural function approximation, manual backpropagation, reinforcement learning and learned dynamical models in one consistent scope. It connects policy optimization to the geometric mechanics of the preceding chapters and provides a reproducible regression example. The pendulum and golf discussions specify models and experiments; they do not report a trained swing controller or new measurements of human performance.

Use q for configuration, v for generalized velocity, x for the complete modeled state, o for observations and u for the chosen physical input. In the learning equations, θ denotes network parameters, never the pendulum angle. A loss ℓ is an optimization criterion; a Lagrangian L is a mechanical function. Their similar notation in the literature does not make them the same quantity.

12.2 The Limits of Analytical Control

Rigid-body algorithms are exact evaluations of declared mathematical models up to numerical error. Real predictions also depend on identified inertias, joint definitions, contact laws, actuator dynamics, measurement errors and the appropriateness of the rigid approximation. Even a manufactured robot is not described perfectly merely because its equations are analytical. Conversely, deformable bodies can have analytical constitutive equations and numerical models. Complexity and nonlinearity alone do not force a choice between physics and learning.

For a selected smooth mechanical mode, one useful model is

$$\begin{aligned}\dot{q} &= N(q)v, \\ M(q)\dot{v} + c(q, v) + g(q) &= B(q)u + J(q)^T\lambda + f_{\text{other}}.\end{aligned}\tag{12.1}$$

Here M is inertia, c contains velocity-dependent inertial forces, g is gravity, B maps inputs to generalized forces, and $J^T\lambda$ represents constraint reactions. N maps generalized velocities to configuration rates; it need not be an identity for quaternion coordinates. Contact

feasibility and actuator or tissue states complete the model. Adding a learned residual to these equations can be more informative than replacing every term with an unconstrained network, provided the residual is identifiable and does not duplicate forces already included.

12.2.1 Prediction, Inversion and Policy Are Different Maps

A forward model predicts $\dot{x} = F_\theta(x, u)$ or a sampled transition $x_{k+1} = T_{\theta,h}(x_k, u_k)$ for a declared sampling interval h and input hold. These maps have different units and different meanings. Fitting x_{k+1} directly does not make the output an acceleration. A learned transition may also depend on the integration scheme and sampling interval used to produce its data.

An inverse model might estimate u from $(q, v, \dot{v}_{\text{desired}})$ and the contact mode. For full actuation and known external loads, a torque-level inverse can be well-defined. Under-actuation, redundant muscles and unknown reactions can make an inverse nonexistent or nonunique. A least-squares torque estimate then depends on an explicitly chosen allocation or regularization criterion. A small residual is not proof that its individual muscle forces were measured correctly.

A policy instead selects an action given information and a task. Let i_k contain the observation history, target and current step. A deterministic policy gives $u_k = \mu_\theta(i_k)$; a stochastic policy specifies a conditional distribution $\pi_\theta(u_k | i_k)$. Identical posture can require different actions when velocity, activation, target, remaining time or disturbances differ. There is no universal hidden function from posture alone to the uniquely correct motor torque.

12.2.2 Why State Choice Matters for Golf

Position and velocity can be a sufficient state for a specified rigid torque-driven model. A muscle-driven swing may also require activation, fiber or tendon variables and shaft deformation. Two swings with the same visible geometry can therefore have different subsequent responses. Observations of joint angles, club motion and ground forces provide incomplete and noisy information about these states.

A partially observed Markov decision process retains a Markov latent state; the observation alone need not be Markov. An estimator, belief distribution or history-dependent network can use past observations to infer missing information. Recurrence does not guarantee recovery of an unobservable state. Sensor placement, synchronization and experimental excitation remain essential.

This distinction links preparation to impact. An earlier action can change configuration, velocity, stored elastic energy and activation before a later interval begins. A policy may exploit that retained state. Explaining its benefit requires controlled comparisons with the retained quantities declared, not a label such as “the network discovered passive dynamics.” Classical optimal control can exploit the same dynamics when its model and objective include them.

12.3 Neural Networks as Function Approximators

With column vectors, a feedforward network composes affine maps and nonlinear activations:

$$h_0 = x, \quad z_j = W_j h_{j-1} + b_j, \quad h_j = \sigma_j(z_j). \quad (12.2)$$

For a real-valued regression output, the last layer may be affine: $\hat{y} = W_L h_{L-1} + b_L$. If every activation is an identity, the whole network is affine, including a composed bias. The nonlinear activations create a richer function class. The weights have shapes determined by the input and output sizes of their respective layers.




Figure 12.1: A Fully Connected Three-Input, Three-Hidden-Unit, Two-Output Network. All Fifteen Inter-Layer Weights Are Shown; the Five Biases Are Omitted.

The network in the figure has three inputs, three hidden units and two outputs: nine input-to-hidden and six hidden-to-output weights, plus five biases. Its architecture alone specifies no task or training data. A torque policy, value function and forward-model residual can share an architecture while having entirely different input definitions, losses and validation requirements.

12.3.1 Universal Approximation and Its Limits

For continuous $f : [0, 1]^n \rightarrow \mathbb{R}$ and any $\epsilon > 0$, Cybenko's theorem permits a finite-width single-hidden-layer network with a continuous sigmoidal activation such that

$$\sup_{x \in [0, 1]^n} \left| f(x) - \sum_{i=1}^H c_i \sigma(w_i^T x + b_i) \right| < \epsilon. \quad (12.3)$$

The c_i are output coefficients and b_i are hidden biases. Width and parameter magnitudes are not fixed in advance. This is an existence result, not an algorithm for finding the coefficients or a bound on the data needed. It does not guarantee uniform approximation of a discontinuous policy by continuous networks. ReLU networks also have universal approximation results with sufficient width; a blanket requirement for extra hidden depth is inappropriate. Depth can help represent particular compositional functions efficiently, but cannot rescue missing information in the inputs.

For an elementary constructive example, $|x| = \max(0, x) + \max(0, -x)$. Two ReLU hidden units and one linear output represent this function exactly. That identity is a statement about representation. Learning those parameters from a finite noisy dataset is a separate optimization and generalization problem.

12.4 Activation Functions

ReLU is $\sigma(z) = \max(0, z)$, with derivative zero for $z < 0$ and one for $z > 0$; an implementation must choose a derivative convention at zero. A unit inactive on every training input supplies no local gradient through that branch. Whether it remains inactive depends on future inputs, upstream changes and the update rule. “Dead forever” is not a theorem for every network.

The logistic sigmoid satisfies

$$s(z) = \frac{1}{1 + e^{-z}}, \quad s'(z) = s(z)(1 - s(z)) \leq \frac{1}{4}, \quad (12.4)$$

with equality at $z = 0$. Tanh has derivative $1 - \tanh^2 z$, reaching one at zero and tending to zero in either tail. Both activations can saturate. GELU is $z\Phi(z)$, where Φ is the standard normal cumulative distribution; its derivative is $\Phi(z) + z\phi(z)$, with ϕ the corresponding density. Smoothness does not imply uniformly large or positive derivatives.

For a stack of layers, the input Jacobian contains products $D_j W_j$, where D_j is the diagonal activation-derivative matrix. Thus

$$\left\| \frac{\partial h_L}{\partial h_0} \right\| \leq \prod_{j=1}^L \|D_j\| \|W_j\|. \quad (12.5)$$

Ten sigmoid derivative factors alone are at most $4^{-10} = 9.5367 \times 10^{-7}$. The actual gradient also depends on weights, inputs and loss derivatives. Large weight norms can counteract that attenuation or cause explosion; multiplying ten sigmoid bounds is not a universal estimate of a first-layer parameter gradient.

12.5 Loss Functions and What They Estimate

For m scalar examples, use the half mean squared error

$$\ell(\theta) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2. \quad (12.6)$$

Under squared loss, the ideal unrestricted predictor is the conditional mean of the target given the input. If two hidden physiological states produce different torques at the same observed posture, averaging their labels need not yield a successful action for either state. This is a representation and information problem, not necessarily inadequate network width. Supervised labels can come from calibrated measurements, a specified inverse-dynamics procedure or simulation; a human need not supply every correct answer manually.

Huber loss with threshold $\delta > 0$ is

$$\ell_\delta(e) = \begin{cases} e^2/2, & |e| \leq \delta, \\ \delta(|e| - \delta/2), & |e| > \delta. \end{cases} \quad (12.7)$$

The threshold has the same units as the residual, unless residuals have been normalized. Linear growth reduces the influence of large residuals relative to squared loss; it does not

identify which observations are erroneous. A contact event, a sensor fault and a rare valid swing can all create a large residual for different reasons.

Multiple outputs require a declared metric. Squared radians, angular velocities and newtons cannot be added meaningfully without scale factors or a covariance model. Normalize using training data only and retain the transformations for deployment. Split evaluation by independent trajectories, sessions or participants as appropriate; random neighboring frames from one swing can leak almost identical states into both training and testing.

12.6 Backpropagation: Chain Rule for Neural Networks

Consider a scalar output $\hat{y} = W_2 h + b_2$, $h = \tanh(W_1 x + b_1)$, and a single-example loss $\ell = (\hat{y} - y)^2/2$. Set $e = \hat{y} - y$. Applying the chain rule gives

$$\begin{aligned} \nabla_{W_2} \ell &= e h^T, & \nabla_{b_2} \ell &= e, \\ d_1 &= (W_2^T e) \odot (1 - h \odot h), & & (12.8) \\ \nabla_{W_1} \ell &= d_1 x^T, & \nabla_{b_1} \ell &= d_1. \end{aligned}$$

For a batch mean, sum the individual gradients and divide by the actual number of examples in that batch. A shorter final batch uses its own size. Compute all derivatives using the same parameter values before applying updates: changing W_2 before computing d_1 would differentiate a different network.

Backpropagation efficiently accumulates these derivatives in reverse order through a computational graph. Gradient descent, momentum and Adam are optimizers that use the derivatives to update parameters. The gradient calculation itself is not a training algorithm, a guarantee of decreasing loss or a method for choosing a physically appropriate loss.

Finite differences provide an independent check. For one parameter component θ_j , compare backpropagation with

$$\frac{\ell(\theta + \varepsilon e_j) - \ell(\theta - \varepsilon e_j)}{2\varepsilon}. \quad (12.9)$$

Choose a step large enough to avoid cancellation and small enough to control truncation error. Check every parameter group, including biases, and avoid ReLU kinks for an ordinary smooth derivative check. Passing this test validates differentiation at the checked inputs; it does not certify generalization or control stability.

12.7 Gradient Descent and Optimization

For a differentiable objective with gradient $g = \nabla \ell(\theta)$, the directional derivative along $-g$ is $-\|g\|^2$. A sufficiently small step decreases the same smooth objective when $g \neq 0$. A fixed large step can overshoot. A mini-batch gradient is a noisy estimate, so even a reasonable stochastic step need not decrease the full-data loss on every iteration.

SGD updates $\theta_{k+1} = \theta_k - \alpha_k \hat{g}_k$. A momentum convention is $m_k = \beta m_{k-1} + \hat{g}_k$, followed by $\theta_{k+1} = \theta_k - \alpha_k m_k$. Other conventions multiply the new gradient by $1 - \beta$; the learning-rate interpretation changes accordingly. Momentum can reduce oscillation in some directions and accelerate progress in others, but it can also overshoot.

Adam maintains coordinatewise first and second moments:

$$\begin{aligned} m_k &= \beta_1 m_{k-1} + (1 - \beta_1) \hat{g}_k, \\ v_k &= \beta_2 v_{k-1} + (1 - \beta_2) \hat{g}_k \odot \hat{g}_k, \\ \hat{m}_k &= m_k / (1 - \beta_1^k), \quad \hat{v}_k = v_k / (1 - \beta_2^k), \\ \theta_{k+1} &= \theta_k - \alpha_k \hat{m}_k / (\sqrt{\hat{v}_k} + \varepsilon). \end{aligned} \tag{12.10}$$

Here $m_0 = v_0 = 0$, operations in the last line are elementwise, and $\varepsilon > 0$ prevents division by zero. Moment adaptation is different from a schedule for α_k . Neither Adam nor momentum universally converges for every nonconvex or changing reinforcement-learning objective. Learning rate, normalization, initialization and optimizer settings remain part of the experiment.

12.8 Markov Decision Processes

For clarity, first use a finite horizon of T decisions. At step t , the agent observes state s_t , samples action a_t , receives reward r_t , and transitions to s_{t+1} . The distribution $P(s_{t+1}, r_t \mid s_t, a_t)$ is the environment model. The reward index here names the action that generated it; some texts use r_{t+1} for the same event.

The Markov property means this conditional distribution does not additionally depend on earlier history when s_t is complete. Include time in the state or use time-indexed policies when the task has a deadline. A terminal state is part of the task definition, not just an accidental interruption of data collection.

12.8.1 Return and Value

Define $G_t = \sum_{k=t}^{T-1} \gamma^{k-t} r_k$ and $J(\theta) = \mathbb{E}_{\pi_\theta}[G_0]$, with $0 \leq \gamma \leq 1$. Rewards must have finite expectations. For an infinite discounted task, bounded rewards and $\gamma < 1$ give a finite return; $\gamma = 1$ requires additional termination or summability conditions. Continuing average-reward control is a different formulation and needs its own assumptions.

Under a specified policy π , define

$$V_t^\pi(s) = \mathbb{E}[G_t \mid s_t = s], \quad Q_t^\pi(s, a) = \mathbb{E}[G_t \mid s_t = s, a_t = a]. \tag{12.11}$$

The advantage $A_t^\pi(s, a) = Q_t^\pi(s, a) - V_t^\pi(s)$ compares an action with the policy's conditional average, not necessarily with doing nothing. Since $\mathbb{E}_{a \sim \pi}[A_t^\pi(s, a)] = 0$, it expresses relative merit under the chosen reward, horizon and future policy.

12.8.2 Policy Evaluation Versus Optimality

With $V_T^\pi = 0$, conditioning on the next transition gives

$$V_t^\pi(s) = \sum_a \pi_t(a \mid s) \mathbb{E}[r_t + \gamma V_{t+1}^\pi(s_{t+1}) \mid s, a]. \tag{12.12}$$

This Bellman equation evaluates the fixed policy. It does not optimize that policy. For finite action sets, the optimal recursion replaces the average with a maximum:

$$V_t^*(s) = \max_a \mathbb{E}[r_t + \gamma V_{t+1}^*(s_{t+1}) \mid s, a]. \tag{12.13}$$

Continuous spaces require integrals, and a supremum replaces the maximum unless an optimizer exists. Policy iteration alternates evaluation and improvement; value iteration applies an optimality operator. Approximate neural versions do not inherit every convergence property of finite tabular algorithms.

For a one-decision example, actions return 0 and 2. The policy choosing each equally has value 1; the optimal value is 2. Solving the evaluation equation more accurately cannot improve the equal-choice policy. This simple distinction prevents a common misunderstanding when interpreting a learned critic.

12.9 Policy Gradient Methods

Let $\zeta = (s_0, a_0, r_0, \dots, s_T)$ be a trajectory. Assume the initial-state and environment distributions are independent of θ , the policy has differentiable positive likelihood on the sampled support, and differentiation can pass through the expectation. Its likelihood factors as

$$p_\theta(\zeta) = p_0(s_0) \prod_{t=0}^{T-1} \pi_\theta(a_t | s_t) P(s_{t+1}, r_t | s_t, a_t). \quad (12.14)$$

The log derivative is a sum of policy scores. Therefore

$$\nabla_\theta J = \mathbb{E} \left[G_0 \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \right]. \quad (12.15)$$

Past rewards have zero expected product with a later score, because the score has conditional mean zero before the action is sampled. Removing those terms yields the return-to-go form

$$\nabla_\theta J = \mathbb{E} \left[\sum_{t=0}^{T-1} \gamma^t \nabla_\theta \log \pi_\theta(a_t | s_t) G_t \right]. \quad (12.16)$$

This derivation accounts for the effect of actions on subsequent state visitation through the trajectory likelihood. It does not simply hold a policy-dependent state distribution constant. If the initial distribution, environment or reward depends directly on θ , additional derivative terms are needed. The explicit γ^t belongs to the stated start-state objective; uniform time-step averages used in practical implementations can correspond to a different weighting convention.

12.9.1 Baselines and REINFORCE

For an action-independent baseline $b_t(s_t)$,

$$\mathbb{E}_{a_t \sim \pi_\theta} [b_t(s_t) \nabla_\theta \log \pi_\theta(a_t | s_t)] = b_t(s_t) \nabla_\theta \int \pi_\theta(a | s_t) da = 0. \quad (12.17)$$

Thus $G_t - b_t(s_t)$ can replace G_t without changing the expected score estimator under these assumptions. Treat the baseline and sampled returns as constants in the actor loss; differentiating through a learned baseline adds unwanted terms. A baseline fitted using the same action's outcome can also introduce finite-sample dependence unless the estimator is designed appropriately.

REINFORCE samples complete returns to estimate the score-weighted update. A useful baseline can reduce variance, but an arbitrary baseline can increase it. The variance-minimizing scalar baseline generally weights returns by score magnitude; it is not always exactly V^π . Reward scale, horizon, rare successes and stochastic transitions affect estimator variance.

12.10 Actor-Critic Architecture

An actor represents the policy and a critic estimates value. For a transition that genuinely terminates the task, let $d_t = 1$; otherwise $d_t = 0$. A one-step target and temporal-difference residual are

$$y_t = r_t + \gamma(1 - d_t)V_\phi(s_{t+1}), \quad \delta_t = y_t - V_\phi(s_t). \quad (12.18)$$

The usual semi-gradient critic step minimizes $(V_\phi(s_t) - \text{stopgrad}(y_t))^2/2$. An actor loss uses a detached advantage estimate, for example $-\log \pi_\theta(a_t | s_t) \text{stopgrad}(\delta_t)$ with the objective's appropriate weighting. A low critic training loss is not a proof that δ_t estimates the true advantage accurately on all visited states.

An artificial rollout cutoff differs from task termination. If the underlying task continues, bootstrap from the final observed state; do not use a reset observation as that state. If a deadline defines the task itself, include remaining time and use zero continuation at its terminal boundary. This choice can materially change a swing policy trained around an impact deadline.

12.10.1 Multi-Step Returns and Bias-Variance Tradeoffs

Before a terminal boundary, an n -step target is

$$G_t^{(n)} = \sum_{j=0}^{n-1} \gamma^j r_{t+j} + \gamma^n V_\phi(s_{t+n}). \quad (12.19)$$

Stop the sum at termination and omit continuation there. A longer return typically relies less on an imperfect bootstrap and uses more sampled rewards, often reducing bootstrap bias while increasing variance. These are tendencies, not universal monotonic laws: reward correlations, critic errors and environment noise matter. With an exact critic under the same policy, the shorter target need not have bootstrap bias.

Generalized advantage estimation forms a truncated weighted sum of TD residuals, $\hat{A}_t = \sum_{j=0}^{K-1} (\gamma\lambda)^j \delta_{t+j}$, with boundary masks and final bootstrap handled consistently. λ changes the mixture of multi-step estimates. Neither this estimator nor nonlinear actor-critic training guarantees convergence without further assumptions.

12.11 Proximal Policy Optimization

PPO gathers data under a frozen old policy and optimizes a surrogate using those data. Define the likelihood ratio $\rho_t(\theta) = \pi_\theta(a_t | s_t) / \pi_{\text{old}}(a_t | s_t)$, keeping the stored old likelihood and advantage fixed during the update. Its clipped sample objective is

$$L_t^{\text{clip}} = \min \left(\rho_t \hat{A}_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right). \quad (12.20)$$

For positive advantage this is $\hat{A}_t \min(\rho_t, 1 + \epsilon)$; for negative advantage it is $\hat{A}_t \max(\rho_t, 1 - \epsilon)$. The first branch stops rewarding a sufficiently large increase in the sampled action's likelihood. The second stops rewarding a sufficiently large decrease. Harmful changes in the opposite direction are still penalized.

With $\epsilon = 0.2$ and $\hat{A} = 2$, ratios 1.2 and 2 both give 2.4. With $\hat{A} = -2$, ratios 0.8 and 0.1 both give -1.6 . These plateaus do not forbid those ratios. Shared parameters, other samples and optimizer steps can move probabilities outside the interval. The clipped quantity is no greater than the corresponding unclipped sampled surrogate, but is not a lower-bound certificate for true future performance. Clipping does not enforce a hard ratio or KL-divergence constraint or guarantee that a policy cannot collapse.

An actor-critic implementation can maximize the average clipped objective minus c_V times a value-error loss plus c_H times an entropy bonus. Equivalently, minimize its negative. Define signs and coefficients explicitly. Monitor actual likelihood ratios, KL estimates, entropy, critic errors and independent rollout performance; a plateau in the surrogate is not sufficient evaluation.

12.11.1 Stochastic Bounded Actions

A deterministic tanh network alone does not provide PPO's sampling likelihood. One possible continuous policy draws $z \sim \mathcal{N}(\mu_\theta(s), \text{diag}(\sigma_\theta(s)^2))$ with positive standard deviations and maps $a_i = a_{\max,i} \tanh z_i$. Inside the bounds, change of variables gives

$$\log \pi_\theta(a | s) = \log p_\theta(z | s) - \sum_i \log [a_{\max,i} (1 - \tanh^2 z_i)]. \quad (12.21)$$

The action scales are positive and fixed here. Use numerically stable log-Jacobian evaluation near saturation. For old and new policies with the same transform and the same stored action, the common Jacobian cancels in their ratio; it still matters for the actual action density and entropy. Sampling an unbounded Gaussian and clipping afterward produces a different distribution, with probability mass at the boundaries, and cannot silently use this tanh density.

Collect new rollouts after a bounded number of optimization epochs and freeze a new old policy for the next update. Excessive reuse of the old data, changed normalization statistics or inconsistent likelihood calculations can invalidate the intended comparison. PPO is an algorithmic design, not a guarantee that these implementation details are automatically correct.

12.12 Value-Based Methods and Continuous Actions

Tabular Q-learning uses the sampled optimality target:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_t + \gamma(1 - d_t) \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]. \quad (12.22)$$

Its standard discounted finite-MDP convergence results require suitable learning-rate schedules and repeated exploration of the relevant state-action pairs. Replacing the table with a neural network changes the problem substantially.

DQN fits a network to targets from a separate target network:

$$\ell_Q = \mathbb{E}_{(s,a,r,s',d) \sim \mathcal{D}} \left[\left(Q_\theta(s, a) - \text{stopgrad}[r + \gamma(1 - d) \max_{a'} Q_{\text{target}}(s', a')] \right)^2 \right]. \quad (12.23)$$

Replay changes how stored experience is reused and can reduce correlations within training batches. A delayed target slows movement of the bootstrap. Neither device eliminates all statistical dependence, approximation error or instability. Exploration still determines what information enters the buffer.

For continuous actions, maximizing a general nonlinear Q can be an expensive nonconvex optimization at every state. It is not categorically impossible; special quadratic or concave forms can be tractable. DDPG uses a deterministic actor μ_θ and a learned critic. A practical actor update follows

$$\hat{g} = \mathbb{E}_{s \sim \mathcal{D}} \left[(D_\theta \mu_\theta(s))^T \nabla_a Q_\phi(s, a) \Big|_{a=\mu_\theta(s)} \right]. \quad (12.24)$$

This is a gradient through a learned critic over replay states, not an exact global argmax or automatically an unbiased start-state performance gradient. The deterministic policy-gradient theorem uses a specified state-visitation weighting and true or suitably compatible action values. Exploration must be supplied during data collection because a deterministic actor itself does not sample alternatives.

TD3 addresses particular actor-critic estimation errors using two critics with a minimum target, delayed actor updates and target-policy smoothing. Taking a minimum can also create underestimation. SAC uses an entropy-regularized objective, schematically

$$J_{\text{soft}} = \mathbb{E} \sum_t \gamma^t [r_t + \alpha_{\text{ent}} \mathcal{H}(\pi(\cdot | s_t))]. \quad (12.25)$$

The entropy coefficient changes the tradeoff being optimized. In continuous action coordinates this is differential entropy, which depends on action scaling and can be negative. These methods differ in policy class, objective and data reuse; there is no universal ranking of their stability or sample efficiency for every swing model.

12.13 Exploration Strategies

In a finite action set, epsilon-greedy exploration chooses a random action with probability ϵ and otherwise a greedy action. A uniform distribution over an unbounded continuous action space does not exist; continuous exploration requires a defined distribution and admissible bounds. For a finite action set, entropy is

$$\mathcal{H}(\pi(\cdot | s)) = - \sum_a \pi(a | s) \log \pi(a | s), \quad (12.26)$$

with $0 \log 0 = 0$ by continuity. Entropy regularization encourages spread according to its chosen density and coordinates, but broad action noise alone does not ensure useful coverage of long-horizon states.

A curiosity bonus can depend on forward-prediction error,

$$r_t^{\text{int}} = \eta \|\hat{T}_\theta(s_t, a_t) - s_{t+1}\|_W^2. \quad (12.27)$$

The weighting W declares the state metric. Large error can indicate a new learnable transition, sensor noise, stochasticity or a deficient model class. Rewarding all such error can repeatedly attract an agent to irreducible noise. Adding the bonus also changes the optimization objective unless its role and limiting behavior are specified.

For golf, define an exploratory simulation task before assigning human significance to its discoveries. A policy that gains club speed by exploiting unphysical contact penetration or unlimited torque has found a weakness in the model or reward, not a coaching principle.

12.14 Worked Example: Pendulum Swing-Up Objective

Use a point mass m on a massless rod of length l , a fixed pivot, viscous rotational damping $b \geq 0$, and torque u . The angle $q = 0$ points downward and $q = \pi$ upward. With $I = ml^2$,

$$I\ddot{q} = u - b\dot{q} - mgl \sin q, \quad |u| \leq u_{\max}. \quad (12.28)$$

The model's mechanical energy above hanging and its rate are

$$E = \frac{1}{2}I\dot{q}^2 + mgl(1 - \cos q), \quad \dot{E} = u\dot{q} - b\dot{q}^2. \quad (12.29)$$

These formulas make signs and input power checkable. A sample implementation must also declare its integration method, step size, torque hold and terminal conditions. A physical pendulum with distributed rod mass would have a different inertia and center-of-mass gravity coefficient.

Let $e_q = \text{atan2}(\sin(q - \pi), \cos(q - \pi))$. One dimensionless upright reward is

$$r(q, \dot{q}, u) = -[e_q^2 + w_v(\dot{q}/\omega_*)^2 + w_u(u/u_{\max})^2], \quad (12.30)$$

where angular measure is in radians, $\omega_* > 0$ is a declared speed scale, and $w_v, w_u \geq 0$ are dimensionless weights. The wrapped error is discontinuous in sign at the opposite orientation, though its square is continuous there; its derivative has a cusp. The smooth periodic alternative $1 + \cos q$ also has its minimum at upright. In either case, the objective must be examined at both equilibria before training.

For illustrative $m = l = 1$, $g = 9.81$, $b = 0$, $u_{\max} = 2$, $\omega_* = 1$, $w_v = 0.1$, and $w_u = 0.01$, upright at rest with zero torque gives reward 0, while hanging gives $-\pi^2$. Replacing e_q with the raw downward angle would reverse the intended preference. At $q = \pi/2$, $\dot{q} = 1$ and $u = 2$, acceleration is -7.81 rad/s^2 and input power is 2 W. The energy difference between the two resting equilibria is 19.62 J.

An action-squared penalty is an effort regularizer, not mechanical work. Work is $\int u\dot{q} dt$; it can be positive or negative, whereas u^2 cannot. Metabolic expenditure requires an additional physiological model. A torque-free interval may retain kinetic energy from earlier actuation; low instantaneous effort does not show that preparation was free.

For a reproducible training experiment, specify an initial-state distribution, sampling interval, finite horizon, policy distribution, optimizer settings and random seeds. Compare against a declared baseline controller and evaluate unseen initial states and parameter perturbations. Report success criteria such as dwell time near upright, speed tolerance, torque feasibility and energy balance in addition to return. No PPO training run is reported here, so no learned swing-up trajectory or success rate is claimed.

12.15 Physics-Informed Neural Networks

A neural network can approximate a solution while a loss penalizes a specified differential-equation residual. For example, a nondimensional viscous Burgers problem is

$$\begin{aligned} w_t + ww_\xi - \nu w_{\xi\xi} &= 0, & \xi \in [-1, 1], \quad t \in [0, 1], \\ w(0, \xi) &= -\sin(\pi\xi), & w(t, -1) = w(t, 1) = 0, \end{aligned} \quad (12.31)$$

with $\nu = 0.01/\pi$ in the cited example. We use w rather than control input u to avoid a notation collision. Given an approximation $w_\theta(t, \xi)$, its residual is $R_\theta = w_{\theta,t} + w_\theta w_{\theta,\xi} -$

$\nu w_{\theta, \xi \xi}$. A training loss combines initial/boundary data errors and squared residuals at collocation points. The nonlinear transport sign and viscosity term must agree with the stated PDE.

This is different from fitting an unknown vector field directly to derivative labels. For $\dot{x} = F_{\theta}(x, u)$, minimizing $\sum_i \|F_{\theta}(x_i, u_i) - \dot{x}_i\|_W^2$ is derivative regression. It becomes physics-informed only through additional known equations, restrictions or priors. Noisy numerical differentiation can bias those labels, and a full-state derivative includes configuration rates as well as accelerations.

For the known forced pendulum, an appropriate trajectory residual is $I\ddot{q}_{\theta} + b\dot{q}_{\theta} + mgl \sin q_{\theta} - u(t)$. Observations and initial conditions constrain the solution as well. A small residual at finitely many points does not establish a global solution error bound, unique parameters or valid behavior at impacts. Enforcing an incorrect conservative equation on a driven dissipative swing can produce a physically misleading fit despite an apparently principled loss.

12.16 Neural ODEs and the Adjoint Method

A Neural ODE parameterizes a continuous vector field $\dot{h} = f_{\theta}(h, t)$ and computes states with an ODE solver. The state can be physical or latent; a latent dimension does not automatically correspond to measured joint coordinates. A policy evaluated at discrete observations is not a Neural ODE merely because it controls a continuous mechanical system.

For a terminal loss $\ell = \Psi(h(T), \theta)$ with fixed endpoints, define a column adjoint $a(t)$ by

$$\dot{a} = -(D_h f_{\theta})^T a, \quad a(T) = \nabla_h \Psi(h(T), \theta). \quad (12.32)$$

Differentiating $a^T \delta h$ gives the parameter contribution after the state-variation terms cancel. The result is

$$\nabla_{\theta} \ell = \partial_{\theta} \Psi + (D_{\theta} h(0))^T a(0) + \int_0^T (D_{\theta} f_{\theta})^T a \, dt. \quad (12.33)$$

If the initial state is parameter-independent and the terminal loss has no direct parameter dependence, only the integral remains. Writing it as $-\int_T^0$ is equivalent; that reversed-limit notation is not itself a sign error. Running costs add source terms to the adjoint and parameter integral, while losses at intermediate observation times introduce adjoint jumps. Contact events need their own event and reset sensitivities.

These equations differentiate the continuous initial-value problem. Differentiating a particular numerical solver produces the derivative of its discrete computation; the two agree only to the relevant numerical accuracy. Backward reconstruction of a forward-dissipative state can be unstable. Checkpointing and solver tolerances affect memory use, accuracy and cost, so a universal constant-memory exact-gradient claim is unwarranted.

As a direct check, let $\dot{h} = \theta h$, $h(0) = 1$, $T = 1$ and $\Psi = h(T)^2/2$. Then $h(T) = e^{\theta}$ and $d\ell/d\theta = e^{2\theta}$. The adjoint integral gives the same value. This verifies the sign convention on a solvable example, not the accuracy of every adaptive-solver implementation.

12.17 Structure-Preserving Neural Networks

12.17.1 Hamiltonian Models

An autonomous Hamiltonian neural network learns a scalar $H_\theta(q, p)$ in canonical coordinates and defines

$$\dot{q} = \nabla_p H_\theta, \quad \dot{p} = -\nabla_q H_\theta. \quad (12.34)$$

Along a smooth exact solution, $dH_\theta/dt = 0$ because the two cross-products cancel. The conserved quantity is the learned Hamiltonian; data and physical calibration must establish whether it equals the real energy. Canonical momentum is not generally equal to velocity. A numerical integrator need not conserve H_θ exactly or preserve the symplectic form merely because its vector field is Hamiltonian.

For an explicit force f_p added to $\dot{p}$, the balance becomes $\dot{H}_\theta = \nabla_p H_\theta^T f_p$, plus $\partial_t H_\theta$ when time dependence is present. Dissipation and actuation therefore call for power balances, not a claim that an actively driven golf swing conserves total mechanical energy.

12.17.2 General Lagrangian Models

For a learned smooth $L_\theta(q, v)$ with $v = \dot{q}$, the forced Euler–Lagrange equation is

$$\frac{d}{dt} \nabla_v L_\theta - \nabla_q L_\theta = \tau. \quad (12.35)$$

Let $H_{vv} = D_v(\nabla_v L_\theta)$ and $H_{vq} = D_q(\nabla_v L_\theta)$. If H_{vv} is invertible,

$$\dot{v} = H_{vv}^{-1} [\tau + \nabla_q L_\theta - H_{vq}v]. \quad (12.36)$$

A generic Lagrangian’s velocity Hessian can depend on both q and v and can be indefinite or singular. It is not automatically a positive-definite mass matrix depending only on configuration. A pseudoinverse may produce a numerical vector, but does not establish that a singular system has unique physically consistent dynamics. Time-dependent Lagrangians require an additional $\partial_t \nabla_v L$ term.

For an autonomous regular model, the energy function is $E_L = v^T \nabla_v L - L$, and the forced balance is $\dot{E}_L = v^T \tau$. Lagrangian gauge freedoms also mean that fitting trajectories need not identify one unique scalar Lagrangian. Do not interpret every learned coefficient as a measured anatomical parameter.

12.17.3 Natural Mechanical Models and DeLaN

A more restricted construction sets

$$M_\theta(q) = R_\theta(q)R_\theta(q)^T + \varepsilon I, \quad L_\theta = \frac{1}{2}v^T M_\theta(q)v - U_\theta(q), \quad (12.37)$$

where R_θ is lower triangular and $\varepsilon > 0$ is defined in appropriately scaled inertia units. This guarantees a positive-definite inertia matrix. Deriving both inertial bias and gravity from this same L_θ gives a coherent natural mechanical model. It does not guarantee identified parameters or accurate contact forces outside the training regime.

The original DeLaN paper parameterizes a triangular inertia factor and a separate gravity-force network. Learning a scalar potential and taking its gradient, as in the construction above, adds an integrability restriction; it should not be attributed to every original

DeLaN implementation. A separately learned force vector is not automatically conservative. Likewise, an independently learned Coriolis matrix must not be assumed consistent with a learned mass matrix without verification.

Friction, muscle activation, tendon storage, contact and input limits require explicit representation beyond the conservative component. For the swing, a useful learned model should expose which quantities exchange power, which dissipate energy and which retain history. Structural priors can reduce certain impossible predictions while still encoding the wrong physiological mechanism.

12.18 Sim-to-Real Transfer and System Identification

A known analytical model or a learned forward model can support model-based planning. Model-free learning does not explicitly require that forward model in its policy/value update, but can still use physical representations, simulation, constrained actions and engineered rewards. Neither category is universally more data-efficient. Model errors, exploration cost and the precision required near impact determine the relevant tradeoffs.

Domain randomization trains over a specified distribution of plausible parameters, delays, sensor errors and disturbances. The distribution is part of the model. Mass and inertia should vary consistently, and correlated subject parameters should not be replaced automatically by independent uniform noise. Terrestrial gravity is usually much better known than muscle strength or sensor bias; arbitrary gravity variation is not a default model of differences among golfers.

Evaluate on withheld parameter combinations and, when claiming transfer, on appropriately governed physical observations. Training on diverse simulations does not substitute for real validation. Report failures and uncertainty rather than only mean reward. A policy may be robust to the sampled uncertainties and fragile to an omitted contact or actuator mechanism.

12.18.1 What Can the Data Identify?

For an unforced ideal point-mass pendulum,

$$\ddot{q} = -\kappa \sin q, \quad \kappa = g/l. \quad (12.38)$$

Any pair (cg, cl) with $c > 0$ gives the same angle dynamics. Angle-only free-motion data therefore cannot identify g and l separately. At an exact resting equilibrium they may provide no information even about κ . A network cannot remove this structural ambiguity by having more parameters.

An independent length measurement, known gravity or a calibrated forced experiment can add information. With measured torque, acceleration depends on both $1/(ml^2)$ and g/l ; which individual parameters become identifiable still depends on what else is known and whether the input sufficiently excites the system. In a golfer, redundant muscle forces and unknown external loads create analogous ambiguities. Design the measurement and intervention before giving a fitted parameter a mechanistic interpretation.

12.19 Python Worked Example: Backpropagation From Scratch

The accompanying example fits the scalar function $y = x^2$ on $[-1, 1]$ using a tanh hidden layer and a linear output. Its implementation is shared between the print and web lessons in the repository module `src.tools.learning_examples`. This is supervised regression, not reinforcement learning, inverse dynamics or a pendulum controller.

The implementation uses row samples, so X has shape $(m, 1)$, the input weights have shape $(1, H)$, hidden activations have shape (m, H) and output weights have shape $(H, 1)$. These are the transposed storage conventions of the earlier column-vector derivation. The actual gradient core is reproduced below; training and validation use that same source.

```

1 prediction = predict(network, features)
2 hidden = np.tanh(features @ network.input_weights + network.hidden_bias)
3 error = prediction - targets
4 output_derivative = error / features.shape[0]
5 hidden_derivative = (output_derivative @ network.output_weights.T) * (1 -
    hidden**2)
6 gradient = Network(
7     features.T @ hidden_derivative,
8     hidden_derivative.sum(axis=0, keepdims=True),
9     hidden.T @ output_derivative,
10    output_derivative.sum(axis=0, keepdims=True),
11 )
12 loss = float(np.mean(error**2) / 2)

```

Mini-batches are shuffled with an explicit random generator. Every nonempty slice is processed, including a partial final batch or one batch larger than the whole dataset. After all updates, the reported epoch loss is recomputed on the full dataset at the final parameters. Averaging batch losses computed at changing parameter values would answer a different question, and dividing by the number of full batches would misweight a remainder or divide by zero for an oversized batch.

```

1 from src.tools.learning_examples import regression_demo
2
3 result = regression_demo()
4 print(result) # Interactive lesson output; the library performs no printing.

```

The declared example uses seed 2026, 12 hidden units, 41 equally spaced training inputs, batch size 16, learning rate 0.05 and 1,200 epochs. It evaluates additional interval midpoints never used as training inputs. The numerical report is:

Initial half MSE: 0.158078703; final training half MSE: 0.001261210; unseen-midpoint half MSE: 0.001021767.

Tests compare manual gradients with centered finite differences for every weight and bias, check partial and oversized batches against the actual final full-data loss, and verify deterministic reproduction. Midpoint accuracy tests interpolation within this example's interval. It is not evidence for extrapolation to arbitrary inputs, a simulation throughput claim, or successful control of a physical system.

12.20 Deep Reinforcement Learning for Robot and Golf Control

The technical chain begins with the task: desired impact position and time, clubhead velocity and orientation, balance or contact requirements, and allowable actuator and tissue loads. Those quantities determine what the policy must know and what its reward should value. A scalar reward is a modeling choice, not an automatic expression of “good golf.” Distinct combinations of speed, accuracy, repeatability and loading can produce different preferred actions.

12.20.1 State Representation and Action Parameterization

Use consistent frames, units and phase conventions. A representation containing pelvis orientation, joint configuration, generalized velocity, club state and target information still may omit activation, shaft vibration or contact history. Test the consequences of these omissions. A periodic angle representation such as $(\sin q, \cos q)$ avoids a raw wrap discontinuity but does not uniquely represent every three-dimensional orientation; rotation encodings need their own constraints and continuity analysis.

Direct joint torque, muscle excitation and desired joint position are different action ports. A desired position normally passes through a feedback controller, so its force response depends on gains and current state. Muscle excitation passes through activation and contraction dynamics. Document these layers and their limits before comparing policies. A low excitation at impact can coexist with retained activation or tendon force, and a bounded position command need not imply a bounded joint torque.

When comparing actions from the same frozen state in a torque-affine mechanical model, instantaneous acceleration increments can add linearly. Policies applied over time change the state, feedback and contact mode, so their full trajectories need not add. Learned strategies should respect this distinction rather than treating trajectory interaction as evidence against force addition.

12.20.2 Reward Shaping and Evidence

Adding speed bonuses, torque penalties or trajectory-tracking terms can change the optimal task. Potential-based shaping has a special telescoping structure: with $r'_t = r_t + \gamma\Phi(s_{t+1}) - \Phi(s_t)$,

$$G'_0 = G_0 - \Phi(s_0) + \gamma^T \Phi(s_T). \quad (12.39)$$

For a fixed initial distribution and an action-independent terminal term, such as zero potential at every terminal state, the added terms do not change policy ranking. In an infinite discounted task a bounded potential makes the terminal term vanish. Arbitrary penalties do not share this property. A finite-horizon impact bonus or a nonzero terminal potential must be checked explicitly.

A large penalty for exceeding a limit is not a guarantee that the limit is never exceeded. Hard constraints, feasibility checks and validated safety mechanisms have separate roles. A reported simulation experiment should name its numerical tolerances, failure conditions, episode resets and contact handling. If an outcome depends on a very narrow impact timing interval, small integration or observation delays can alter the apparent strategy.

To explain a learned swing, compare matched models and controlled interventions: retain the initial state, define the altered input port, preserve unrelated parameters and measure the specified downstream outcome. Assess whether the explanation survives plausible modeling errors and withheld observations. A zero-torque counterfactual removes the declared torque channel; it does not erase prior work, stored energy, gravity or other active channels. Sensitivity analysis can reveal interactions, while identifiability analysis tells us which of those interactions the available measurements can distinguish.

12.21 Chapter Summary and Further Reading

Function approximation supplies a flexible representation; differentiation supplies parameter sensitivities; optimization uses those sensitivities to change a stated objective. Reinforcement learning adds the consequences of decisions over time. Learned mechanics adds structural assumptions about how states evolve. Each layer can be mathematically correct while the overall physical explanation remains unsupported because the state, input, objective or evidence is wrong.

The preceding geometry and recursive-dynamics chapters provide coordinate maps and consistent mechanics. This chapter adds tools for fitting uncertain components and optimizing policies. Their connection is especially useful in golf because earlier actions change the conditions under which later actions operate. Technical rigor comes from making those dependencies testable, not from assigning a universal sequence or physiological interpretation to a network's behavior.

The references below identify methods and assumptions, not empirical validation of a human golf policy. Cybenko's Theorem 2 supports the bounded approximation statement. Goodfellow, Bengio and Courville's chapters on feedforward networks and optimization distinguish gradient computation from training. Sutton and colleagues' policy-gradient paper explains the state-visitation weighting and compatible-critic conditions. The PPO paper's clipped surrogate and actor-critic algorithm motivate the update described here; they do not prove hard clipping or safety.

Cybenko (1989), Theorem 2 [8]. Uniform approximation on the unit cube; no training guarantee.

Goodfellow, Bengio and Courville (2016), Chapters 6 and 8 [15]. Feedforward models, loss functions, backpropagation and optimization; see the book navigation for Chapter 8.

Sutton and Colleagues (1999), Policy Gradients [42]. State-visitation weighting and compatible function approximation; not a blanket deep actor-critic convergence theorem.

Schulman and Colleagues (2017), PPO [36]. Sections 3 and 5 define the clipped surrogate and rollout/update procedure.

Mnih and Colleagues (2015), DQN. Replay and target-network method; its game benchmarks are not golf evidence.

Silver and Colleagues (2014), Deterministic Policy Gradients. The deterministic policy-gradient framework and exploration distinction.

Lillicrap and Colleagues, Continuous Control. The DDPG actor, critic, replay and target-network implementation.

Fujimoto, Van Hoof and Meger (2018), TD3. Twin critics, delayed updates and target-policy smoothing; Algorithm 1.

Haarnoja and Colleagues (2018), SAC. Entropy-regularized objective and approximate soft actor-critic; theoretical assumptions differ from practical approximation.

Raissi, Perdikaris and Karniadakis, PINNs [34]. The authors' specified Burgers problem and collocation residual; no general solution-error certificate.

Chen and Colleagues (2018), Neural ODEs [6]. Section 2, numerical considerations and Appendix B explain adjoint calculation and its numerical limits.

Greydanus, Dzamba and Yosinski (2019), HNNs [16]. Section 2 constructs a learned conserved quantity in canonical coordinates.

Cranmer and Colleagues (2020), LNNs [7]. Euler–Lagrange derivative construction and its distinction from a natural mechanical model.

Lutter, Ritter and Peters (2019), DeLaN [23]. Sections 3–4 parameterize inertia and a separate gravity force; distinguish a scalar-potential variant.

12.22 Exercises

1. **Representation:** Construct $|x|$ from two ReLU units. Then explain why continuous networks cannot uniformly approximate a jump of nonzero height to arbitrarily small error on a domain containing the jump.
2. **Backpropagation:** Derive all four parameter gradients for the tanh network. Compare them with centered finite differences at three nonzero inputs. State the loss normalization and array shapes.
3. **Activation Functions:** Plot sigmoid, tanh, ReLU and GELU and their derivatives. Identify saturation regions and nondifferentiability. Discuss advantages and disadvantages for a specified regression problem without asserting one activation is universally best.
4. **Gradient Bounds:** Derive the product-of-Jacobians bound for ten sigmoid layers. Evaluate the bound when every weight norm is one and when every weight norm is eight. Explain why neither is a measured parameter gradient.
5. **Pendulum Task:** Specify an MDP with the downward angle convention, torque bounds, input hold, integrator, observation, horizon and upright reward. Check reward and energy at both resting equilibria and at the numerical state in the text.
6. **Evaluation and Improvement:** In a 5-by-5 grid, movements are deterministic; attempted off-grid moves leave the state unchanged. The upper-right cell is absorbing with zero future reward; every preceding action earns -1 , including the action entering the goal. Use $\gamma = 0.9$. Compare uniform-random policy evaluation with optimal value iteration from all-zero initial values.
7. **Policy Gradient:** Derive the finite-horizon score estimator for the stated discounted start-state objective. Show explicitly where the γ^t factor and the cancellation of past rewards enter.
8. **Bootstrap Boundaries:** For rewards $(1, 2, 3)$, $\gamma = 0.9$, critic values $V(s_1) = 4$, $V(s_2) = 5$, and true termination at s_3 , compute the one-, two- and three-step targets from s_0 . Explain what repeated trajectories would be needed to compare their variance.

9. **Actor-Critic:** Write pseudocode that keeps bootstrap targets and actor advantages detached. Distinguish a genuine terminal transition from a truncated continuing roll-out and specify the final observation used for bootstrapping.
10. **PPO:** Plot the clipped sample objective for $\epsilon = 0.2$ and advantages 2 and -2 . Give ratios outside the interval that lie on its plateaus. Explain why the plot does not establish a hard probability bound or prevent policy collapse.
11. **Physics Residual:** Write a trajectory-fitting loss for the forced damped pendulum with known input and initial conditions. Give a residual that would be incorrect if actuation or damping were omitted. Propose held-out residual checks.
12. **Domain Randomization:** Define a physically consistent uncertainty distribution for the pendulum's mass, length, torque calibration and observation delay. Explain how inertia changes with mass and length, and separate training draws from held-out evaluation cases.
13. **Identification:** Show that scaling both g and l by the same positive factor leaves free angle trajectories unchanged. Identify an additional measurement or calibrated experiment and explain which ambiguity it removes.
14. **Adjoint and Conservation:** Verify the scalar Neural ODE derivative in the text analytically and with finite differences. Separately show that explicit Euler applied to $H = (q^2 + p^2)/2$ multiplies H by $1 + h^2$ per step, despite the exact Hamiltonian flow conserving it.
15. **Reproducible Learning:** Run the shared regression example with batch sizes that divide the dataset, leave a remainder and exceed its size. Then design, but do not claim to have executed, a multi-seed PPO pendulum comparison with explicit success and failure criteria.
16. **Golf Interpretation:** A simulated policy uses less commanded torque near impact after a stronger early action. List the retained states and power flows needed to interpret that result. Design a matched counterfactual and identify which conclusions would still require independent human measurements.

Glossary of Important Concepts

These definitions distinguish geometric objects, physical models, and the guarantees obtained from a calculation. A definition does not establish that a particular golf model satisfies its assumptions.

Articulated Body Algorithm

A recursive forward-dynamics algorithm for a rigid-body tree. With bounded joint dimension, its arithmetic cost scales linearly with the number of bodies. Closed-loop constraints and contact require additional treatment; the tree algorithm alone does not solve every constrained contact problem.

Configuration Space

The set of configurations allowed by a model, including positions, orientations, and declared geometric constraints. It is a smooth manifold when suitable regularity conditions hold; redundant coordinates and constraint singularities require care. Velocities are additional state information.

Contraction Analysis

Analysis of convergence between trajectories under stated common inputs or closed-loop dynamics, in a specified metric and domain. A differential contraction inequality needs regularity, metric bounds, and domain/existence conditions to imply the corresponding finite-distance result. Contraction does not by itself prove input feasibility or a desired contact outcome.

Control Contraction Metric (CCM)

A positive-definite differential metric satisfying control-dependent contraction conditions that support feedback construction for feasible reference trajectories. The result depends on the theorem's domain, regularity, metric bounds, and controller construction. Actuator saturation and state constraints require explicit treatment.

Degrees of Freedom (DOF)

The number of locally independent coordinates needed to specify a configuration. Independent constraints can reduce it; redundant parameters can exceed it. The number of actuators and the dimension of the full state space are different quantities.

Drift

The evolution of the declared effective plant when the declared control channel is zero. It includes the configuration-velocity kinematics and all retained uncontrolled dynamics, such as gravity, inertial coupling, passive elasticity, and damping. Its meaning depends on the chosen state, control channel, and environment; zero control does not mean zero acceleration or zero velocity.

Drift-Control Ratio (DCR)

A model-based comparison of drift and bounded control authority in a declared acceleration

or task-projected space:

$$\text{DCR}_{W,\mathcal{U}}(x) = \frac{\|W a_{\text{drift}}(x)\|_2}{\sup_{u \in \mathcal{U}(x)} \|W B_a(x)u\|_2 + \varepsilon}.$$

State the projection/scaling W , admissible input set $\mathcal{U}$, acceleration input map B_a , and regularizer ε , with compatible units. This scalar magnitude ratio omits directional alignment and does not by itself establish controllability, cancellation feasibility, or a universal coaching threshold.

Euler Angles

Three parameters describing an orientation using a declared sequence of rotations. Different sequences have different singular sets. Angle rates are related to angular velocity by a configuration-dependent map and are not generally its Cartesian components.

Exponential Map

For a Lie group, the map from a Lie-algebra element to the group element reached at parameter one along its generated one-parameter subgroup; for a matrix Lie group it is the matrix exponential. It need not be globally one-to-one. It is not general integration of an arbitrary time-varying velocity. The Riemannian exponential defined by geodesics is a related but distinct construction requiring a connection or metric.

Flow

The solution map of an ordinary differential equation where solutions exist uniquely. For an autonomous smooth vector field it forms a local one-parameter family with the composition property. Time-dependent systems use a two-time solution map; hybrid resets need not belong to a single smooth flow.

Gimbal Lock

A singularity of an Euler-angle rate-to-angular-velocity map. The parameterization loses local rank; the rigid body's physical rotational freedom does not disappear. Quaternions or overlapping coordinate charts handle orientation differently and have their own representation conventions.

Jacobian

The derivative of a map expressed in coordinates. For a task map $y = h(q)$, $\dot{y} = J(q)\dot{q}$. Acceleration also contains $\dot{J}\dot{q}$. Spatial and body Jacobians use different frame conventions, which must match the represented velocities and forces.

Kinematics

Description of position, orientation, velocity, and acceleration relationships without using forces to determine the motion. A kinematically allowed path is not automatically dynamically feasible under available actuation.

Lagrangian Mechanics

A formulation using a Lagrangian, often $L = T - U$, together with applied generalized forces and any constraints. For unconstrained coordinates, $d(\partial L/\partial \dot{q})/dt - \partial L/\partial q = Q$. Dissipation, actuation, nonconservative loads, and contact must be specified; kinetic and potential energy alone do not determine a forced system's motion.

Lie Algebra

A vector space equipped with a bilinear antisymmetric bracket satisfying the Jacobi identity. For a Lie group, it is the tangent space at the identity with the induced bracket. Brackets of smooth vector fields are another important example; accessibility arguments require the appropriate system and theorem assumptions.

Lie Group

A smooth manifold that is also a group, with smooth multiplication and inversion. Examples include the rotation group $SO(3)$ and rigid-transformation group $SE(3)$. A choice of coordinates represents these objects without changing the group itself.

Lyapunov Stability

For an equilibrium, the property that sufficiently small initial perturbations remain small for all future times of the stated system. It does not require return to equilibrium. Asymptotic stability additionally requires attraction; orbital and finite-horizon stability statements refer to different targets and must be defined separately.

Manifold

A space locally described by Euclidean coordinates, with the usual topological regularity conditions. A smooth manifold also has smoothly compatible charts. Its dimension is local and does not depend on how many coordinates are used to embed or redundantly represent it.

Screw Axis

For a rigid-body twist with nonzero angular velocity, an instantaneous axis and pitch describing rotation together with translation along that axis. Pure translation is a limiting case, often represented as an axis at infinity; a finite unique rotational axis is then unavailable.

State Space

The set of variables sufficient to determine future evolution under specified inputs and a model. Mechanical models often include configuration and velocity or momentum, and may additionally include actuator states, flexible modes, estimators, or discrete contact modes.

Tangent Bundle

The union of tangent spaces $T_q Q$ over configurations $q \in Q$, with its bundle and smooth-manifold structure. For an n -dimensional smooth configuration manifold, TQ has dimension $2n$. It represents configuration-velocity states for a basic mechanical model; momentum states instead lie in the cotangent bundle T^*Q .

Tangent Space

The vector space of instantaneous derivatives of smooth curves through a point on a manifold. Its vectors are local velocities, not finite displacements between arbitrary configurations. Applied differential constraints can restrict the admissible velocities further.

Twist

A six-component representation of an instantaneous rigid-body velocity field, combining angular and linear terms with a declared reference point and frame. In one frame, $v(r) = v_0 + \omega \times (r - r_0)$. Spatial and body representations are related by the appropriate adjoint transformation; the linear term must not be confused with the velocity of an unspecified point.

Wrench

The dual force representation consisting of a moment about a declared reference point and a resultant force in a declared frame. With matching twist conventions, mechanical power is $\tau \cdot \omega + F \cdot v_0$. Changing point or frame requires the corresponding moment shift and dual transformation so that power is unchanged.

Zero-Torque Counterfactual (ZTCF) Family

Model interventions setting the declared applied generalized-control channel to zero while preserving the declared effective plant. Pointwise samples, stitched traces, forward trajectories, and branched trajectories answer different questions. A stitched pointwise acceleration trace is not automatically a feasible unforced trajectory.

Zero-Velocity Counterfactual (ZVCF)

The instantaneous acceleration after the declared velocity and control intervention:

$$a_{\text{ZVCF}}(q, z) = -M(q)^{-1}h(q, 0, z_0; p).$$

Specify the auxiliary-state reset z_0 , retained loads, parameters, and contact mode. This separates the retained configuration-dependent loads under that intervention; it is not necessarily gravity alone, and it is not a forward simulation of a stationary body.

Further Reading and References

The following books develop complementary parts of the framework used here. Choose a source for the particular claim or calculation being investigated, and consult its assumptions, notation, edition, and errata. A robotics or control theorem does not become evidence about human golf performance merely because it can be applied to a simplified swing model.

Geometry, Kinematics, and Mechanical Models

- **A Mathematical Introduction to Robotic Manipulation**, Murray, Li, and Sastry [31], treats rigid-body motion, manipulator and hand kinematics, dynamics, and nonholonomic motion planning. The [authors' companion site](#) provides chapter information and corrections. Use it to connect Lie groups and Jacobians to explicitly defined frames and mechanical systems.
- **Geometric Control of Mechanical Systems**, Bullo and Lewis [5], develops differential-geometric modeling and control of mechanical systems. Its [companion page](#) includes contents, selected material, examples, and errata. This is a suitable route into connections, mechanical control structures, and the hypotheses behind geometric analysis.
- **Robot Modeling and Control**, Spong, Hutchinson, and Vidyasagar [40], covers kinematics, Jacobians, trajectory planning, dynamics, joint and multivariable control, force control, and vision-based methods. The [publisher's chapter listing](#) identifies the scope of the cited edition. Its robot models still require adaptation and validation before supporting a biomechanical interpretation.

Computational Rigid-Body Dynamics

- **Rigid Body Dynamics Algorithms**, Featherstone [12], develops spatial-vector conventions and recursive dynamics algorithms. The [author's teaching materials](#) and [software resources](#) support implementation. Distinguish the linear-cost articulated-body calculation for a tree from the additional work needed for closed loops, constraints, and contacts.

Nonlinear Stability and Feedback

- **Applied Nonlinear Control**, Slotine and Li [39], develops nonlinear stability and feedback methods. [Slotine's accompanying course materials](#) identify the relevant book sections and separate the later contraction lectures and papers. Check the domain and model conditions of a stability argument before extending it to saturated or hybrid systems.
- **Contraction Theory for Dynamical Systems**, Bullo [4], develops contractivity and incremental stability using specified norms and metrics. The [author's book page](#) provides version information and materials. The bibliography cites the 2024 edition; later revisions may reorganize or extend the treatment. Contraction claims must retain their metric, input, domain, and existence assumptions.

The chapter-level primary references supply more specialized results for transverse dynamics, phase constraints, nonlinear funnels, stochastic models, and learning. Keep mathematical derivations, numerical verification, model calibration, and experimental validation identifiable when combining those sources.

Index

- 2R arm
 - manipulability, 179
- 7-DOF arm
 - null-space optimization, 179
- A Matrix, 37
- Adjoint, 103
- angular velocity, 81
 - body frame, 81
 - spatial frame, 81
- articulated inertia
 - floating base, 190
- Asymptotic Stability, 40
- Axis-Angle, 125
- axis-angle representation, 82
- B Matrix, 37
- Baker-Campbell-Hausdorff, 125
 - higher-order terms, 115
- Basis, 5
- bilateral constraints, 138
- Center, 40
- centroidal momentum, 190
- Characteristic Equation, 12
- Chasles Theorem, 103
- Chasles' theorem
 - proof, 95
- Cholesky Decomposition, 16
- Cholesky decomposition, 26, 235
- complementarity condition, 196
- condition number, 25
- Configuration Space, 53
- congruence transformation, 23
- constraint
 - bilateral, 138
- constraint handling, 138
- contact dynamics, 196
- contact forces, 138
- Controllability, 32, 41
- Controllable System, 41
- coordinate frame
 - body frame, 68
 - spatial frame, 68
- coordinate-free formulation, 23
- Coulomb friction, 138, 196
- Cross Product, 10
- curiosity-driven exploration, 232
- DDPG, 232
- Deep Deterministic Policy Gradient, 232
- Deep Lagrangian Networks, 235
- Deep Q-Network, 231
- Degrees of Freedom, 53
- DeLaN, 235
- Determinant, 7
- diffeomorphism, 210
- Dimension, 5
- discrete-time systems, 45
- distribution
 - smooth, 61
- Dot Product, 8
- DQN, 231
- dual space, 95
- duality
 - wrench-twist, 96
- dyad, 20
- dyadic, 21
- Eigenvalue, 11
- Eigenvector, 11
- energy error
 - bounded, 216
- entropy regularization, 232
- epsilon-greedy, 232
- Equilibrium Point, 38
- Euler angles, 72
- experience replay, 231
- exploration strategies, 232
- exploration-exploitation tradeoff, 232
- Exponential Map, 125
- floating-base dynamics, 189
- Focus, 40
- force at offset, 95
- frame invariance

- power, 96
- friction cone, 138, 196
- Frobenius norm, 24
- Frobenius theorem, 61
- Geodesics, 125
- geodesics
 - Riemannian proof, 117
- gimbal lock, 73
- Hamiltonian Neural Network, 235
- HNN, 235
- humanoid robot
 - dynamics, 190
- inertia tensor, 22
- infinitesimal generator, 210
- Instability, 40
- integrability, 61
- intrinsic reward, 232
- Jacobian Matrix, 39
- kinematic redundancy, 179
- kinetic energy
 - spatial form, 159
- Lagrangian Mechanics, 201
- Lagrangian mechanics
 - RNEA equivalence, 137
- Lagrangian Neural Network, 235
- LCP, *see* linear complementarity problem
- leapfrog integrator, 216
- Lie algebra
 - definition, 74
- Lie bracket, 61
- Lie group
 - definition, 74
- Linear Complementarity Problem, 139
- linear complementarity problem, 196
- Linear Independence, 4
- Linear State Space, 37
- Linear Transformation, 6
- LNN, 235
- LU decomposition, 26
- Lyapunov
 - direct method, 49
- Lyapunov function, 48
- Lyapunov stability, 48
- manipulability, 177
- manipulability ellipsoid
 - force, 178
 - velocity, 178
- manipulability measure, 178
- Marginal Stability, 40
- mass matrix
 - column extraction, 138
- Mass-Spring-Damper, 43
- Matrix, 2
- matrix decomposition, 26
- Matrix Exponential, 40, 125
- matrix exponential
 - rotation, 75
- Matrix Logarithm, 125
 - numerical stability, 113
- matrix norm, 24
- maximum entropy, 232
- Mechanical Power, 103
- mechanical power, 96
- Noether charge, 210
- Noether's theorem
 - rigorous proof, 210
- non-holonomic constraint, 63
- Null Space, 7
- null space
 - of Jacobian, 179
- null-space motion, 179
- null-space projector, 179
- Observability, 32, 42
- Observable System, 42
- Orthogonality, 9
- outer product, 20
- Output Equation, 38
- pendulum
 - Lyapunov analysis, 49
- Phase Portrait, 39
- Positive Definite, 15
- power
 - inner product, 96
- pure couple, 95
- pure force, 95
- Python
 - quaternion, 82
- Q-learning, 231
- QR decomposition, 26

- Quadratic Form, 15
- quaternion
 - algebra, 76
 - definition, 76
 - double cover, 78
 - rotation representation, 77
 - unit quaternion, 77
- Rank, 6
- Rank-Nullity Theorem, 7
- reachable workspace, 178
- Reciprocal Screws, 103
- redundancy, 179
- Riemannian metric
 - $SO(3)$, 117
- RNEA, 137
 - Lagrangian equivalence, 137
- Rodrigues Formula, 125
- Rodrigues' rotation formula, 75, 76
- rotation
 - history
 - Euler, 68
 - Hamilton, 68
- rotation matrix
 - definition, 69
 - quaternion conversion, 78
 - R_x , 70
 - R_y , 70
 - R_z , 71
- rotation representations
 - comparison, 79
- SAC, 232
- Saddle Point, 40
- screw
 - pitch, 95
- Screw Axis, 103
- $SE(3)$, 103
- $se(3)$ Lie Algebra, 125
- Signorini conditions, 196
- Singular Value, 17
- Singular Value Decomposition, 16
- Singularity, 6
- Skew-Symmetric Matrix, 10
- SLERP, 80, 125
 - Python example, 82
- $SO(3)$
 - definition, 71
- $so(3)$ Lie Algebra, 125
- Soft Actor-Critic, 232
- spatial inertia, 159
 - kinetic energy derivation, 159
 - quadratic form, 159
- spectral norm, 25
- Spectral Theorem, 14
- Spiral, Stable, 40
- Spiral, Unstable, 40
- Stability, 39
- stability
 - discrete-time, 45
- Stable Node, 40, 44
- Stable Spiral, 43
- State Space, 31
- State Space Representation, 37
- stiffness matrix, 22
- Stormer-Verlet, 216
- stress tensor, 22
- structure-preserving neural networks, 235
- SVD
 - of Jacobian, 178
- symplectic Euler, 216
- symplectic form, 216
- symplectic integrator, 216
- target network, 231
- TD3, 232
- time-stepping methods, 197
- Trajectory, 36
- Transfer Function, 32, 44
- Twist, 103
- unilateral constraints, 138
- Unstable Node, 40
- value-based methods, 231
- Vector, 2
- Vector Space, 3
- virtual joint, 190
- Virtual Work, 103
- workspace analysis, 177
- Wrench, 103
 - geometric object, 95
 - offset force, 95
 - physical interpretation, 95
 - pure couple, 95
 - pure force, 95
- Yoshikawa manipulability, 178
- zero-order hold, 45

Bibliography

- [1] Vladimir I. Arnold. *Mathematical Methods of Classical Mechanics*. Springer, 2nd edition, 1989.
- [2] Robert Stawell Ball. *A Treatise on the Theory of Screws*. Cambridge University Press, 1900. Legacy duplicate key retained for backwards compatibility.
- [3] R. W. Brockett. Robotic manipulators and the product of exponentials formula. *Mathematical Theory of Networks and Systems*, pages 120–129, 1984. Legacy duplicate key retained for backwards compatibility.
- [4] F. Bullo. *Contraction Theory for Dynamical Systems*. Kindle Direct Publishing, 2024.
- [5] F. Bullo and A. D. Lewis. *Geometric Control of Mechanical Systems*. Springer, 2004.
- [6] Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David Duvenaud. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- [7] Miles Cranmer, Sam Greydanus, Stephan Hoyer, Peter Battaglia, David Spergel, and Shirley Ho. Lagrangian neural networks. *ICLR Workshop on Integration of Deep Neural Models and Differential Equations*, 2020. Neural networks that learn Lagrangian structure for physical systems.
- [8] George Cybenko. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems*, 2:303–314, 1989.
- [9] Jacques Denavit and Richard S. Hartenberg. A kinematic notation for lower-pair mechanisms based on matrices. *Journal of Applied Mechanics*, 22(2):215–221, 1955. Legacy duplicate key retained for backwards compatibility.
- [10] R. Featherstone. The calculation of robot dynamics using articulated-body inertias. *International Journal of Robotics Research*, 2(1):13–30, 1983.
- [11] R. Featherstone. *Robot Dynamics Algorithms*. Kluwer Academic Publishers, 1987. Legacy duplicate key retained for backwards compatibility.
- [12] R. Featherstone. *Rigid Body Dynamics Algorithms*. Springer, 2008.
- [13] Roy Featherstone. Spatial vector and rigid-body dynamics software: Spatial_v2, 2012. Reference algorithms and coordinate-transform documentation; accessed 2026-09-07.
- [14] Herbert Goldstein, Charles Poole, and John Safko. *Classical Mechanics*. Addison-Wesley, 3rd edition, 2002.
- [15] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. Legacy duplicate key retained for backwards compatibility.
- [16] Sam Greydanus, Misko Dzamba, and Jason Yosinski. Hamiltonian neural networks. *NeurIPS*, 2019. Neural networks that learn to conserve Hamiltonian structure.

- [17] Brian C. Hall. *Lie Groups, Lie Algebras, and Representations*. Springer, 2nd edition, 2015. Legacy duplicate key retained for backwards compatibility.
- [18] R. E. Kalman. A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1):35–45, 1960.
- [19] H. K. Khalil. *Nonlinear Systems*. Prentice Hall, 3rd edition, 2002.
- [20] O. Khatib. A unified approach for motion and force control of robot manipulators: The operational space formulation. *IEEE Journal on Robotics and Automation*, 3(1):43–53, 1987.
- [21] W. Lohmiller and J.-J. E. Slotine. On contraction analysis for non-linear systems. *Automatica*, 34(6):683–696, 1998.
- [22] J. Y. S. Luh, M. W. Walker, and R. P. C. Paul. On-line computational scheme for mechanical manipulators. *Journal of Dynamic Systems, Measurement, and Control*, 102(2):69–76, 1980. Legacy duplicate key retained for backwards compatibility.
- [23] Michael Lutter, Christian Ritter, and Jan Peters. Deep lagrangian networks: Using physics as model prior for deep learning. In *International Conference on Learning Representations*, 2019.
- [24] Kevin M. Lynch and Frank C. Park. Modern robotics: Controllability of wheeled mobile robots, part 4. Official companion video transcript; accessed 2026-09-07.
- [25] Kevin M. Lynch and Frank C. Park. Modern robotics: Numerical inverse kinematics, part 1. Official companion video transcript; accessed 2026-09-07.
- [26] Kevin M. Lynch and Frank C. Park. Modern robotics: Task space and workspace. Official companion video transcript; accessed 2026-09-07.
- [27] Kevin M. Lynch and Frank C. Park. *Modern Robotics: Mechanics, Planning, and Control*. Cambridge University Press, 2017.
- [28] I. R. Manchester and J.-J. E. Slotine. Control contraction metrics: Convex and intrinsic criteria for nonlinear feedback design. *IEEE Transactions on Automatic Control*, 62(6):3046–3053, 2017.
- [29] Jerrold E. Marsden and Tudor S. Ratiu. *Introduction to Mechanics and Symmetry*. Springer, 2nd edition, 1999. Legacy duplicate key retained for backwards compatibility.
- [30] MuJoCo Contributors. Computation: MuJoCo documentation, 2026. Online algorithm and contact-model documentation; accessed 2026-09-07.
- [31] R. M. Murray, Z. Li, and S. S. Sastry. *A Mathematical Introduction to Robotic Manipulation*. CRC Press, 1994.
- [32] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer, 2nd edition, 2006. Legacy duplicate key retained for backwards compatibility.
- [33] Pinocchio Contributors. Pinocchio: Rigid body dynamics algorithms and their analytical derivatives, 2026. Software description and feature documentation; accessed 2026-09-07.

- [34] Maziar Raissi, Paris Perdikaris, and George Em Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- [35] Olinde Rodrigues. Des lois géométriques qui régissent les déplacements d’un système solide. *Journal de Mathématiques Pures et Appliquées*, 5:380–440, 1840. Legacy duplicate key retained for backwards compatibility.
- [36] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. Legacy duplicate key retained for backwards compatibility.
- [37] A. S. Shiriaev, L. B. Freidovich, and S. V. Gusev. Transverse linearization for controlled mechanical systems with several passive degrees of freedom. *IEEE Transactions on Automatic Control*, 55(4):893–906, 2010.
- [38] Bruno Siciliano, Lorenzo Sciavicco, Luigi Villani, and Giuseppe Oriolo. *Robotics: Modelling, Planning and Control*. Springer, 2009. Legacy duplicate key retained for backwards compatibility.
- [39] J.-J. E. Slotine and W. Li. *Applied Nonlinear Control*. Prentice Hall, 1991.
- [40] M. W. Spong, S. Hutchinson, and M. Vidyasagar. *Robot Modeling and Control*. Wiley, 2006.
- [41] Gilbert Strang. *Linear Algebra and Its Applications*. Wellesley-Cambridge Press, 5th edition, 2019. Legacy duplicate key retained for backwards compatibility.
- [42] Richard S. Sutton, David McAllester, Satinder Singh, and Yishay Mansour. Policy gradient methods for reinforcement learning with function approximation. In *Advances in Neural Information Processing Systems*, volume 12, 1999.
- [43] Emanuel Todorov, Tom Erez, and Yuval Tassa. Mujoco: A physics engine for model-based control. In *IEEE/RSJ International Conference on Intelligent Robots and Systems*, pages 5026–5033, 2012. Legacy duplicate key retained for backwards compatibility.